

# Affordable Formal Verification for Practical Distributed Systems.

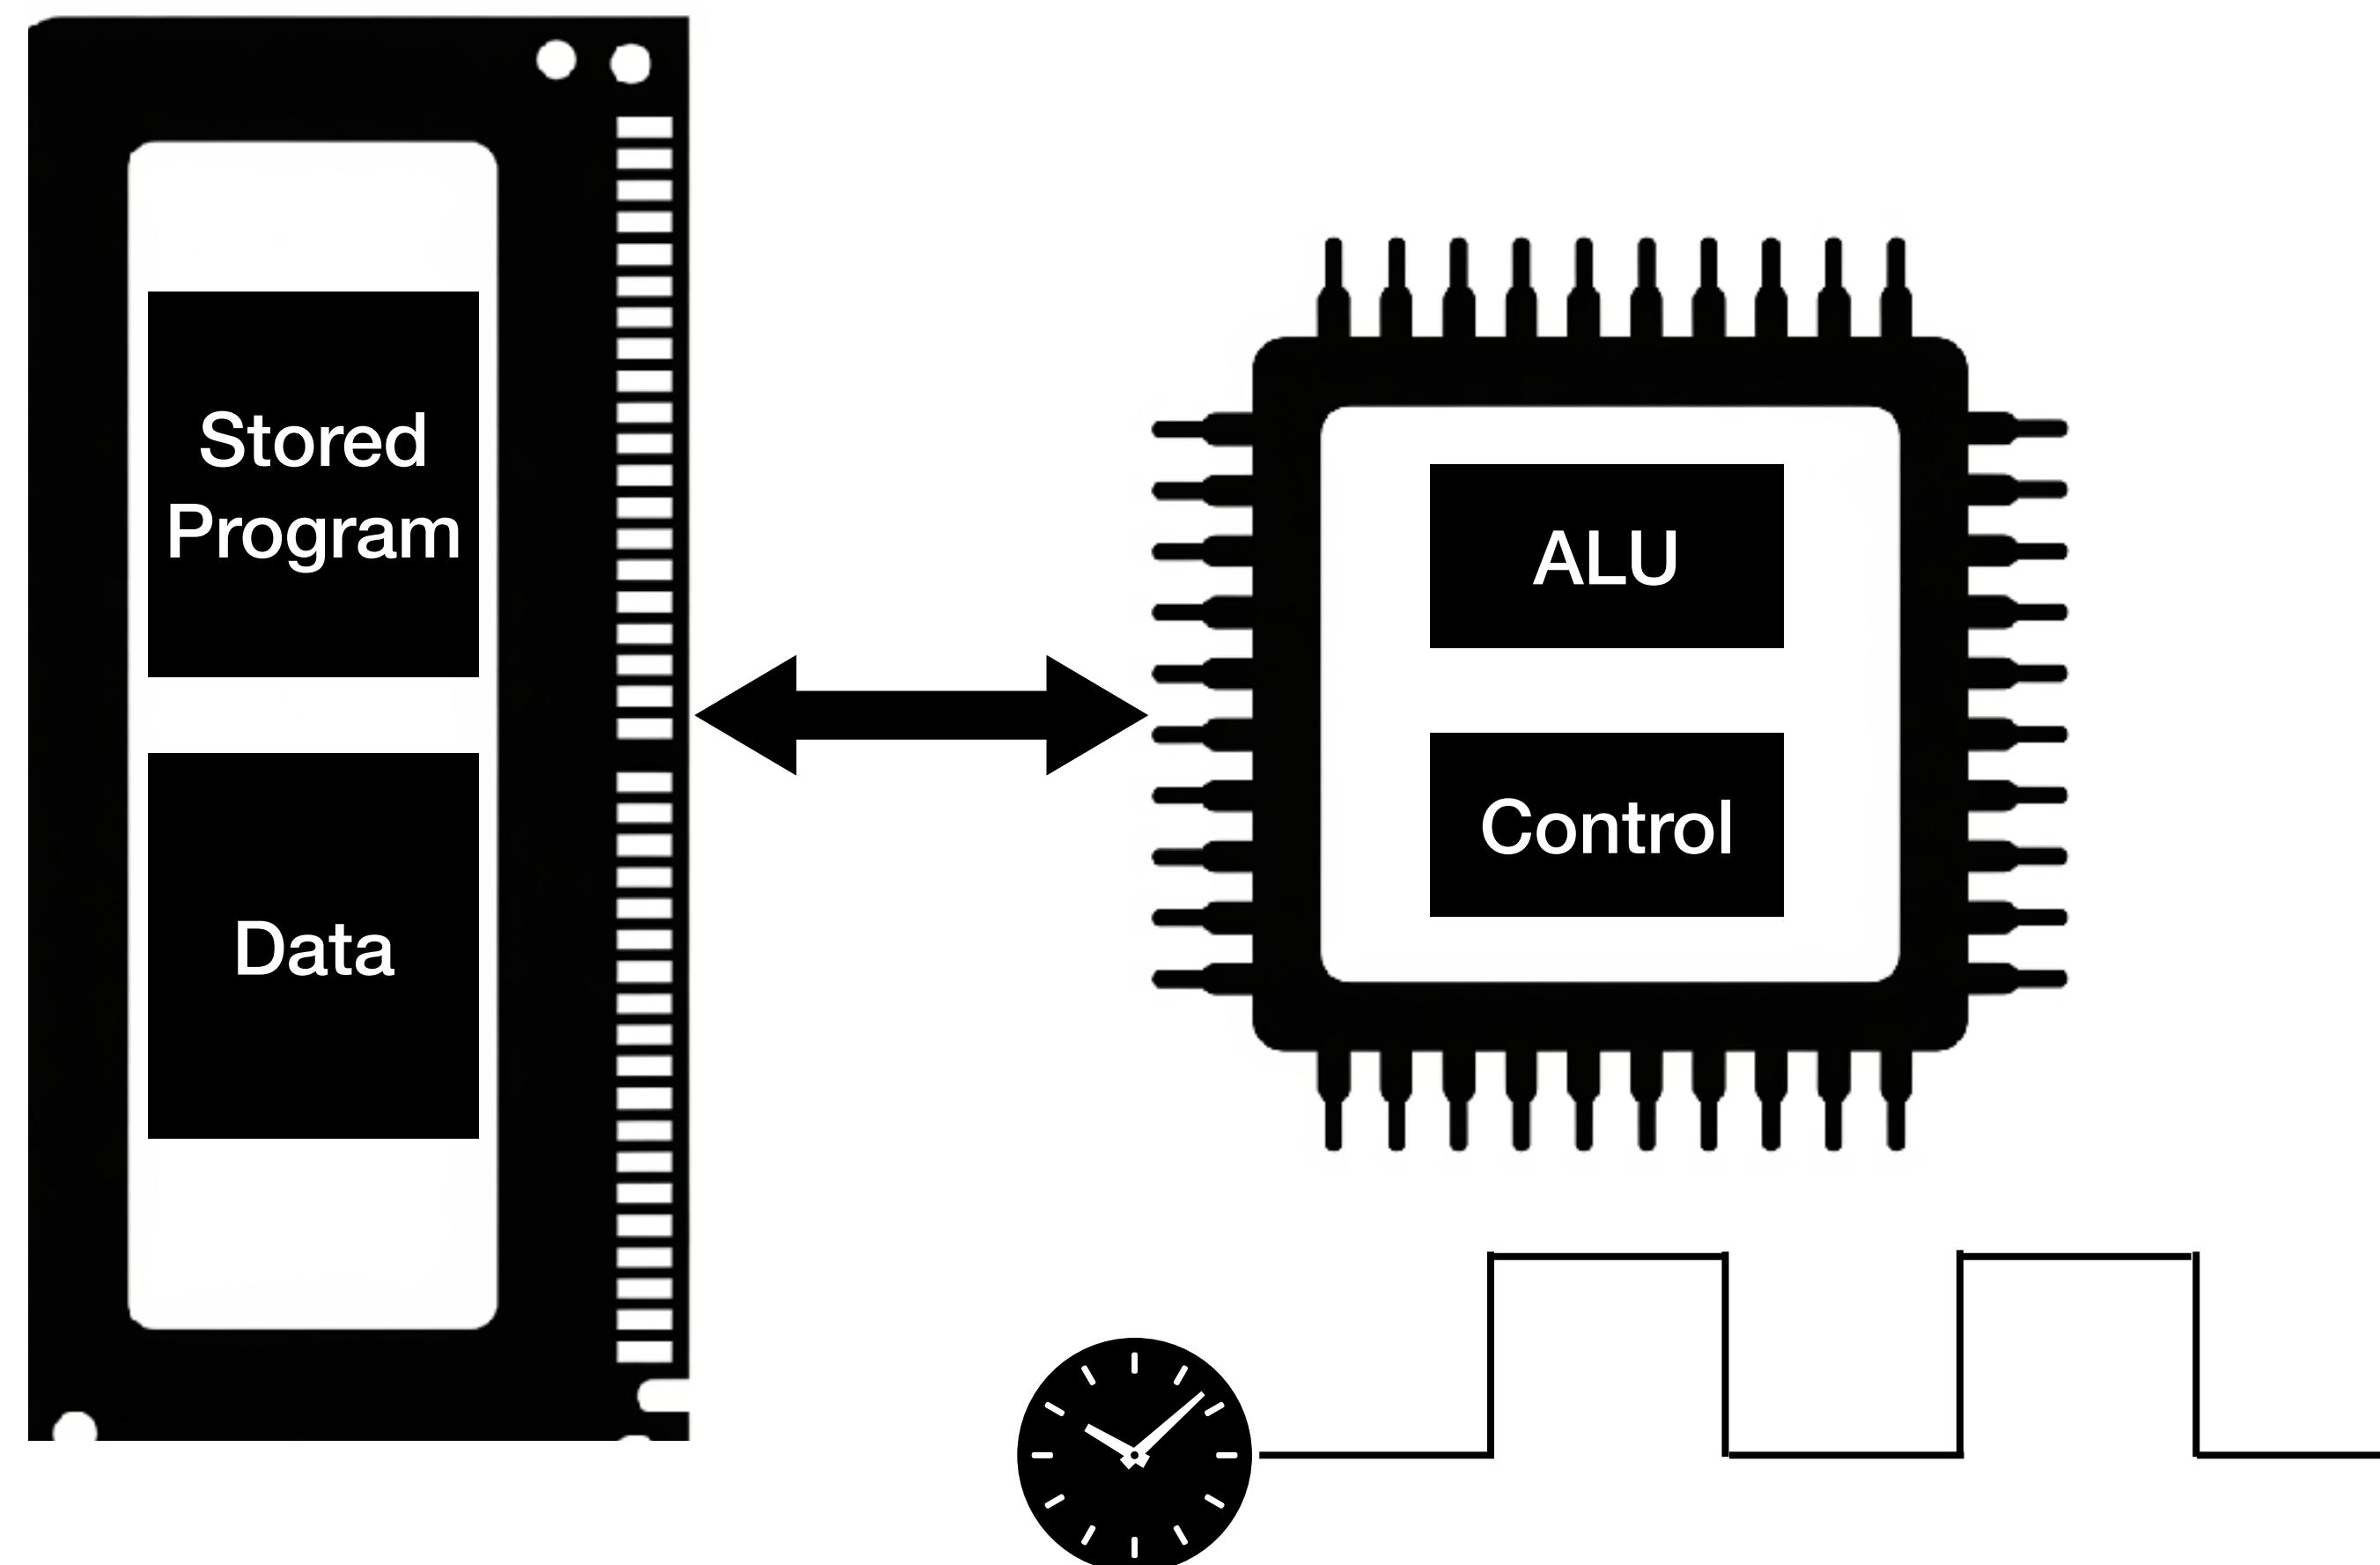
Gowtham Kaki  
University of Colorado Boulder



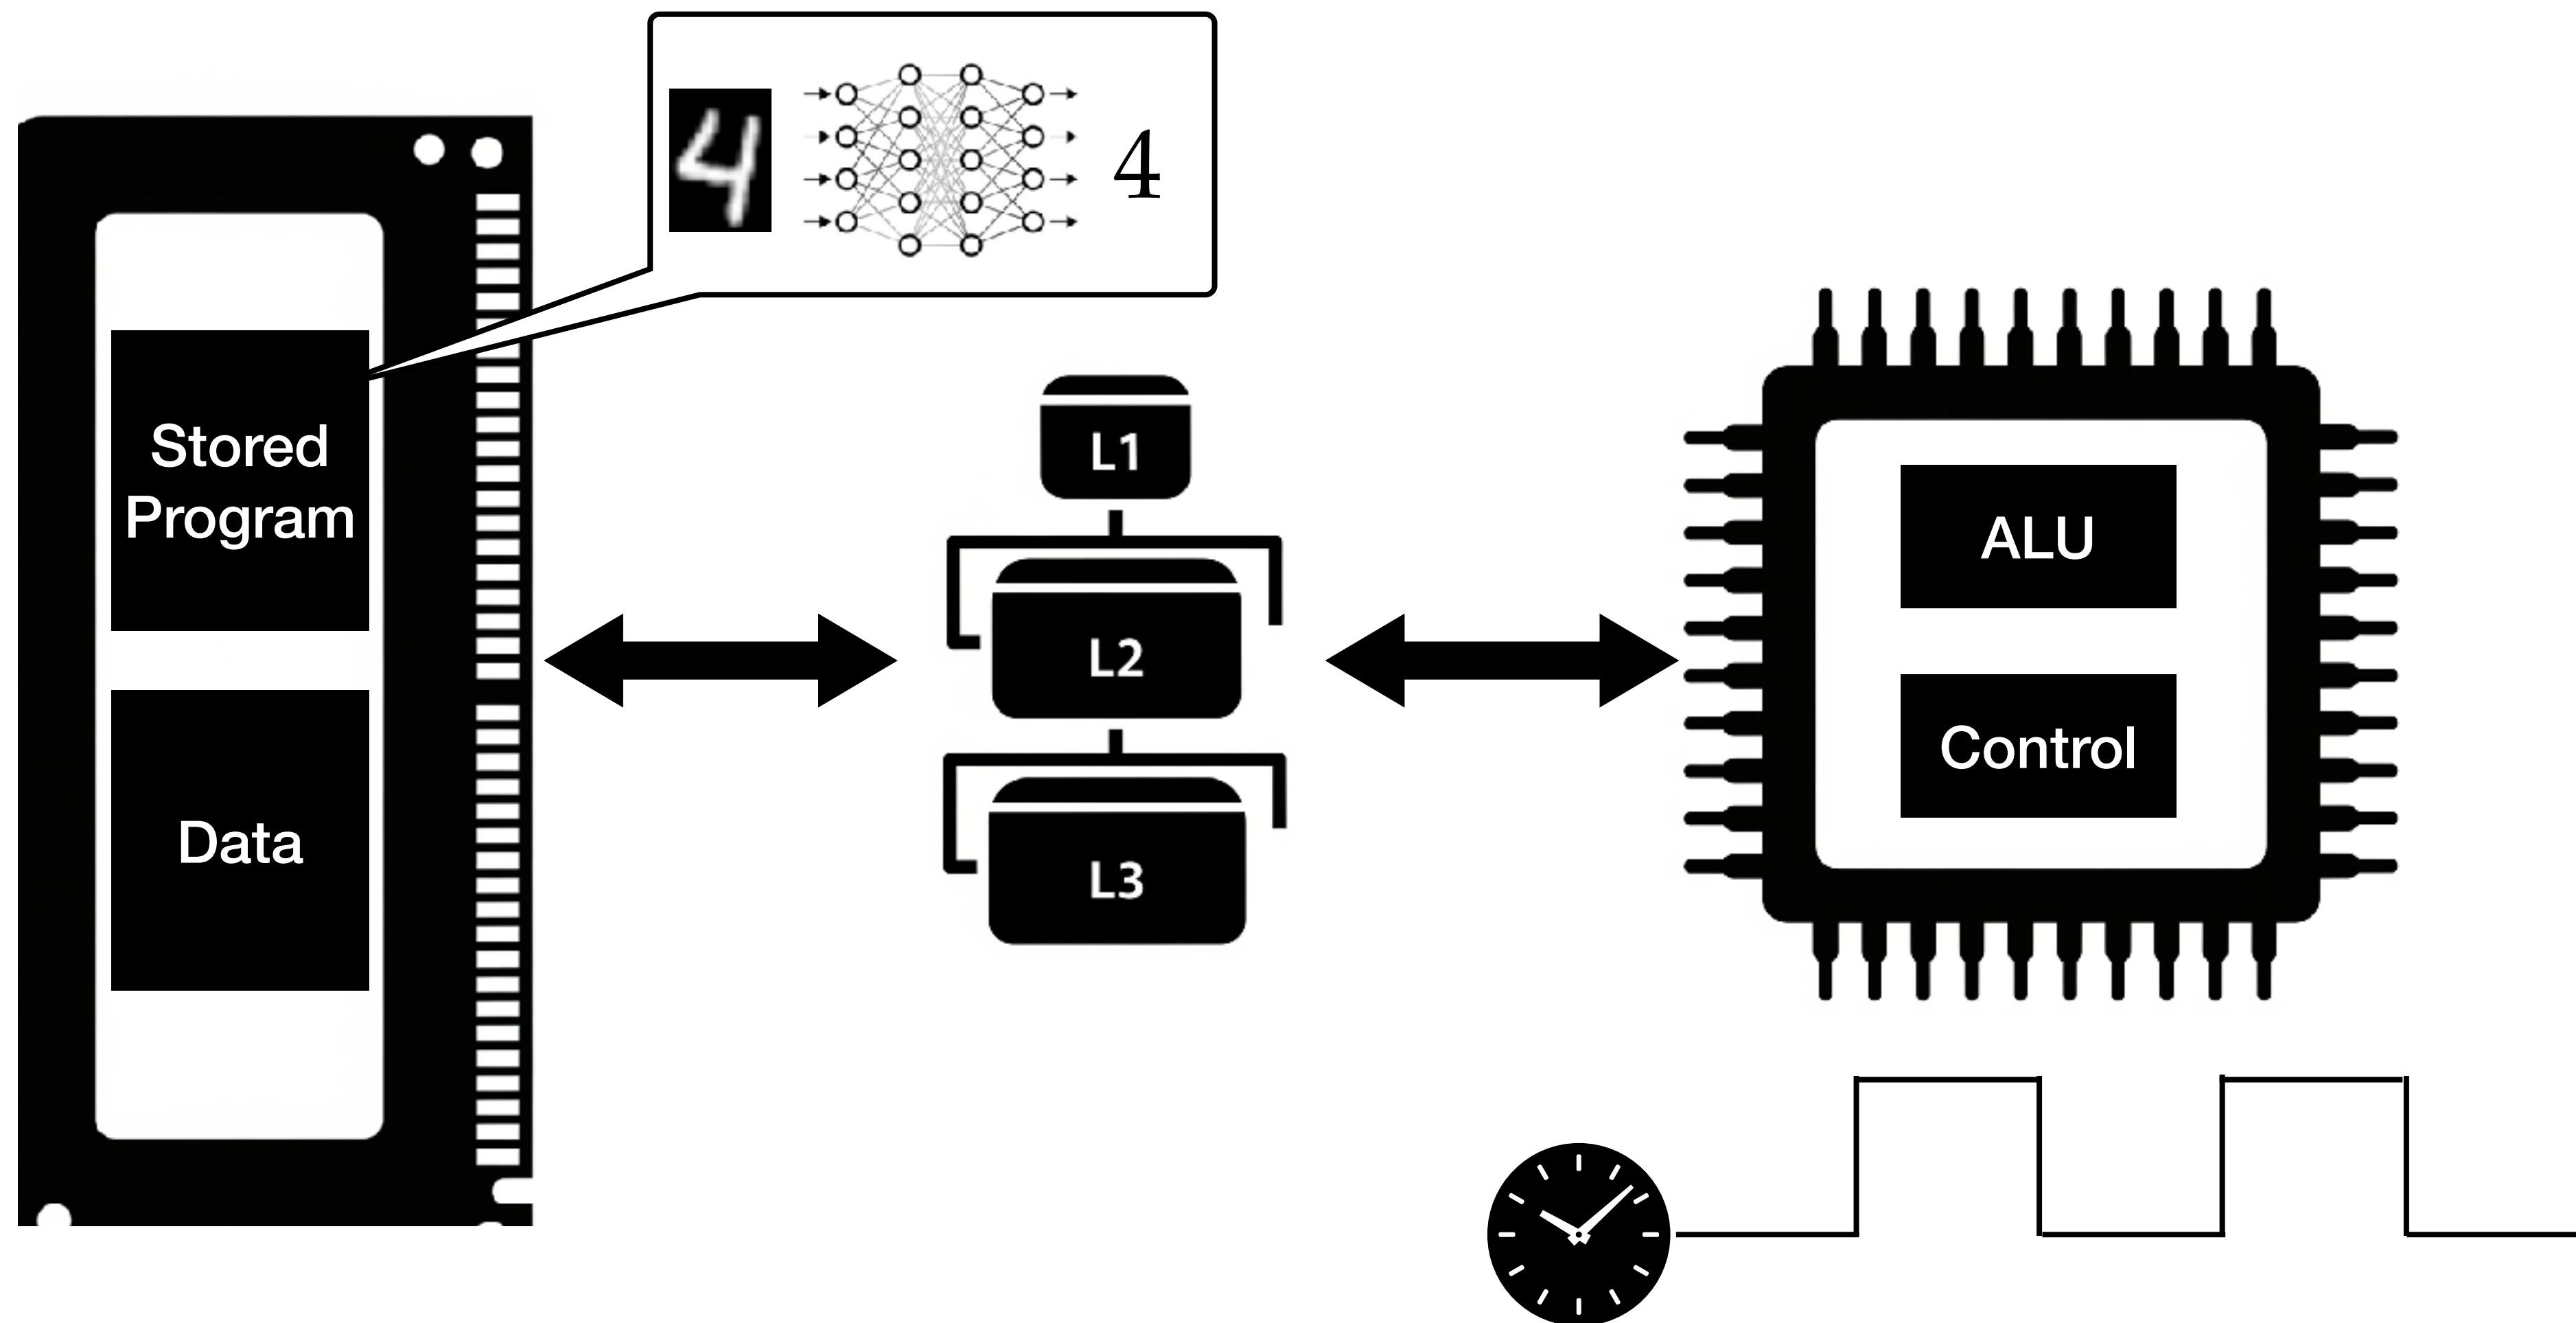
CU Programming Languages  
& Verification

# A von Neumann Computer

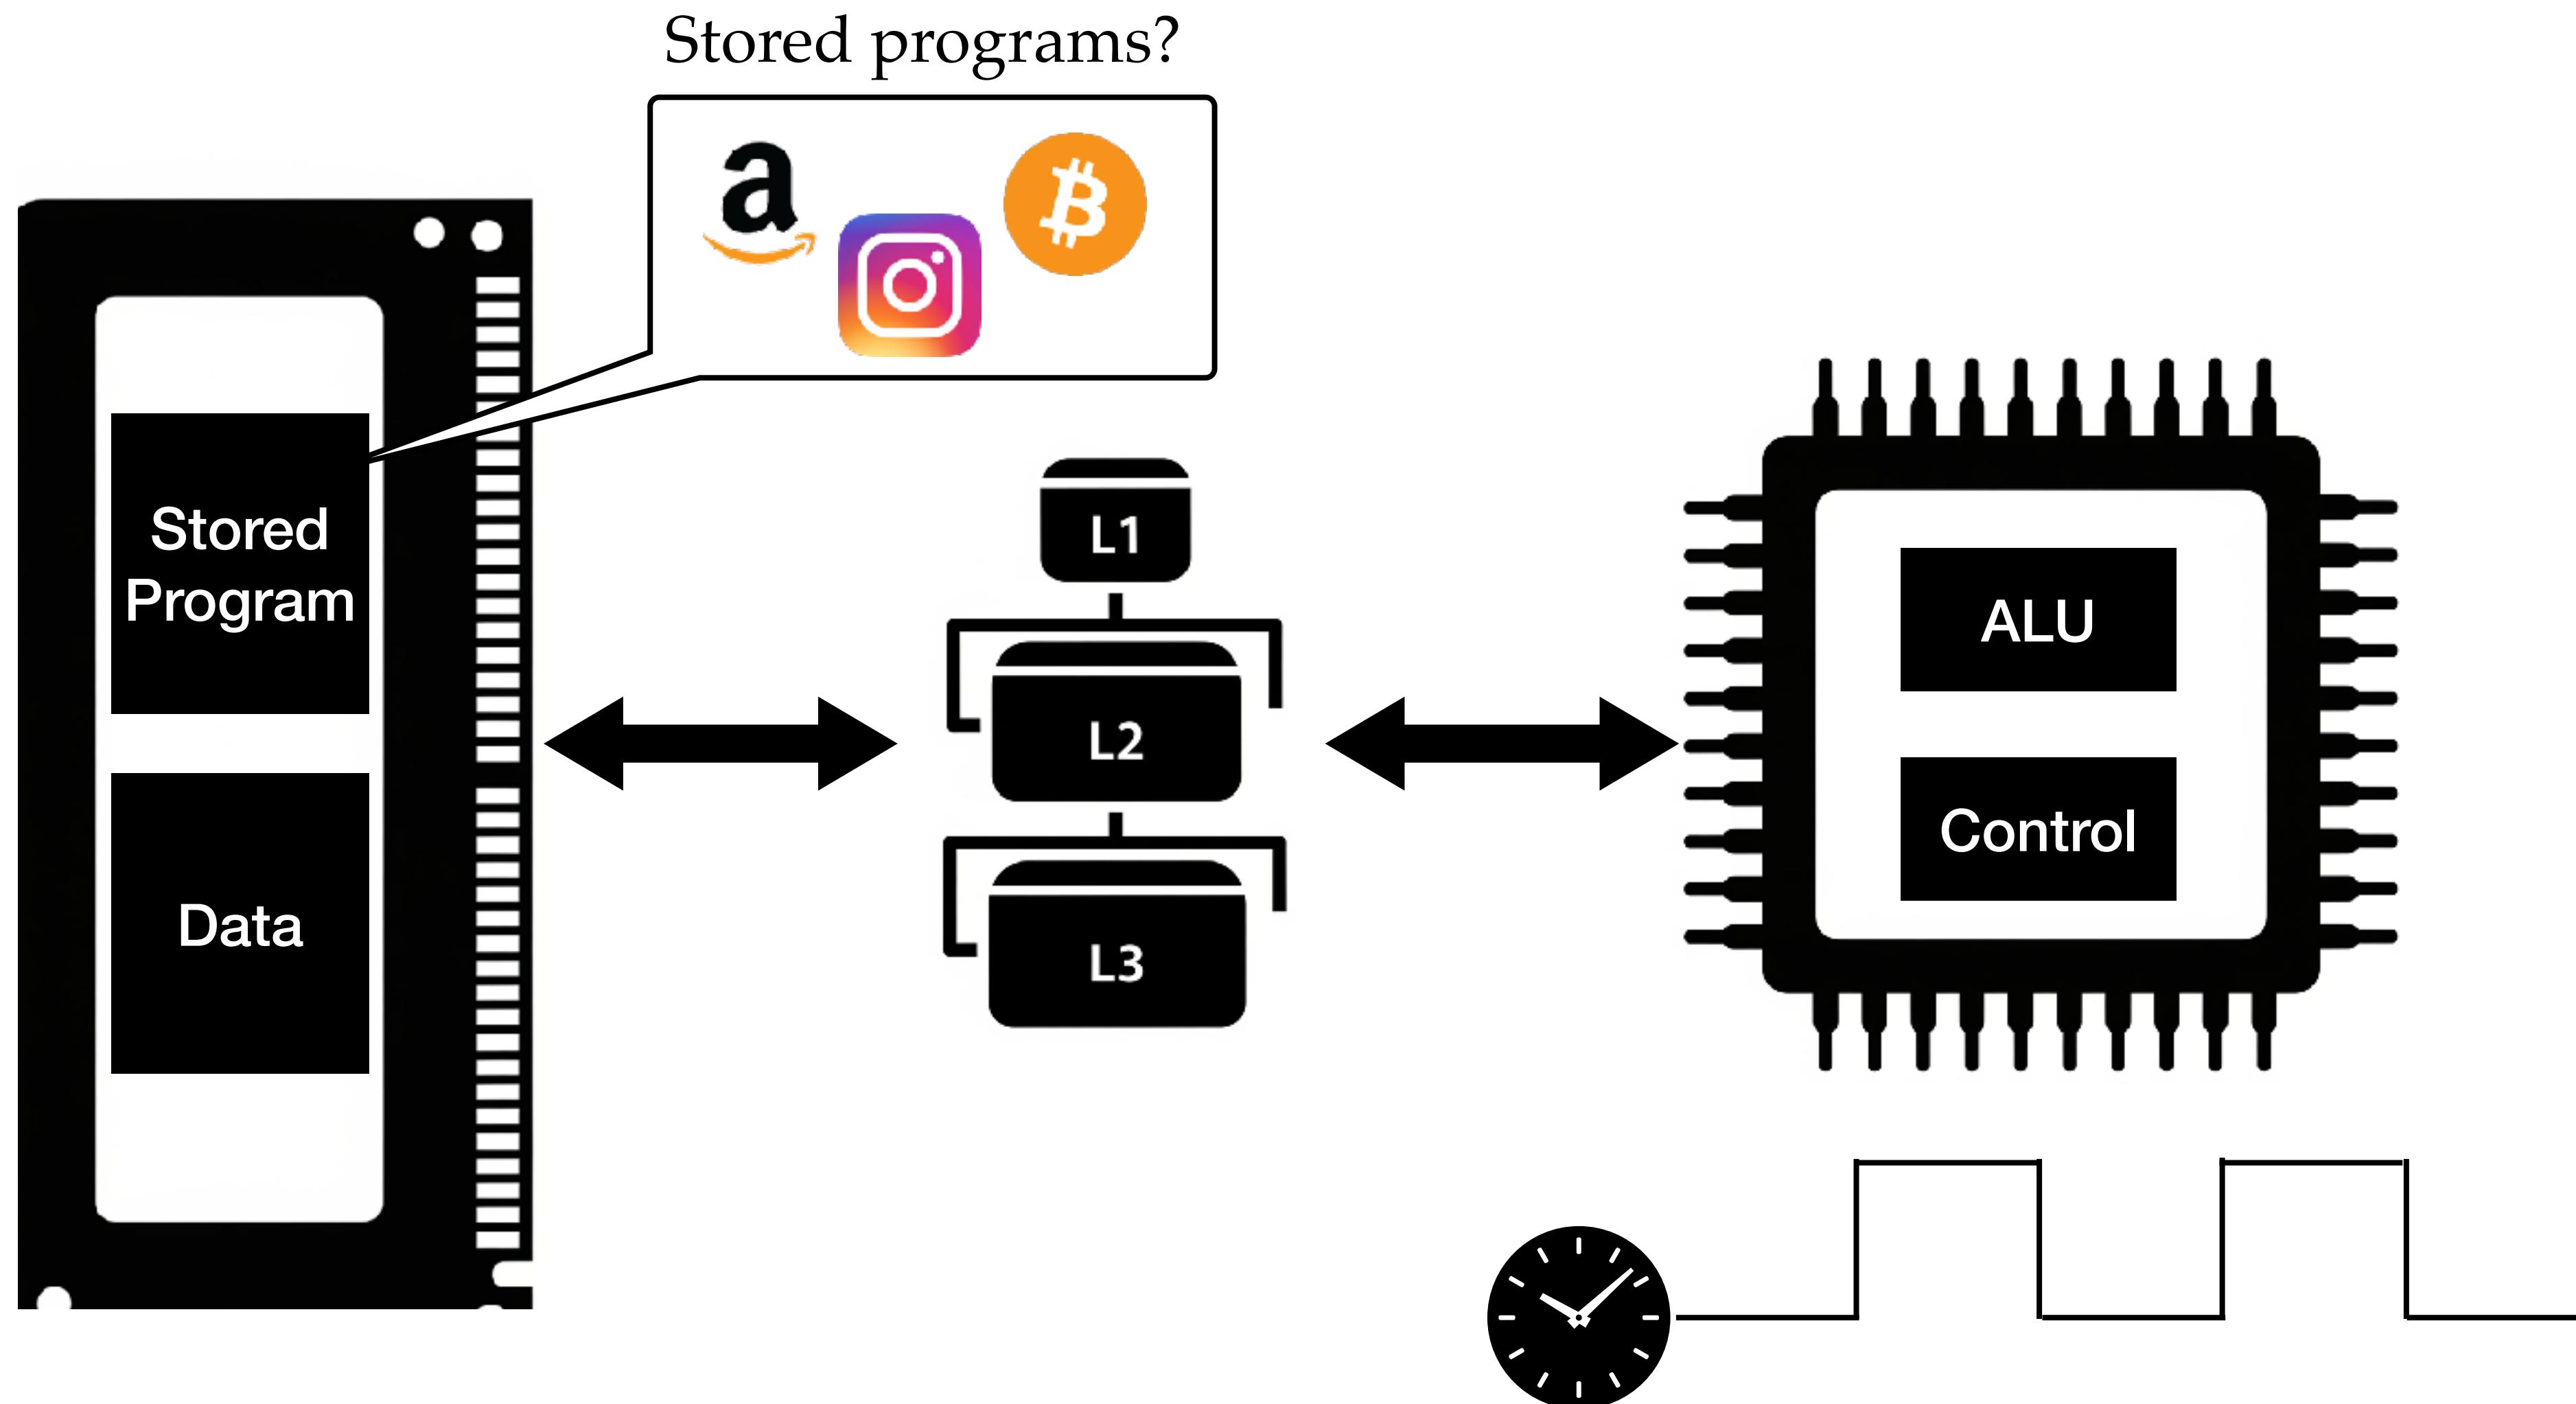
“A stored-program computer”



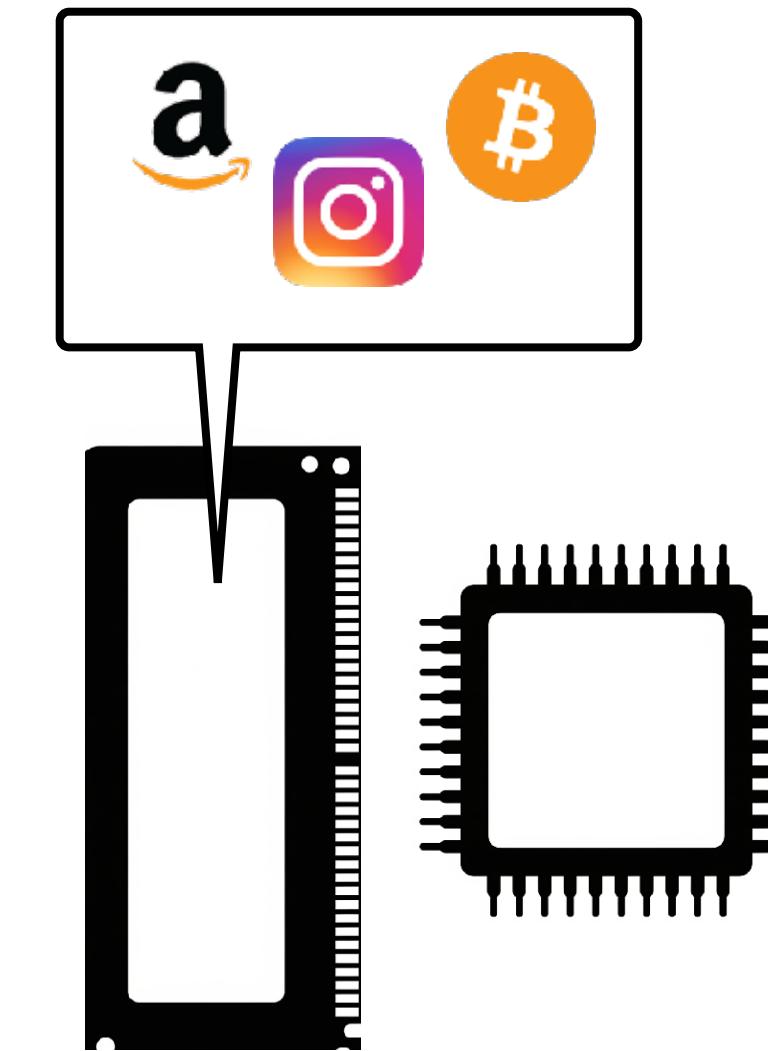
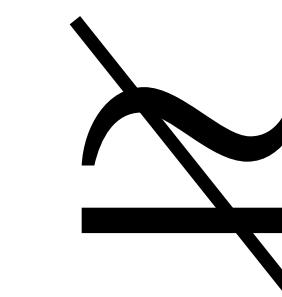
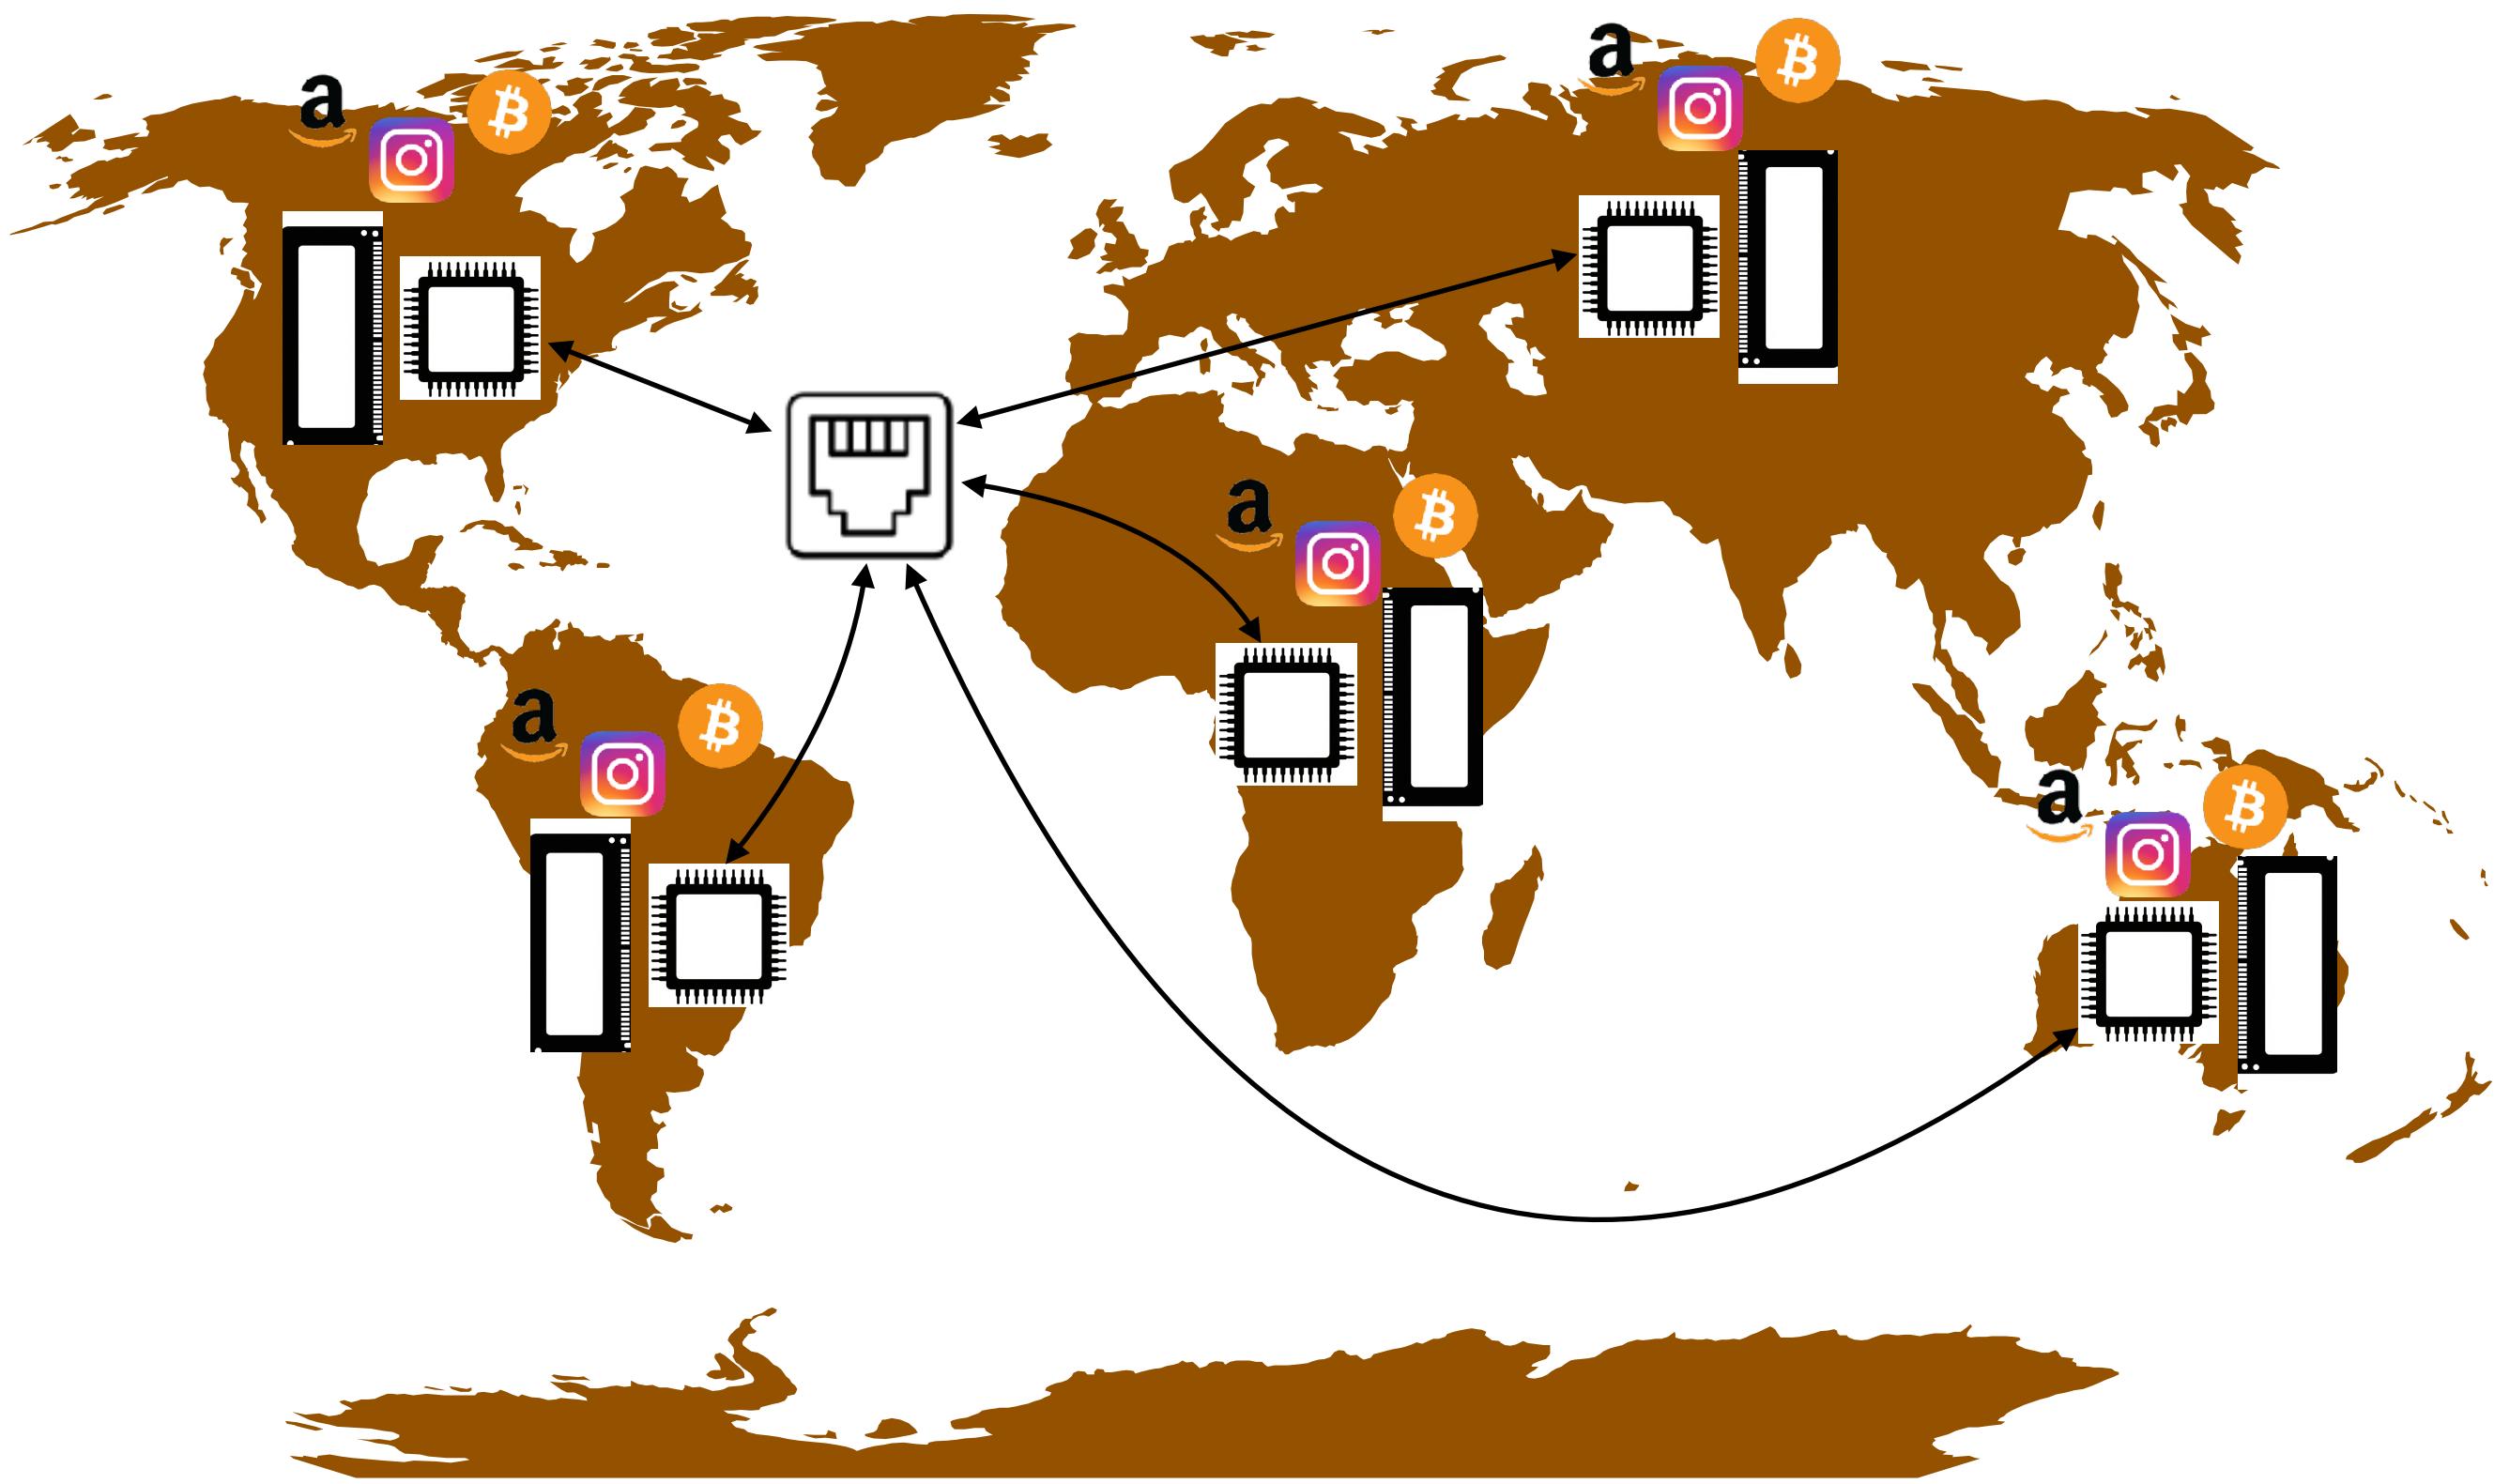
# A von Neumann Computer



# A von Neumann Computer

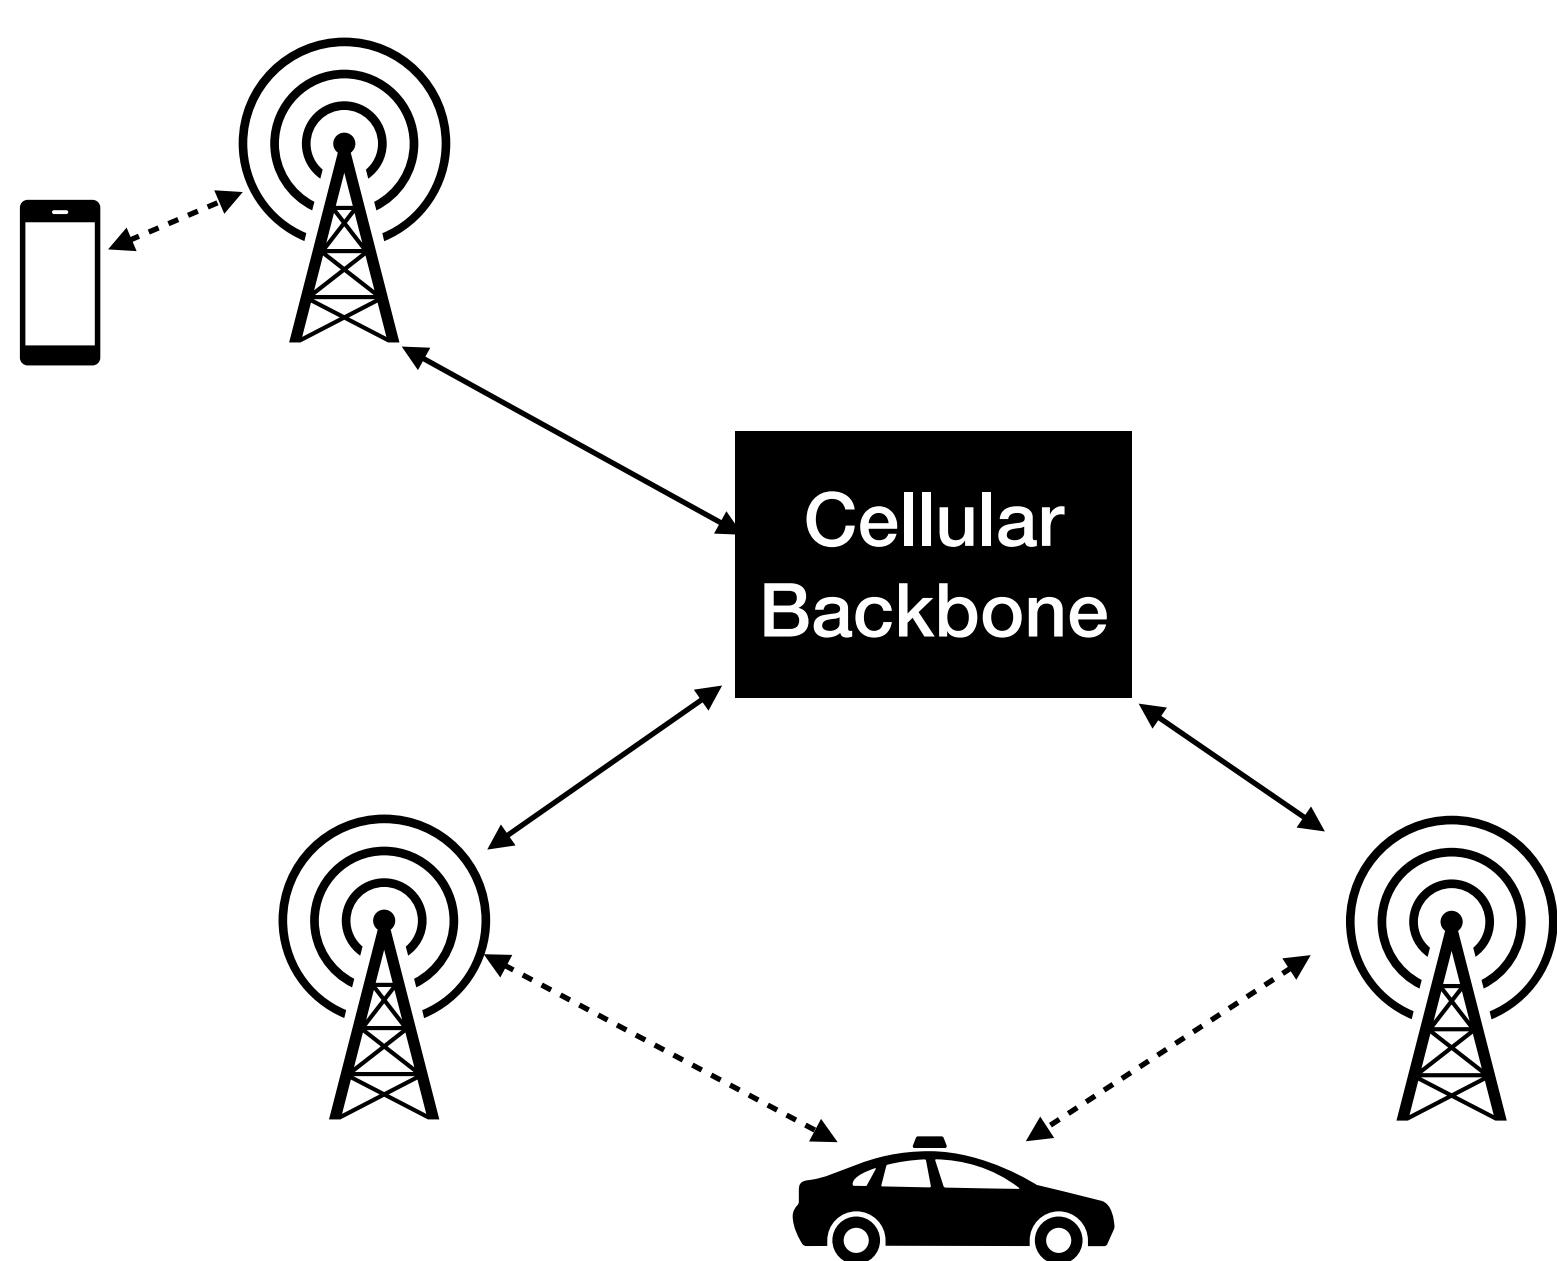


# A Distributed System

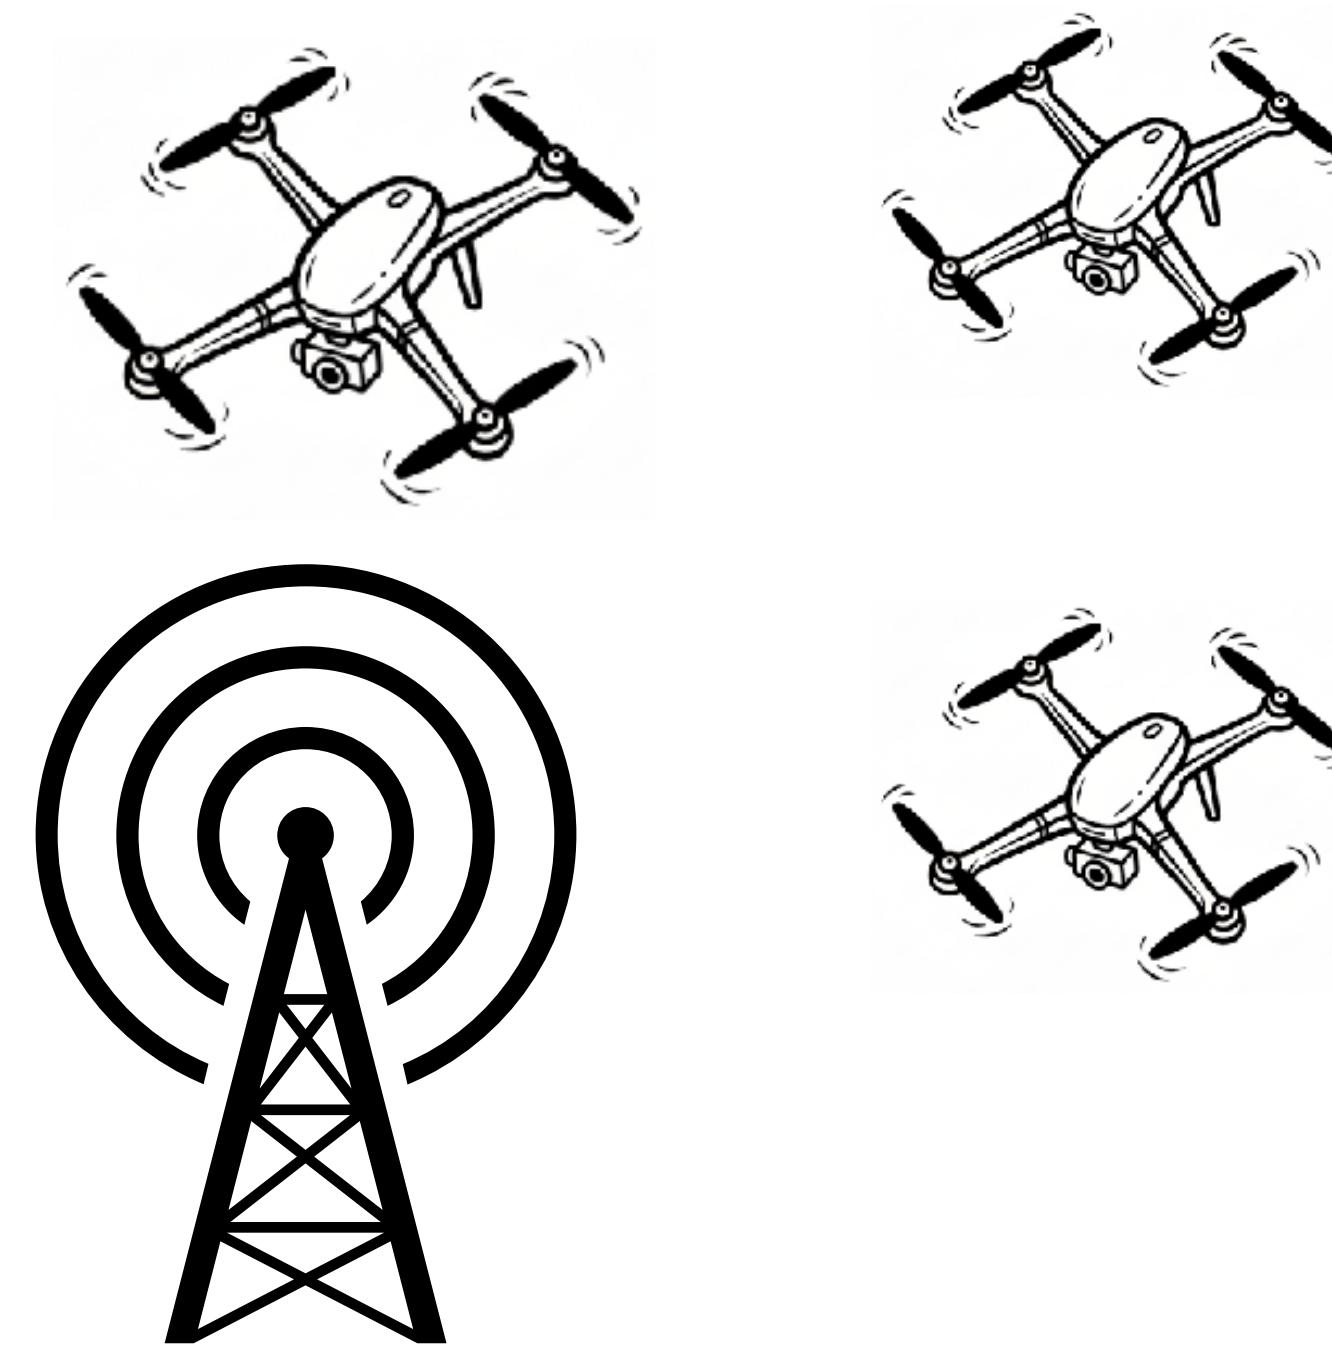


- The speed of light is too low: ~6in per clock cycle
  - $\Rightarrow$  Control and data flow is *not* instantaneous.
- Internet is a shared network
  - $\Rightarrow$  Performance is *not* predictable.
  - $\Rightarrow$  System is *not* isolated.

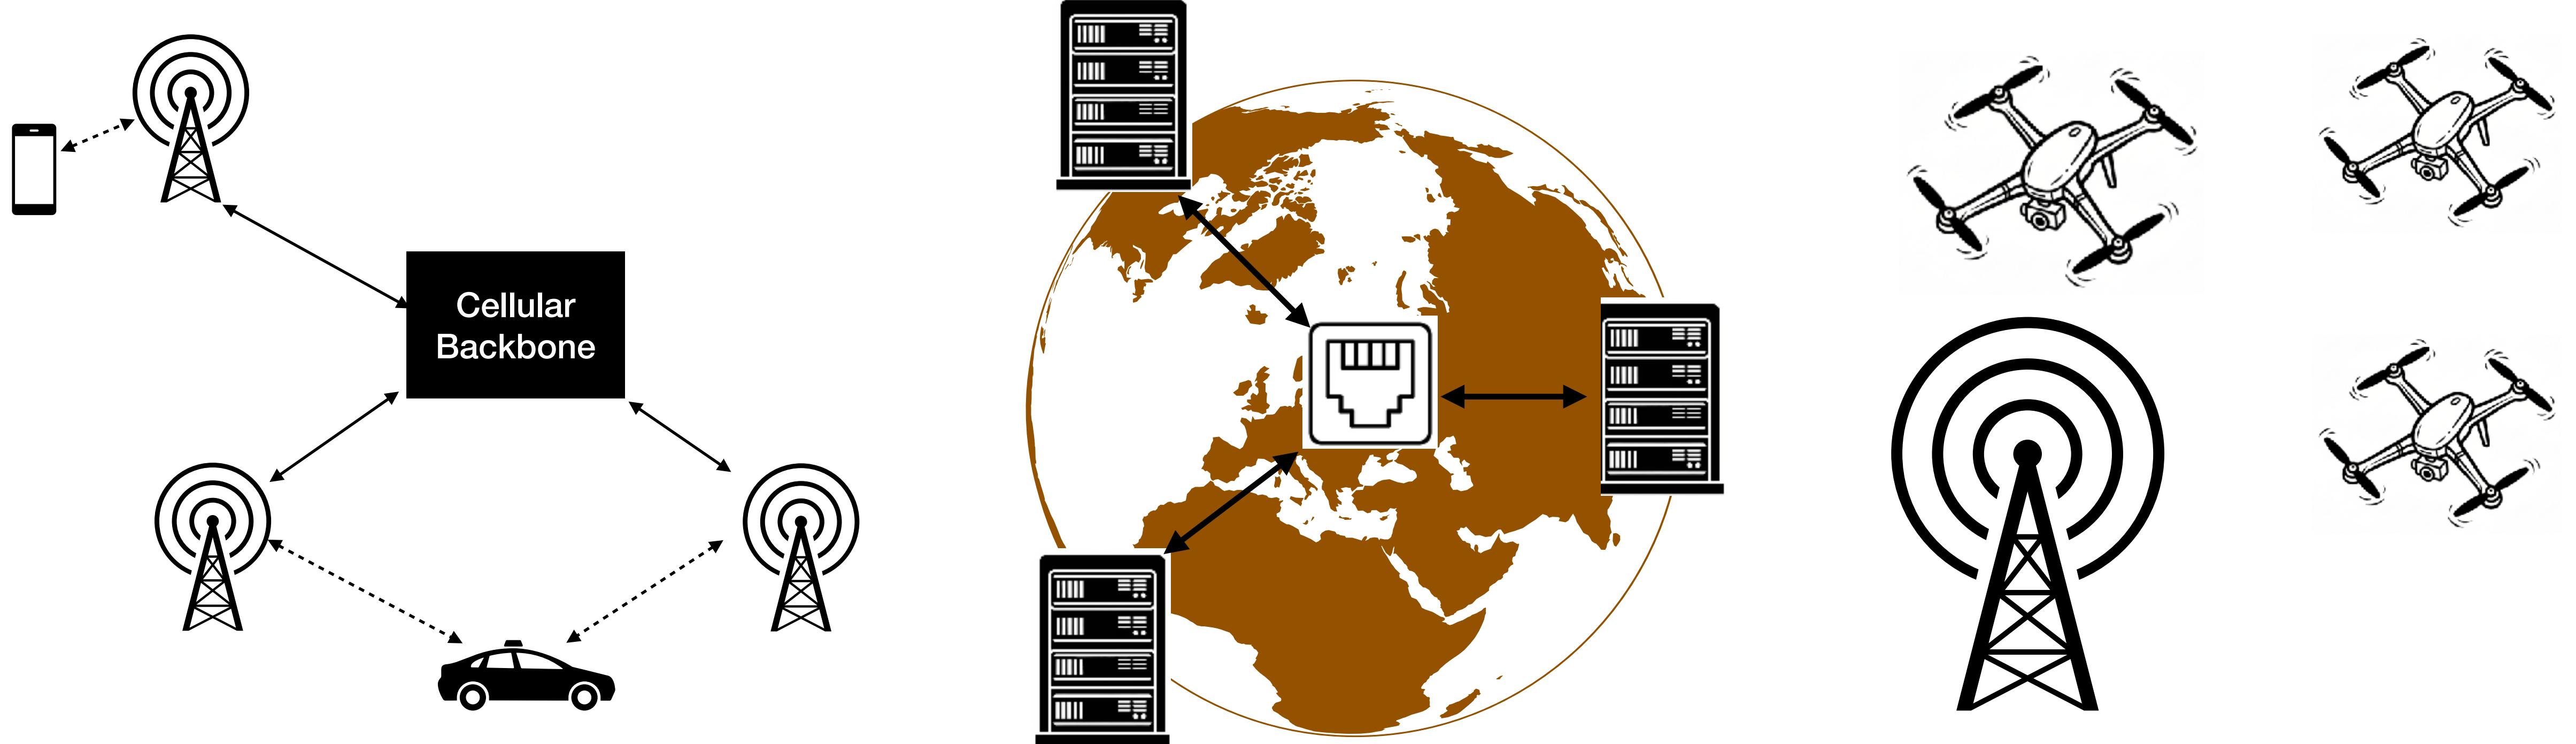
# Also Distributed Systems



5G Cellular Network



A fleet of unmanned aerial vehicles (UAVs)



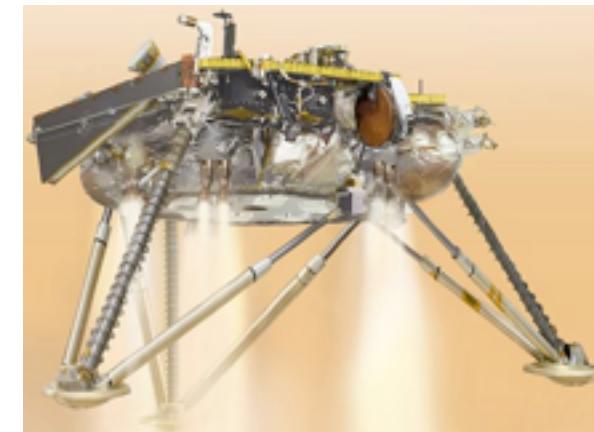
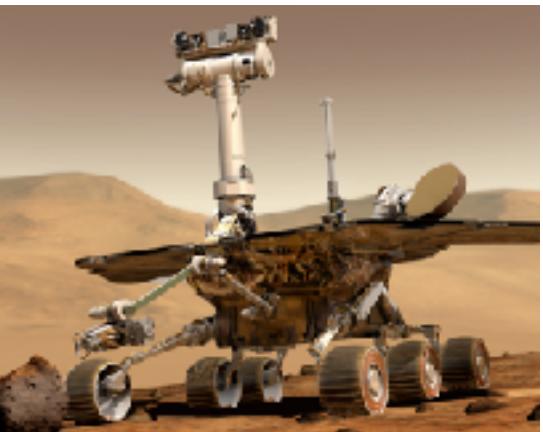
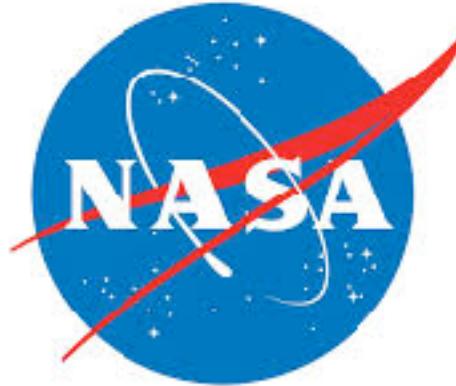
How do we program massive distributed systems to ensure fail-safe operation and resilience against adversarial attacks?

- A question motivating Programming Language and Formal Methods research.

# Formal Methods: Recent Success Stories



- “As AWS has grown, formal methods have become increasingly valuable” — Brooker and Desai, ACM Queue, 02/04/2025.
- “Project Everest produced provably correct and secure software deployed in a range of production systems” — Project Everest Collaborators, Microsoft.



- Kornfeld et al, Verification and Validation of the Mars Science Laboratory/Curiosity Rover Entry, Descent, and Landing System, JSR Vol. 51, 2014.



- “.. DARPA successfully applied the HACMS tools to other military systems, including satellites and self-driving convoy trucks. Each successful project has validated the advantages of applying a formal methods approach to software security” — DARPA HACMS Project Report.
- “DARPA touts ‘formal methods’ for nipping cyber disasters in the bud” — Defense News, 02/21/2025.



Ethereum



Tezos



SOLANA

- “Formal verification is one of the recommended techniques for improving smart contract security” — Ethereum Official Documentation.

# Formal Methods: Impediments to Widespread Adoption



## Poor scalability

- Formal verification is mostly confined to abstract models.
- Verifying real distributed systems requires enormous computational resources and great manual effort.



## Insufficient generality

- Formal methods research is almost entirely focused on safety properties.
- Emerging internet-native distributed systems are most concerned about privacy and security.

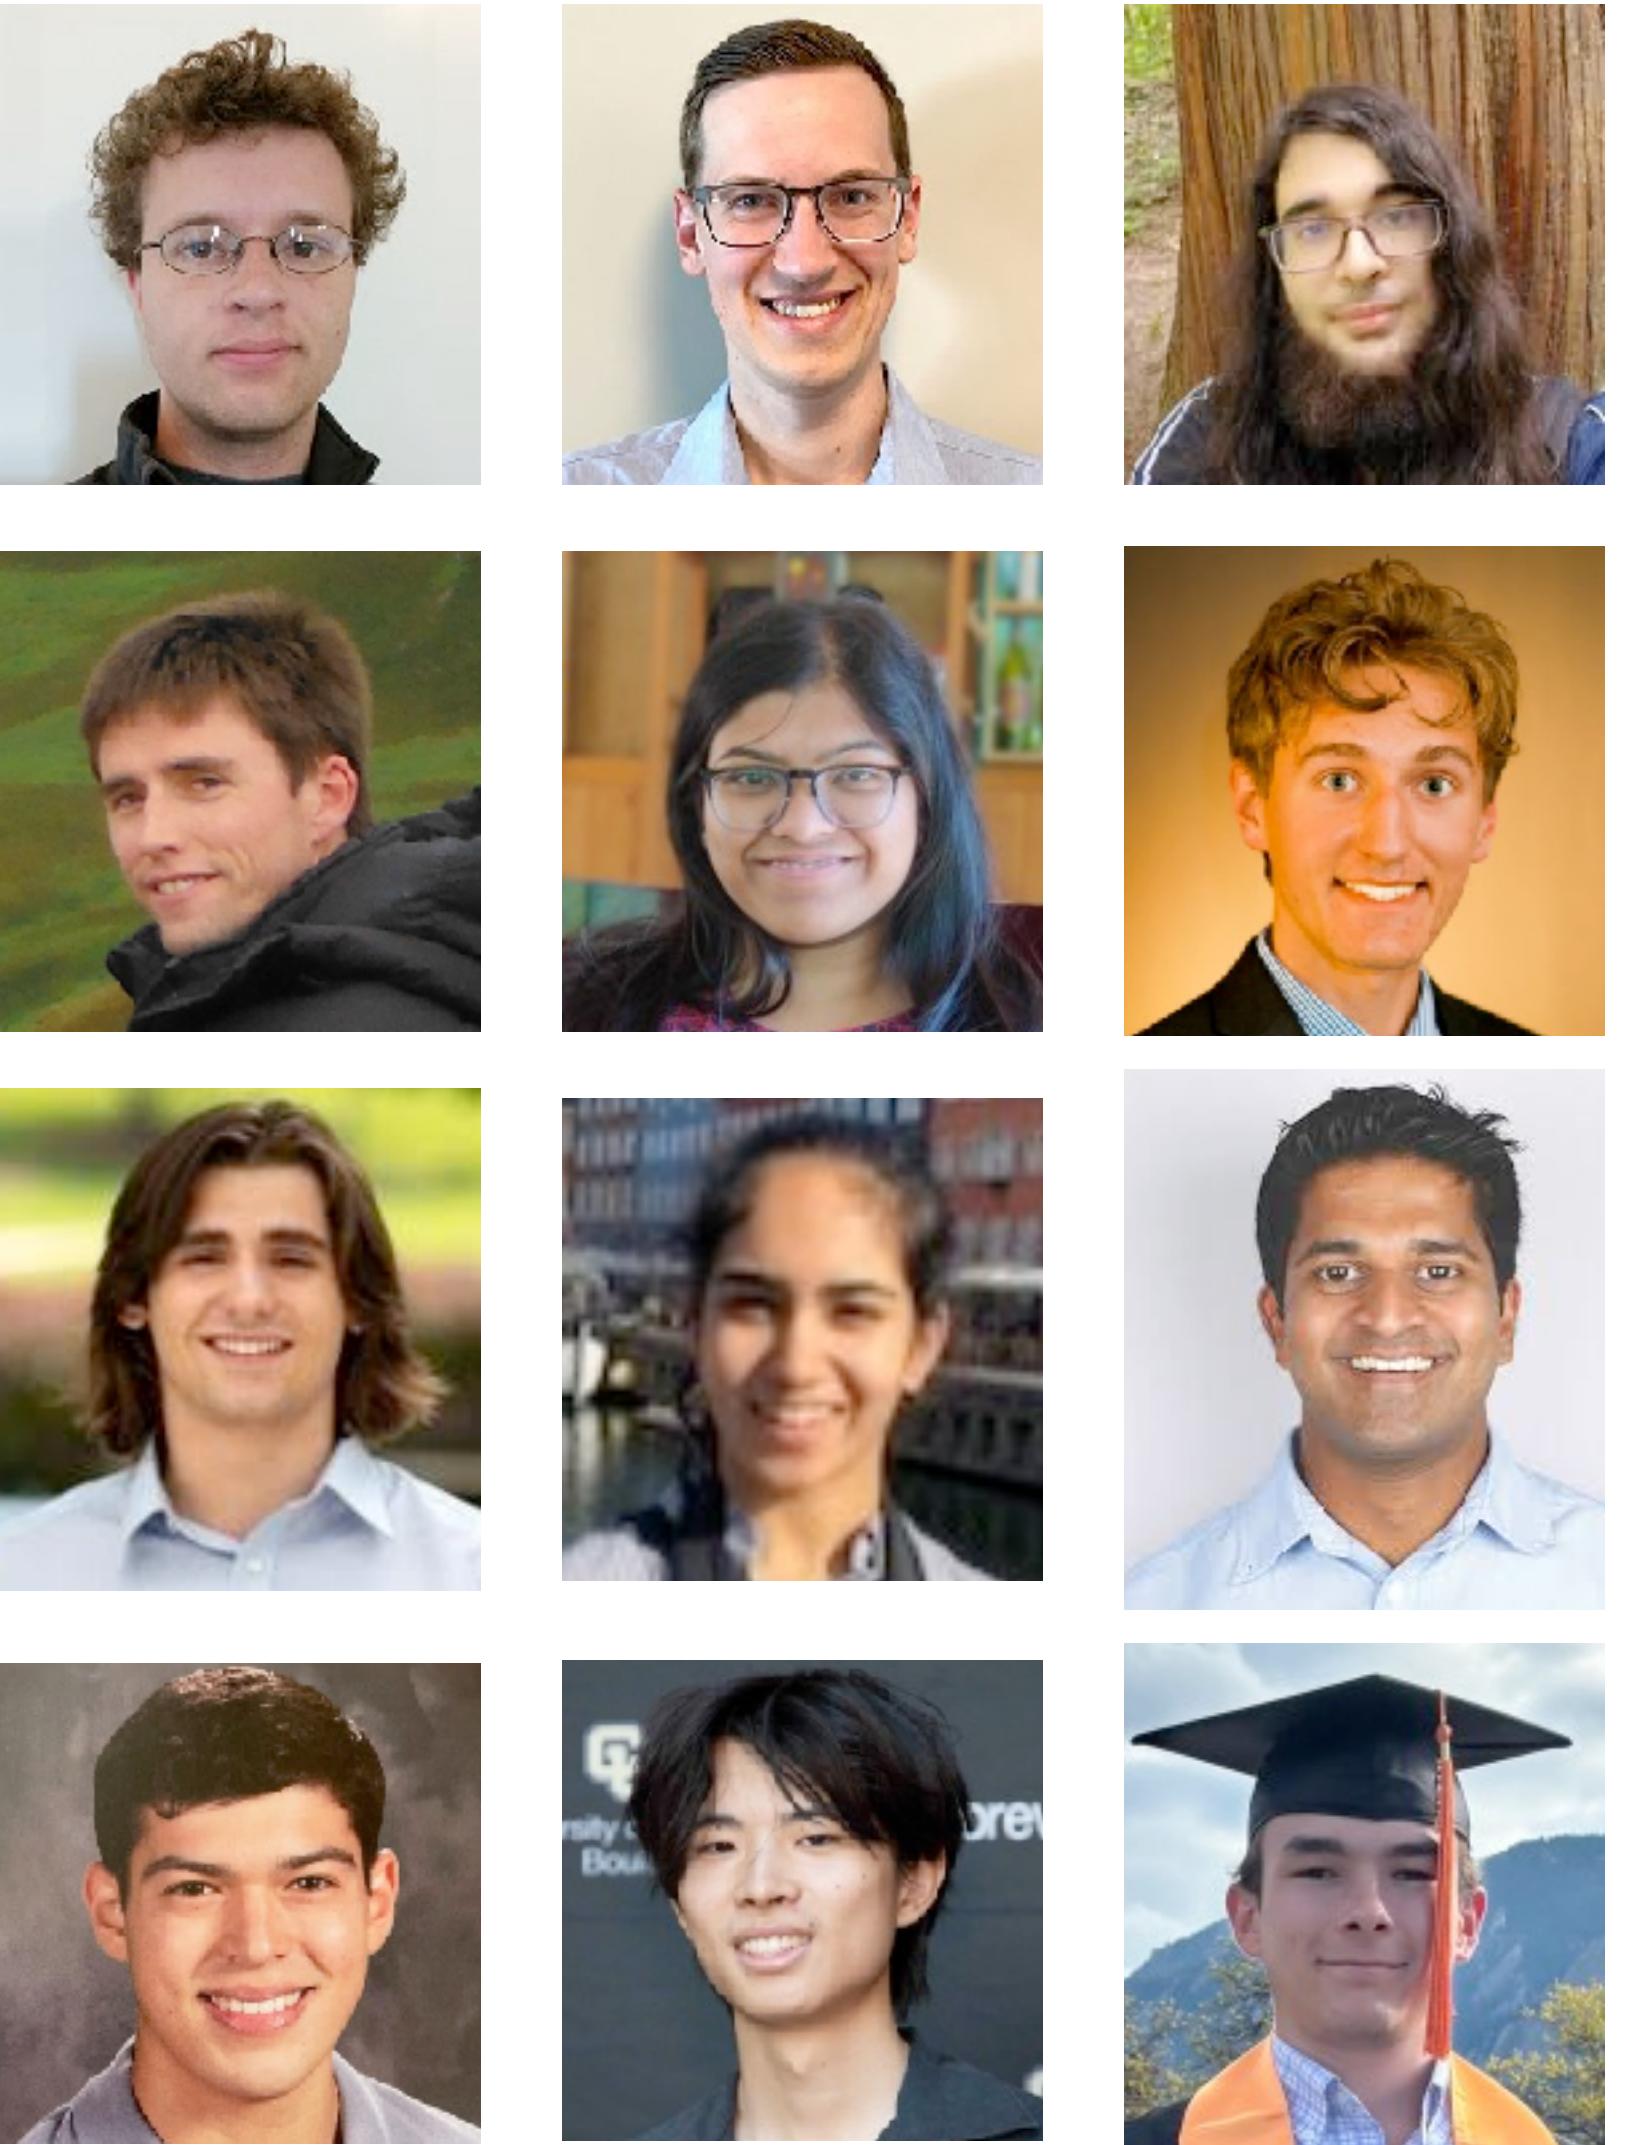
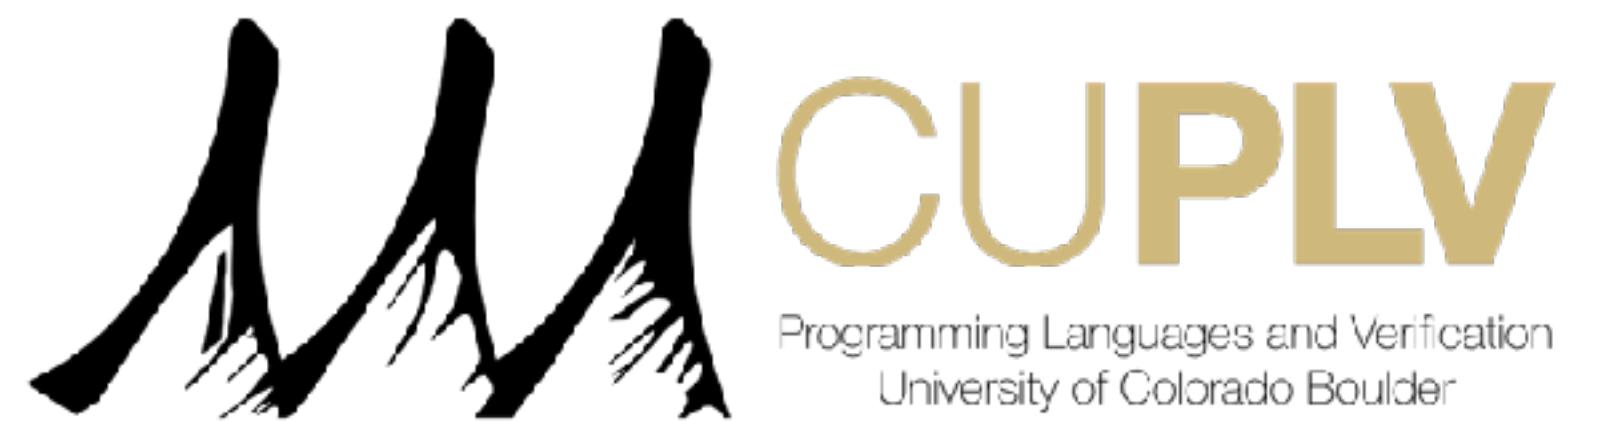


## Steep Learning Curve

- Writing specifications, invariants, and proofs requires training in formal logic.
- Formal methods need to work with what's available: unit tests, assertions, and natural language documentation.

# My Research

Programming abstractions and automated reasoning techniques to bring affordable formal verification to practical distributed systems.



# This Talk



Poor scalability

- Formal methods research is based on a single formalism.
- **1. Improve Scalability:  
from *models* to *implementations***
  - requires enormous computational resources and great manual effort.



Insufficient generality

- Formal methods research is based on a single formalism.

**2. Improve Generality:  
from *safety* to *security & privacy***

distributed systems are most concerned about privacy and security.



Steep Learning Curve

- Writing specifications,

**3. Improve Usability:  
from *experts* to *programmers***

with what's available: unit tests, assertions, and natural language documentation.

# 1. Improving Scalability: From *models* to *implementations*

## Contributors



Nicholas Lewchenko



Prasanth Prahladan



Evan Chang

## Tools



<https://github.com/cuplv/Quark>

**Super-V**

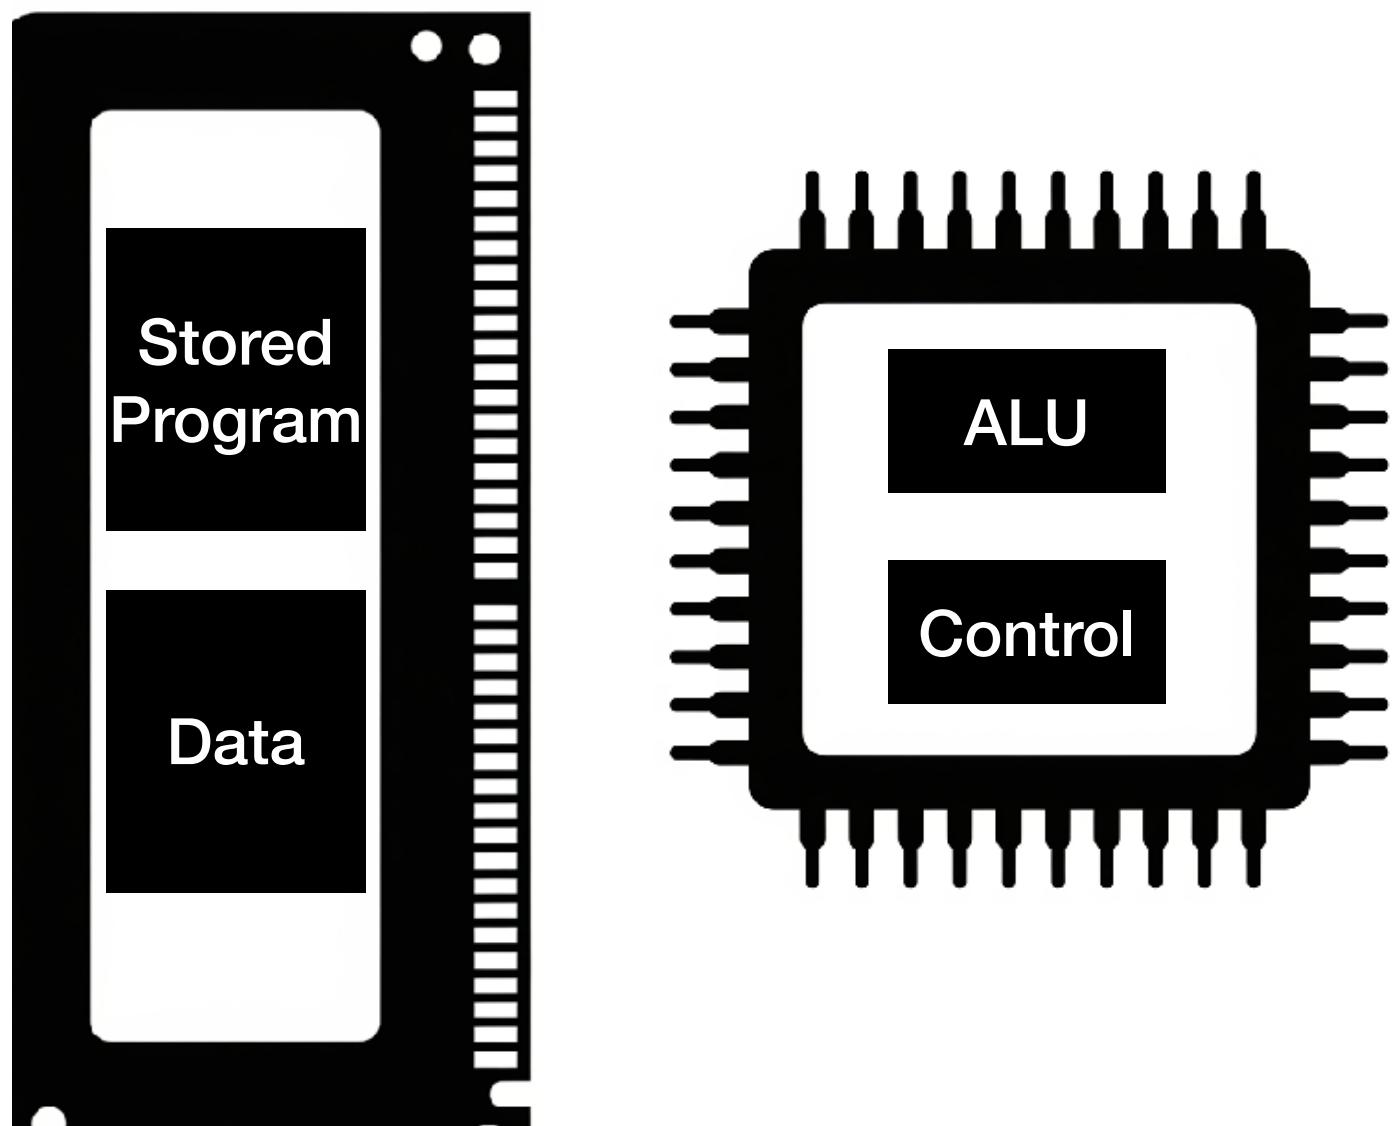
<https://github.com/cuplv/super-v>

## Publications

- Nicholas Lewchenko, Gowtham Kaki, and Bor-Yuh Evan Chang. **Bolt-On Strong Consistency: Specification, Implementation, and Verification**. Proceedings of the ACM on Programming Languages (PACMPL), issue **OOPSLA1 2025**, Article No.: 137, Pages 1604 - 1631.
- Gowtham Kaki, Prasanth Prahladan, and Nicholas Lewchenko. **RunTime-Assisted Convergence in Replicated Data Types**. Proceedings of 43rd ACM SIGPLAN Programming Language Design and Implementation (PLDI 2022), Pages 364 - 378.
- Nicholas Lewchenko and Gowtham Kaki. **Distributed Consensus Algorithms as Replicated State Applications**. ACM SIGOPS Principles and Practice of Consistency for Distributed Data (PaPoC 2023).
- Gowtham Kaki, Prasanth Prahladan, and Nicholas Lewchenko. **Bolt-On Convergence in Mergeable Replicated Data Types**. ACM SIGOPS Principles and Practice of Consistency for Distributed Data (PaPoC 2022).

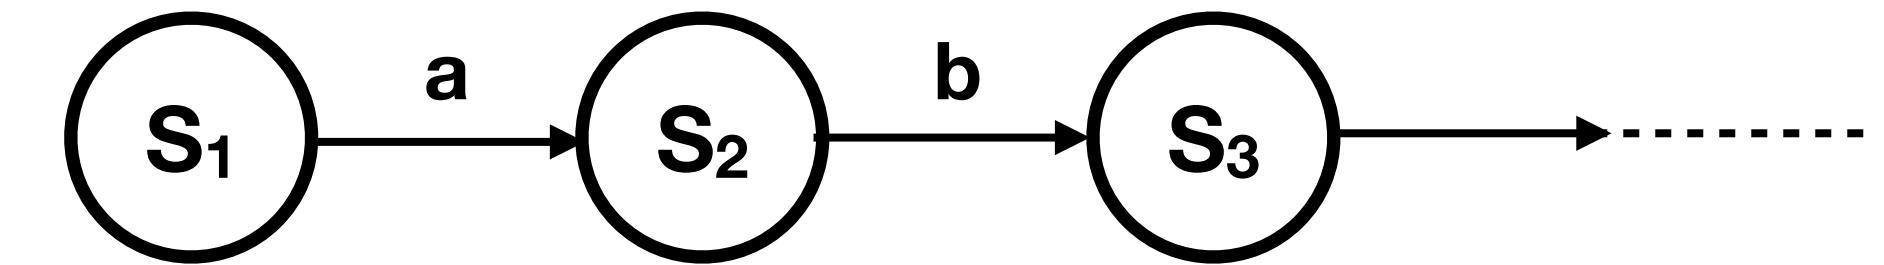
# State Machine Abstraction

Micro-architectural abstraction

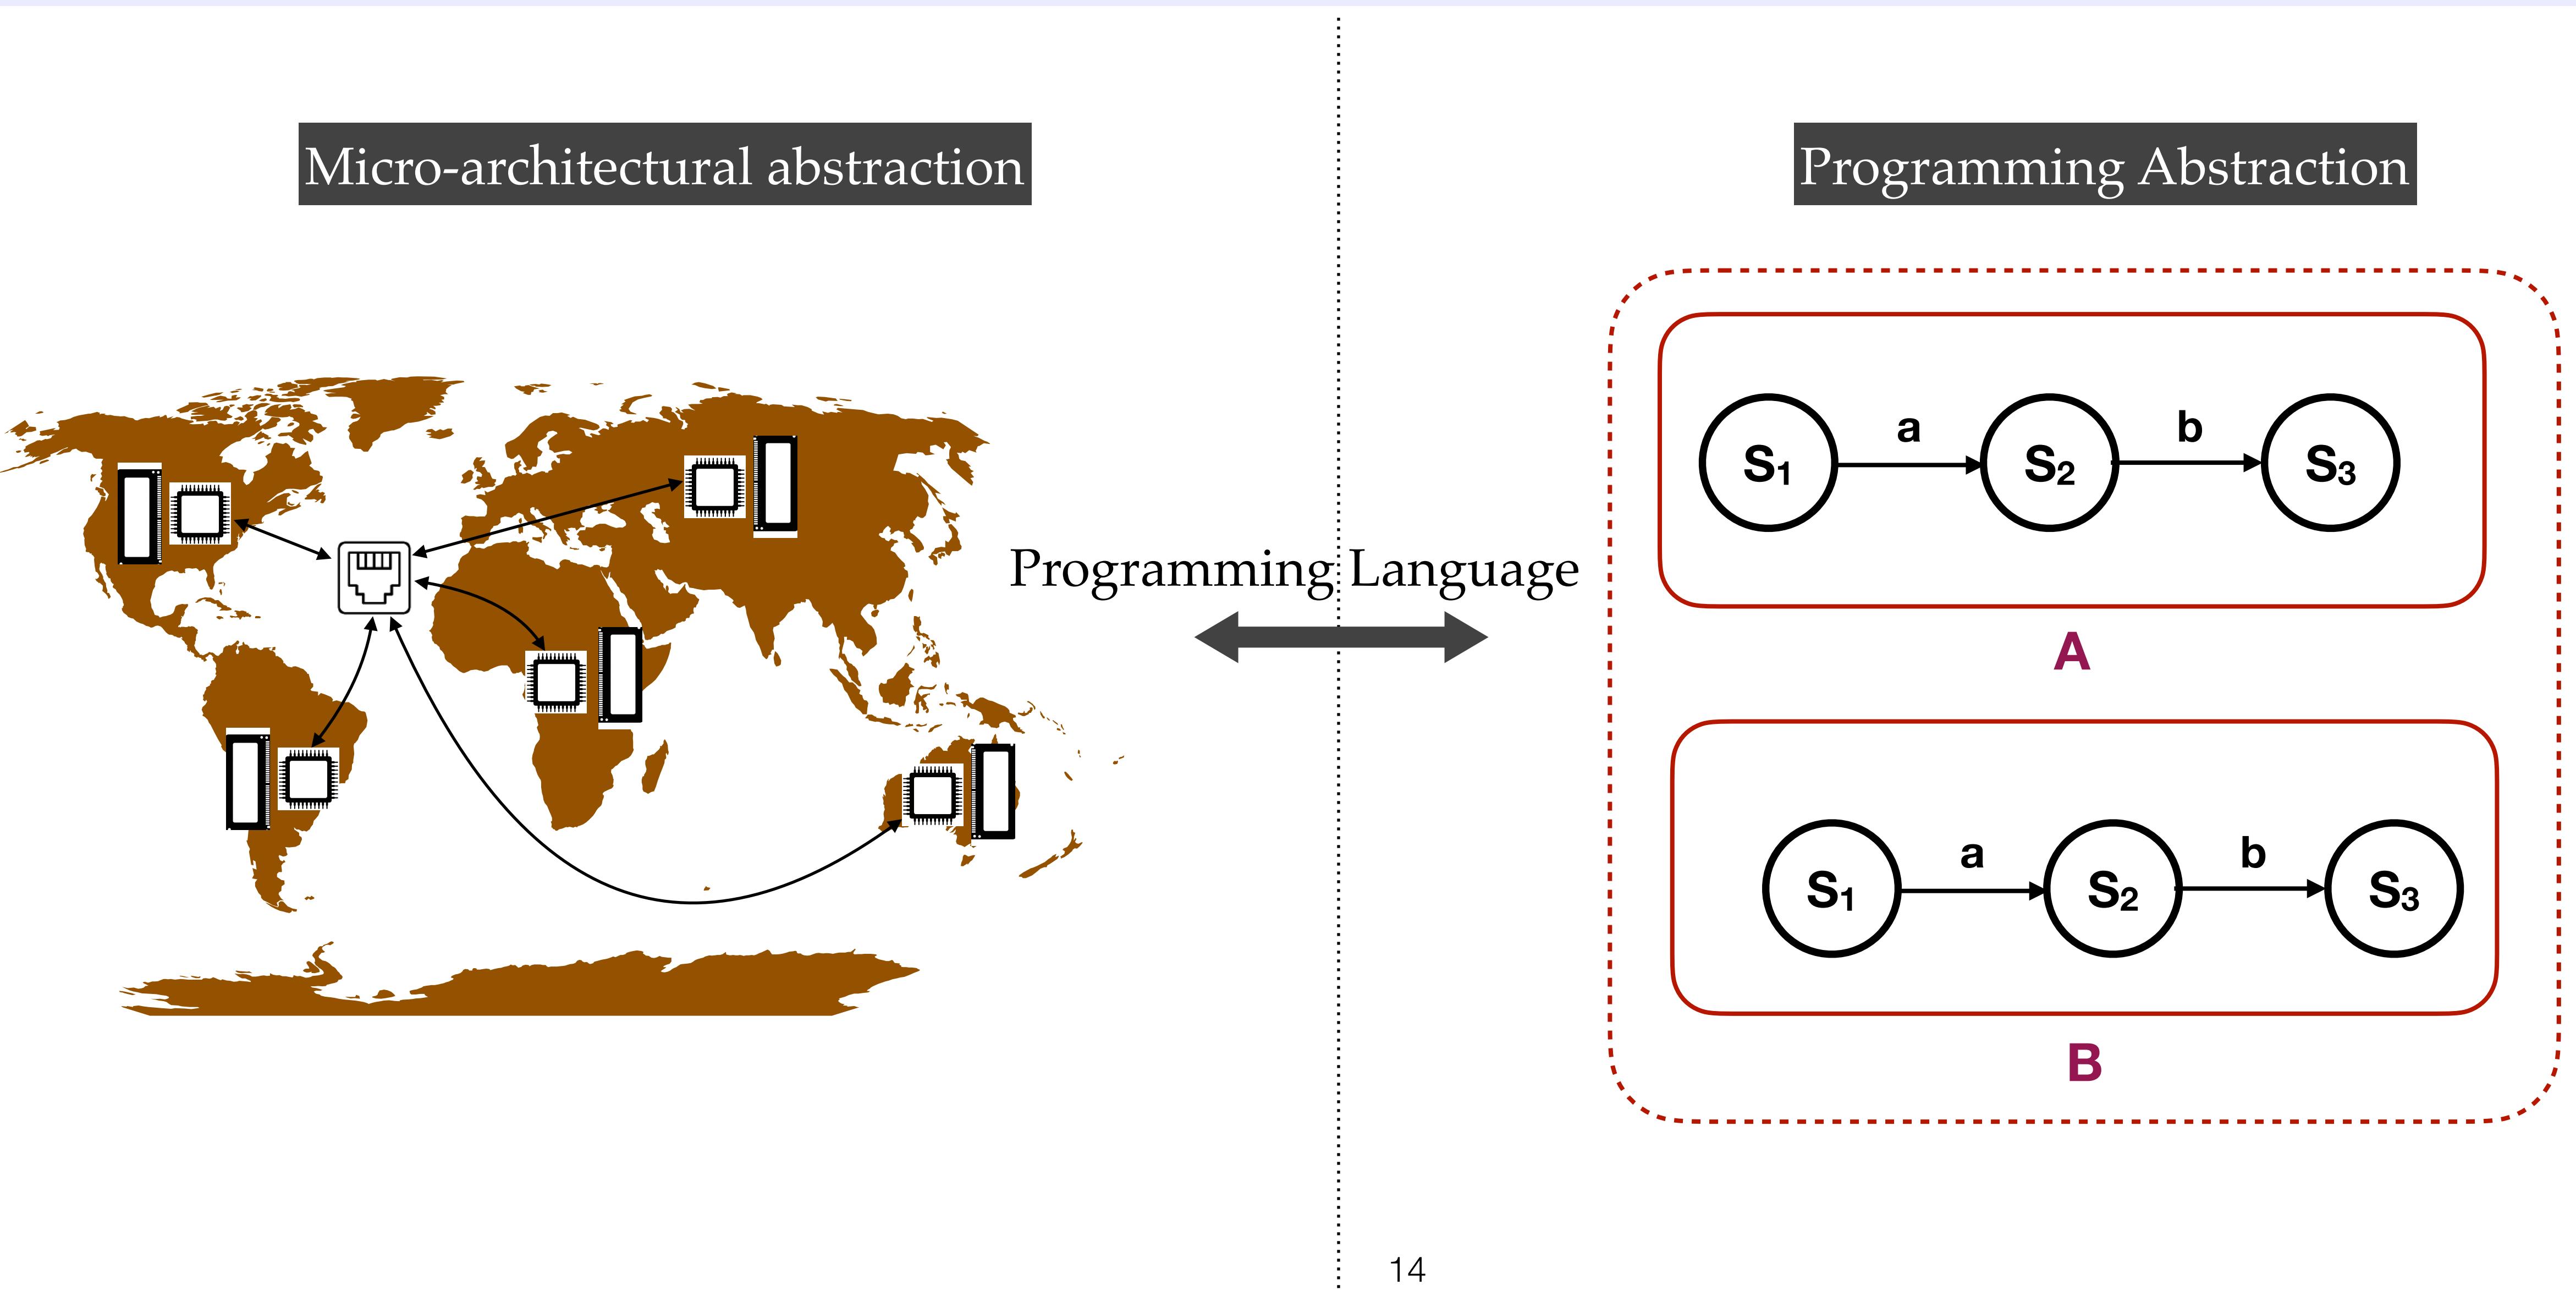


Programming Abstraction

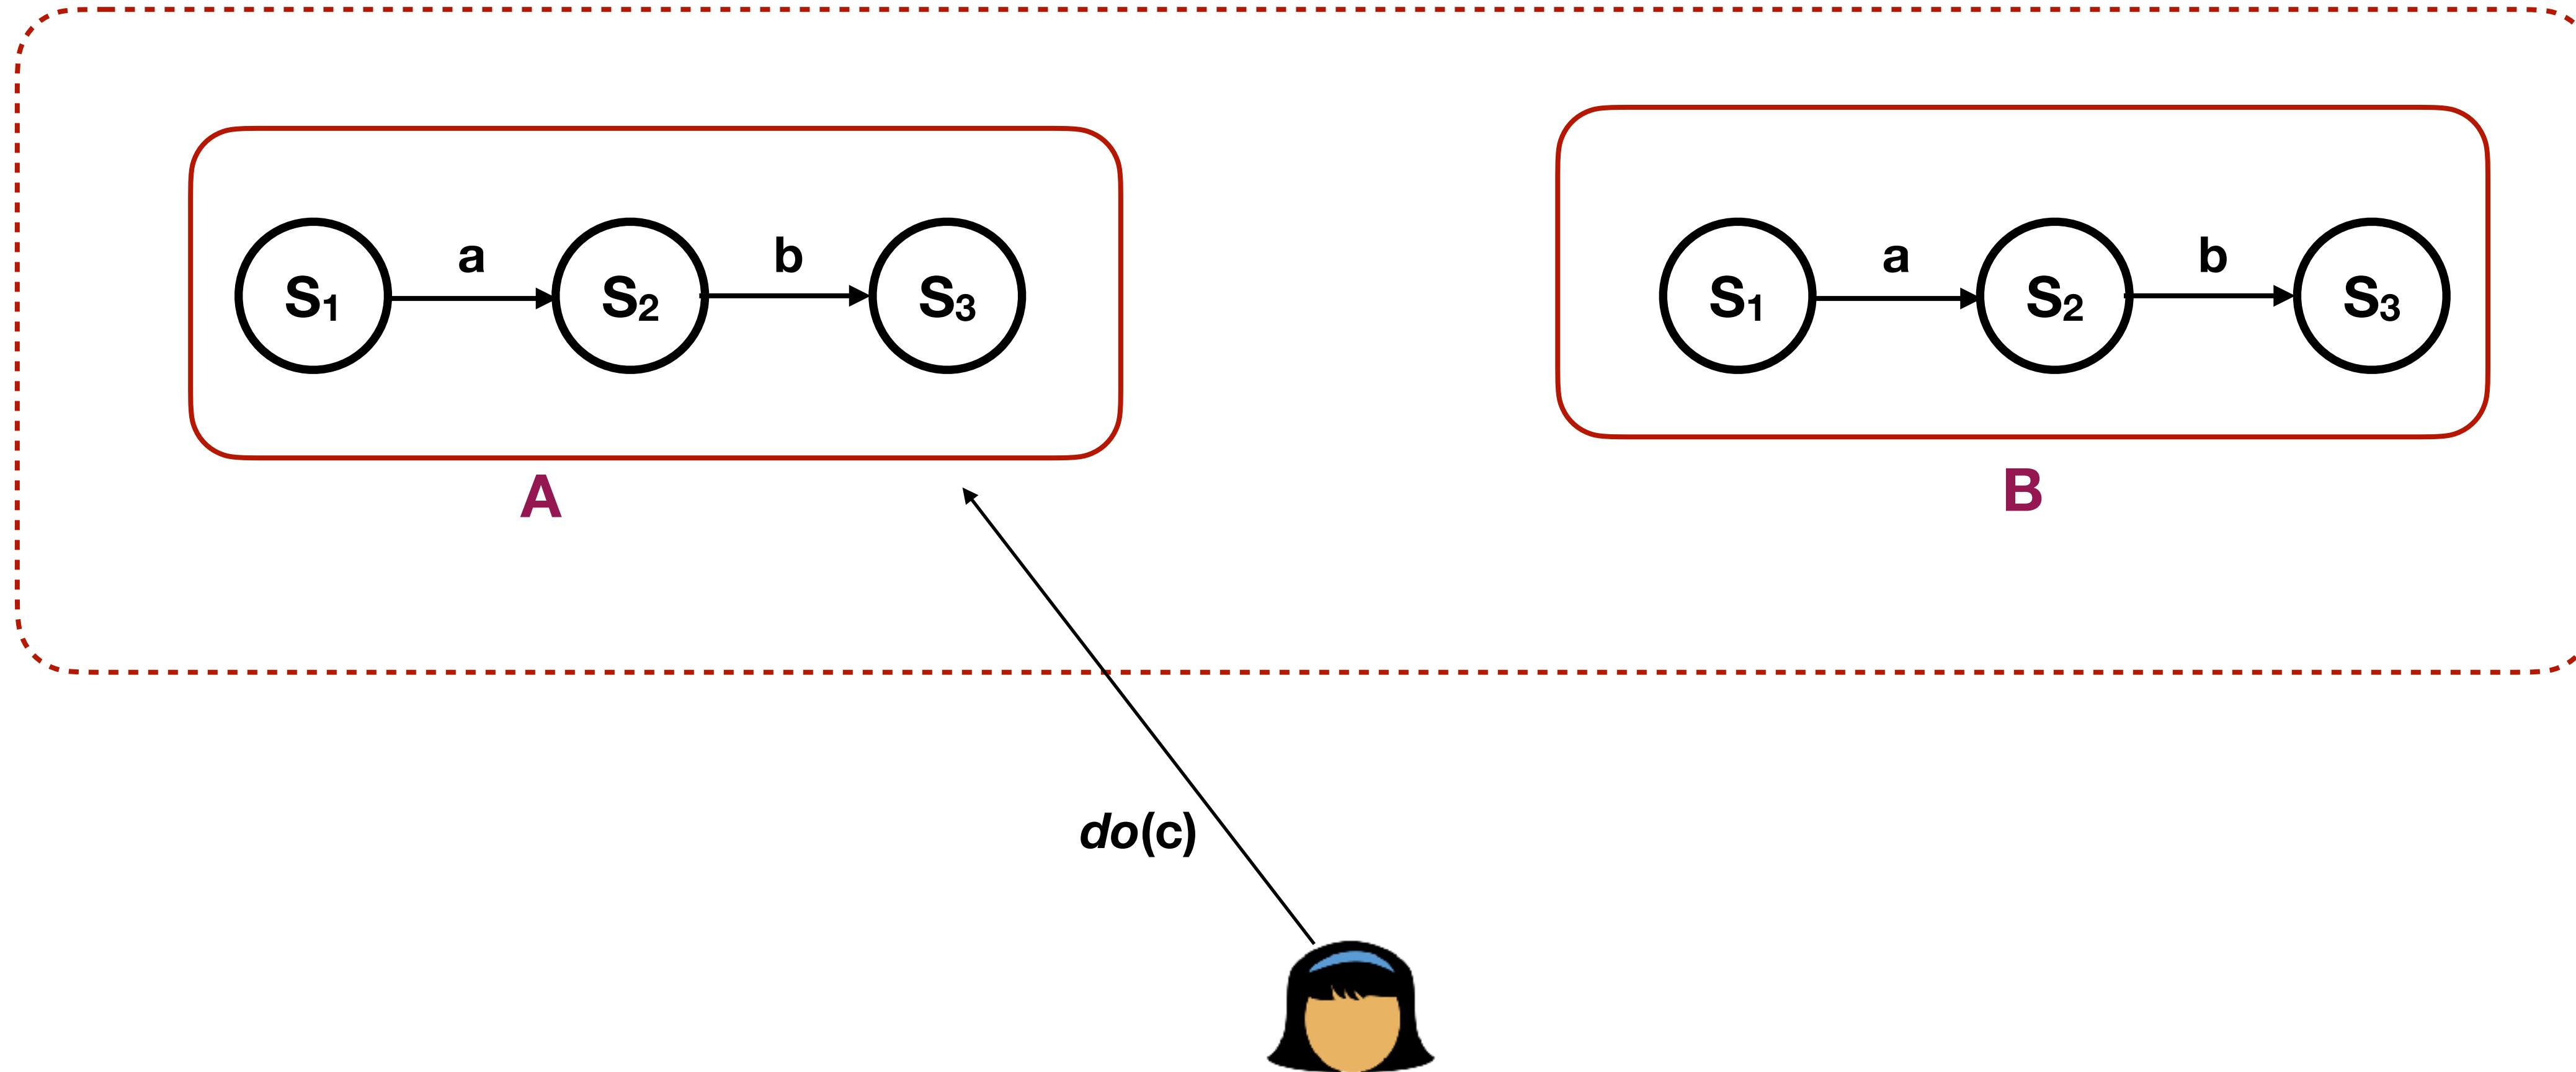
Programming Language



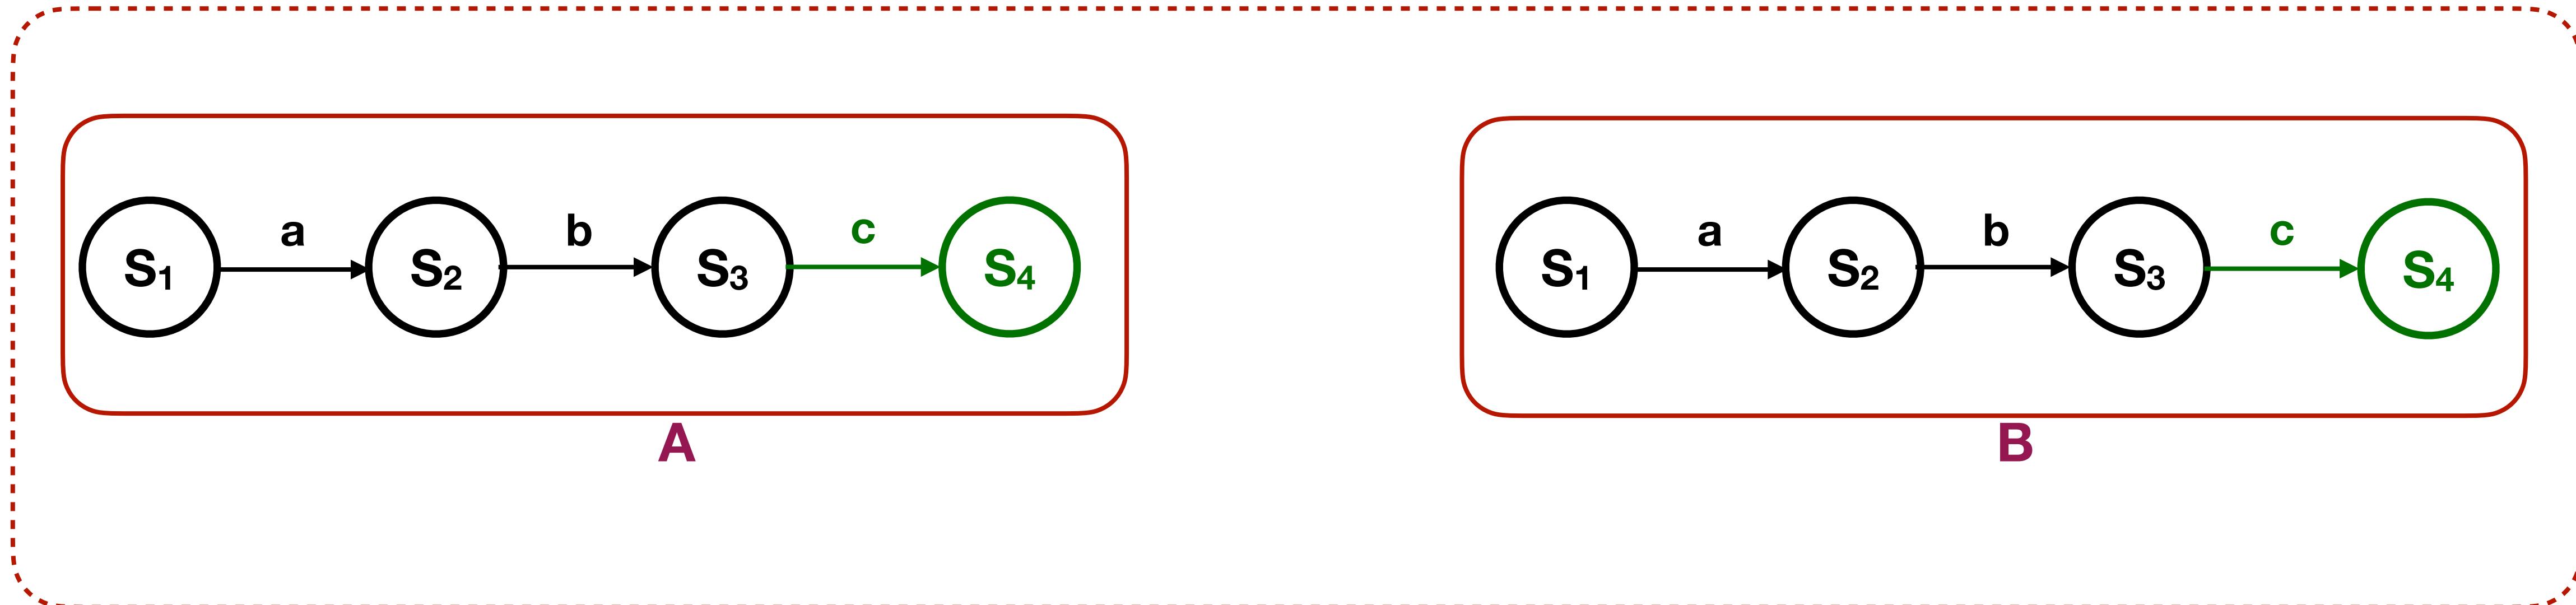
# Replicated State Machine (RSM) Abstraction



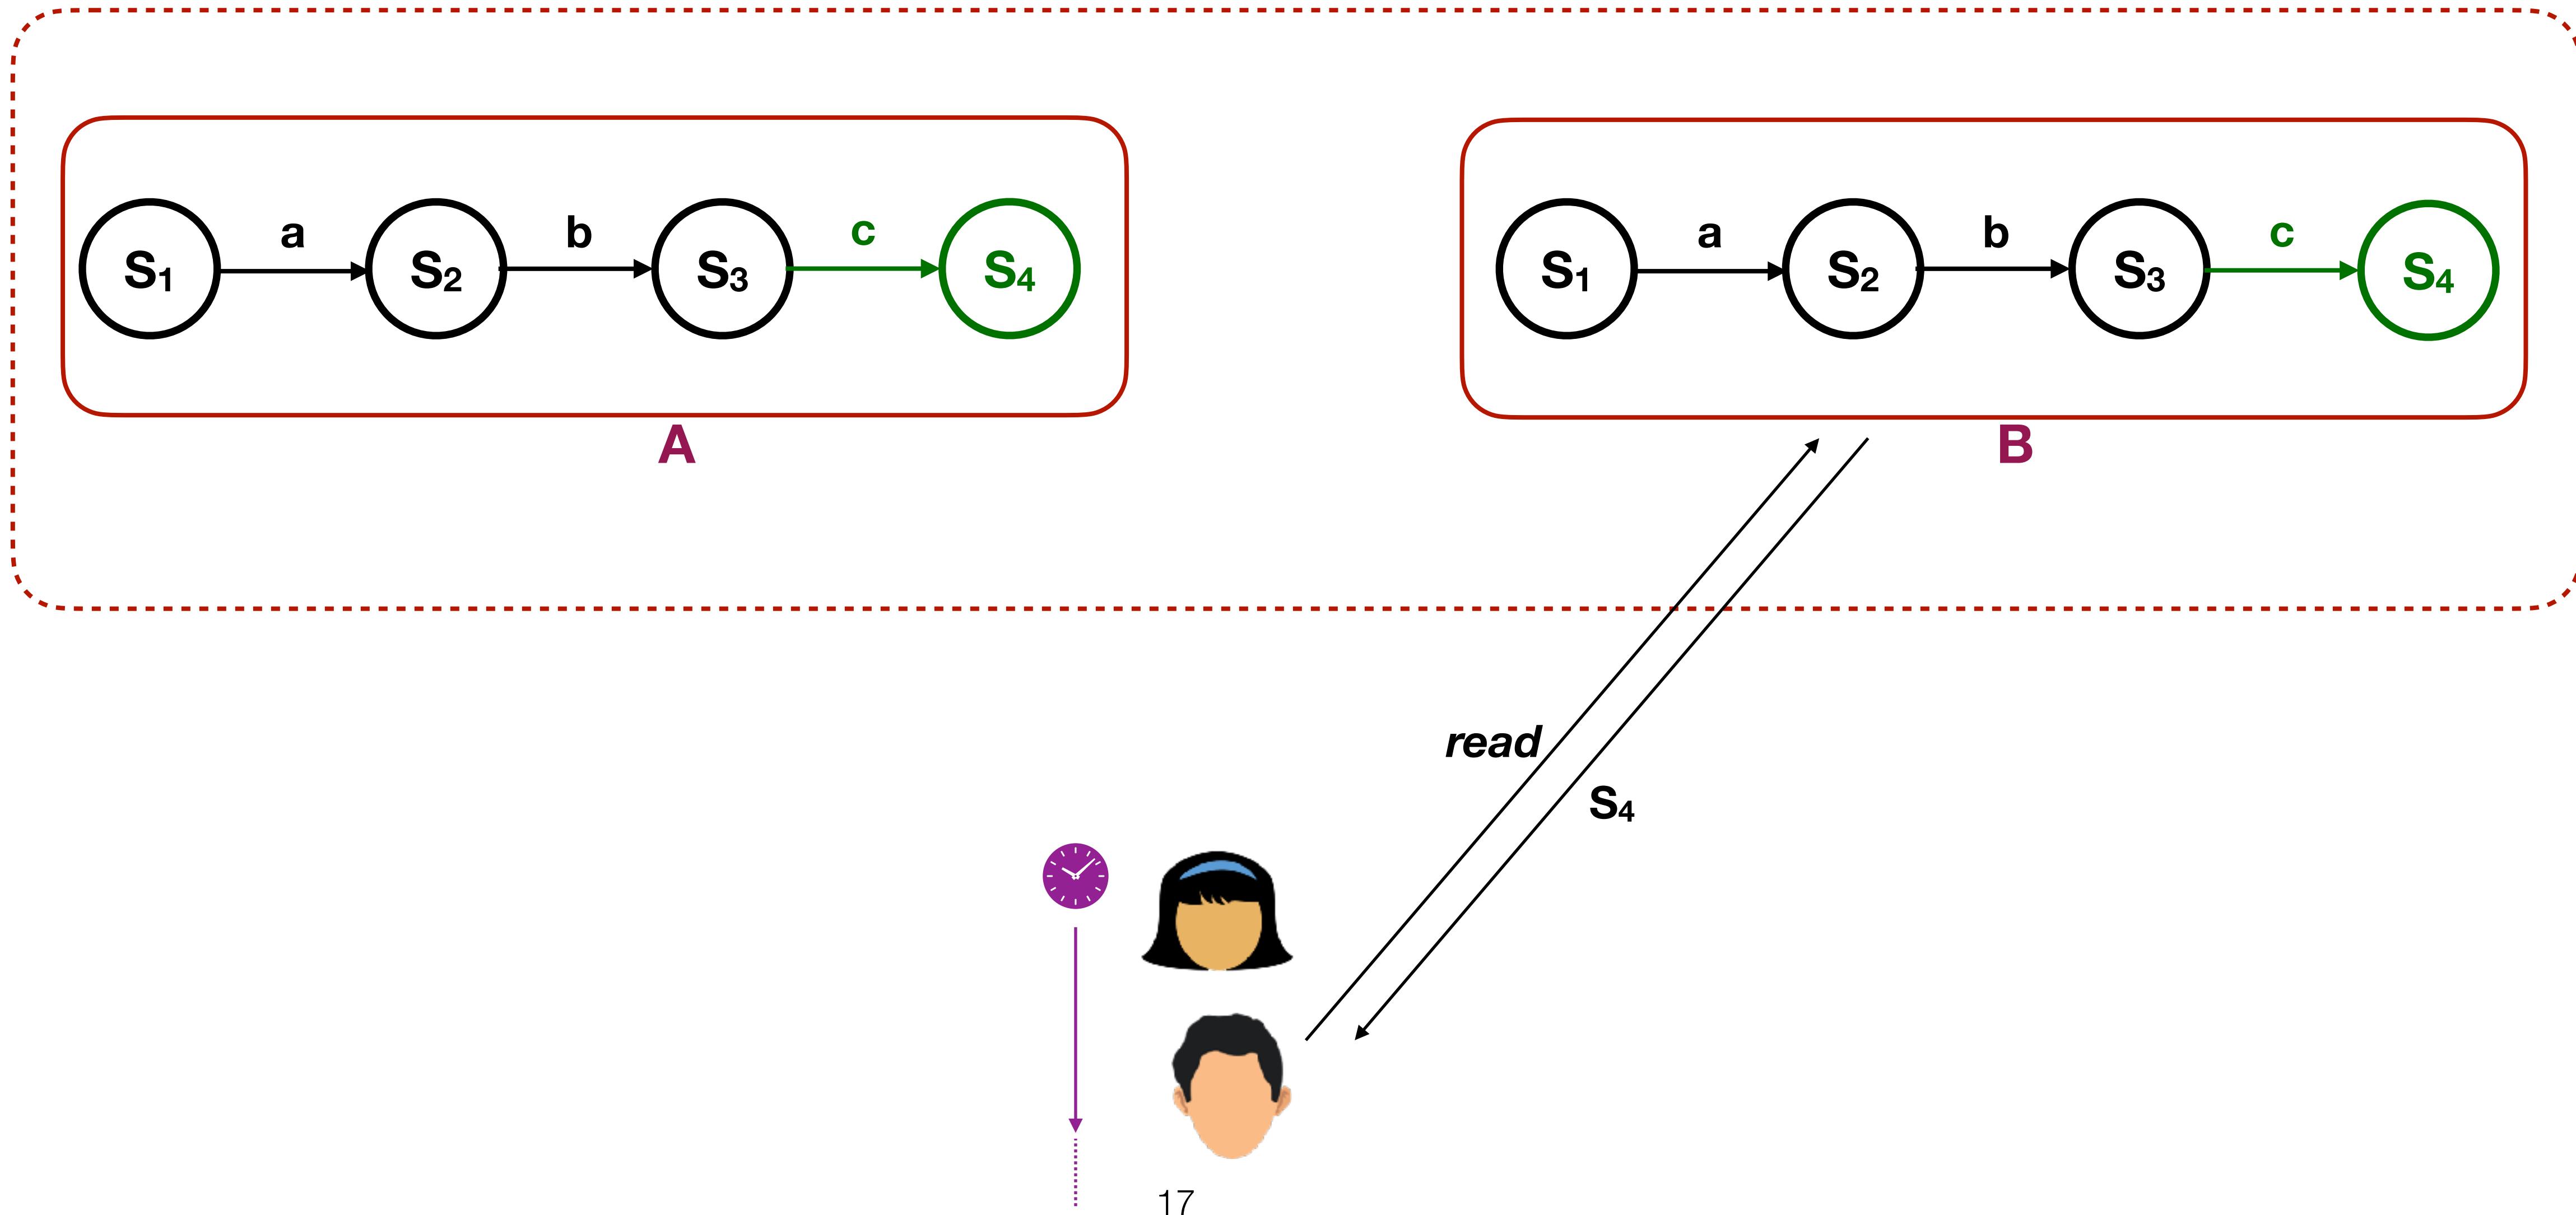
# Replicated State Machines



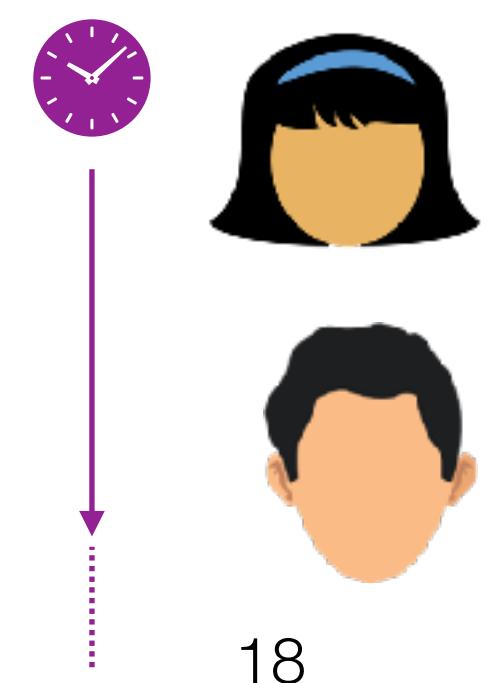
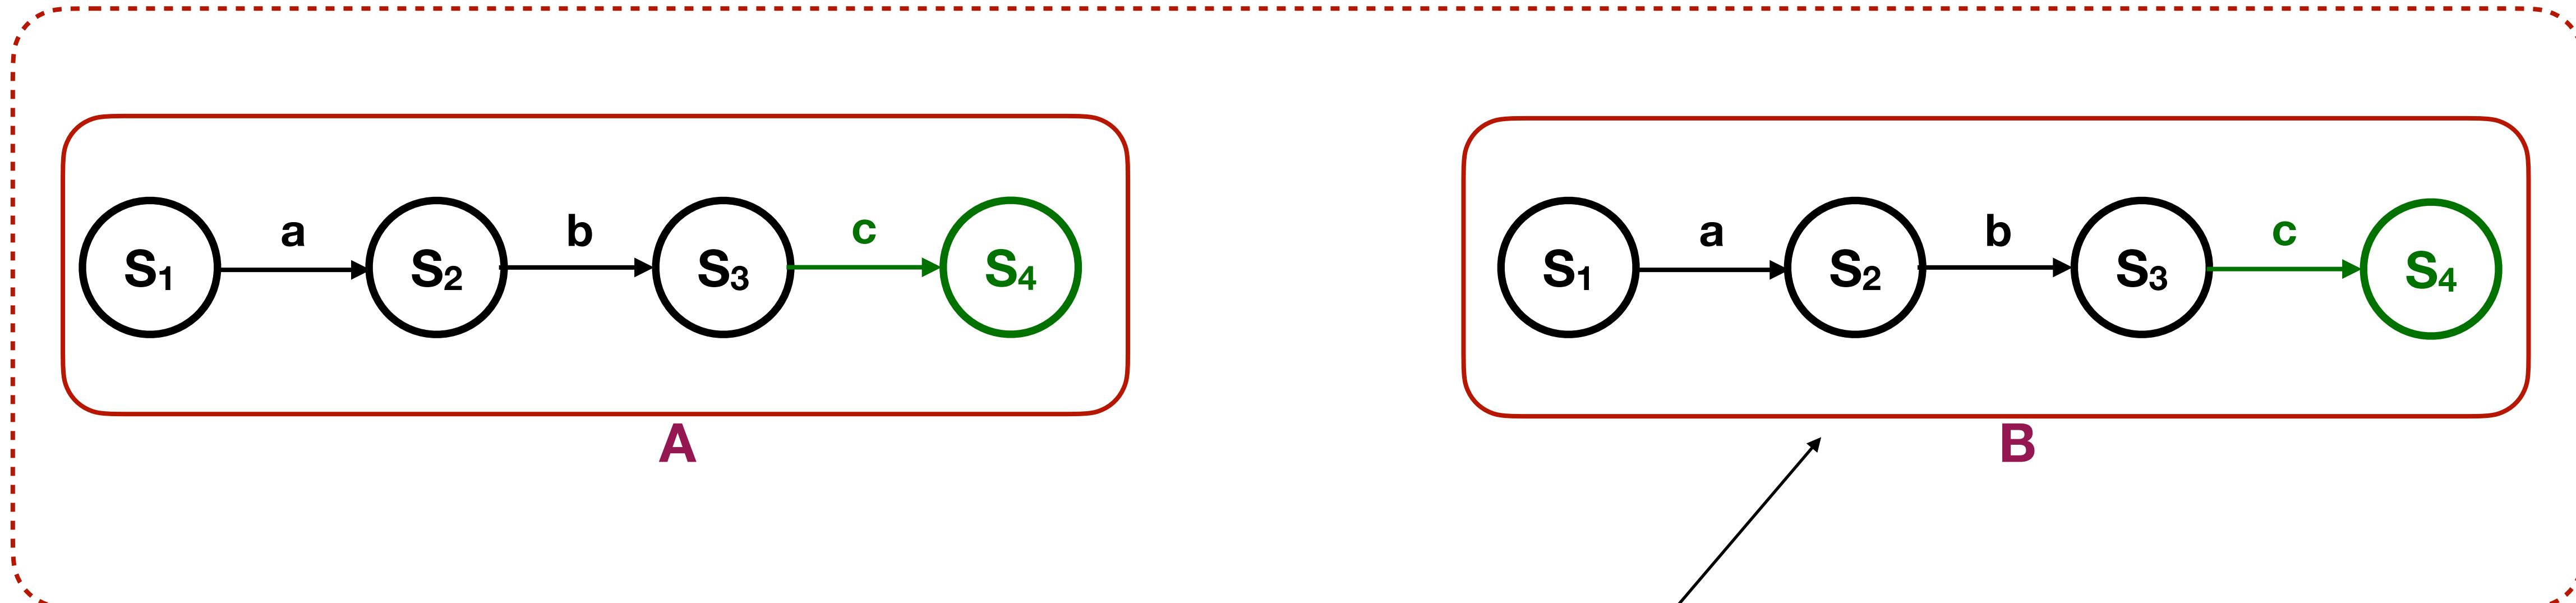
# Replicated State Machines



# Replicated State Machines

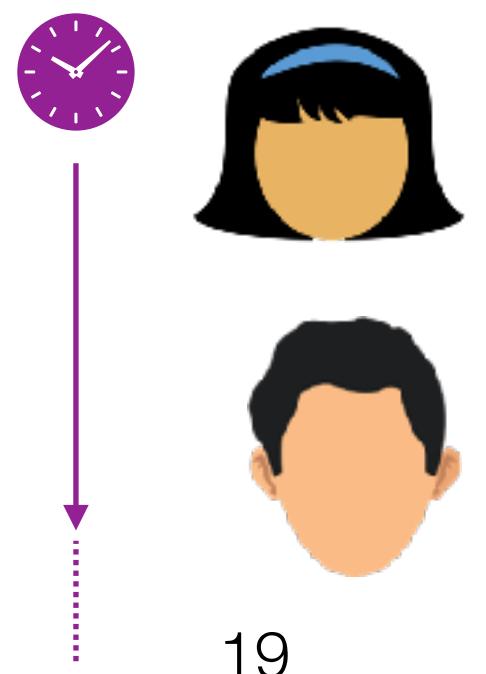
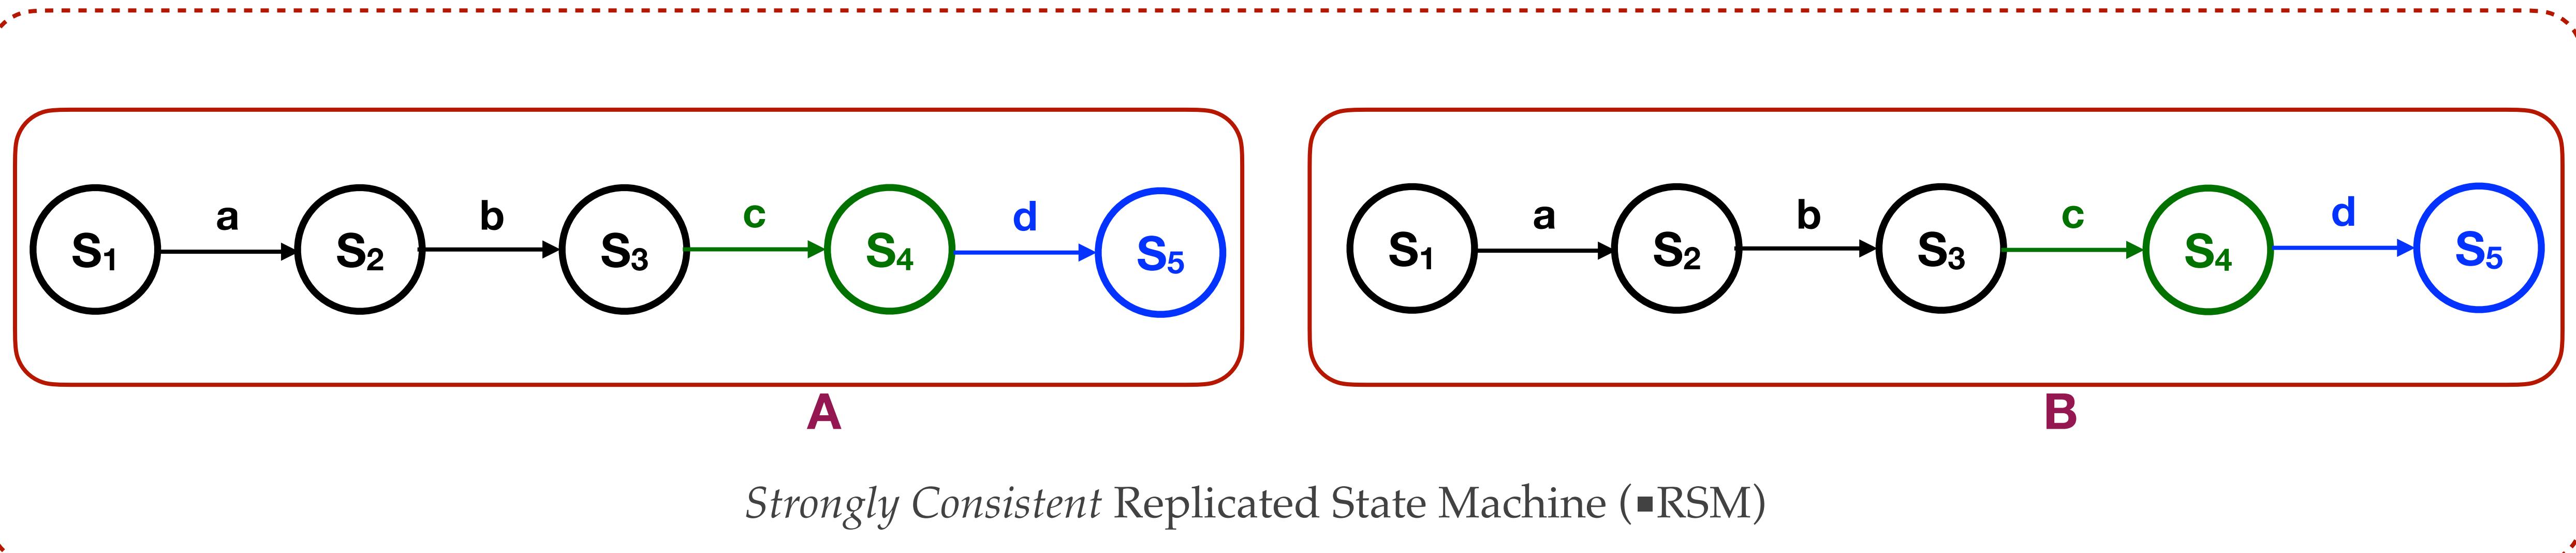


# Replicated State Machines

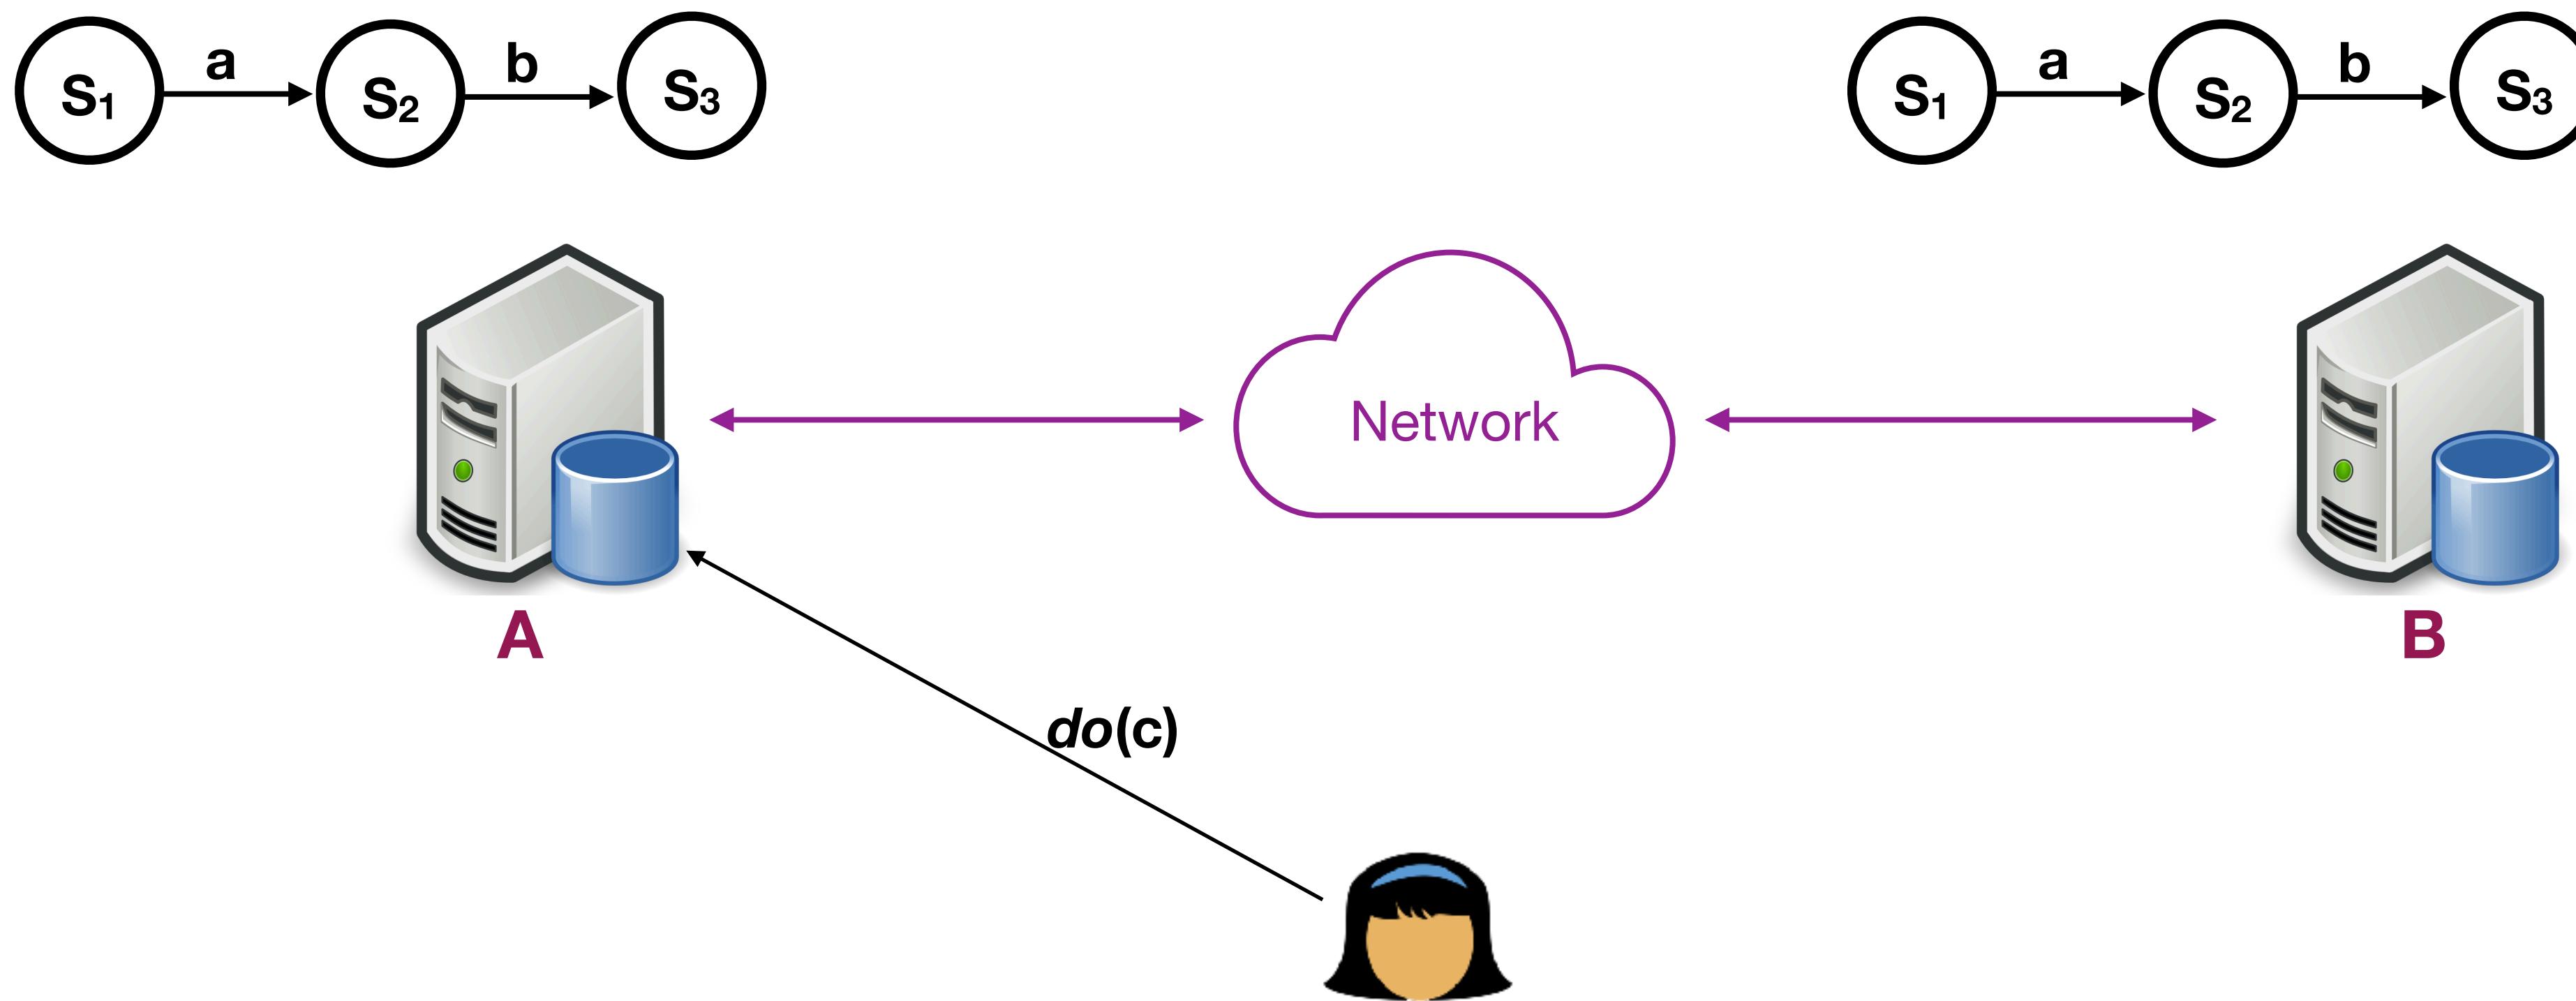


*do(d)*

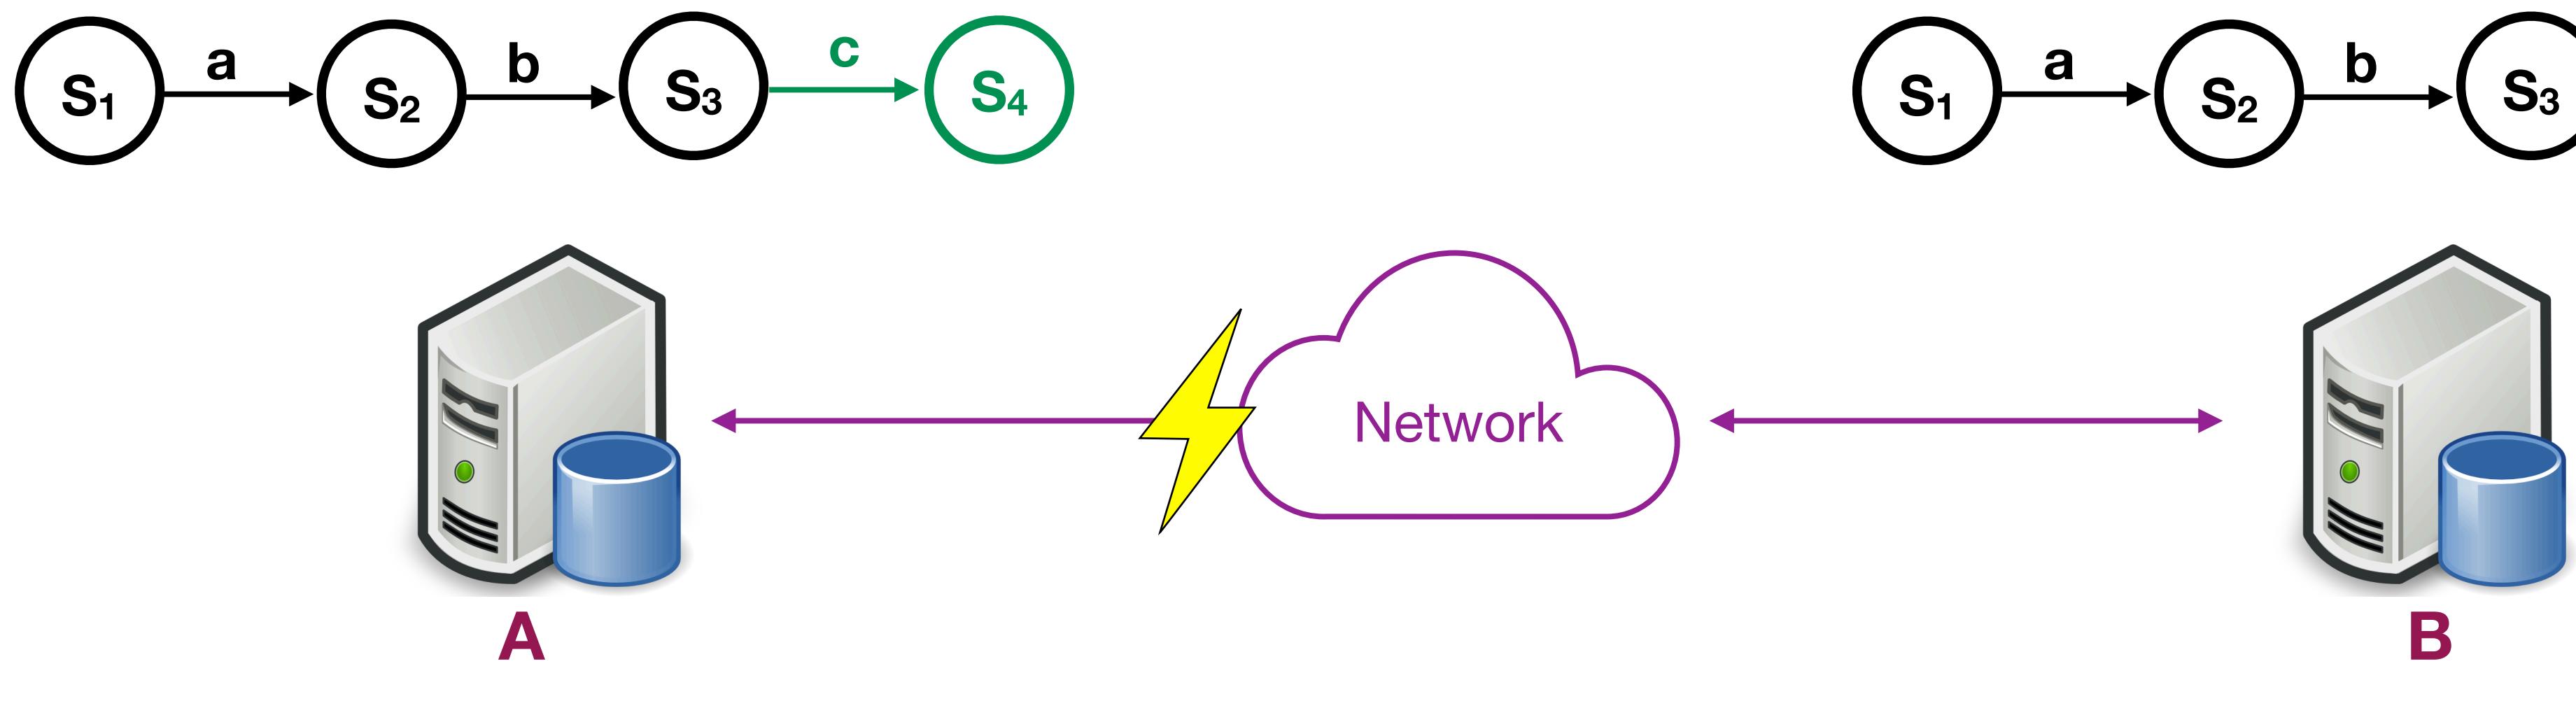
# Replicated State Machines



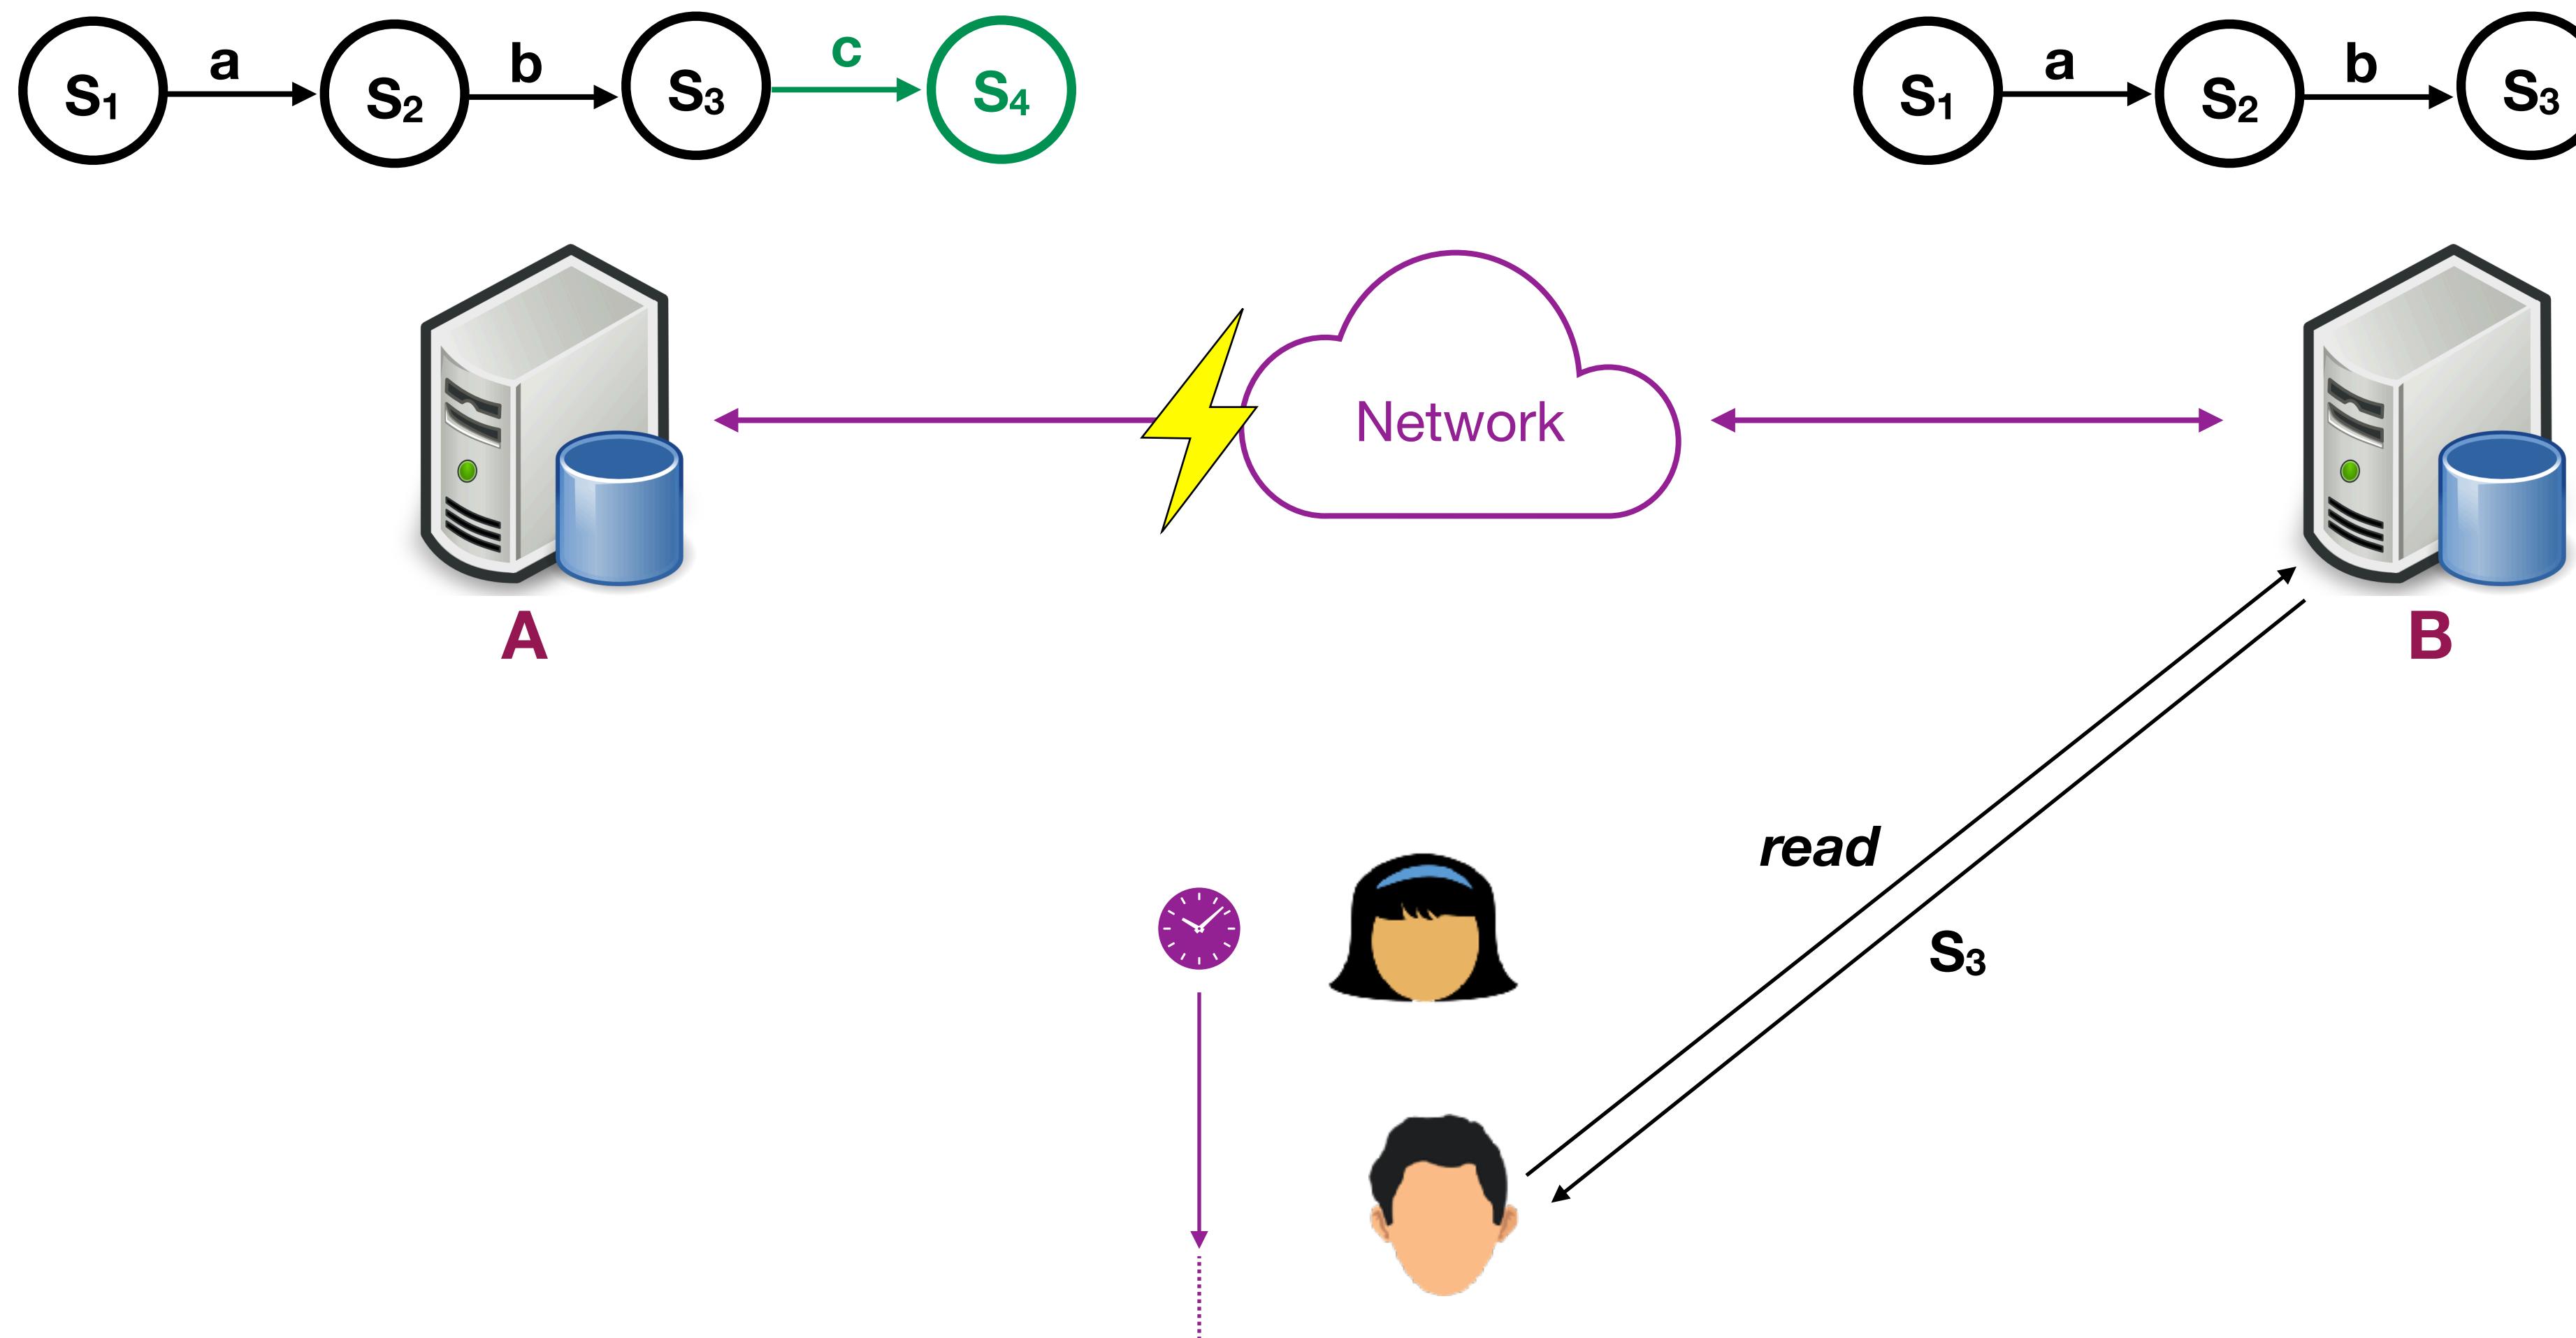
# Distributed Systems In Practice



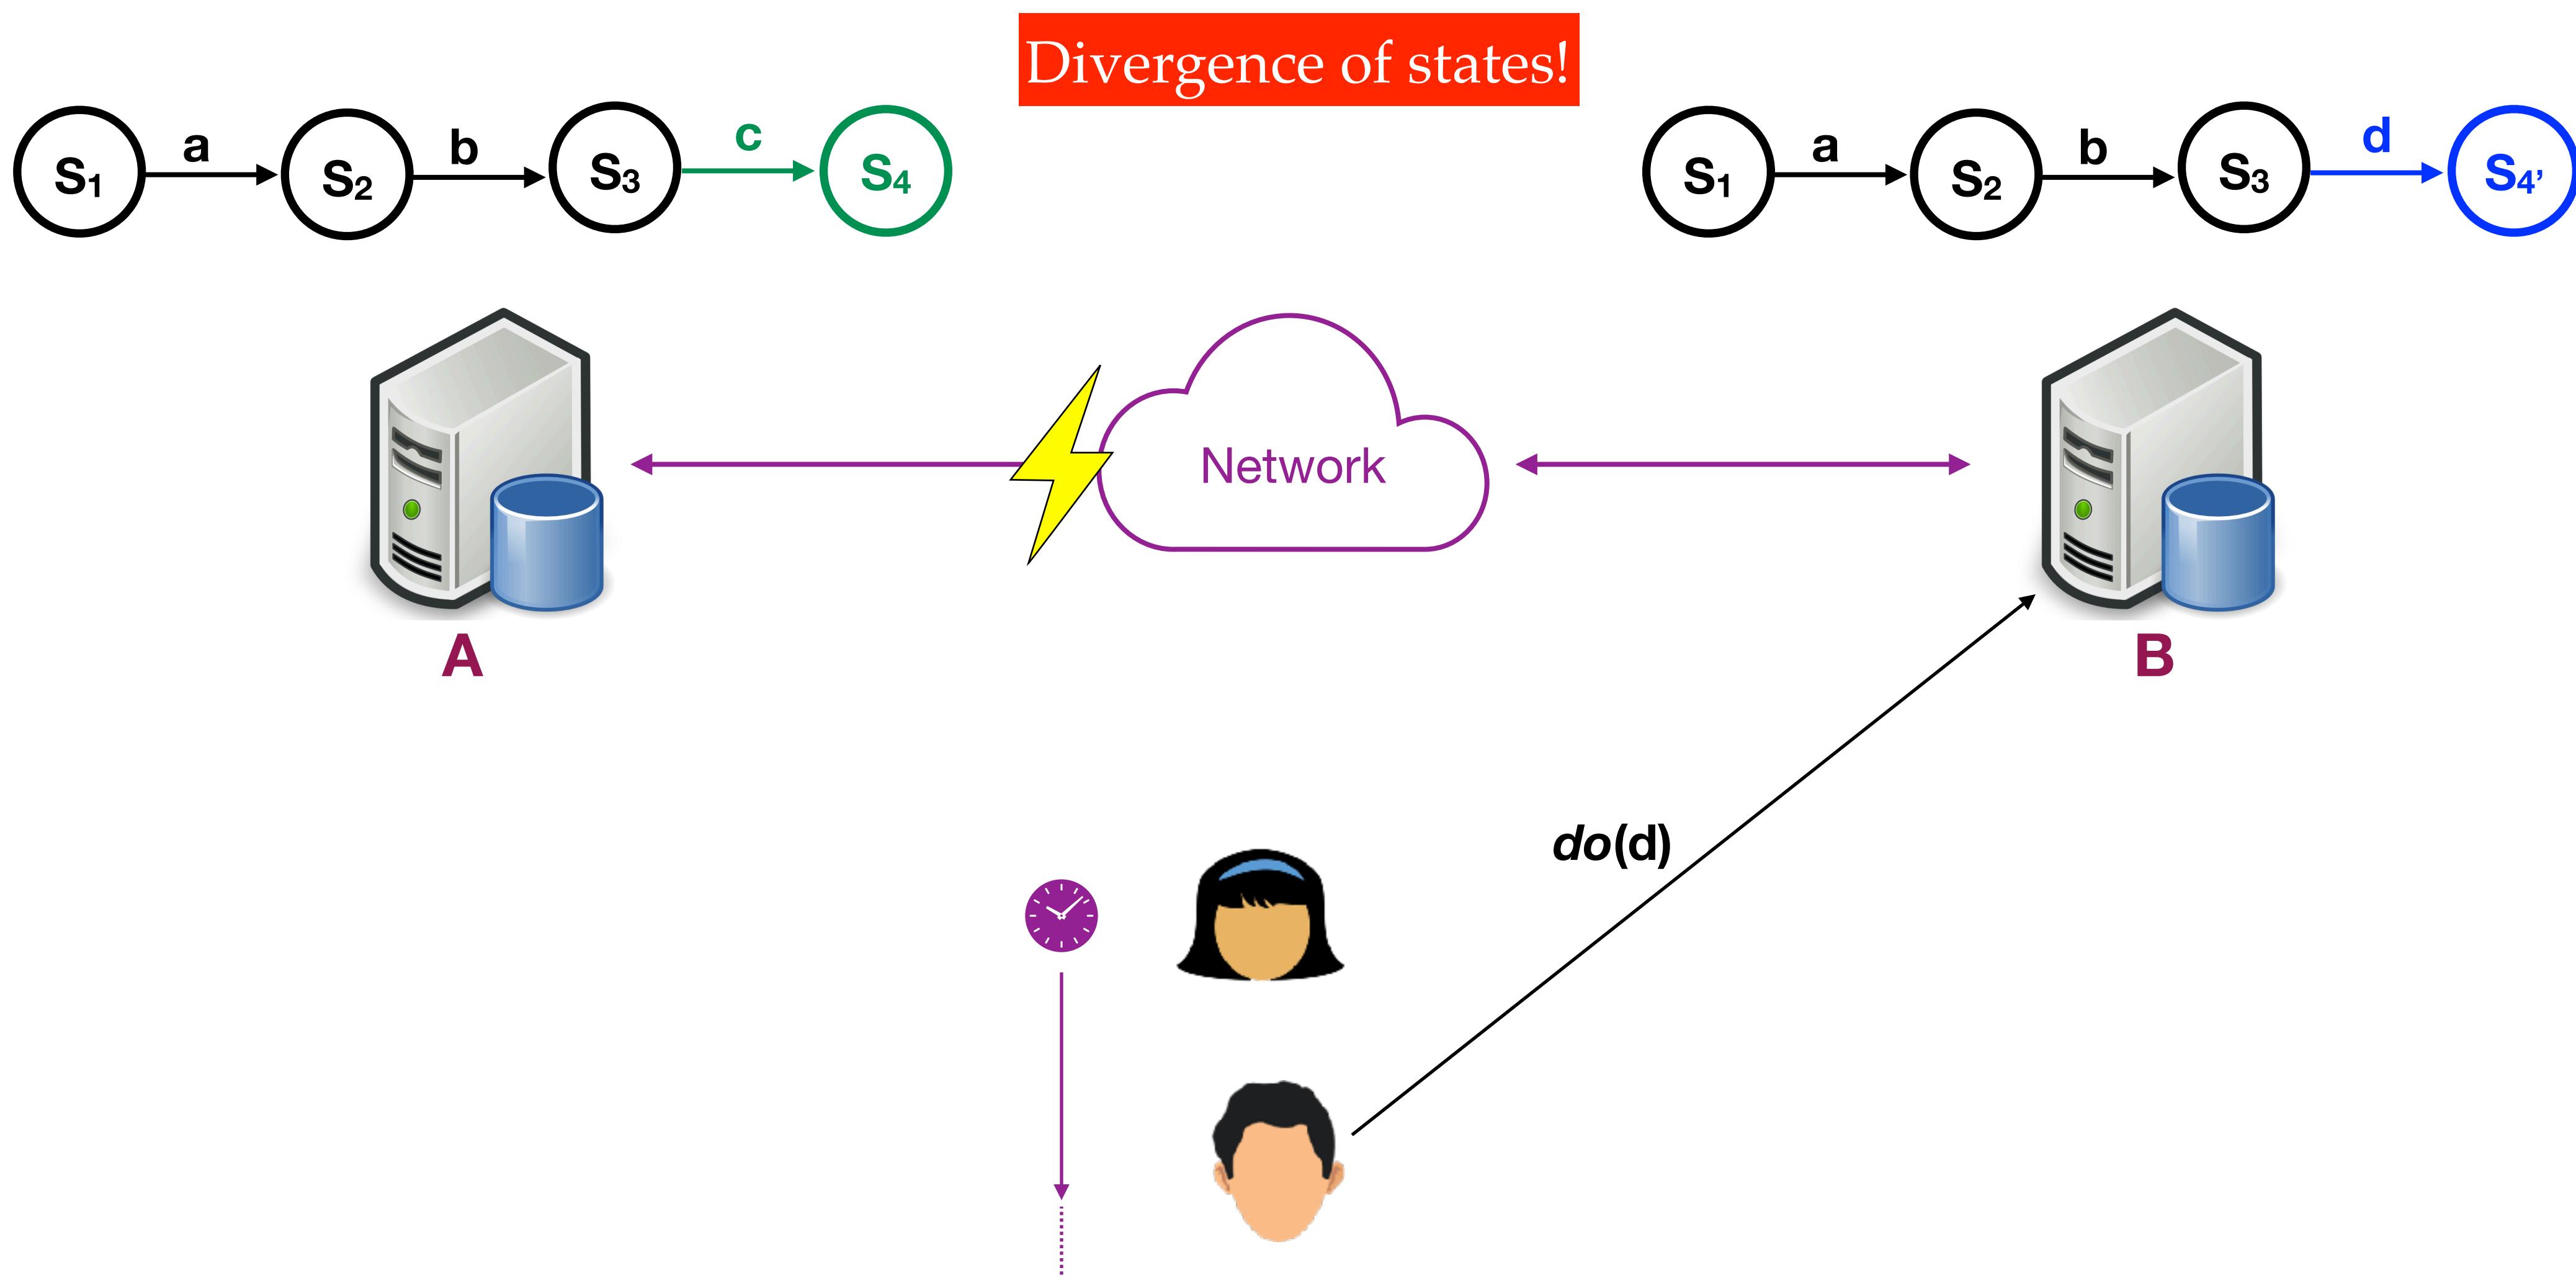
# Distributed Systems In Practice



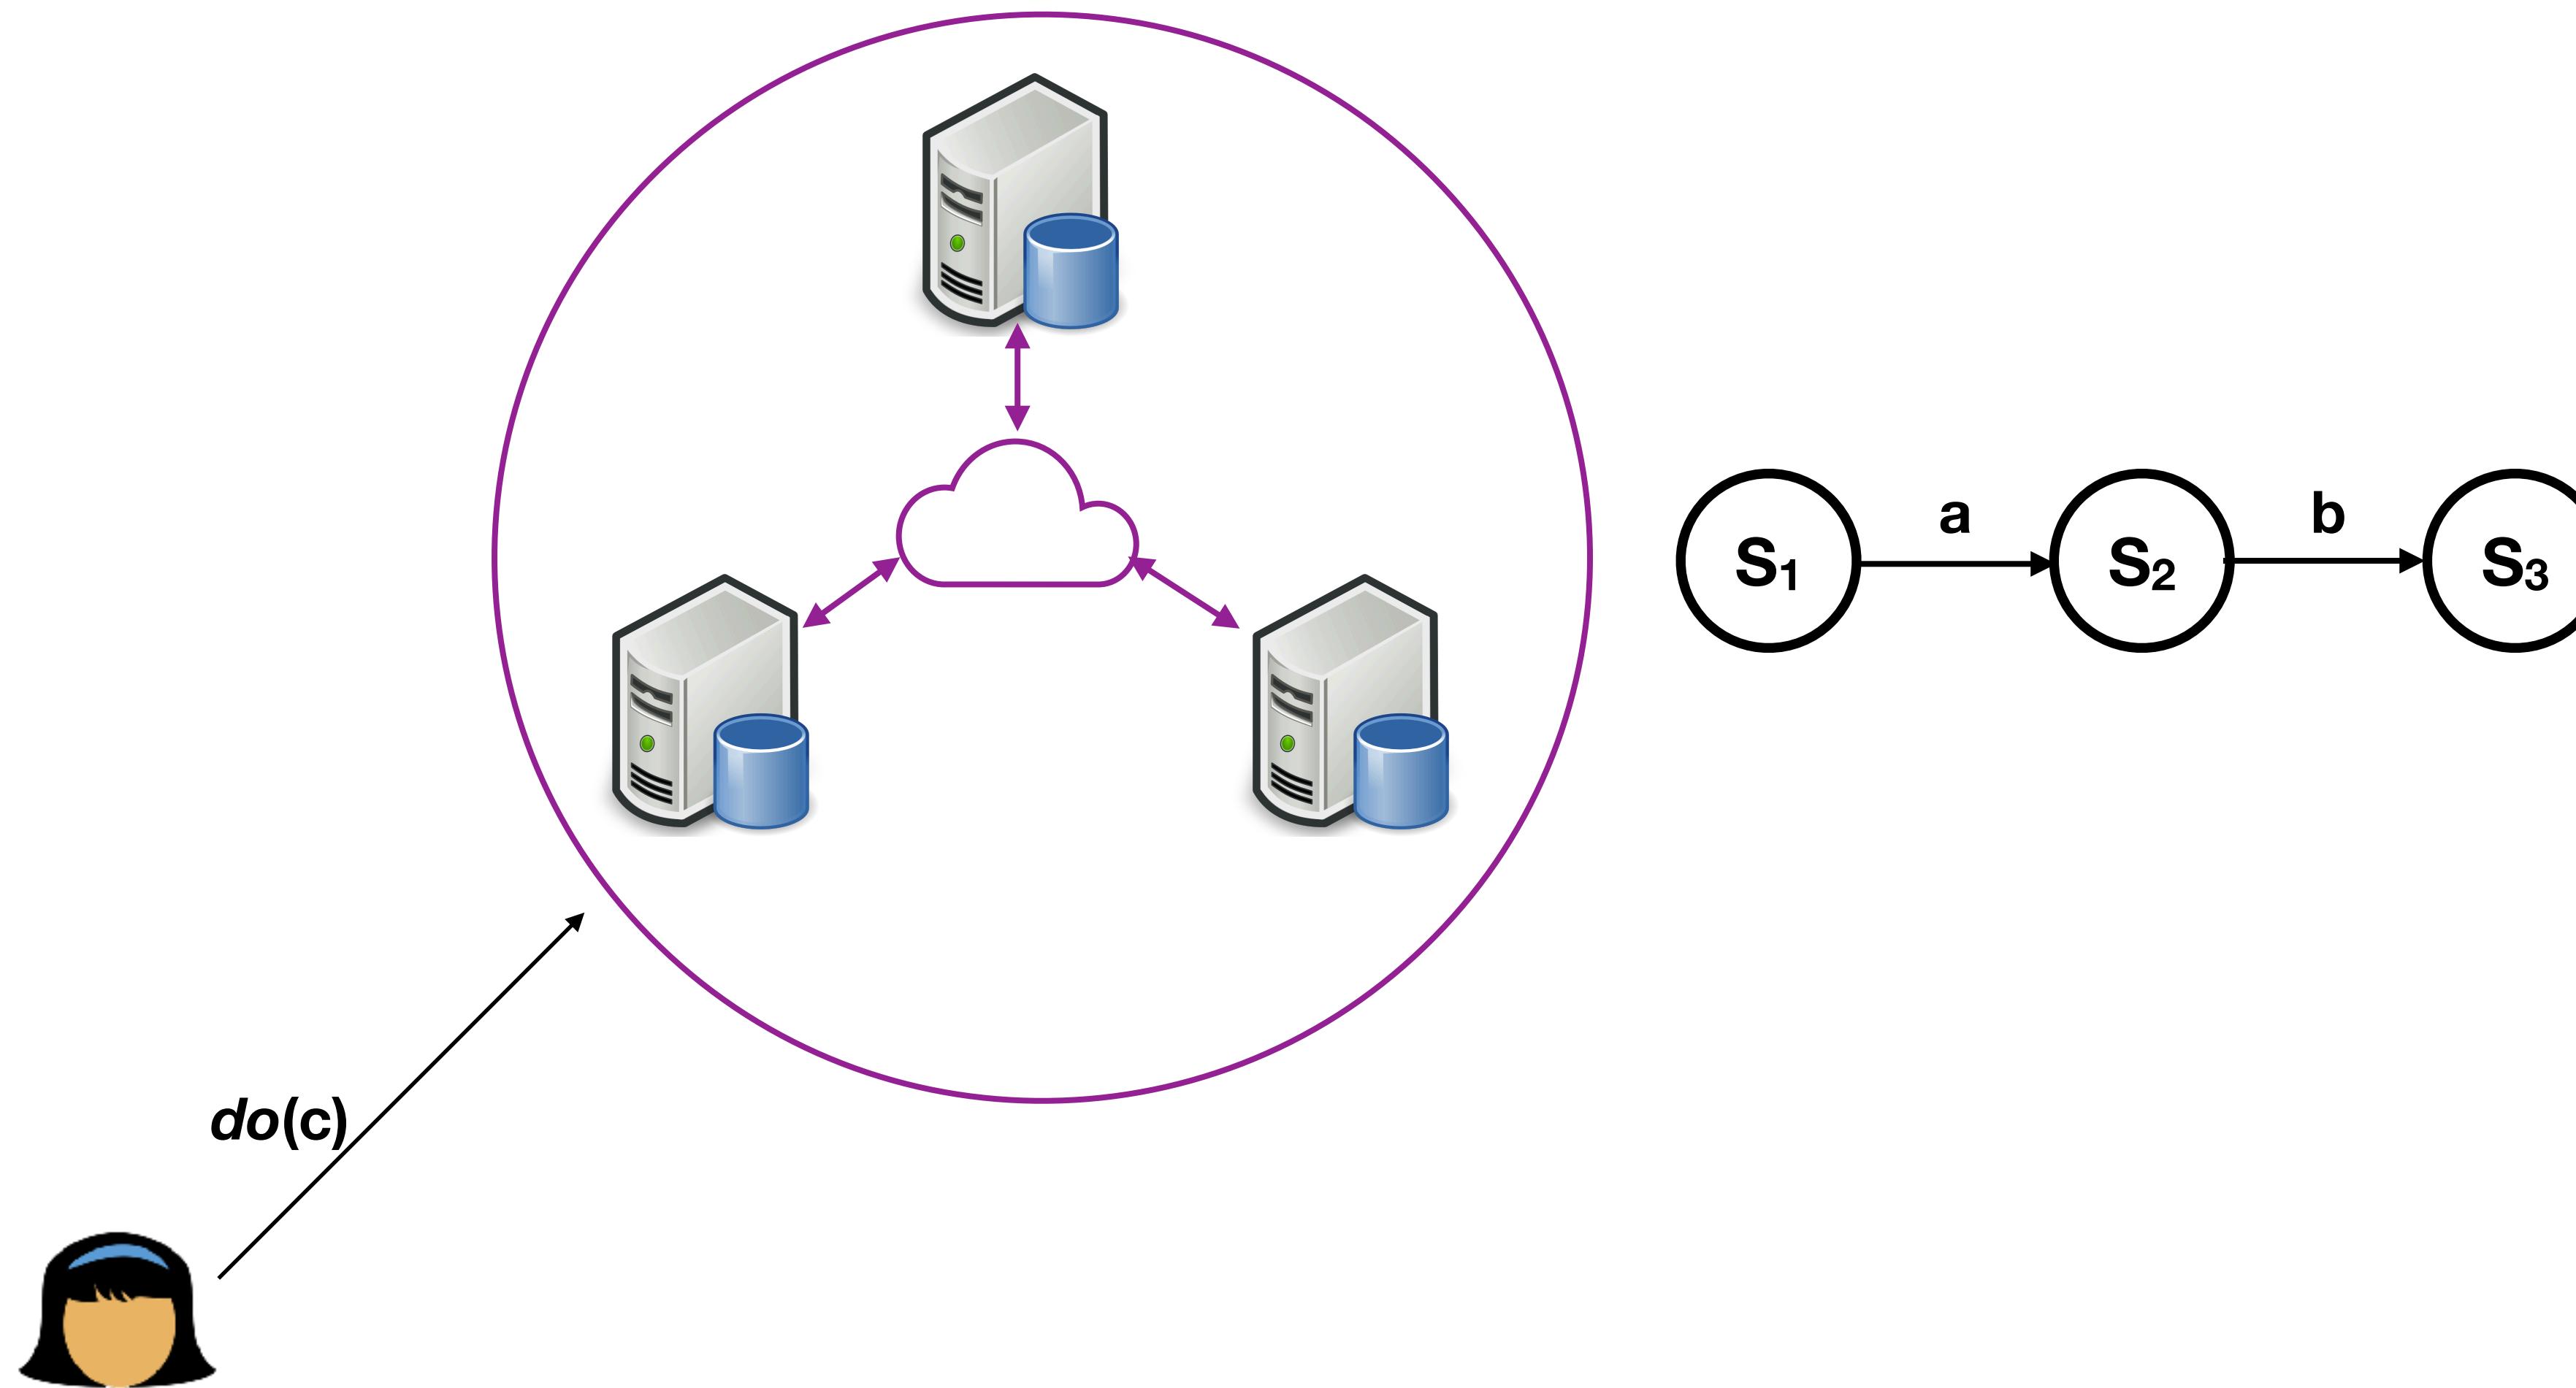
# Distributed Systems In Practice



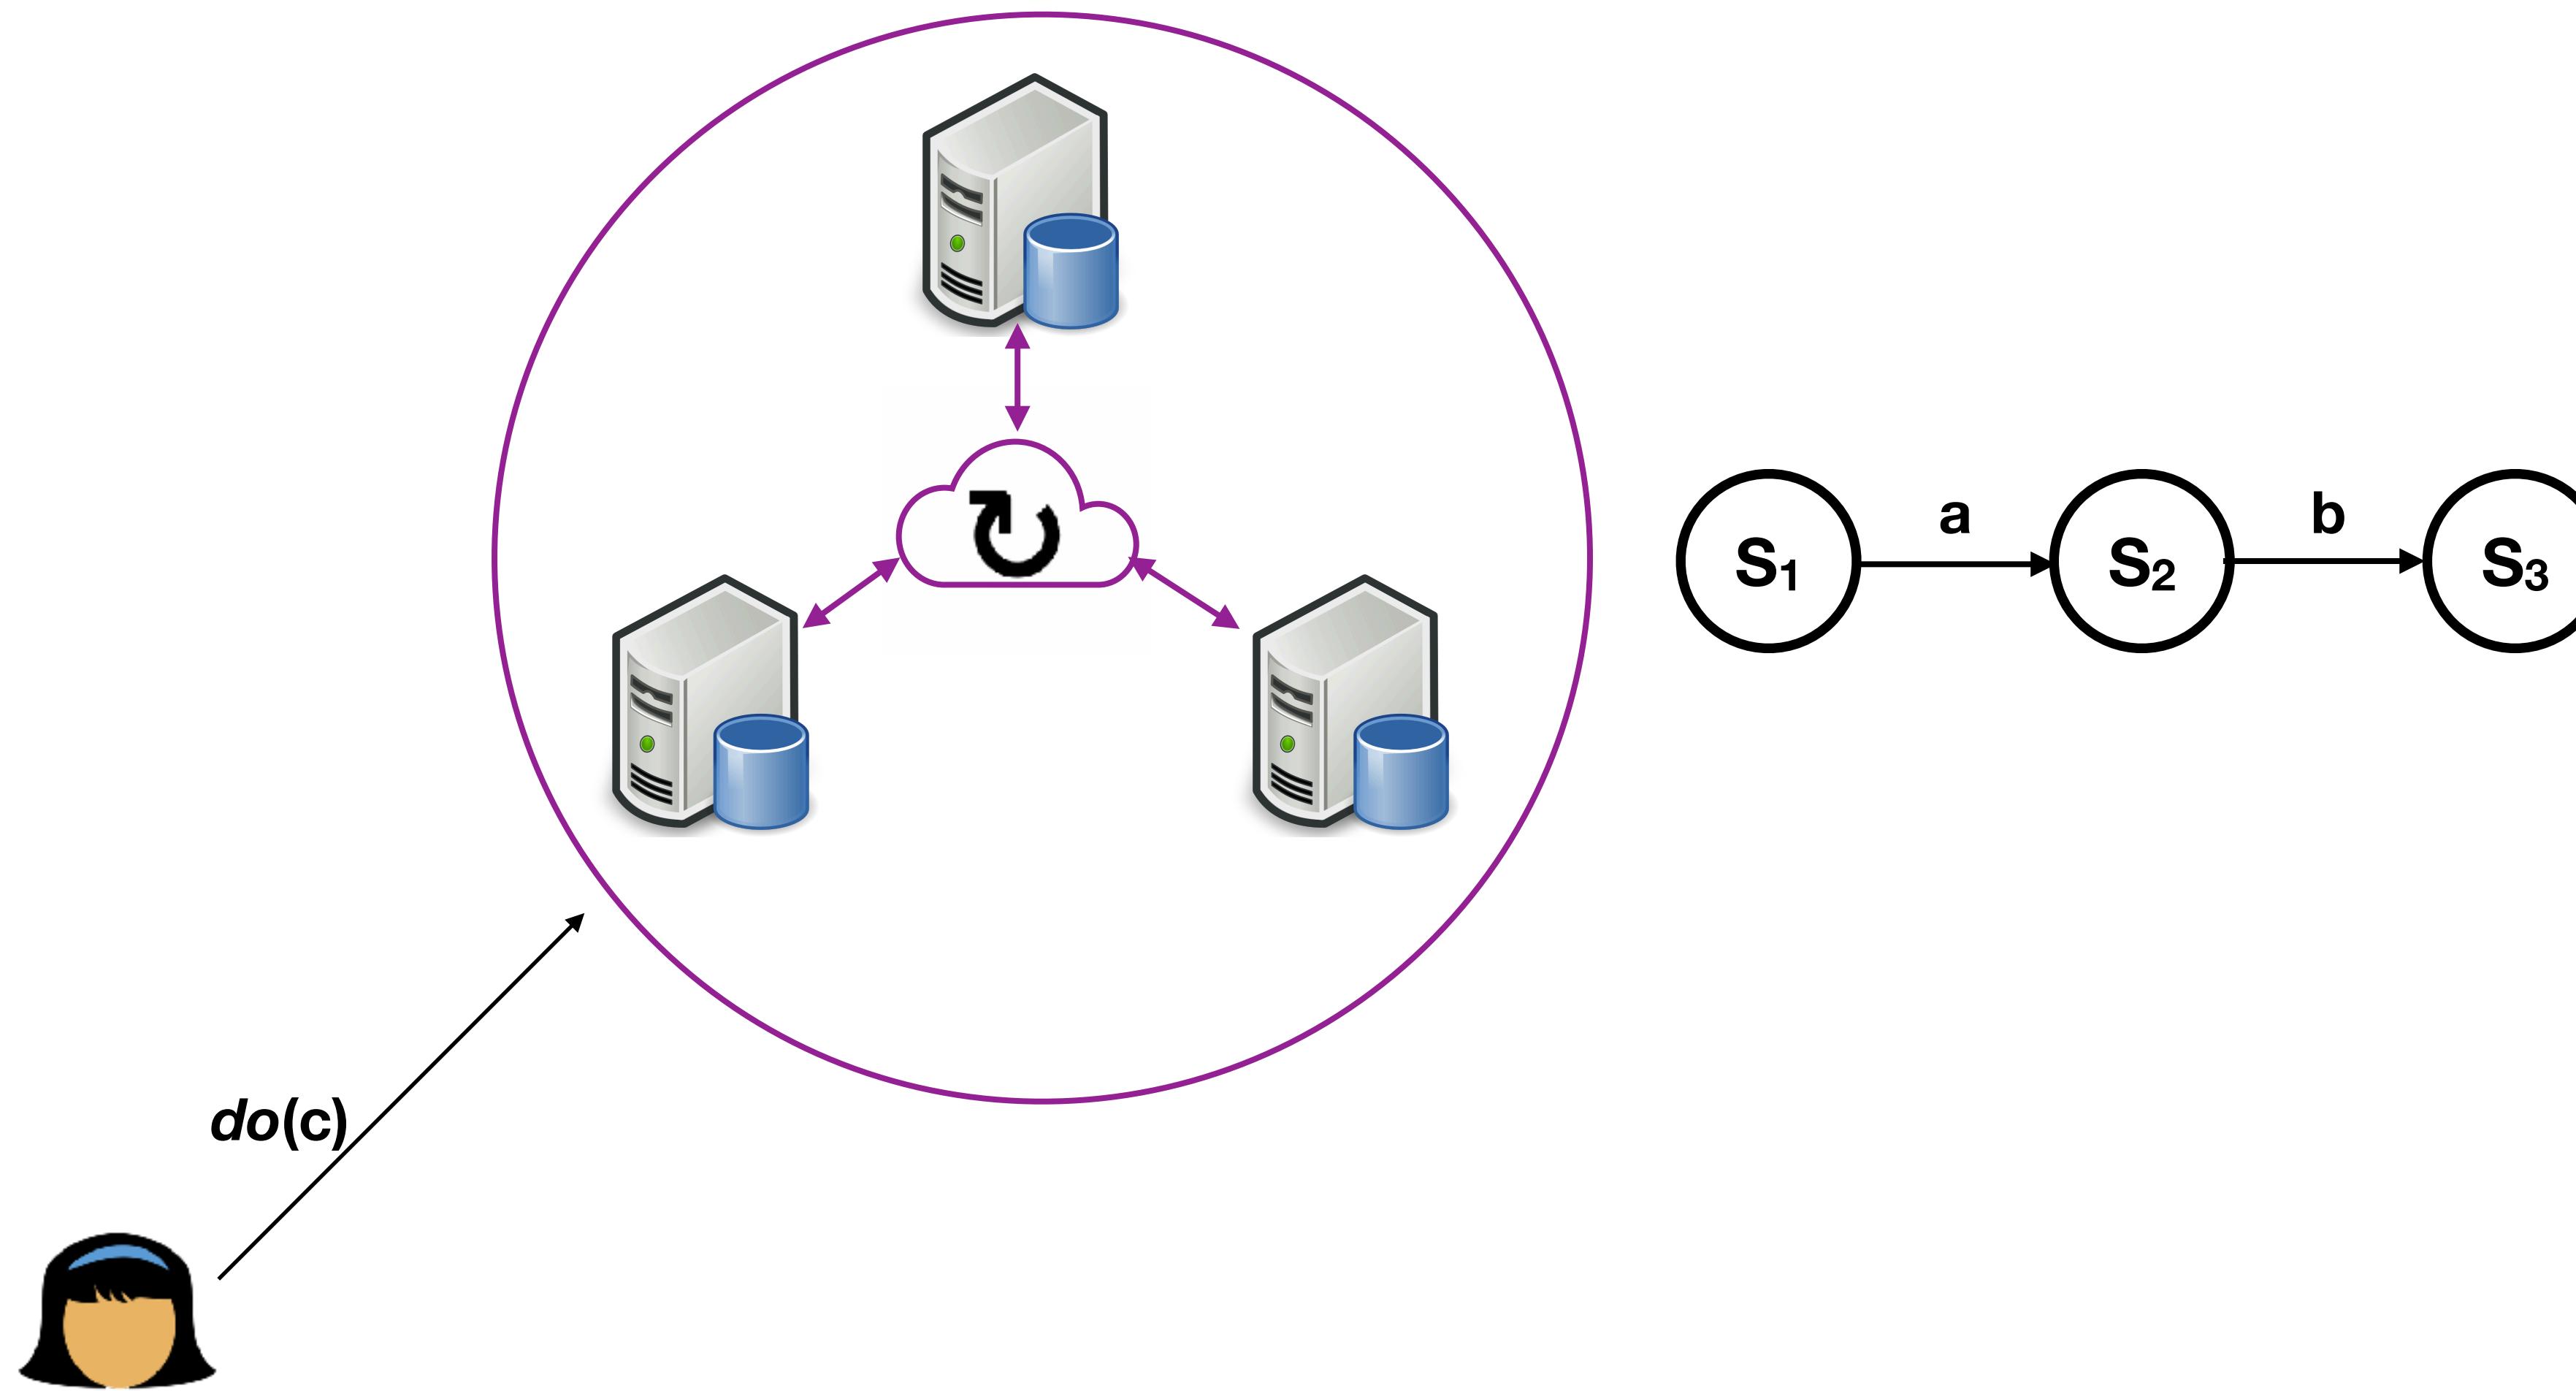
# Distributed Systems In Practice



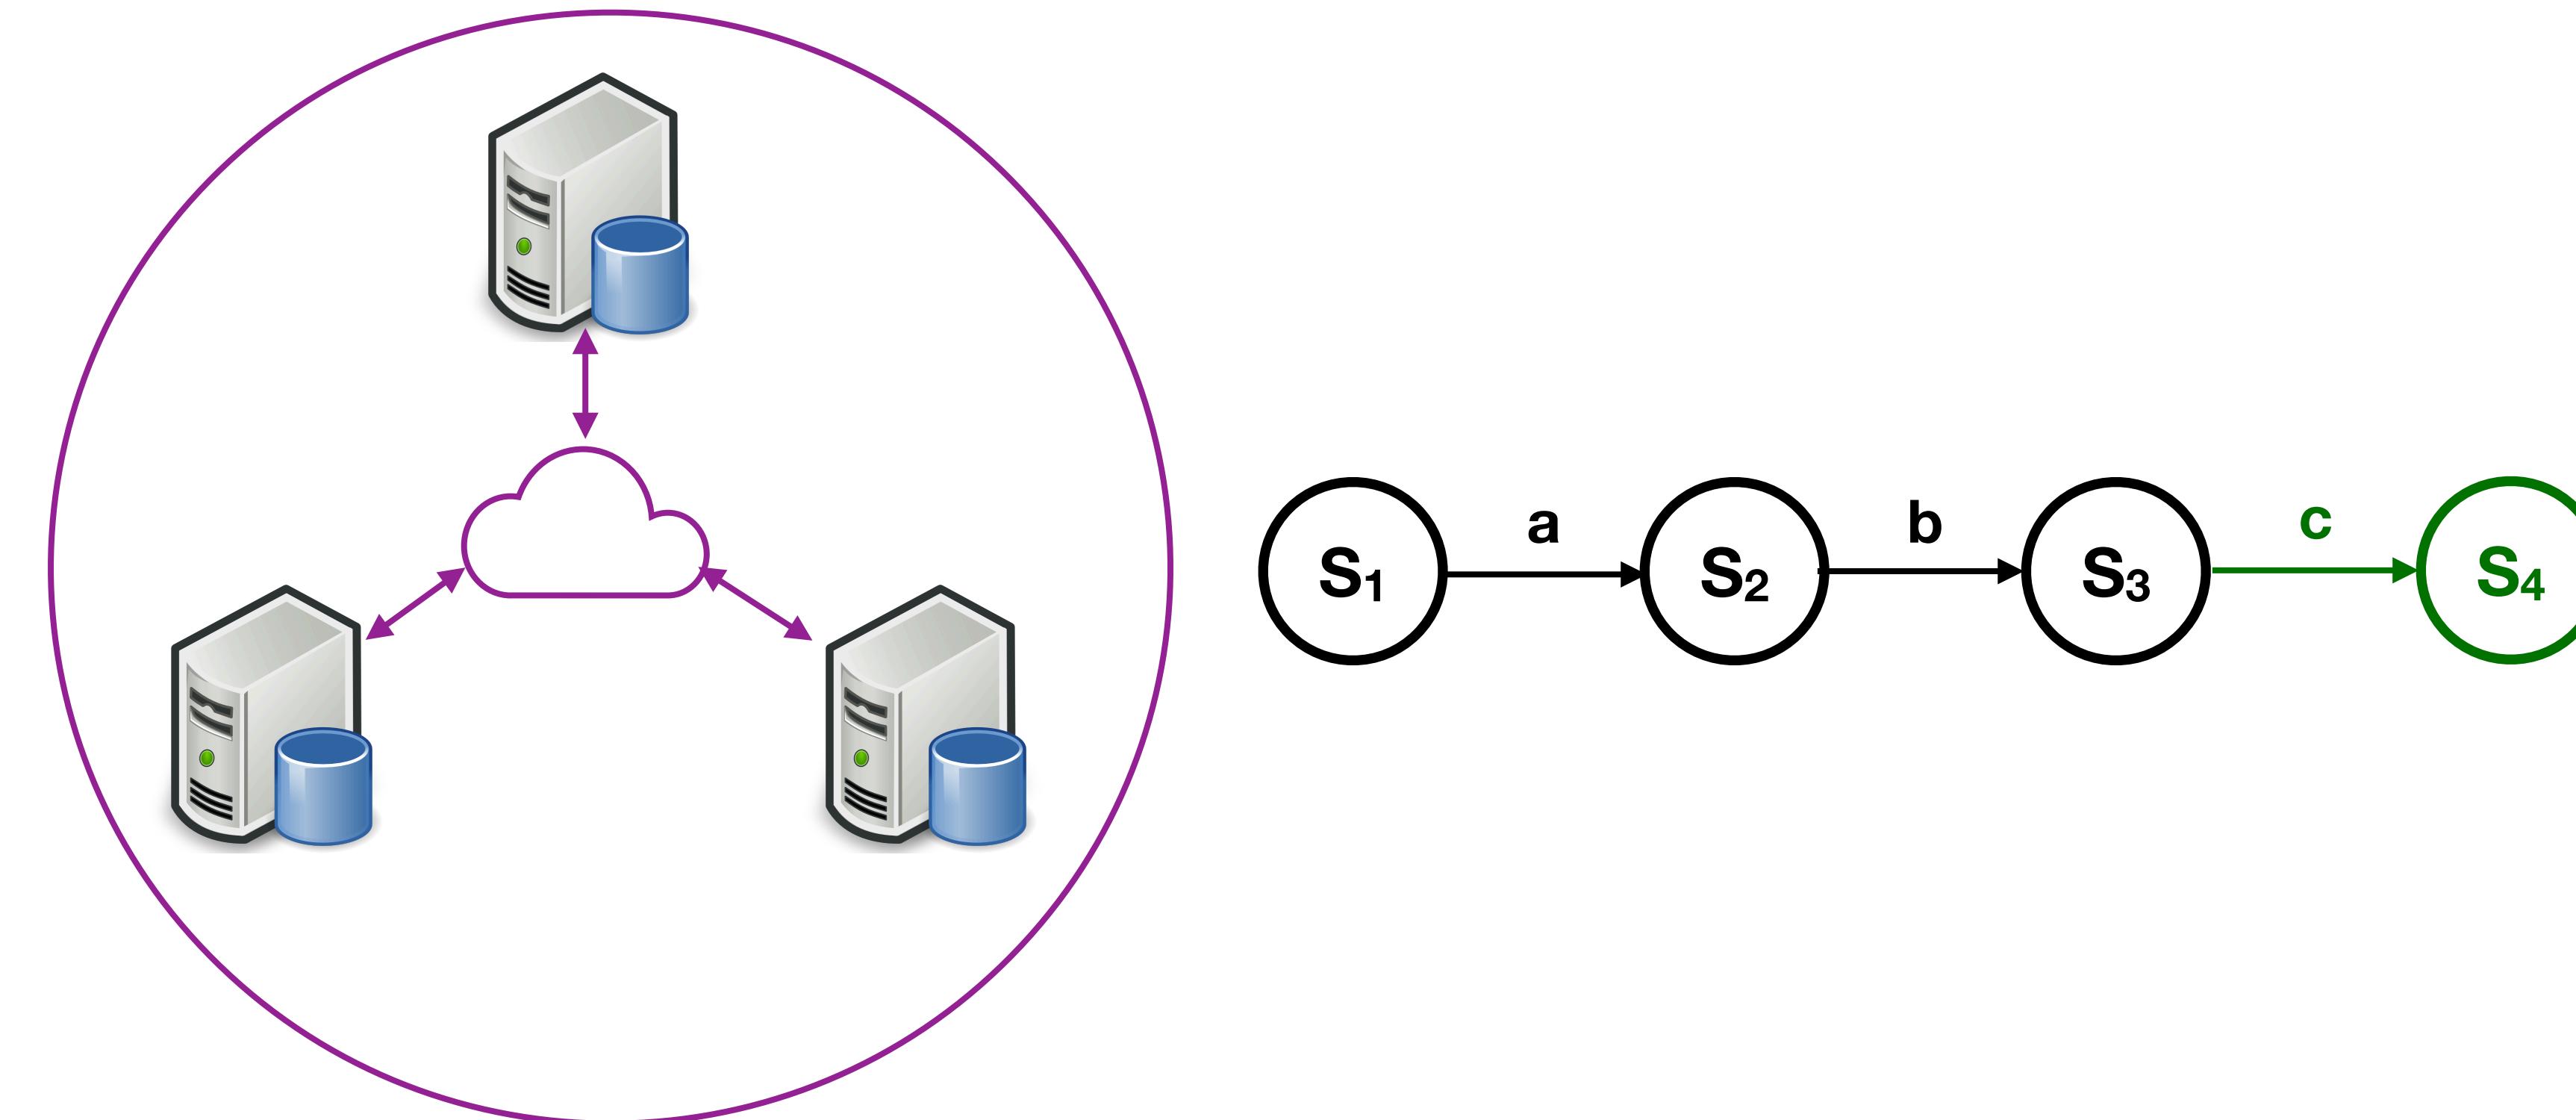
# Distributed System + Consensus = ■RSM



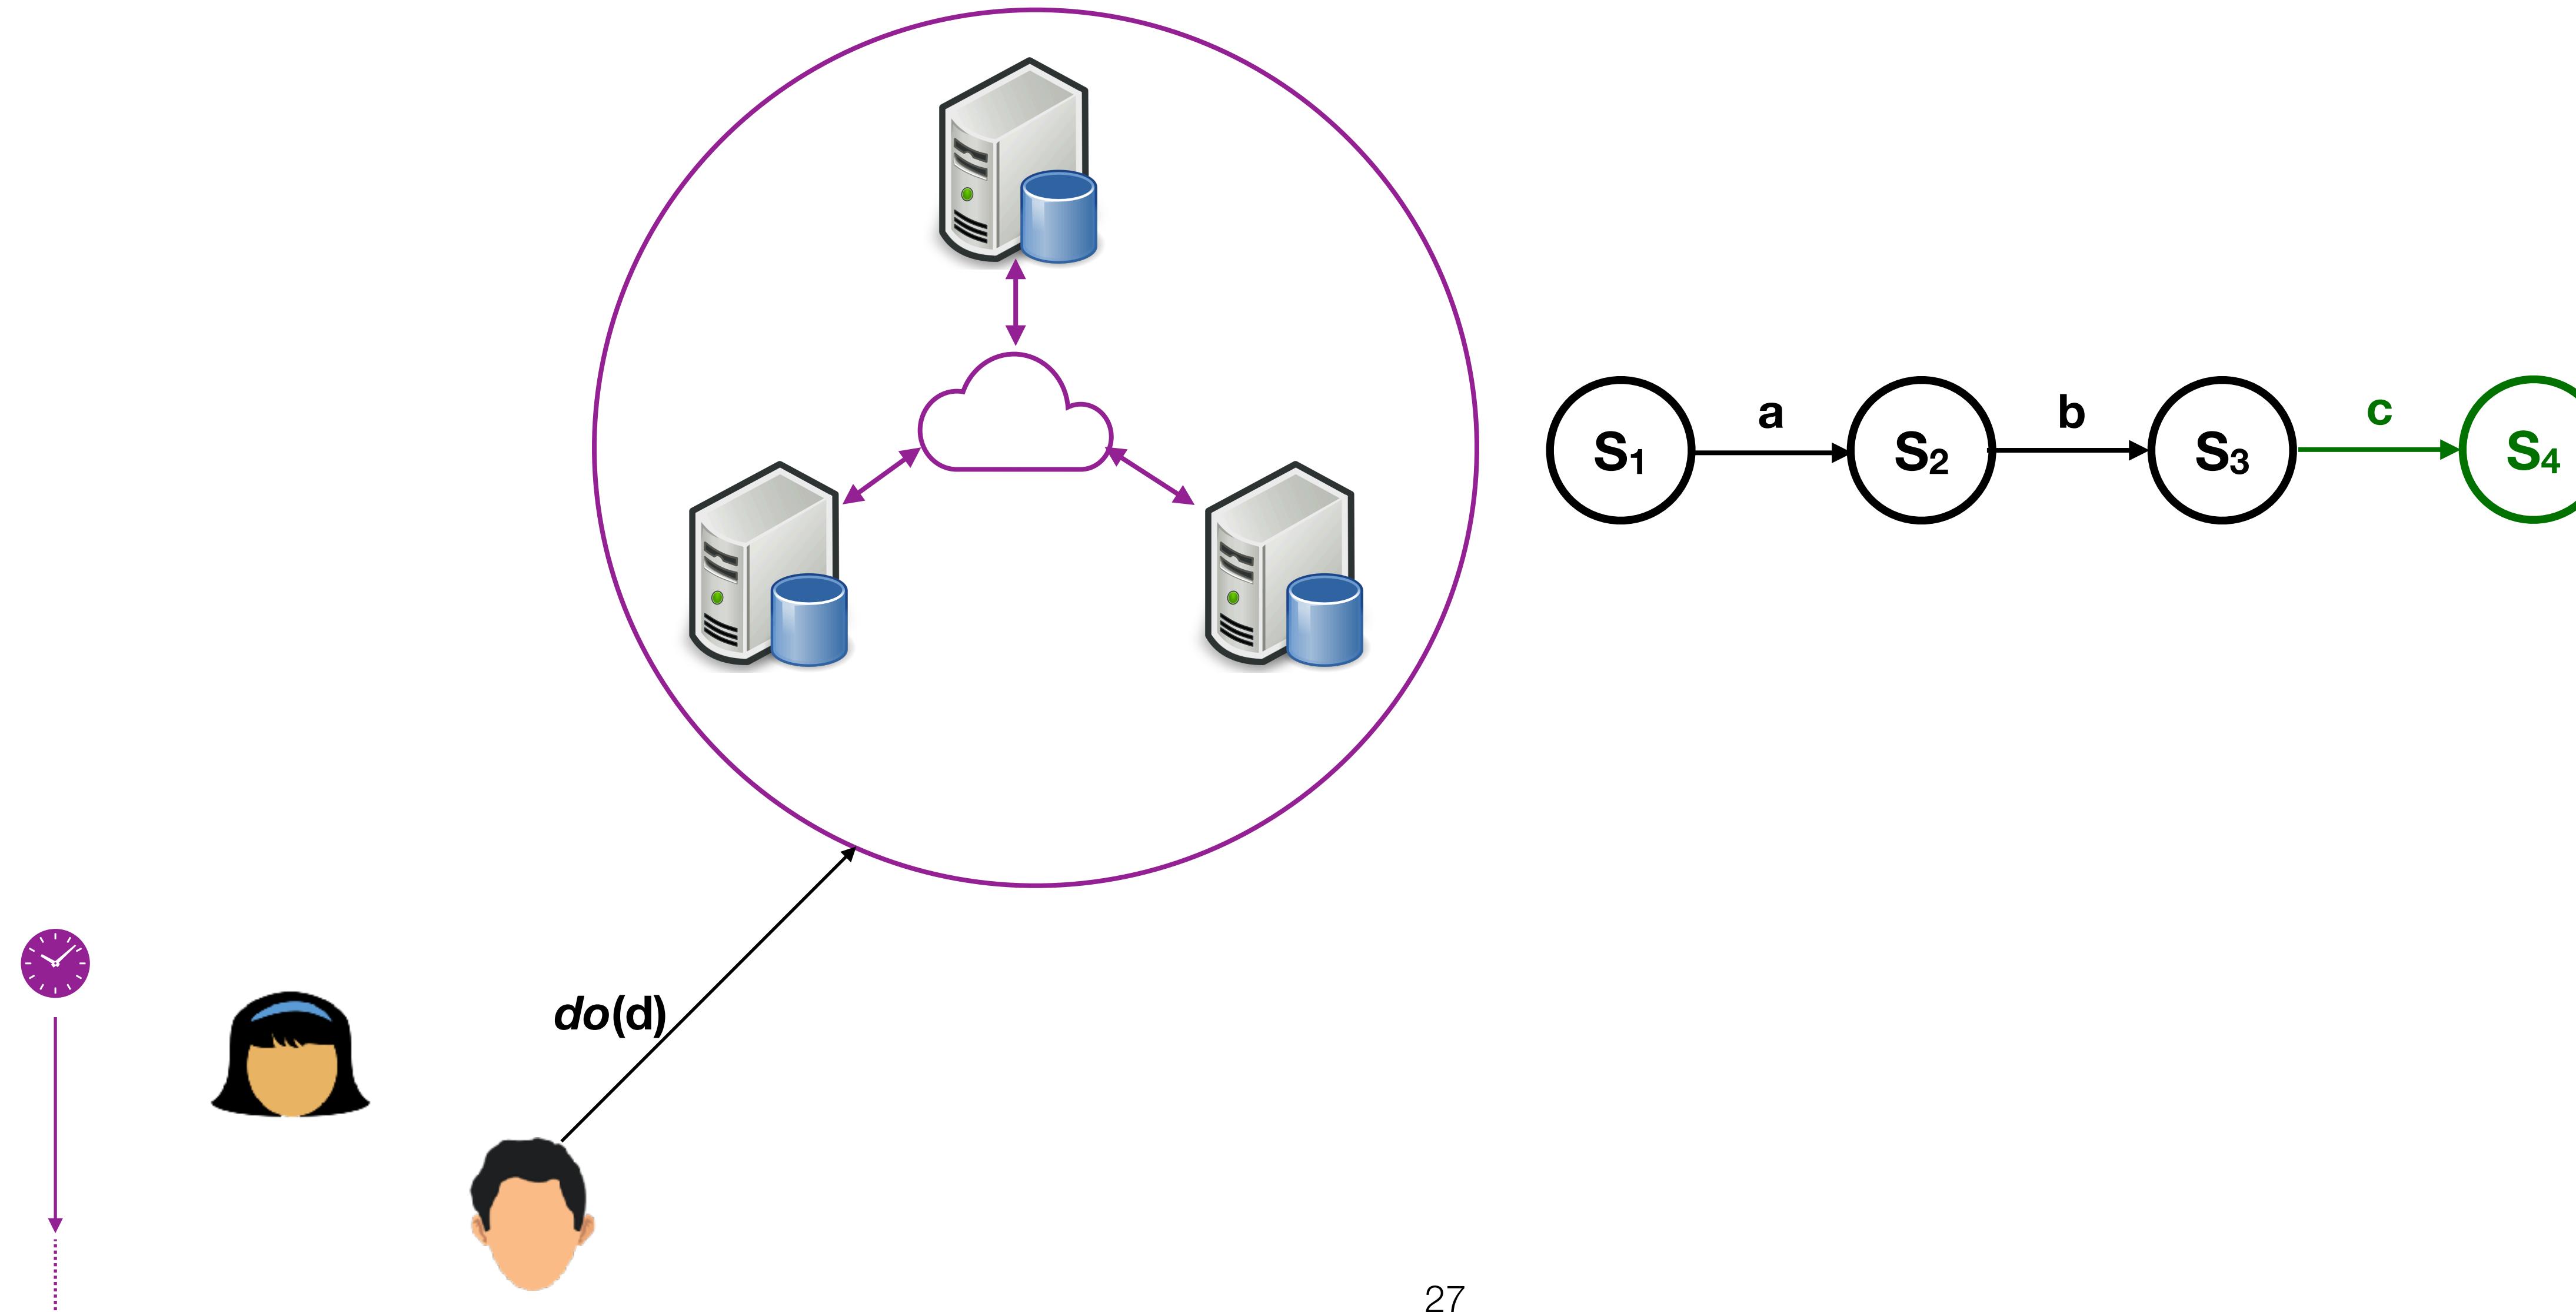
# Distributed System + Consensus = ■RSM



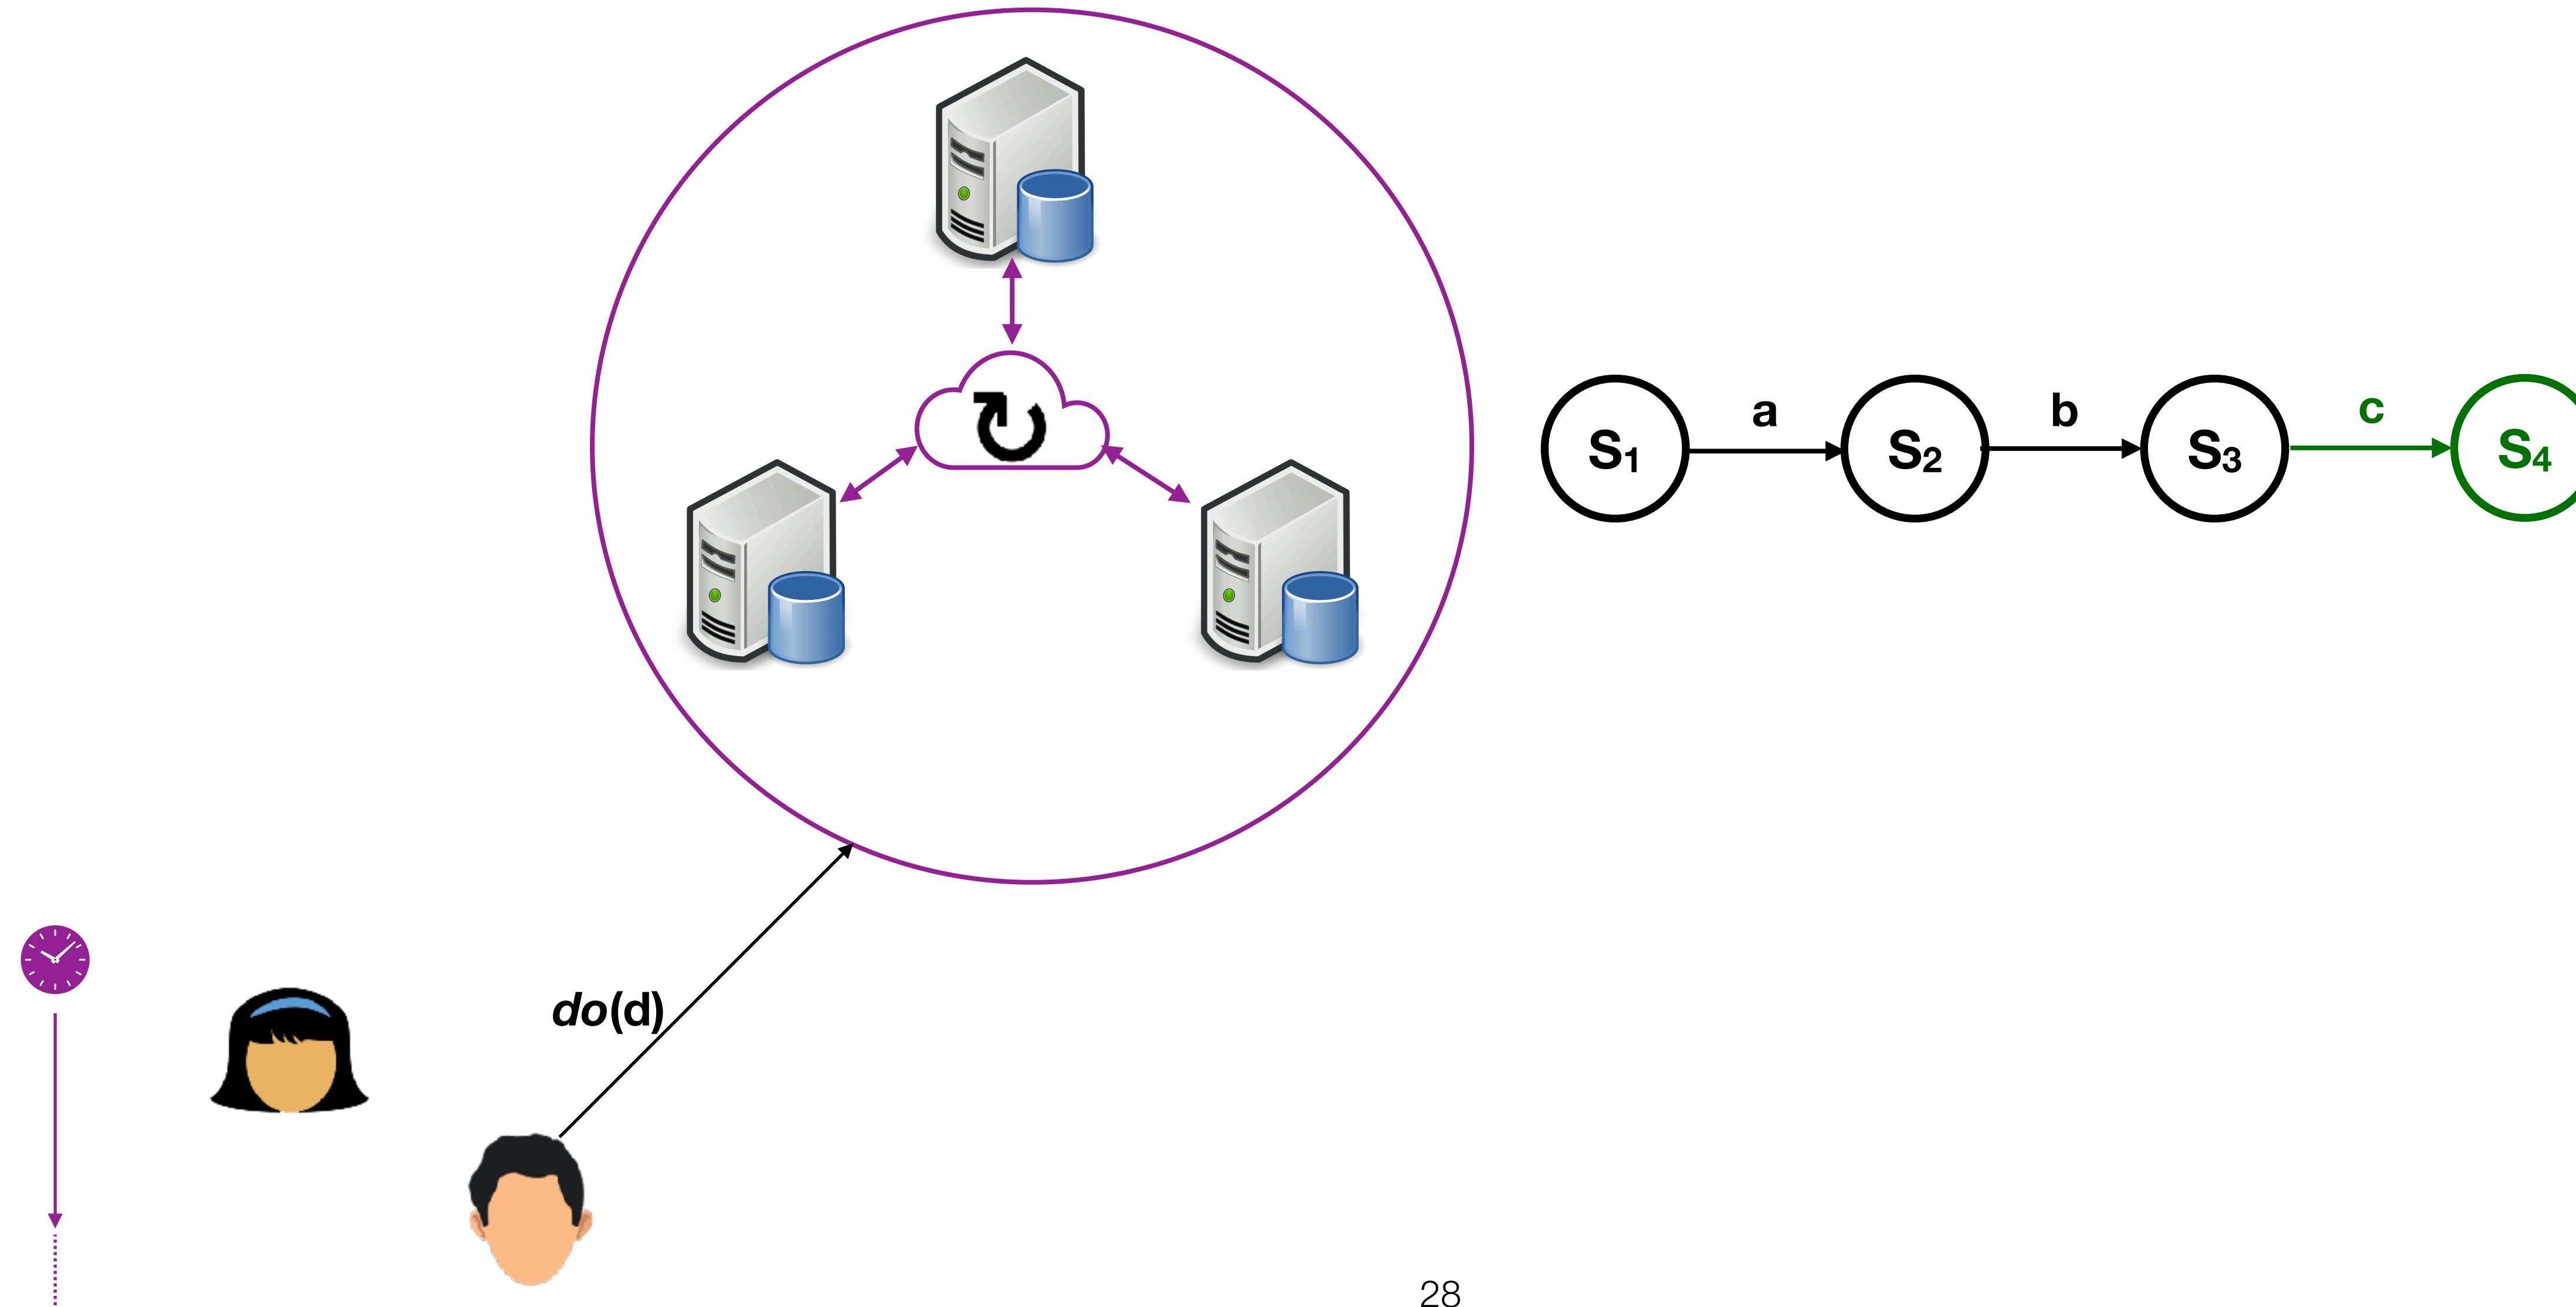
# Distributed System + Consensus = ■RSM



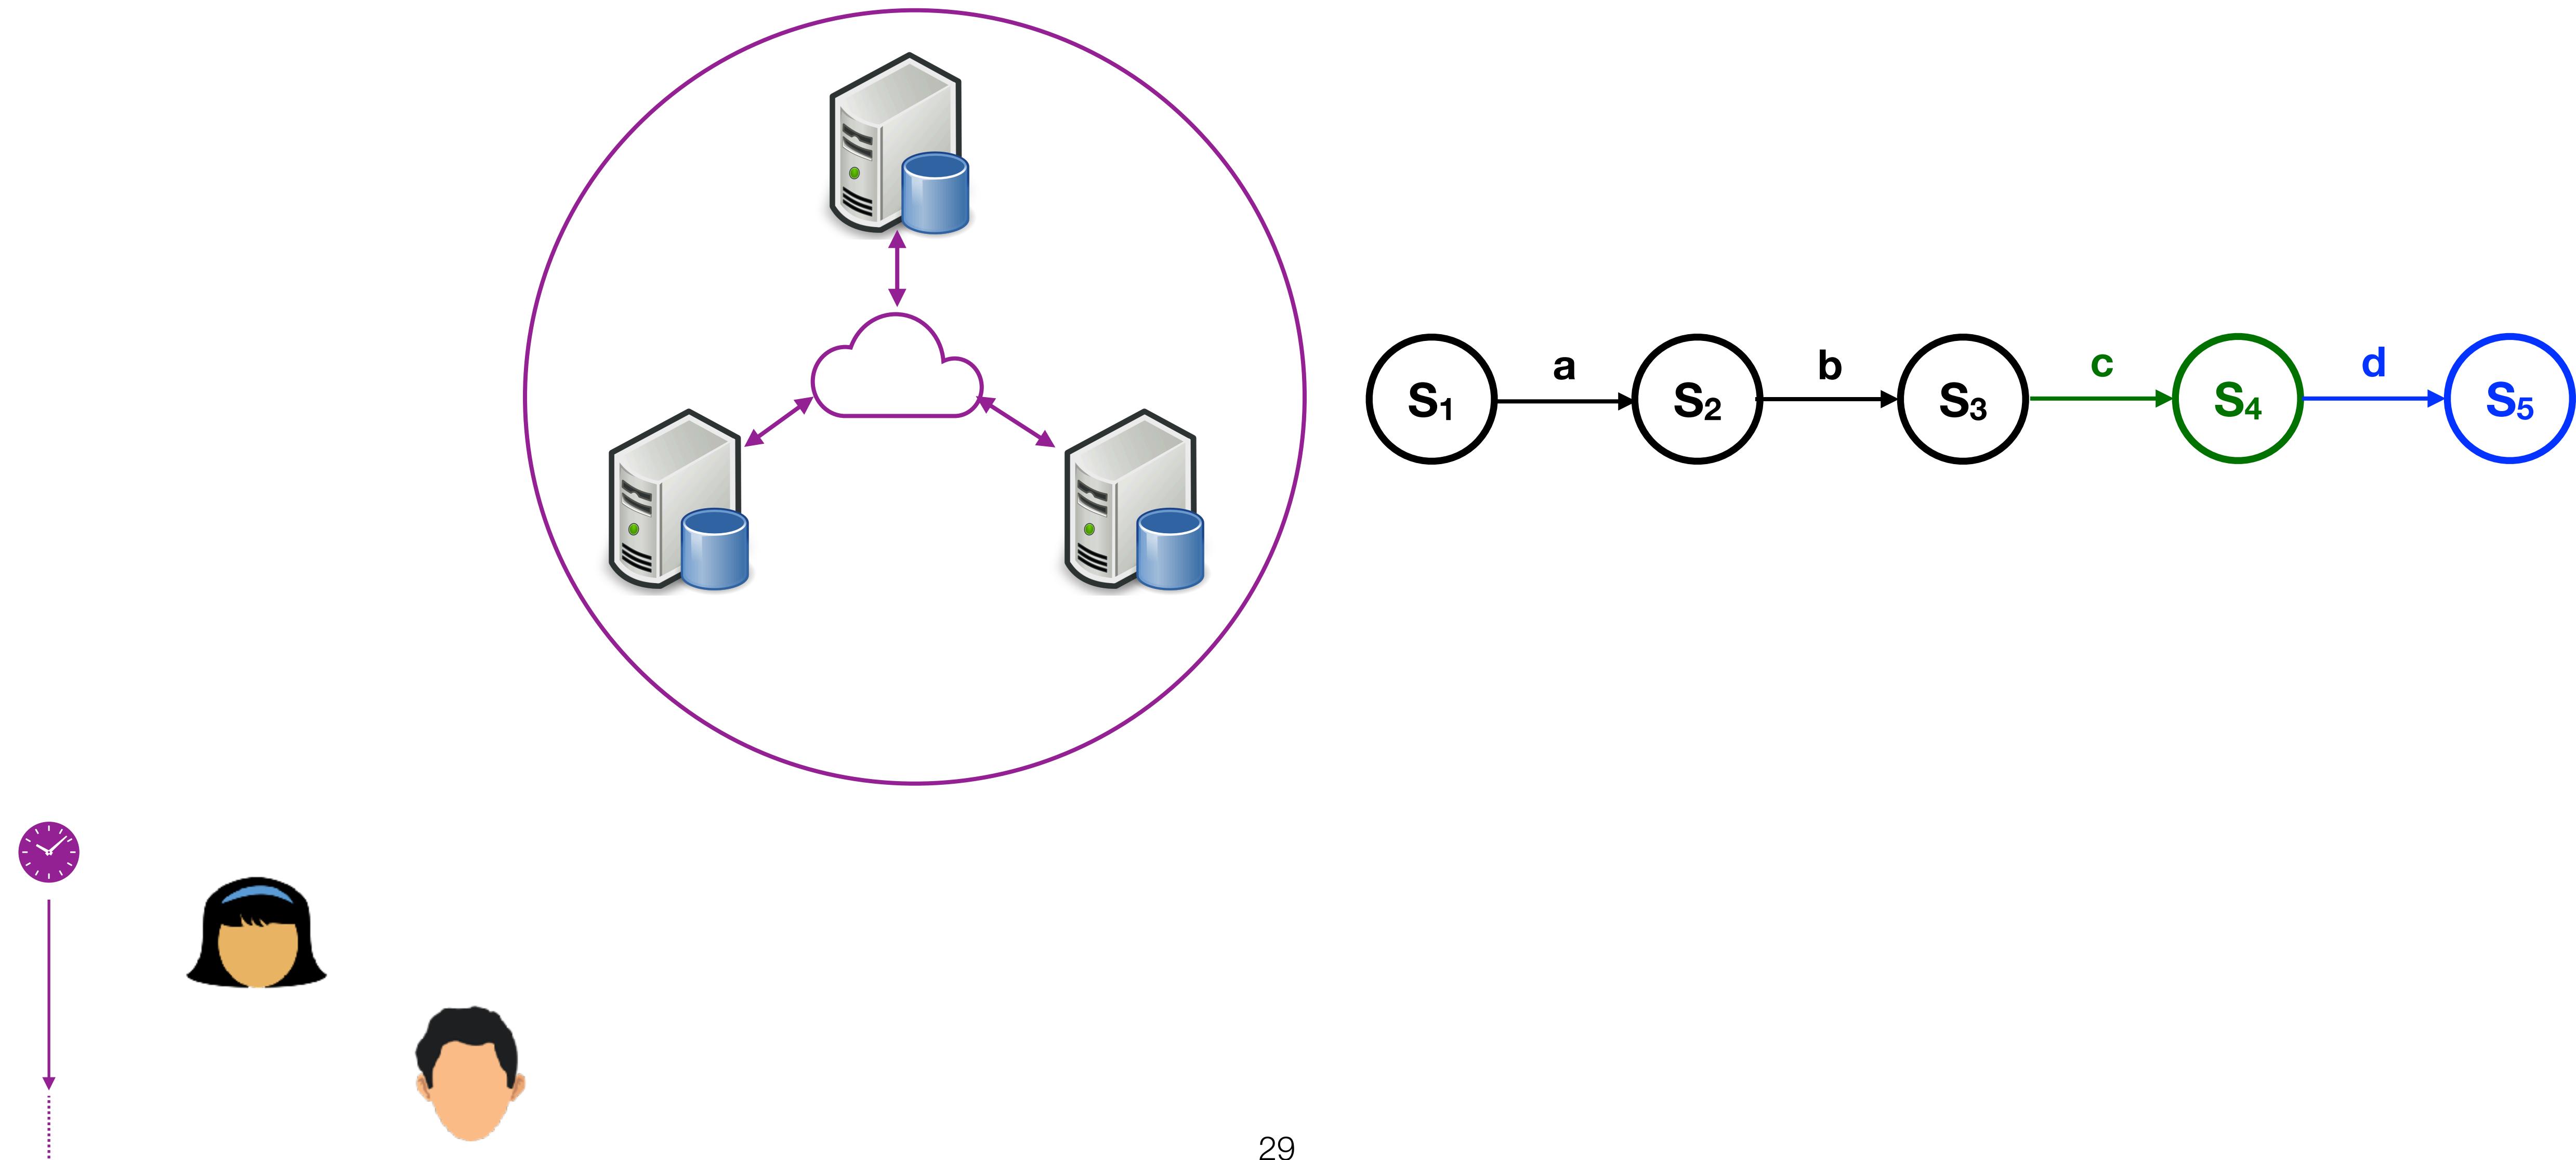
# Distributed System + Consensus = ■RSM



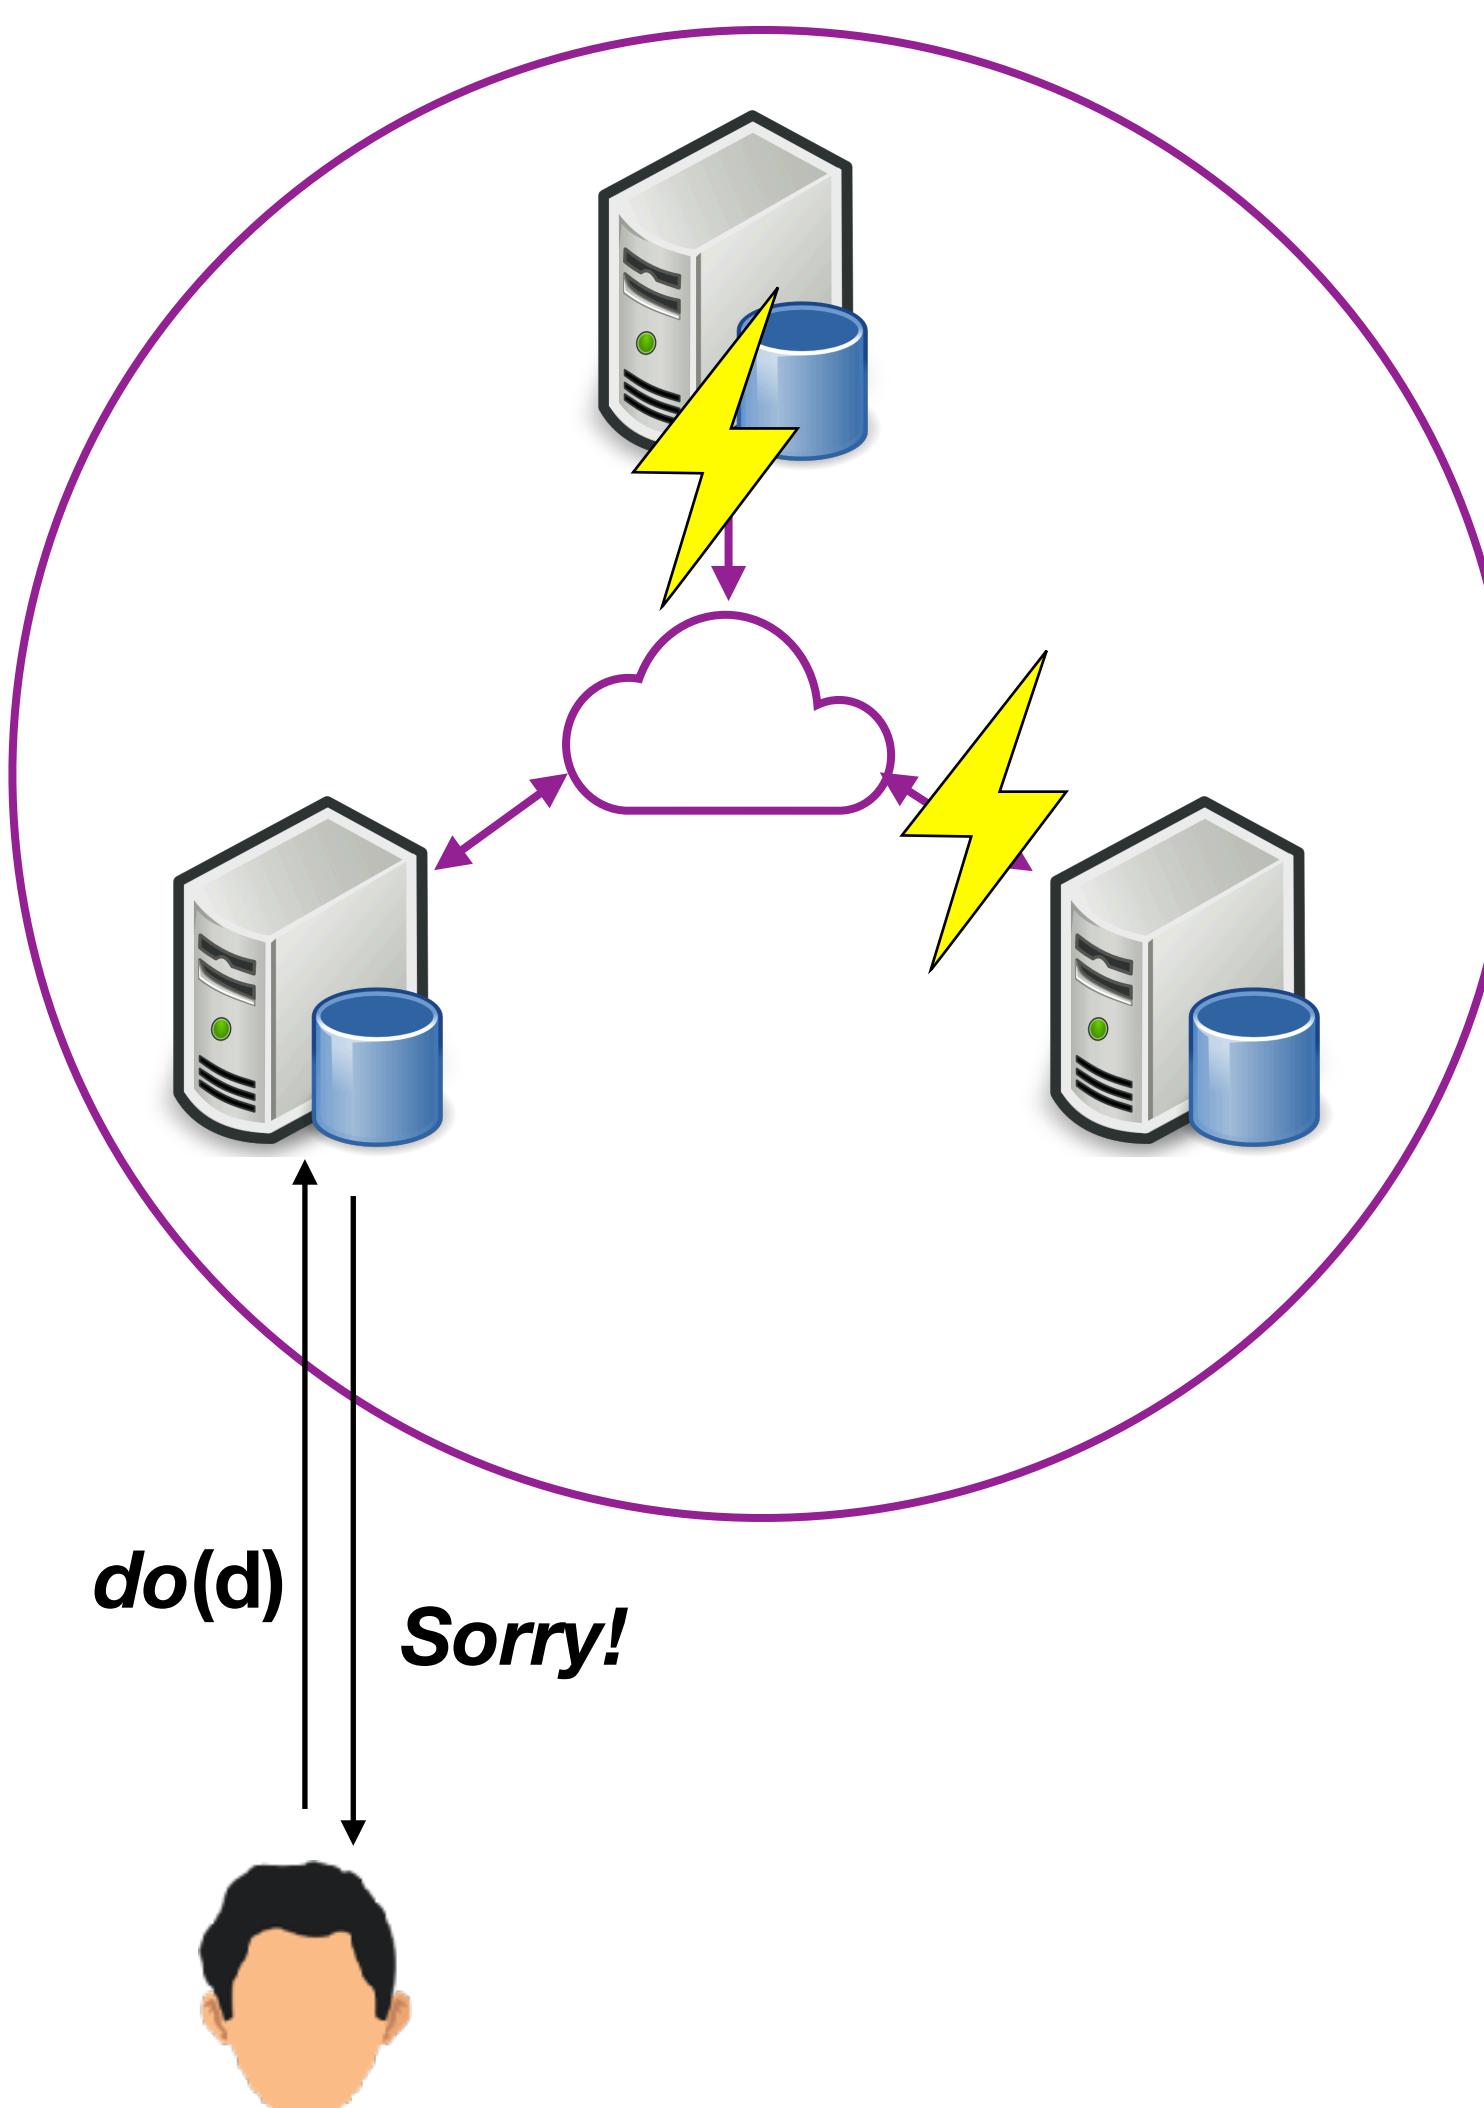
# Distributed System + Consensus = ■RSM



# Distributed System + Consensus = ■RSM



# Distributed System + Consensus = ■RSM



## Consensus Protocols

- Paxos [Lamport 1998]
  - Multi-Paxos [Lamport 2001]
  - Disk Paxos [Gafni et al 2002]
  - Fast Paxos [Lamport 2006]
  - Stoppable Paxos [Malkhi et al 2008]
  - ...
- Raft [Ongaro et al 2014]
  - Muti-Region Raft [Cockroach DB, TikV DB]
  - FlexiRaft [Yadav et al 2023]
- Viewstamped Replication [Oki et al 1988]
- Zookeeper [Junqueira et al 2011]
- Chubby [2006]
- ...

*A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.* – Leslie Lamport (Turing award 2013).

- Negatively affect the *availability* (CAP theorem).
- Extremely complex  $\Rightarrow$  Error-prone implementations

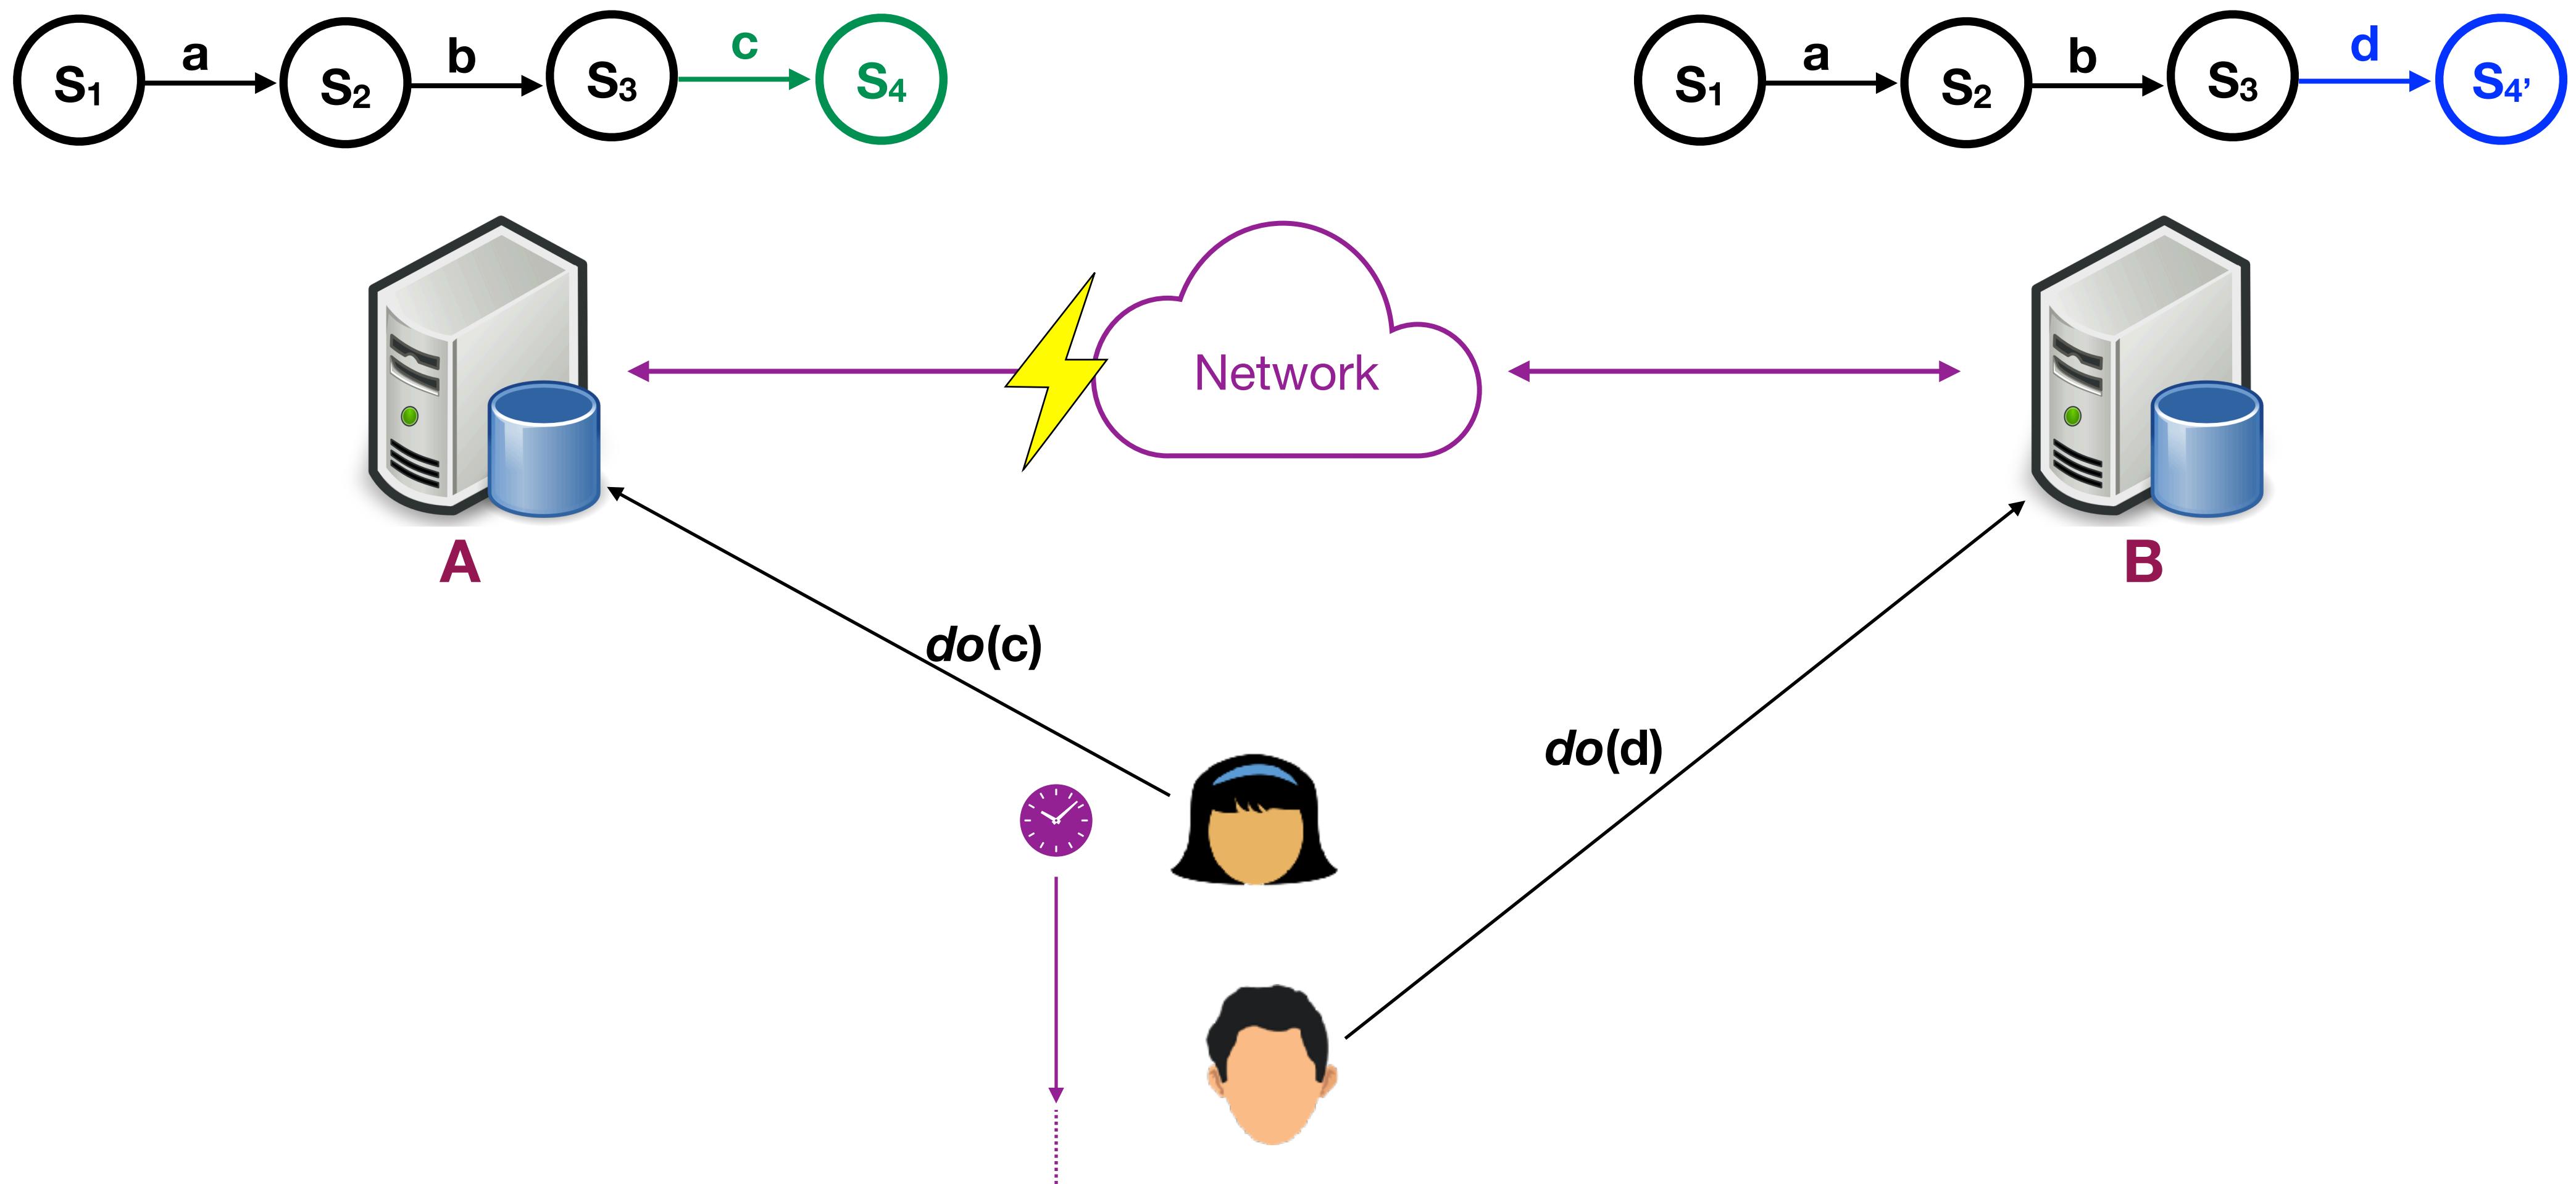
# Key Questions

Can we avoid strong consistency / consensus in practice?

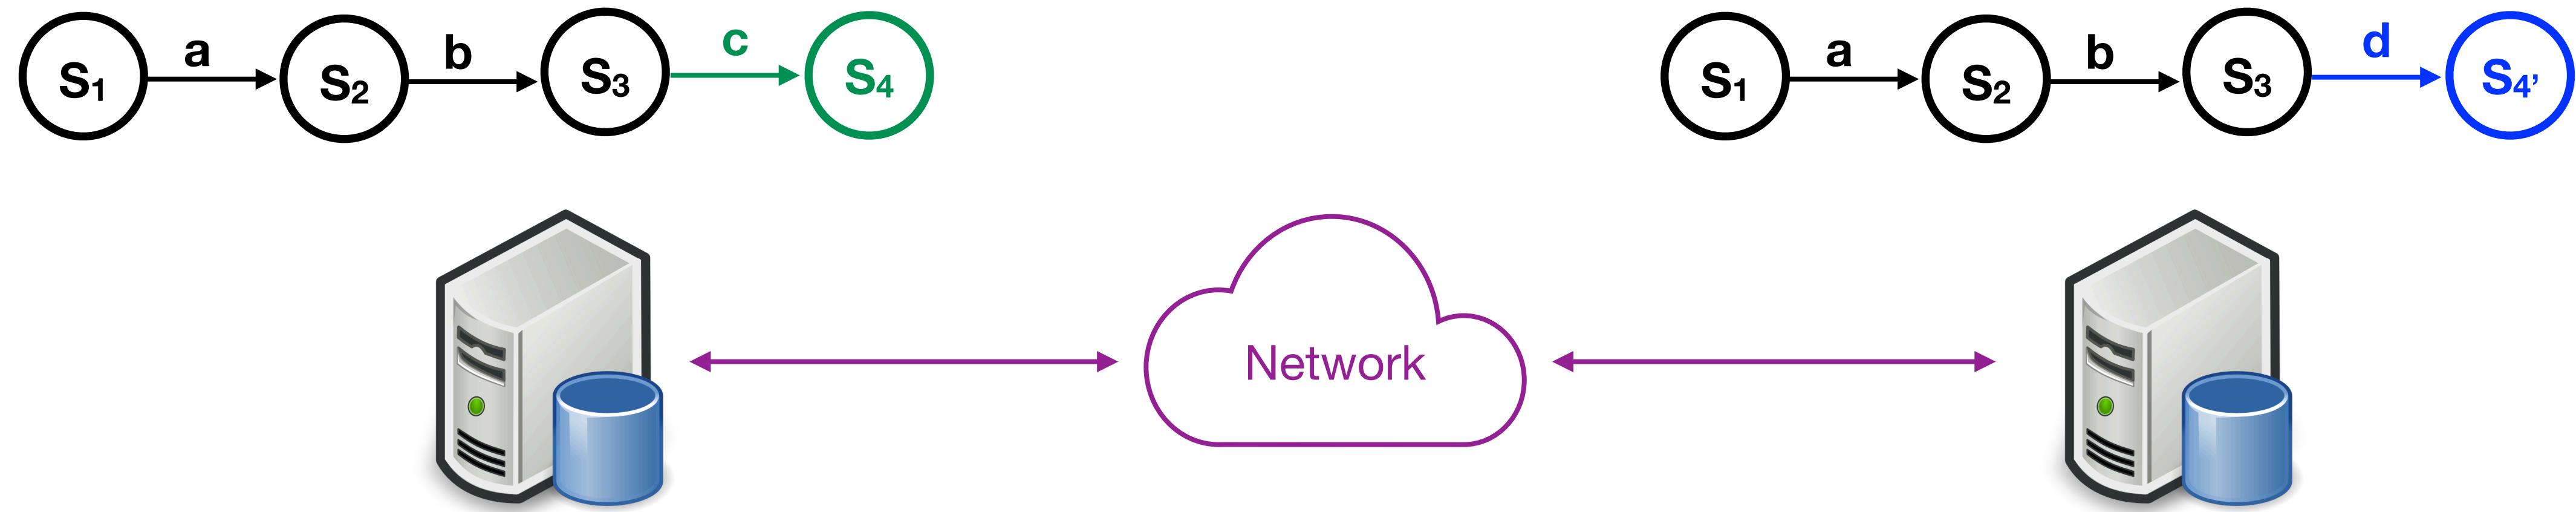
- Negatively affect the *availability* (CAP theorem).
- Extremely complex  $\Rightarrow$  Error-prone implementations

If not, can we make formal verification scale to implementations?

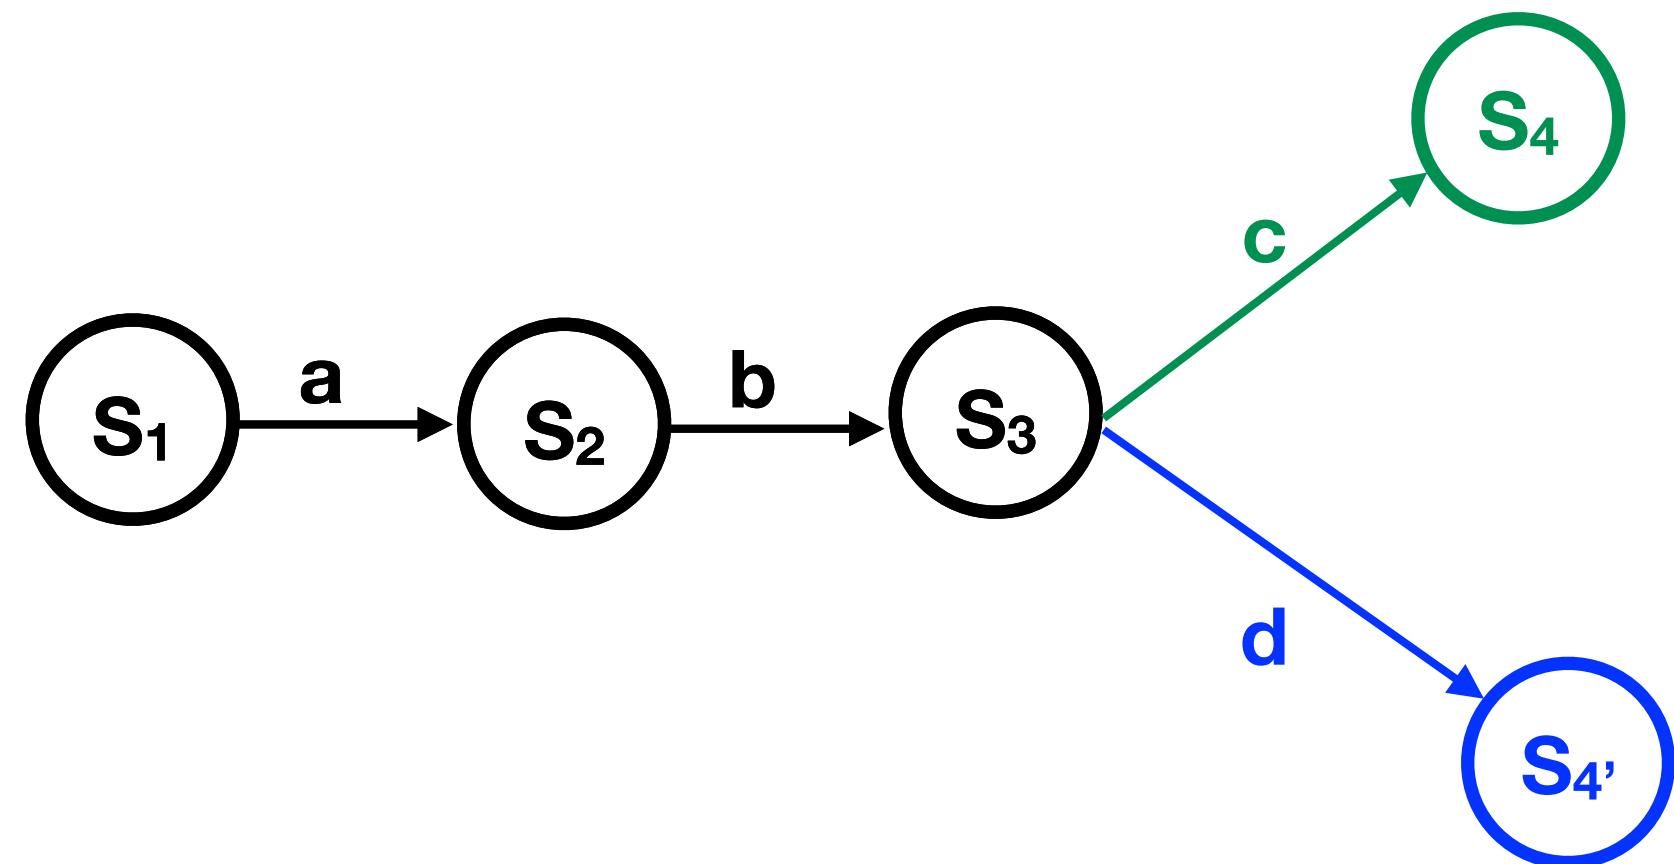
# Divergence re-visited



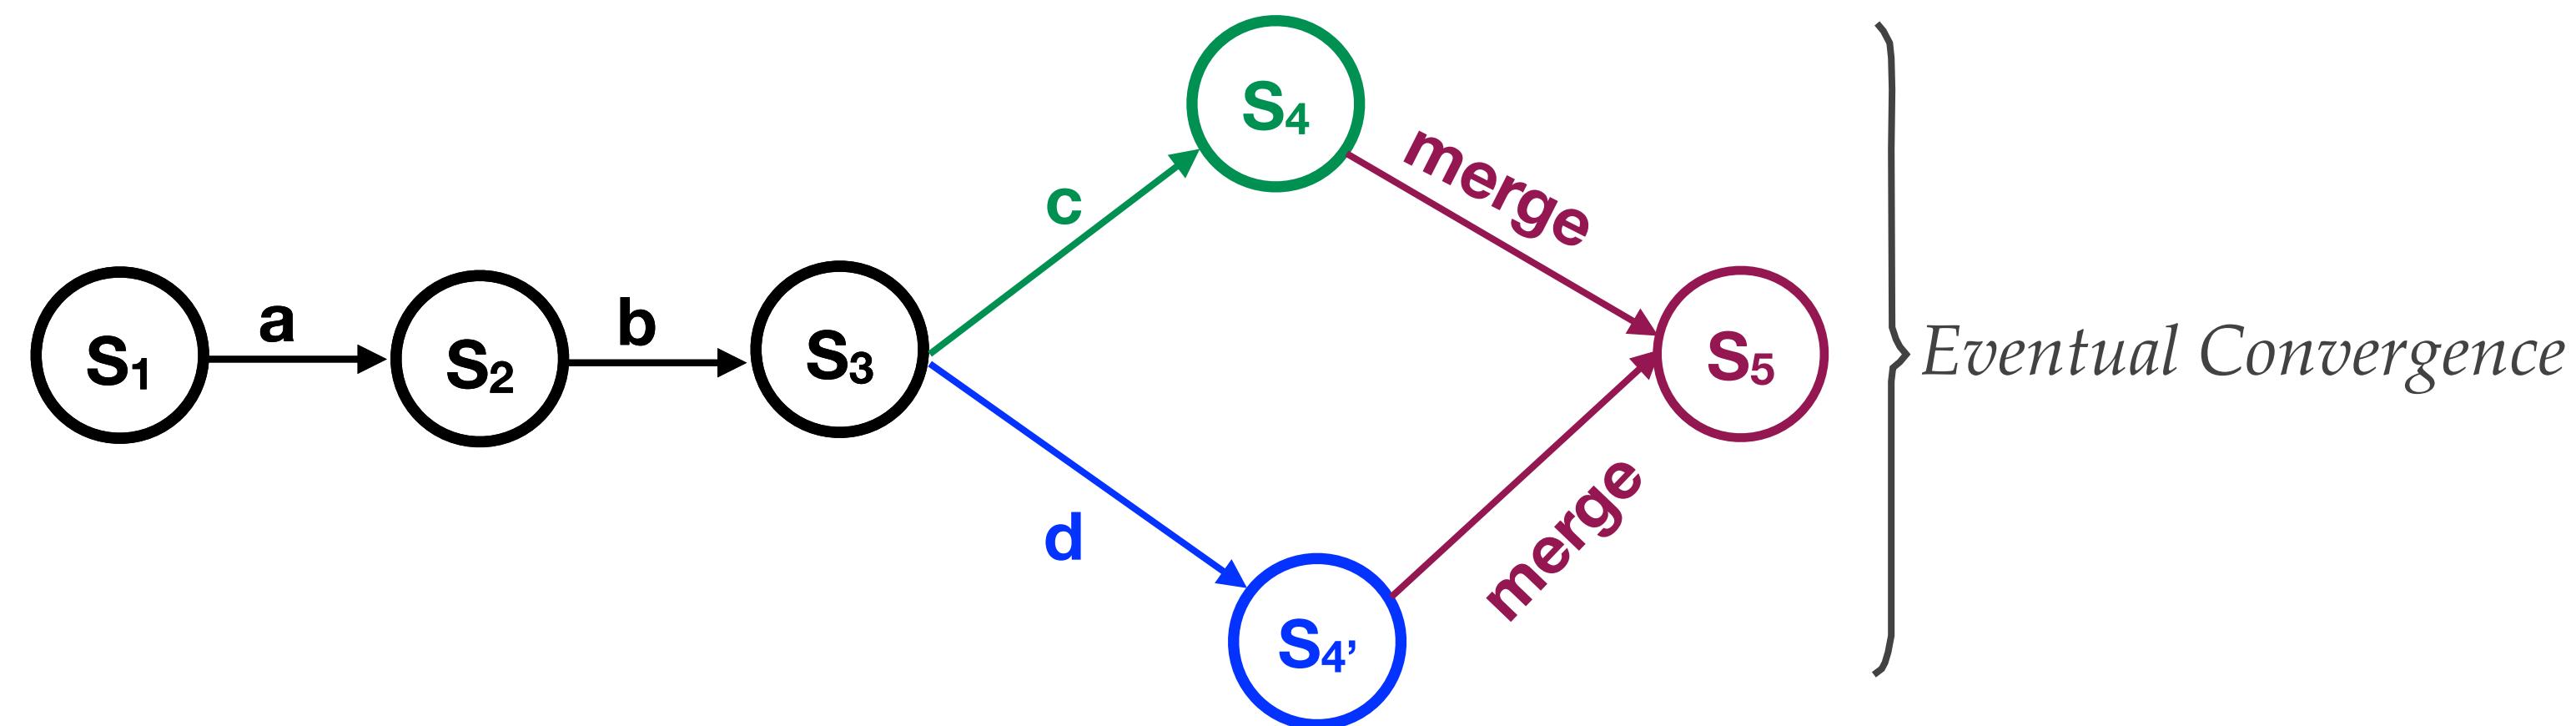
# Divergence re-visited



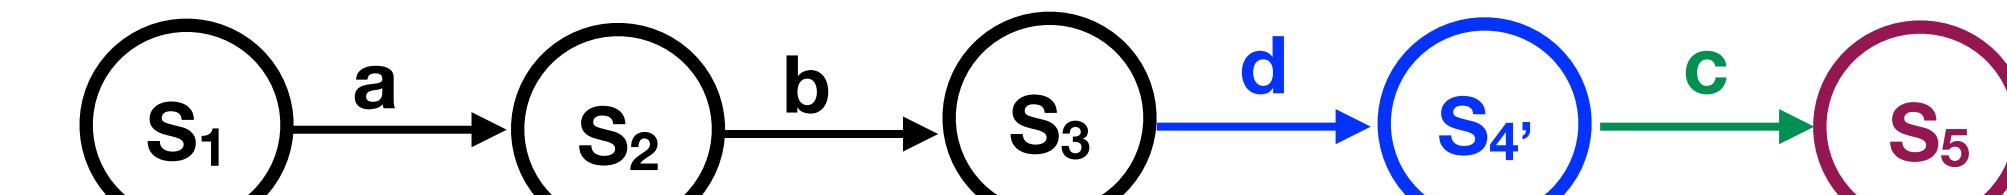
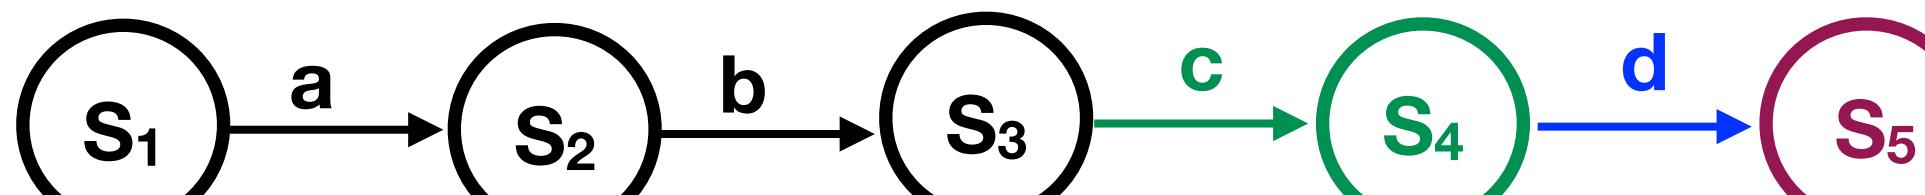
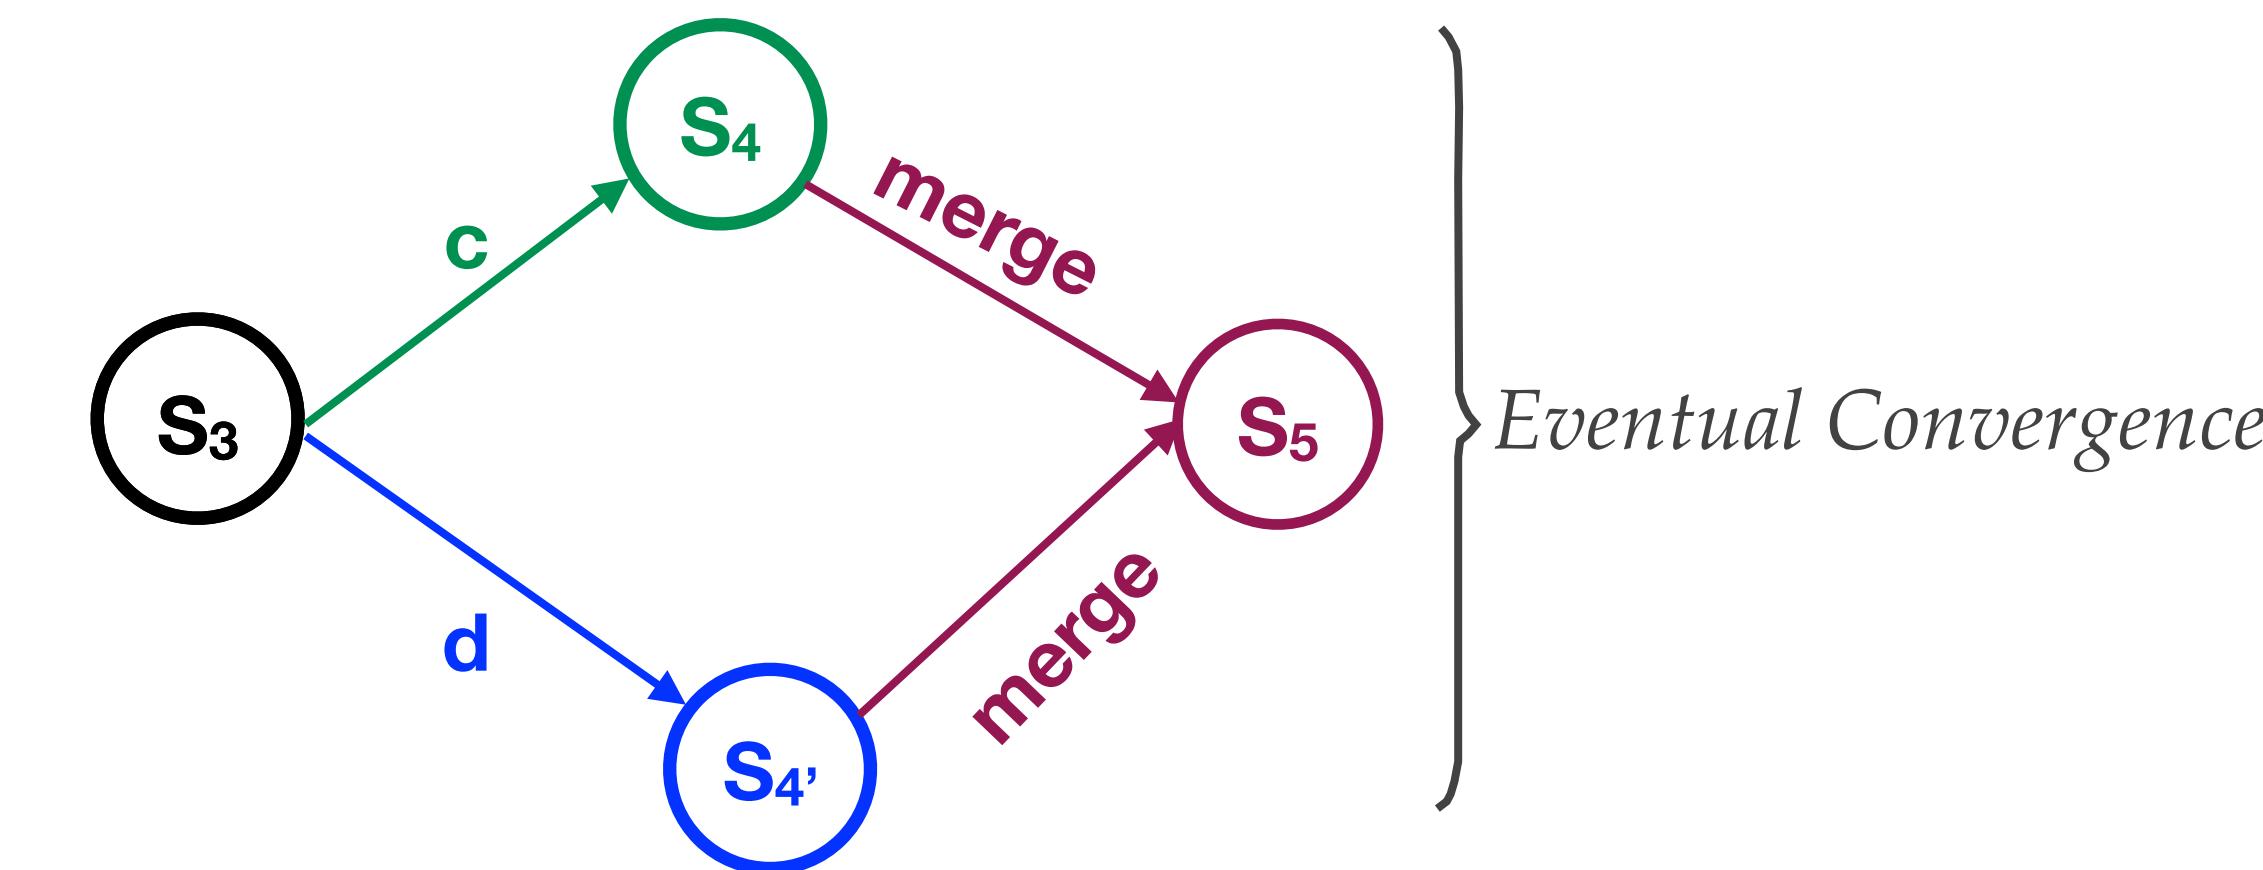
# Divergence: Zoomed In



# Merging Divergent States



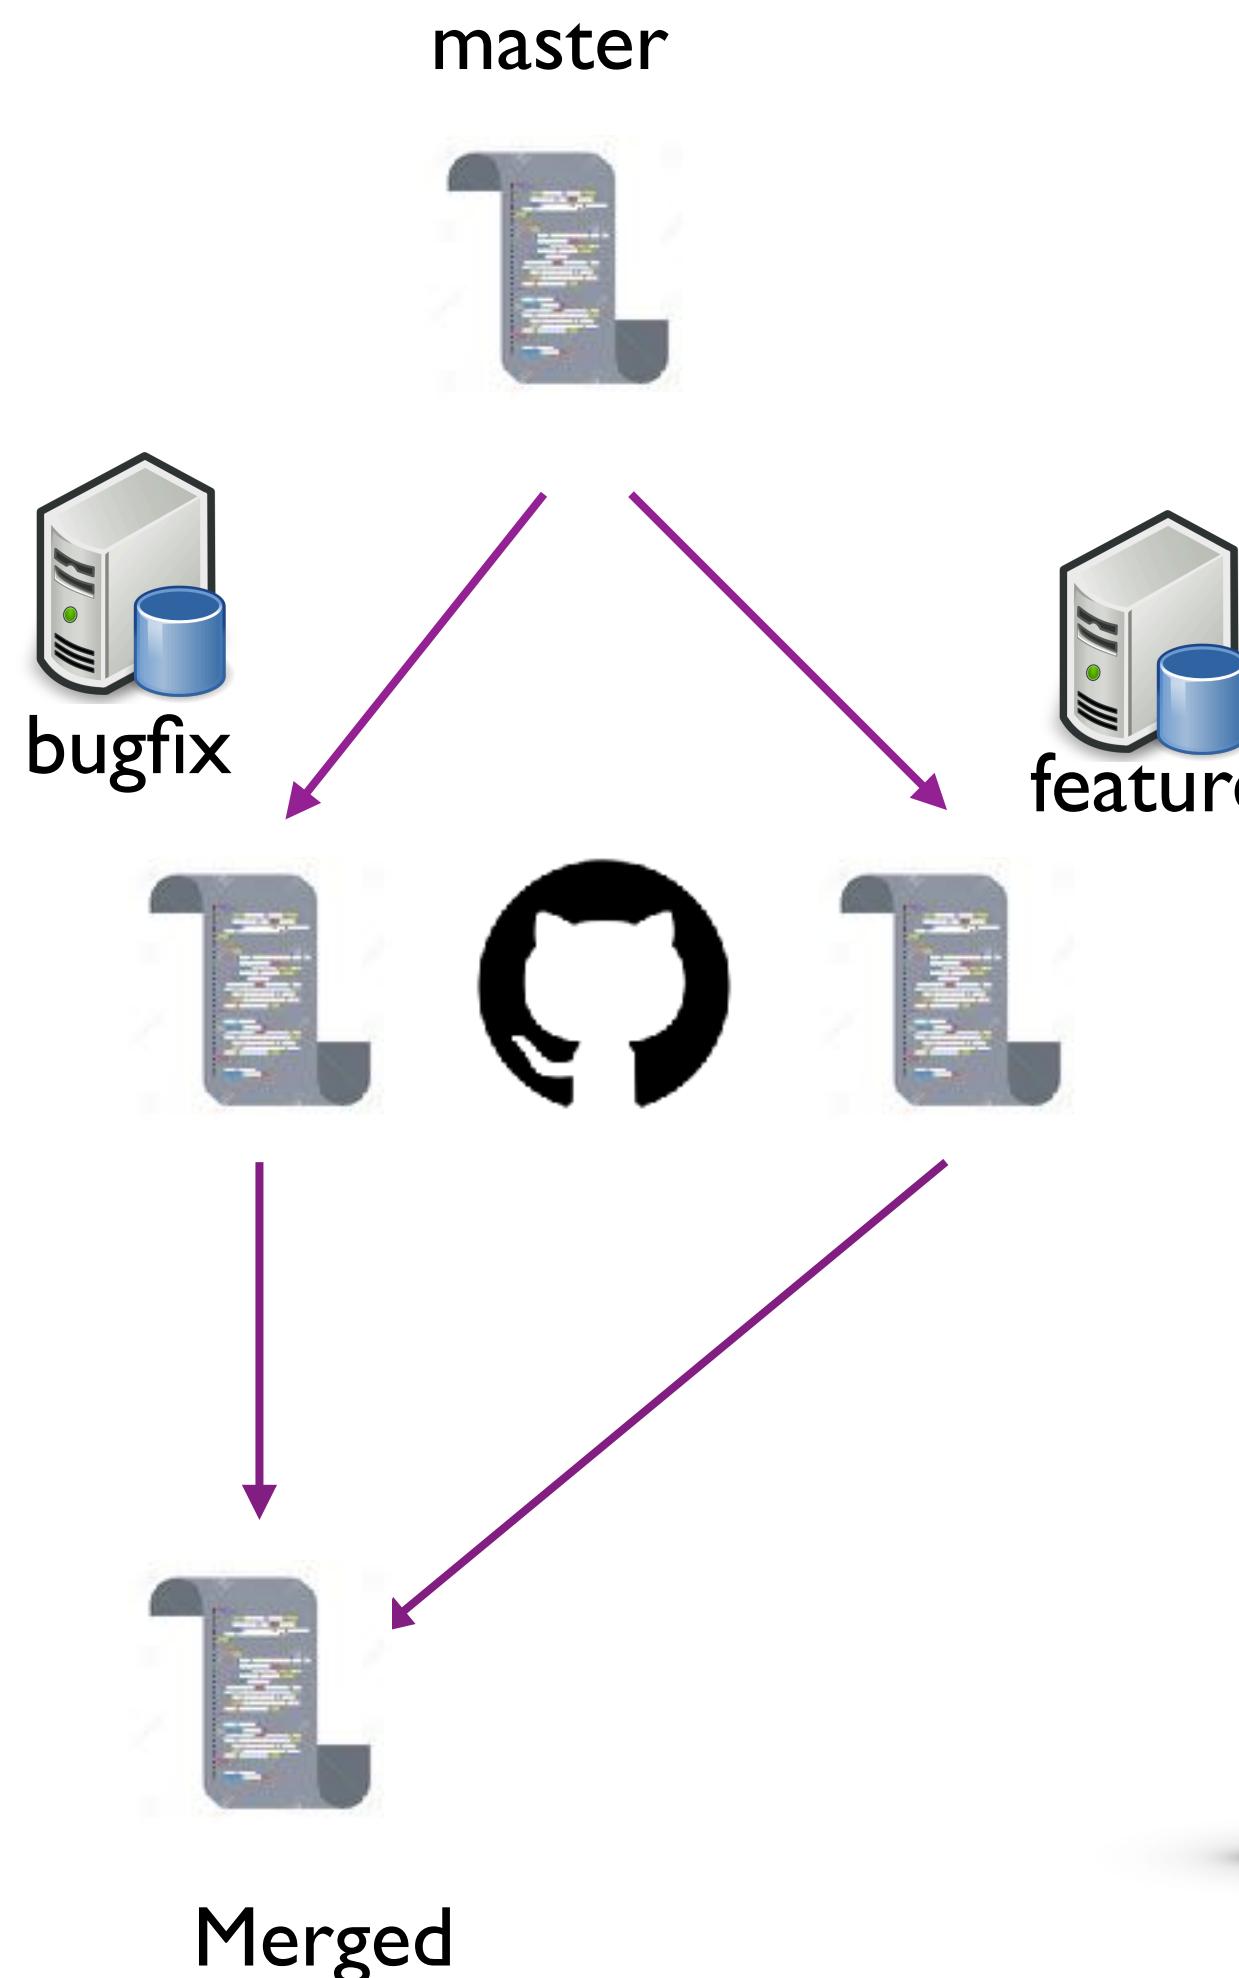
# Merging Divergent States



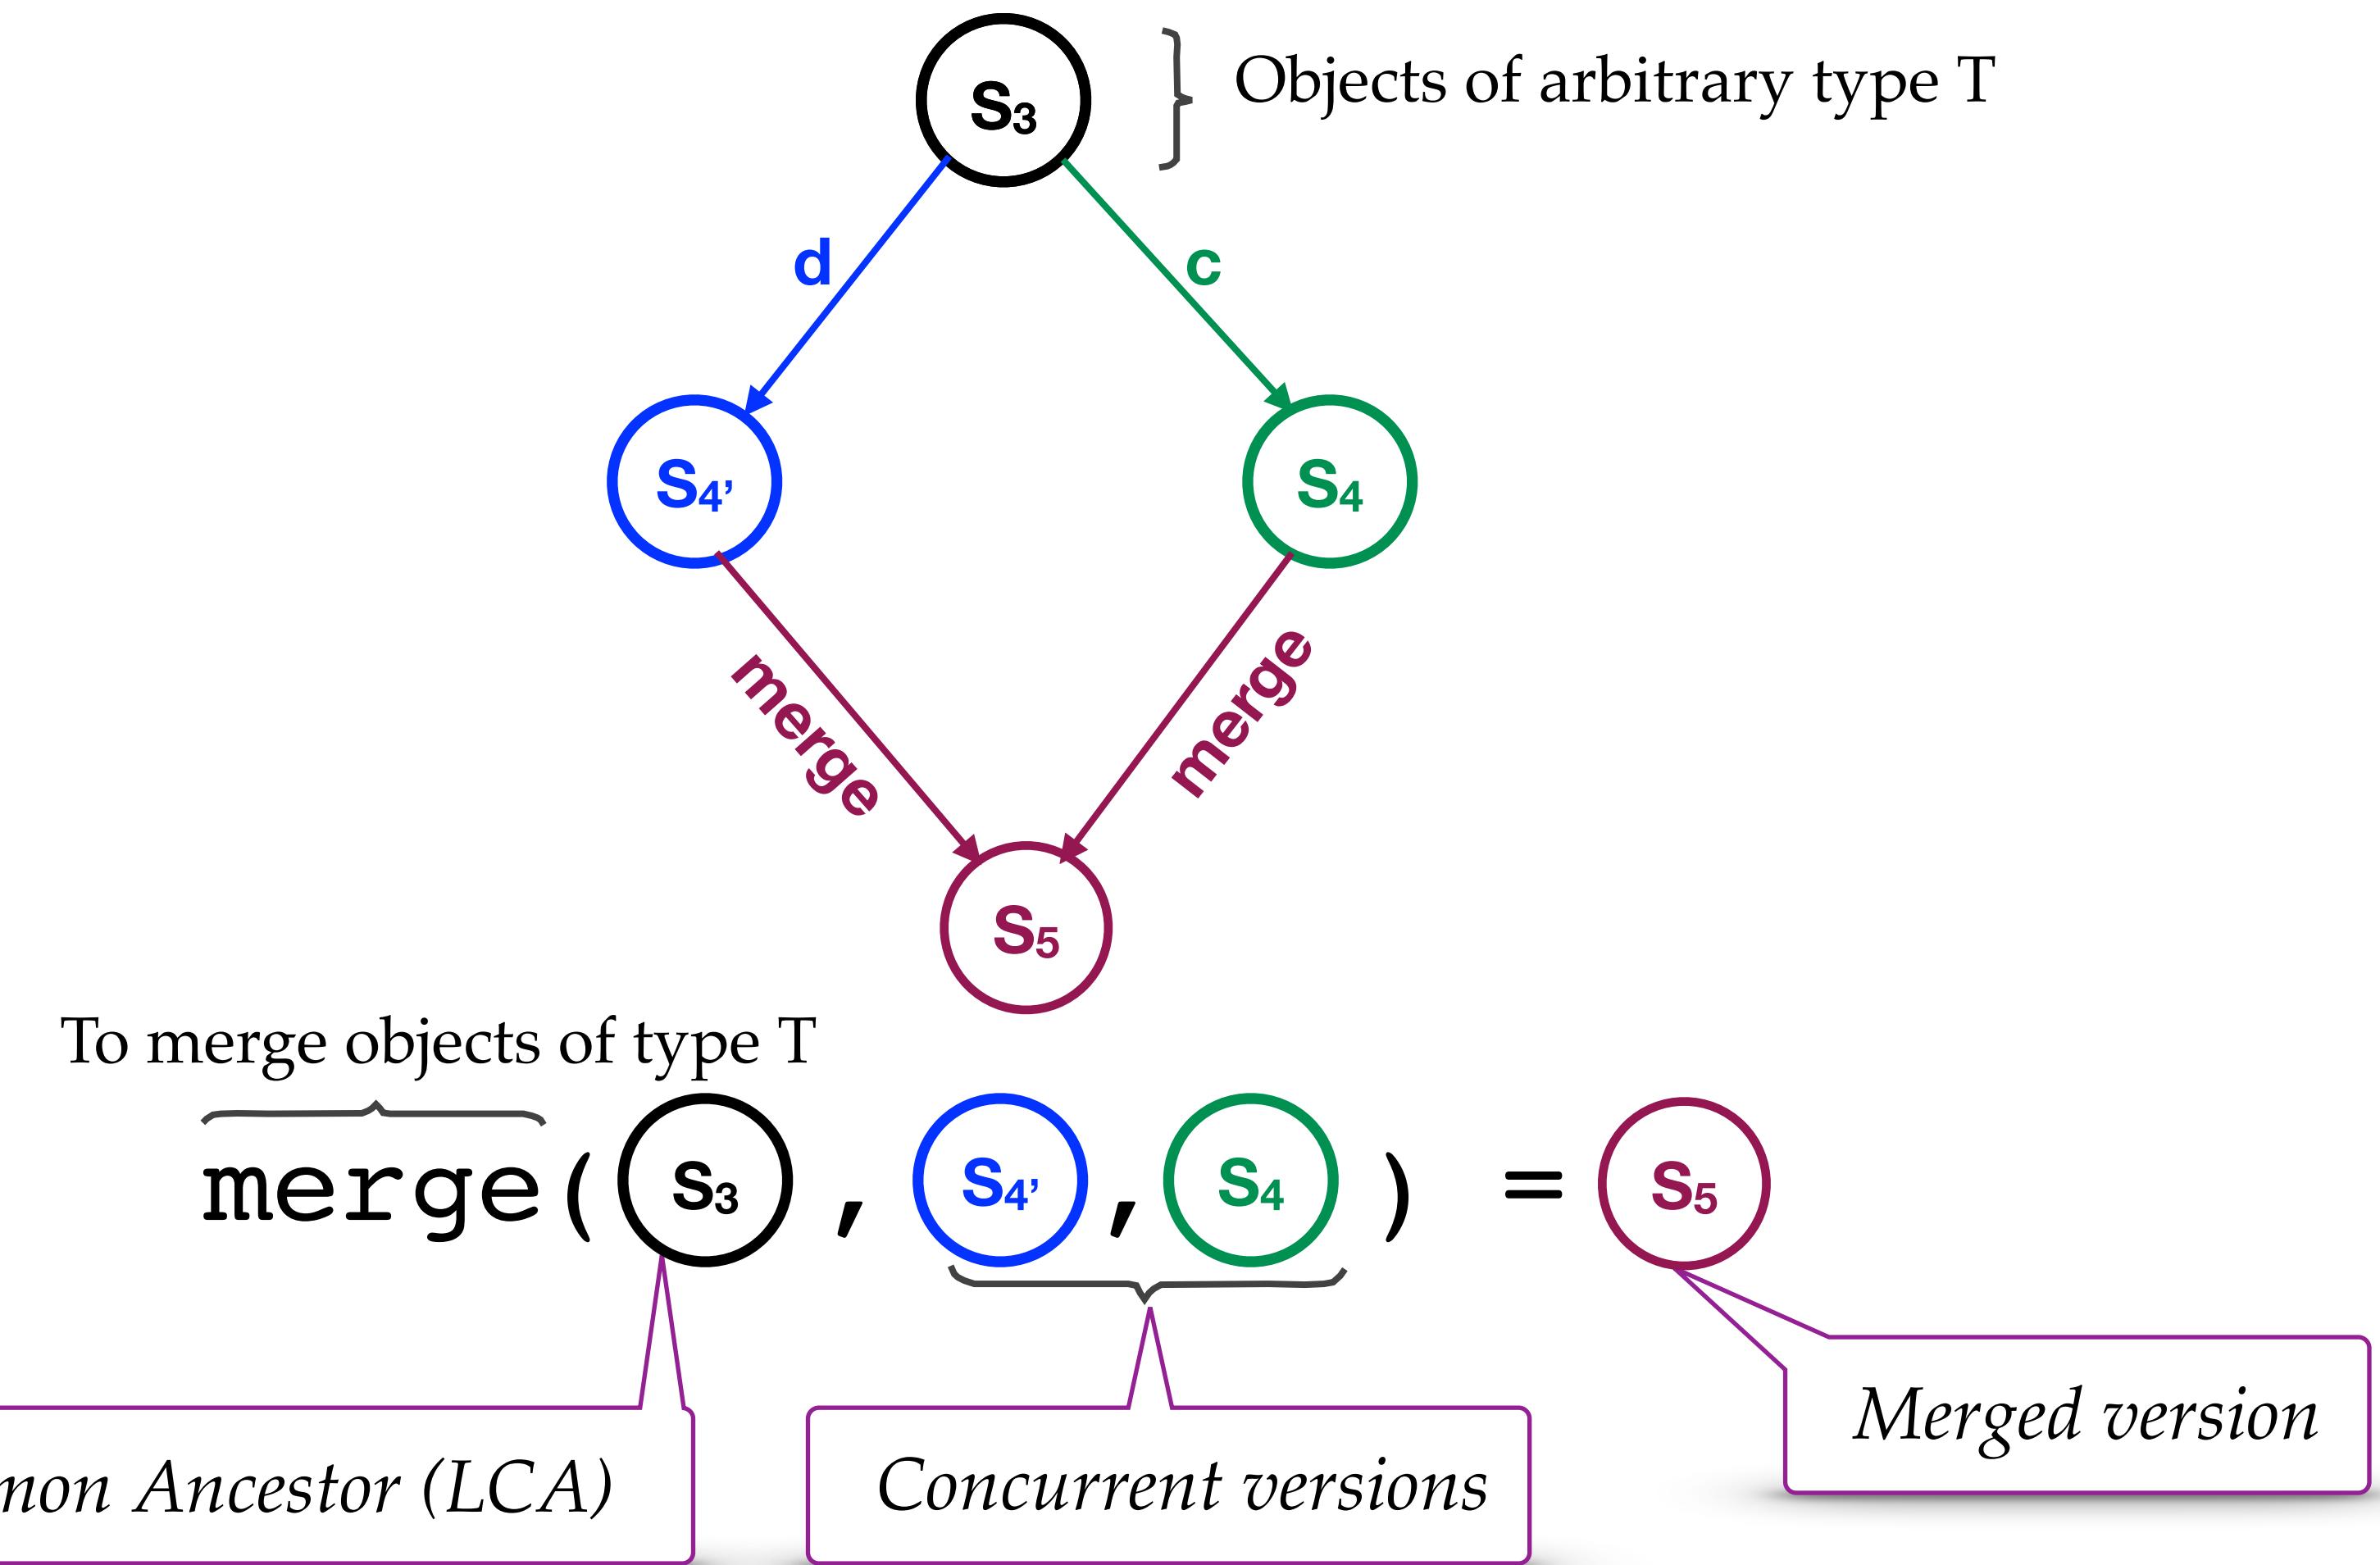
*Eventually Consistent Replicated State Machine (♦RSM)*

# ◆RSM in Practice

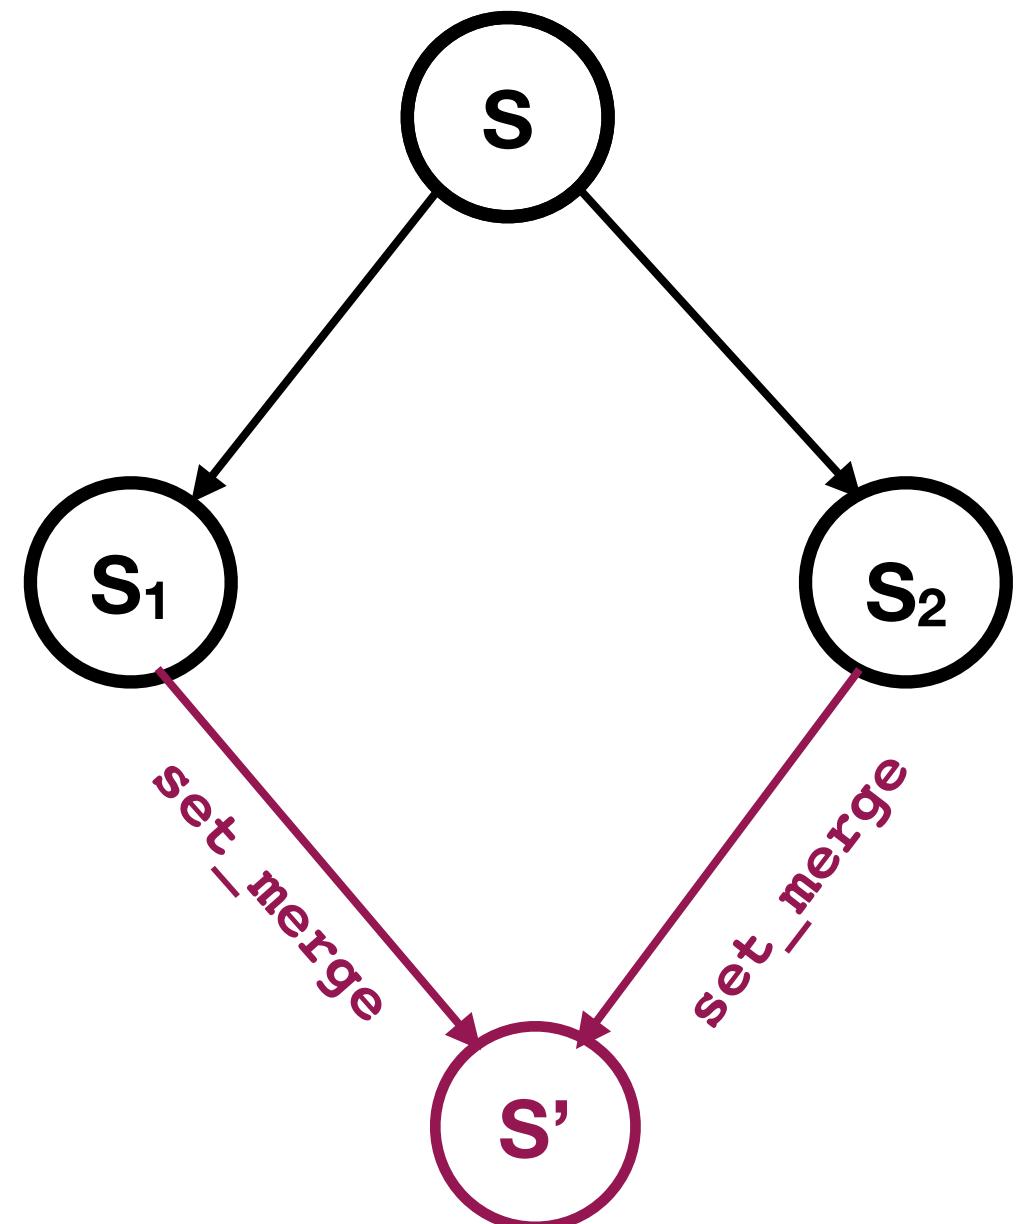
A familiar example: Git



In general:



# A merge function for Set data type



$$\text{set\_merge}(s, s_1, s_2) = (s \cap s_1 \cap s_2) \cup (s_1 - s) \cup (s_2 - s)$$

- Merged set contains:
  - Elements in  $S$  that are not deleted in  $S_1$  and  $S_2$
  - Elements newly added in  $S_1$
  - Elements newly added in  $S_2$



Stdlib Set data type

$$+ \lambda \text{ set\_merge} =$$



Mergeable Replicated Data Type (MRDT)

# Generality & Practicality of MRDTs

Gowtham Kaki, Swarn Priya, KC Sivaramakrishnan, and Suresh Jagannathan. **Mergeable Replicated Data Types**, OOPSLA 2019.

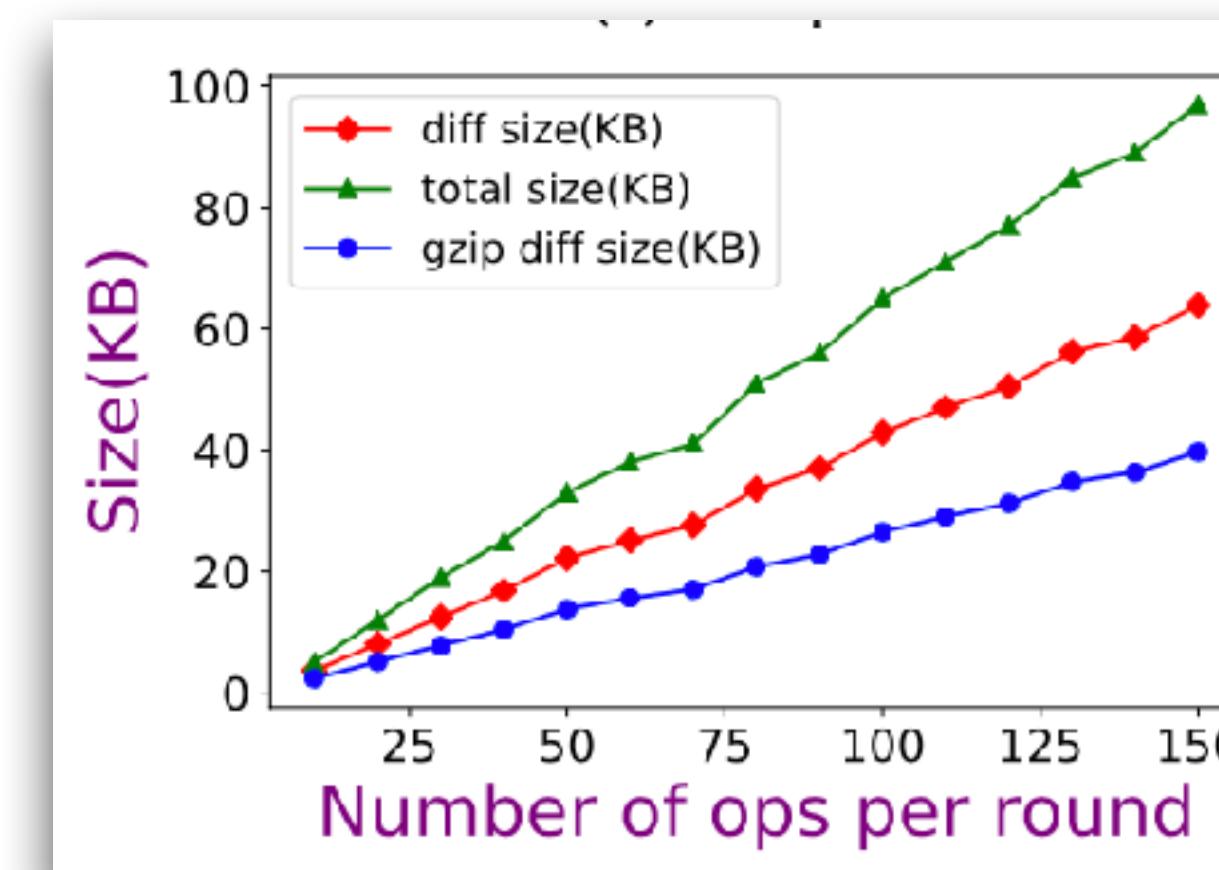
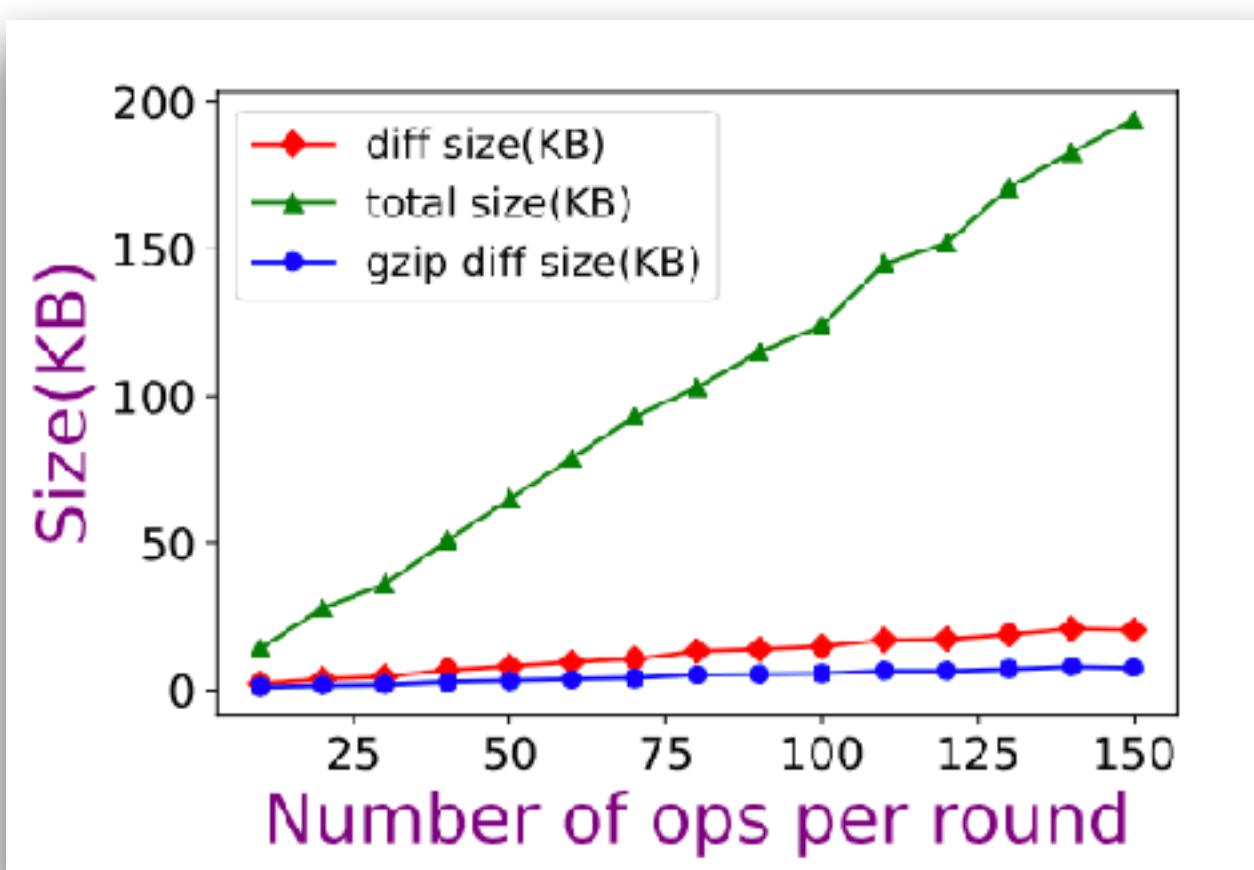
Table 1. Characteristic relations for various data types

Data Type	Characteristic Relations
Binary Heap	Membership ( $R_{mem}$ ), Ancestor ( $R_{ans} \subseteq R_{mem} \times R_{mem}$ )
Priority Queue	Membership ( $R_{mem}$ )
Set	Membership ( $R_{mem}$ )

Hacker News new | threads | past | comments | ask | show

Mergeable replicated data types – Part I (acolyer.org)

104 points by telotortium 75 days ago | hide | past | web | un-favorite | 21 comments



Sens: data them

## Mergeable replicated data types – Part I

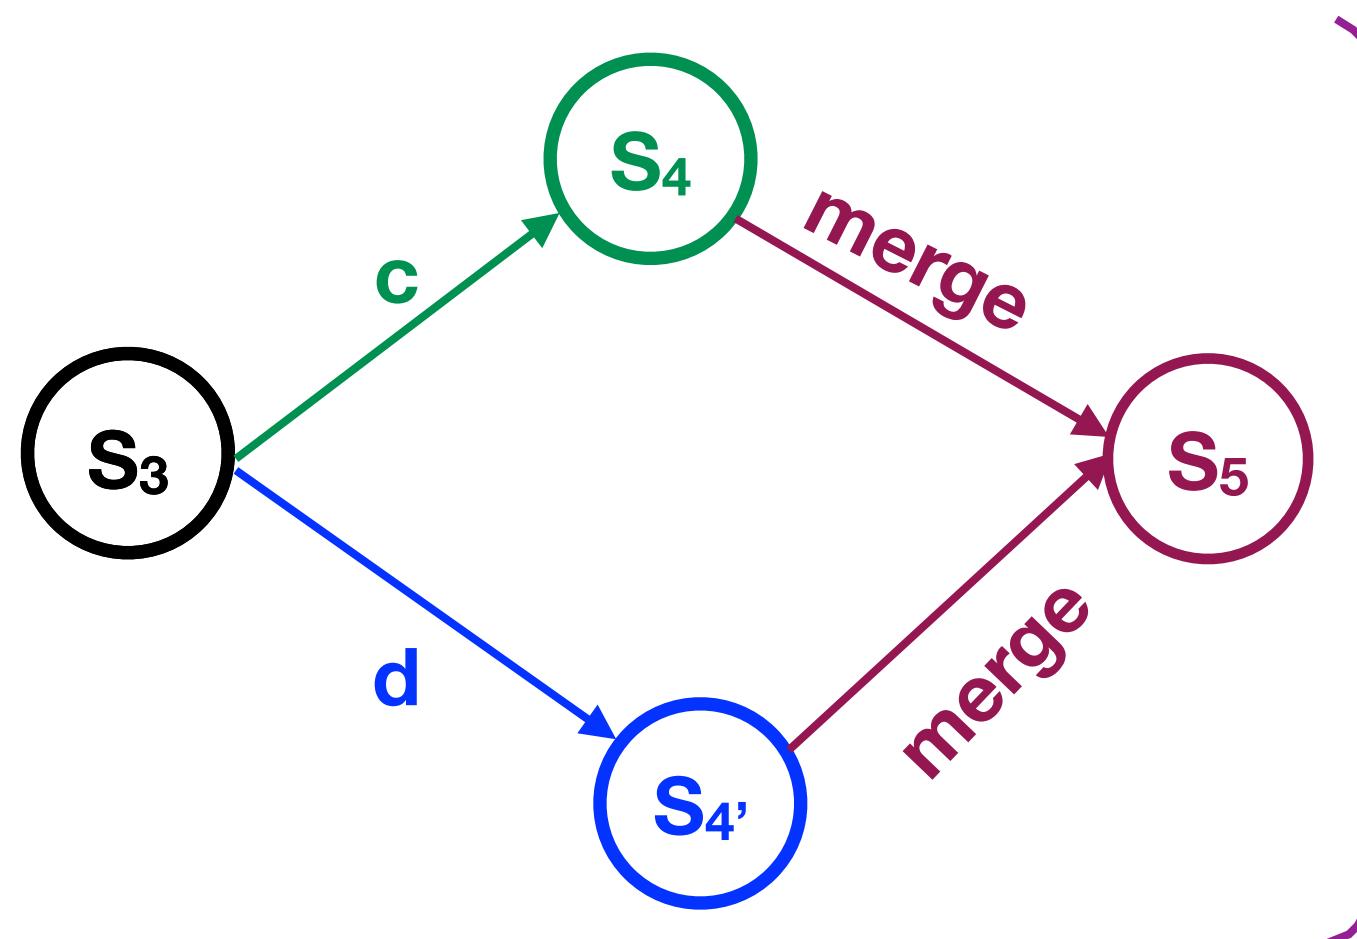
NOVEMBER 25, 2019

Mergeable replicated data types Kaki et al., OOPSLA'19

This paper was published at OOPSLA, but perhaps it's amongst the distributed systems community that I expect there to be the greatest interest.

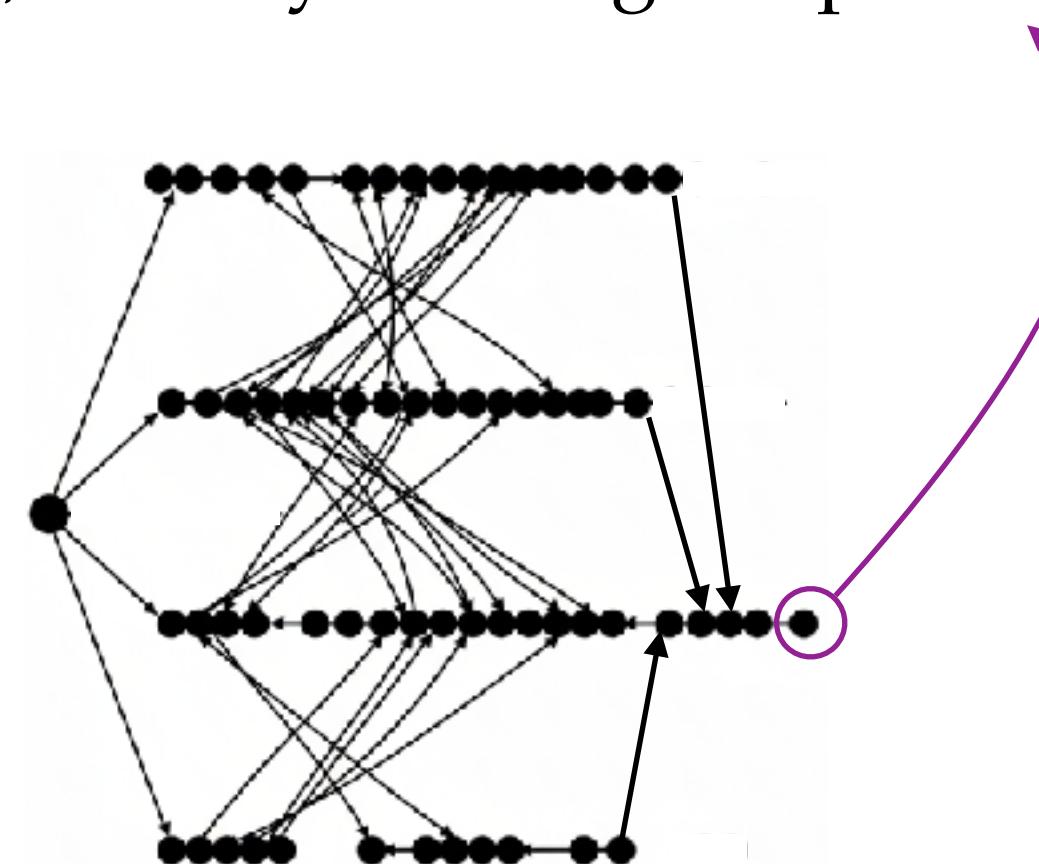
Network overhead imposed by state merges can be kept low by efficiently computing *diffs* (à la Git)

# Theoretical Guarantees of MRDTs



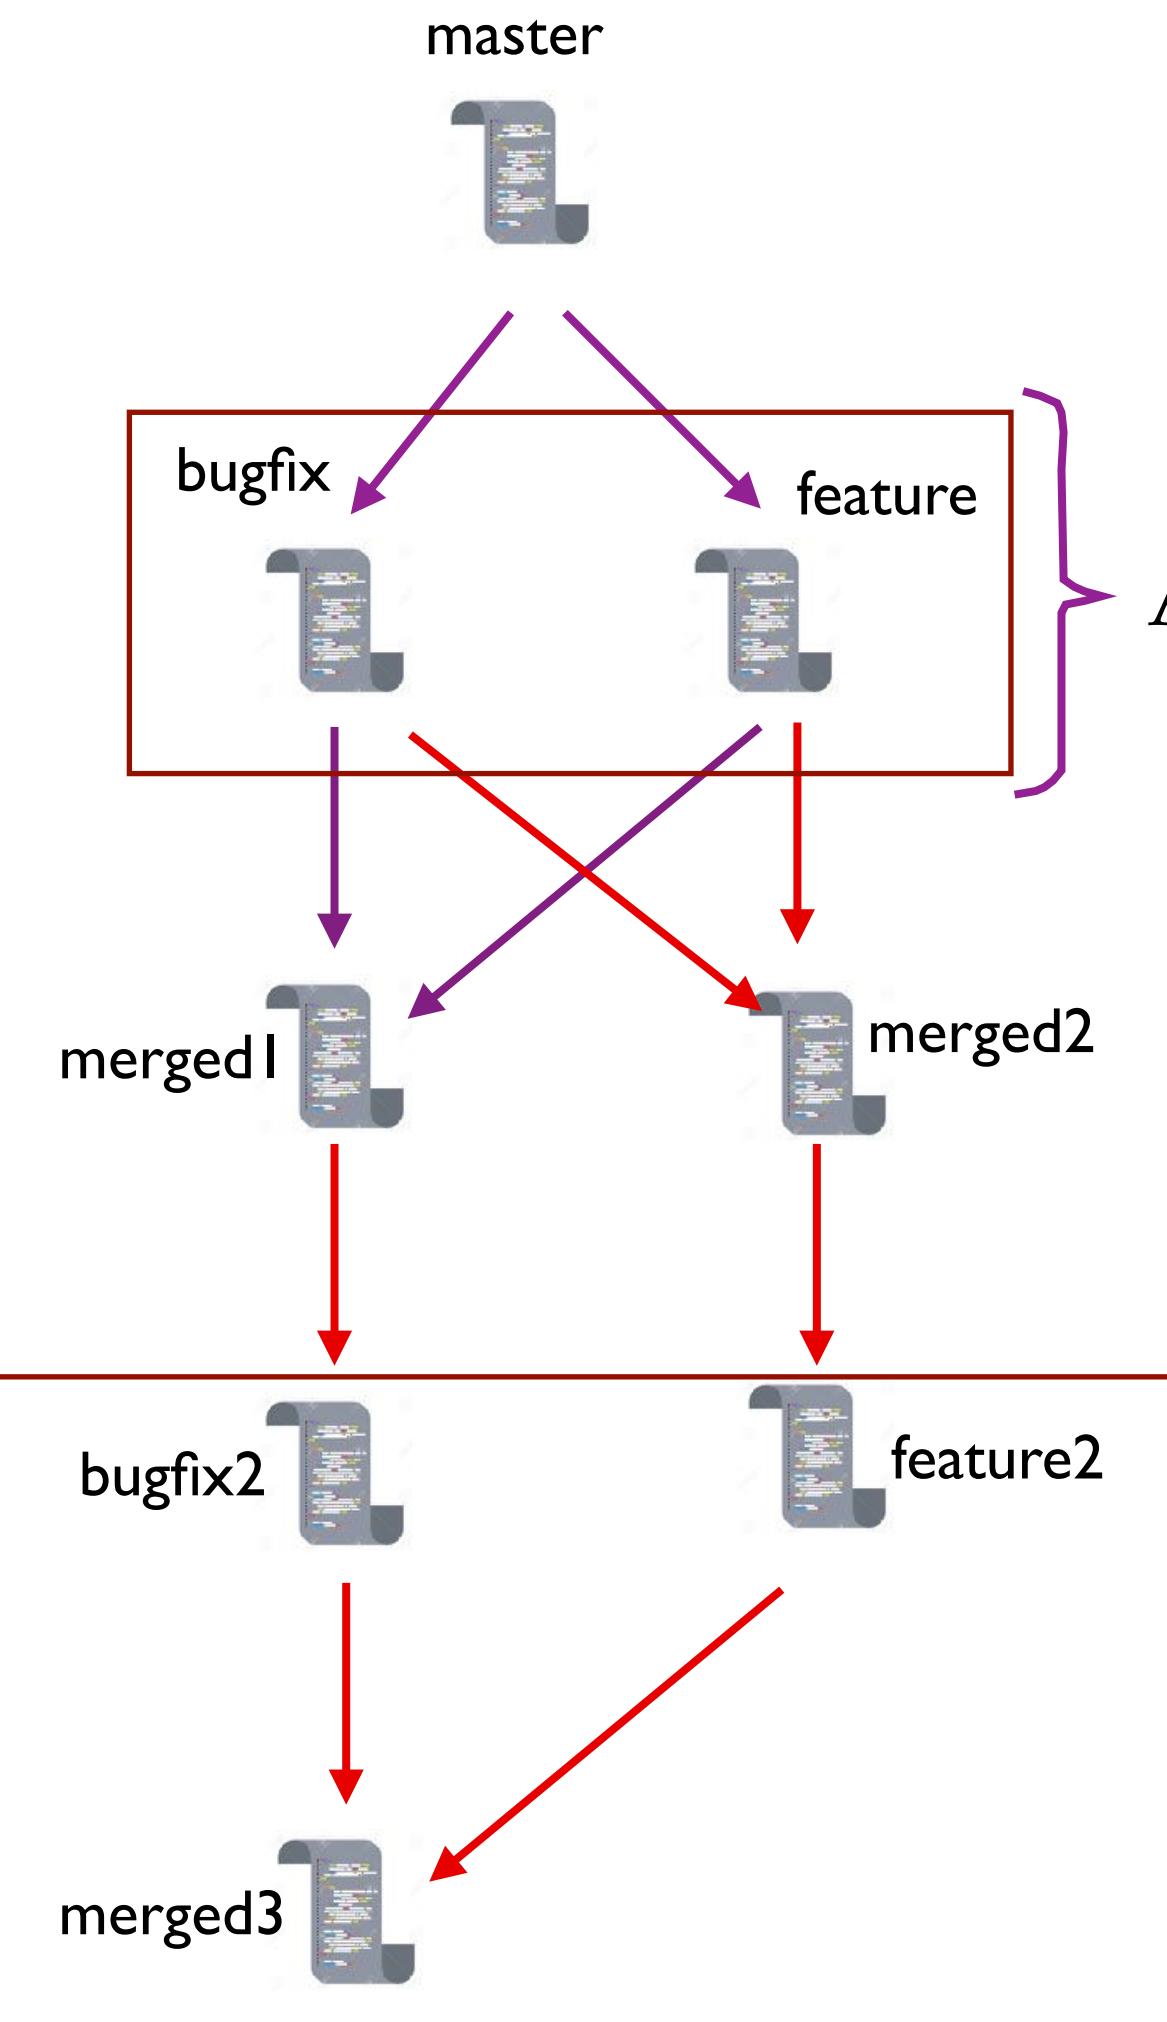
States have to converge.

But, do they converge in practice?

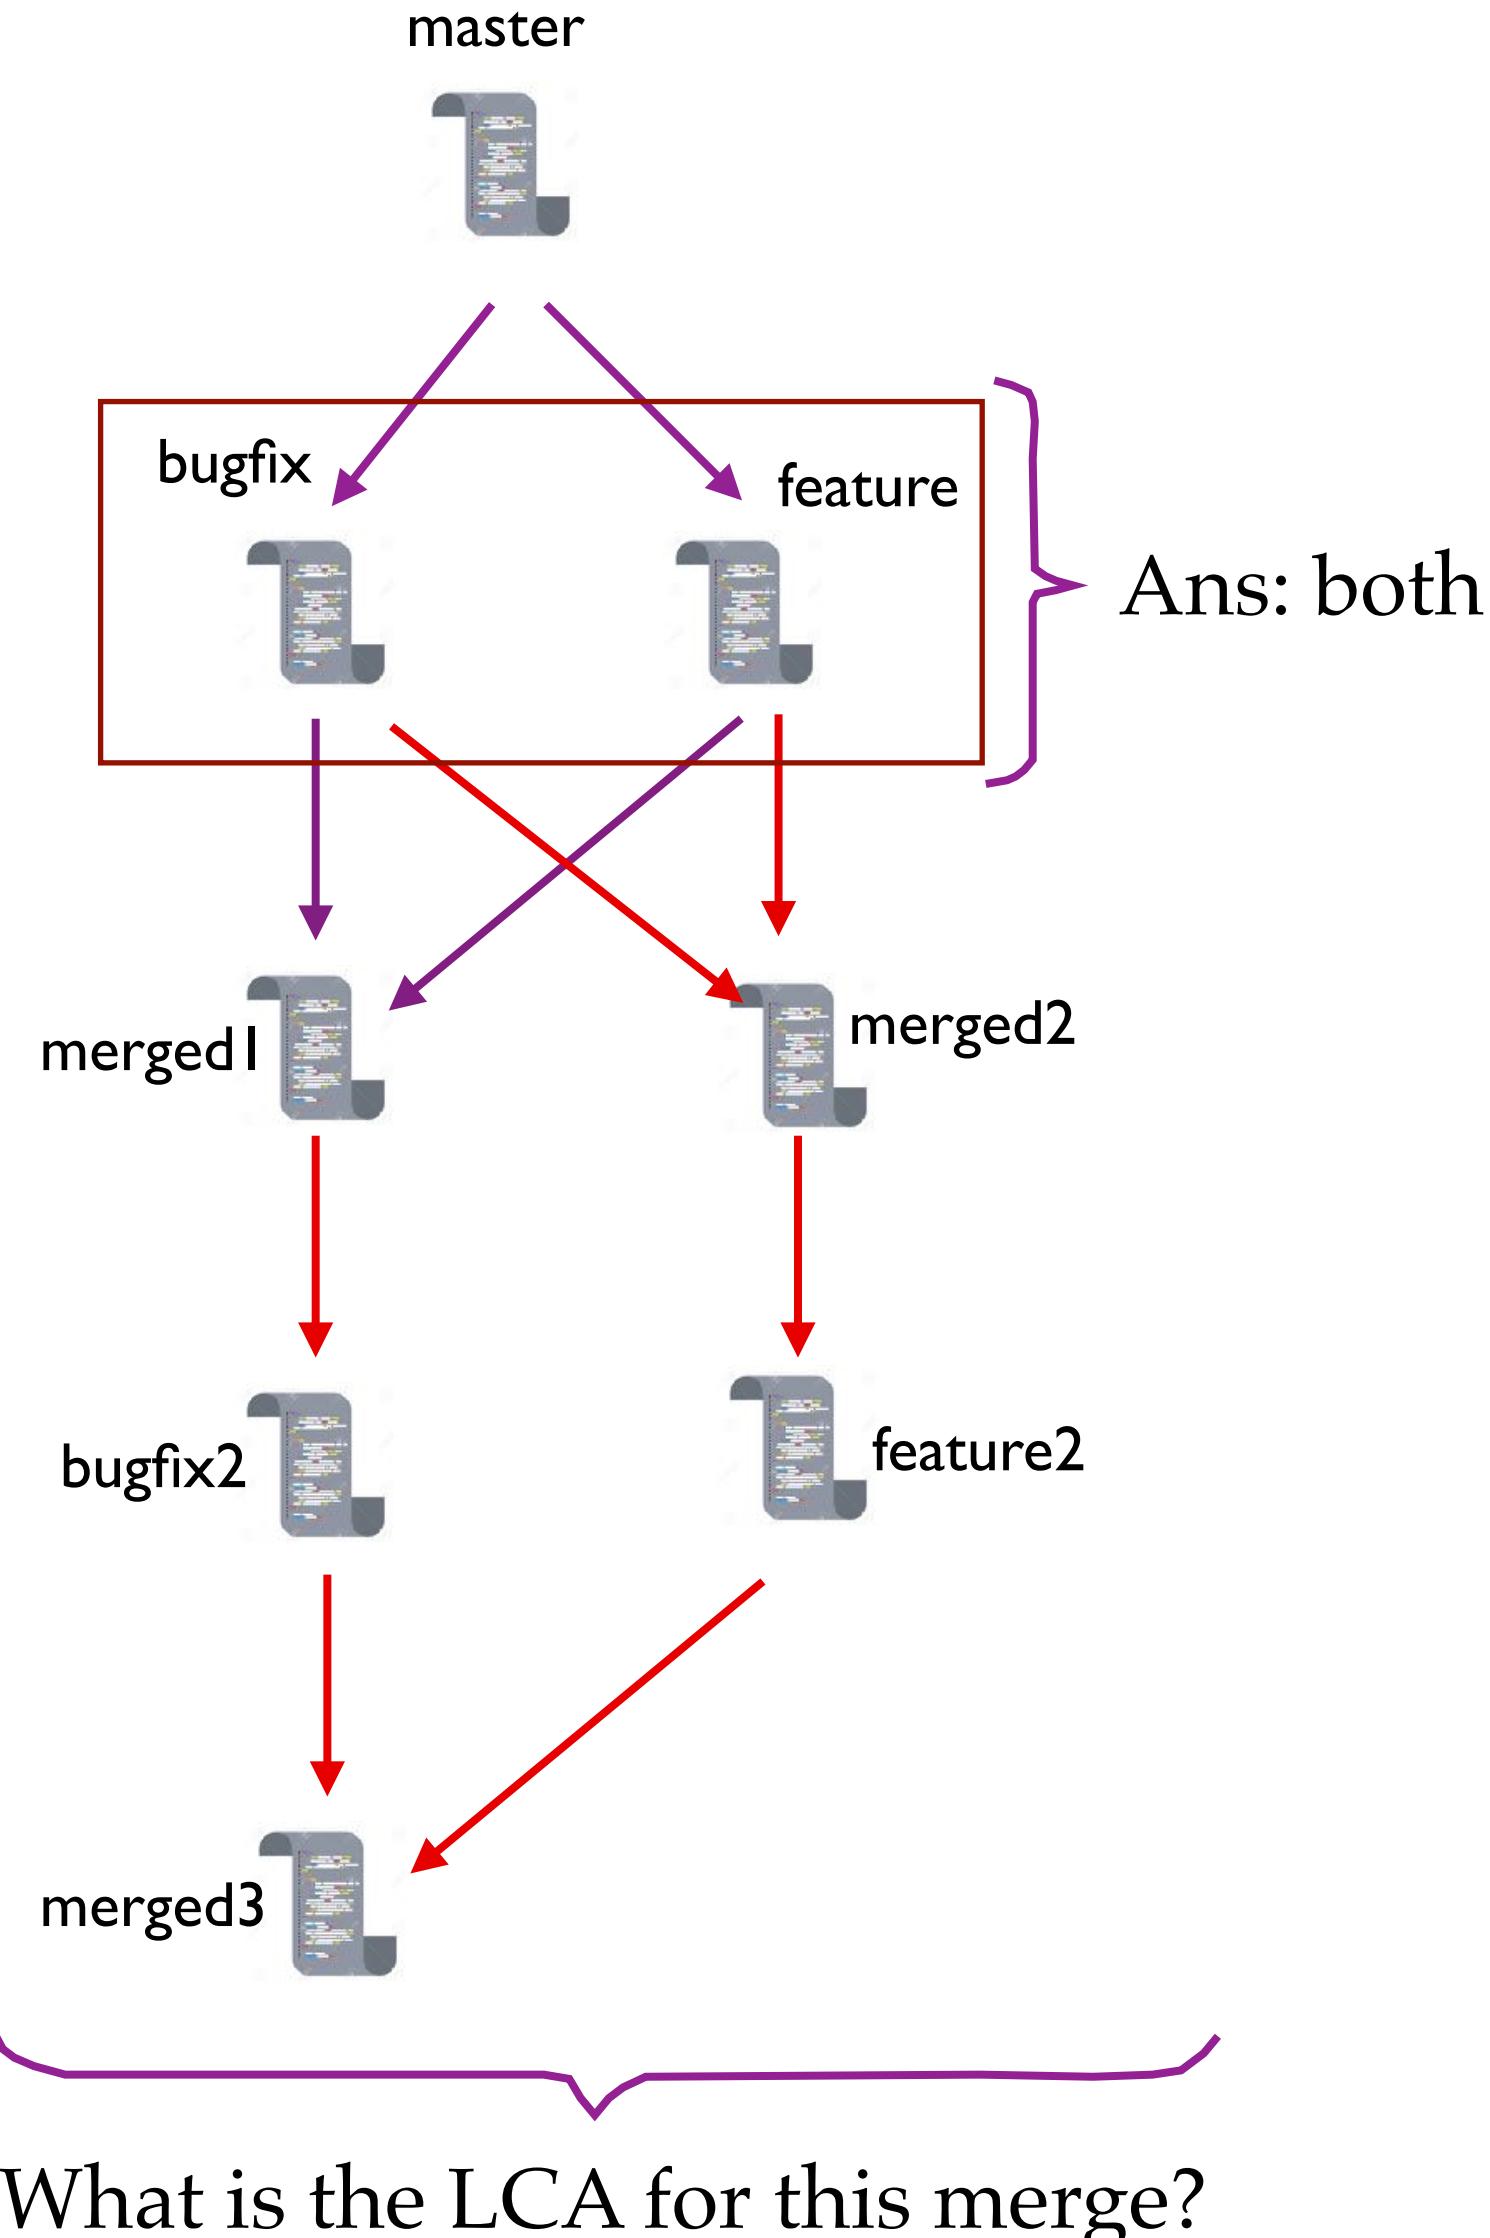


What is the LCA for this merge?

A fly in the ointment:



# Ill-formed Version Graphs



From my blog (<https://tycon.github.io/git-inconsistencies.html>)

## Vagaries of Git Merge

Git's merge algorithm seems to have inexplicable semantics leading to some interesting cases. I describe a couple of examples below. The bottom line is that:

1. Git merge is inconsistent: Merge result seems to depend not just on *what* versions were merged, but also *how* they were merged.
2. Git merge is unintuitive: The result of a successful (auto) merge may not necessarily be what we expect.

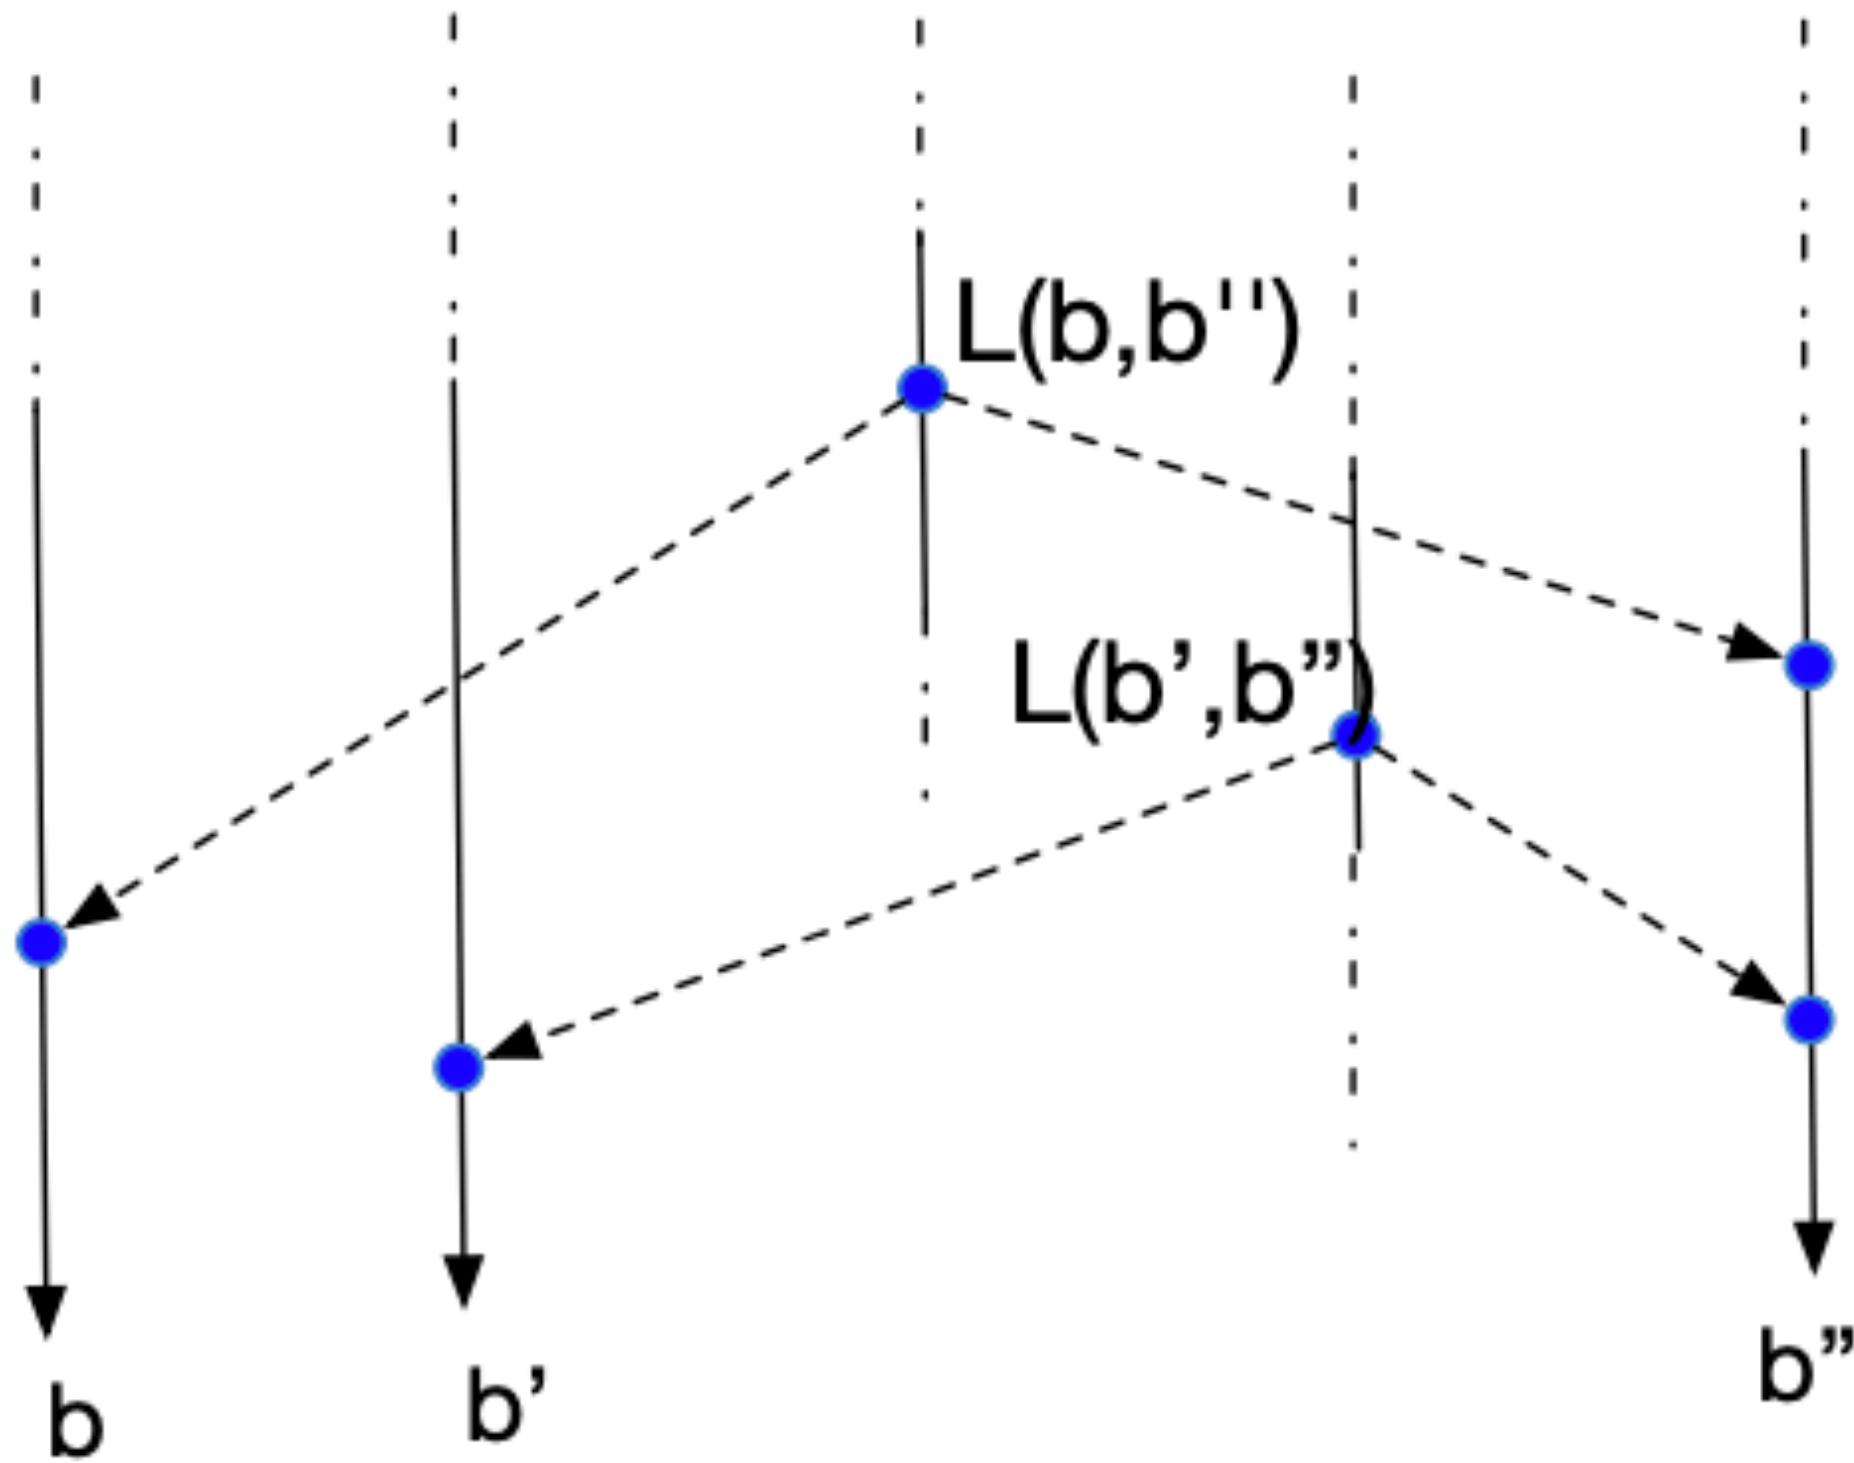
Git's recursive merging strategy™

What do we do? Options:

- Constrain merge function s.t. Git-style recursive merging is safe, i.e., guarantees convergence.
- Constrain version graphs s.t. ill-formed graphs never manifest.

# Constraining Version Graphs

Key idea: Enforce a semi-lattice structure



For every branch  $b''$  in the system, the LCA of  $b''$  with merging branches  $b$  and  $b'$  should be causally related, i.e:

$$L(b, b'') \rightarrow^* L(b', b'') \vee L(b', b'') \rightarrow^* L(b, b'').$$

# Constraining Version Graphs

## RunTime-Assisted Convergence in Replicated Data Types

Gowtham Kaki  
University of Colorado Boulder  
Boulder, USA  
gowtham.kaki@colorado.edu

Prasanth Prahladan  
University of Colorado Boulder  
Boulder, USA  
prasanth.prahladan@colorado.edu

Nicholas V. Lewchenko  
University of Colorado Boulder  
Boulder, USA  
nile1033@colorado.edu

[PLDI 2022]

A fault-tolerant distributed algorithm to orchestrate  
*well-formed* distributed executions of MRDTs.

**Theorem 3.7 (Convergence).** *In every reachable state  $\Delta = ((V, E), N, C, H, L)$  of the abstract machine: For all distinct  $b_1, b_2 \in \text{dom}(H)$ , and  $v_1, v_2 \in V$  such that  $v_1 = H(b_1)$  and  $v_2 = H(b_2)$ , the following is true:  $C(v_1) = C(v_2) \Rightarrow N(v_1) = N(v_2)$ .*



$$(G, N, C, H, L) \longrightarrow (G', N', C', H', L')$$

$$\frac{b \in \text{dom}(H) \quad v \notin V \quad i \notin \text{codom}(C)}{((V, E), N, C, H, L) \longrightarrow (V \cup \{v\}, E \cup \{H(b) \xrightarrow{c} v\}, N[v \mapsto n], C[v \mapsto \{i\} \cup C(H(b))], H[b \mapsto v], L)}$$
 [COMMIT]
  

$$\frac{b \in \text{dom}(H) \quad b' \notin \text{dom}(H) \quad v \notin V}{((V, E), N, C, H, L) \longrightarrow (V \cup \{v\}, E \cup \{H(b) \xrightarrow{f} v\}, N[v \mapsto N(H(b))], C[v \mapsto C(H(b))], H[b' \mapsto v], L[(b', b) \mapsto H(b)][\{(b', b'') \mapsto L(b, b'') \mid b'' \neq b\}])}$$
 [FORK]
  

$$\frac{b, b' \in \text{dom}(H) \quad C(H(b)) \supset C(L(b, b')) \quad C(H(b')) \supset C(L(b, b')) \quad \forall (b'' \in \text{dom}(H)). L(b, b'') \xrightarrow{*} L(b', b'') \vee L(b', b'') \xrightarrow{*} L(b, b'') \quad n = \text{merge}(N(L(b, b')), N(H(b)), N(H(b'))) \quad v \notin V}{((V, E), N, C, H, L) \longrightarrow (V \cup \{v\}, E \cup \{H(b) \xrightarrow{m} v, H(b') \xrightarrow{m} v\}, N[v \mapsto n], C[v \mapsto C(H(b)) \cup C(H(b'))], H[b \mapsto v], L[(b, b') \mapsto H(b)][\{(b, b'') \mapsto L(b, b'') \mid L(b, b'') \xrightarrow{*} L(b', b'')\}])}$$
 [MERGE]
  

$$\frac{b, b' \in \text{dom}(H) \quad C(H(b)) = C(L(b, b')) \quad C(H(b')) \supset C(L(b, b')) \quad \forall (b'' \in \text{dom}(H)). L(b, b'') \xrightarrow{*} L(b', b'') \vee L(b', b'') \xrightarrow{*} L(b, b'') \quad v \notin V}{((V, E), N, C, H, L) \longrightarrow (V \cup \{v\}, E \cup \{H(b) \xrightarrow{m} v, H(b') \xrightarrow{m} v\}, N[v \mapsto N(H(b'))], C[v \mapsto C(H(b'))], H[b \mapsto v], L[(b, b') \mapsto H(b)][\{(b, b'') \mapsto L(b, b'') \mid L(b, b'') \xrightarrow{*} L(b', b'')\}])}$$
 [FASTFWD]

**Theorem 3.8 (Progress).** *Every reachable state  $\Delta = ((V, E), N, C, H, L)$  of the abstract machine is either:*

- A “quiescent” state, where:  
 $\forall b_1, b_2 \in \text{dom}(H). C(H(b_1)) = C(H(b_2))$ , Or
- An “unstuck” state, where there exist  $b_1, b_2 \in \text{dom}(H)$  that satisfy the pre-conditions of MERGE or FASTFWD rules, i.e.,  $b_1$  and  $b_2$  are mergeable.

# Quark

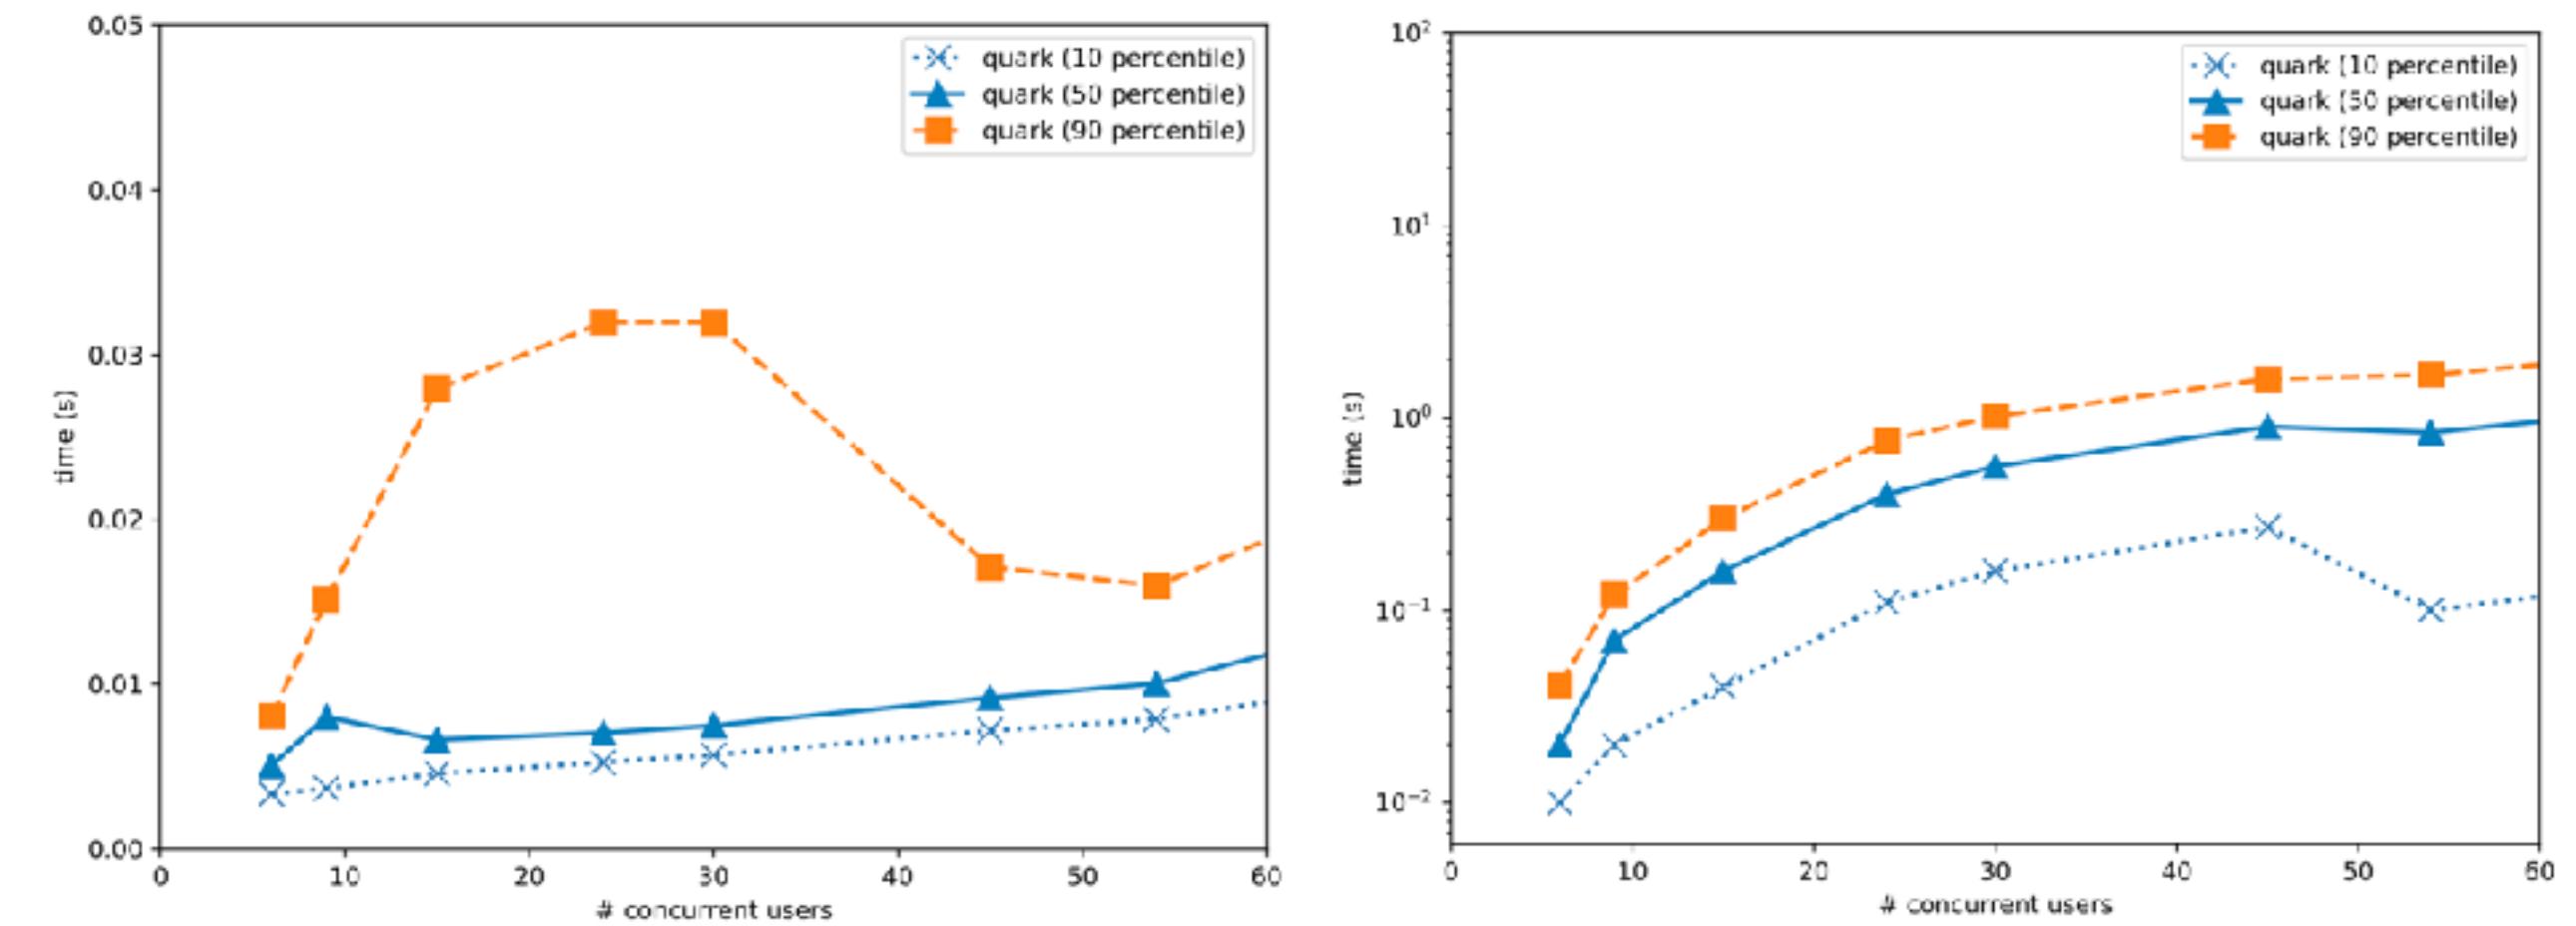
## RunTime-Assisted Convergence in Replicated Data Types

Gowtham Kaki  
University of Colorado Boulder  
Boulder, USA  
gowtham.kaki@colorado.edu

Prasanth Prahladan  
University of Colorado Boulder  
Boulder, USA  
prasanth.prahladan@colorado.edu

Nicholas V. Lewchenko  
University of Colorado Boulder  
Boulder, USA  
nile1033@colorado.edu

[PLDI 2022]



Implemented in a distributed runtime called:



<https://github.com/cuplv/Quark>

## Key Takeaway

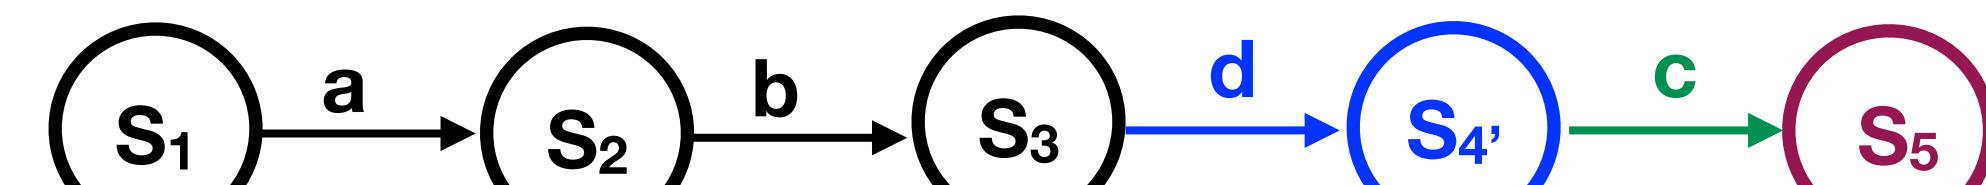
- For any data type  $T$ , Quark lets the programmer implement a  $T \diamond RSM$  that is:
- Fault-tolerant and scalable, and
  - Guaranteed to converge without needing a type-specific convergence proof.

# But, Convergence $\neq$ Consensus

Convergence is about *eventually* agreeing



A



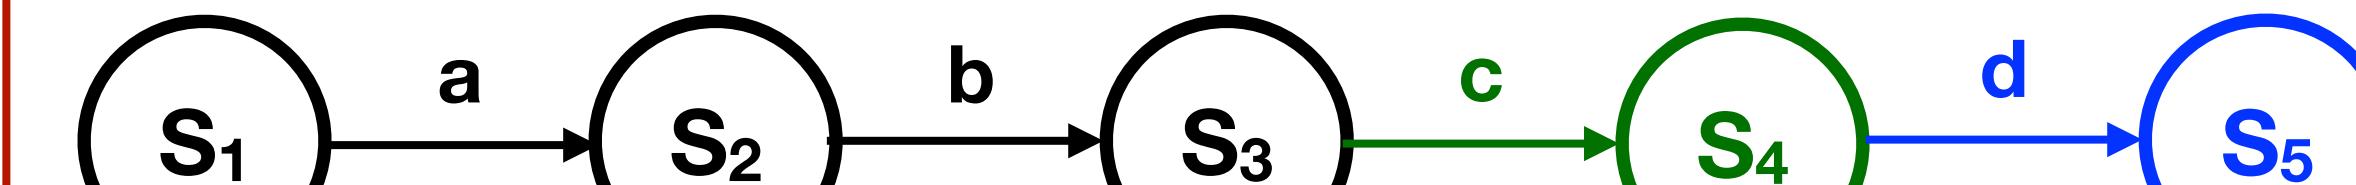
B

Eventually Consistent Replicated State Machine ( $\diamond$ RSM)

Consensus is about *never* disagreeing



A



B

Strongly Consistent Replicated State Machine ( $\blacksquare$ RSM)

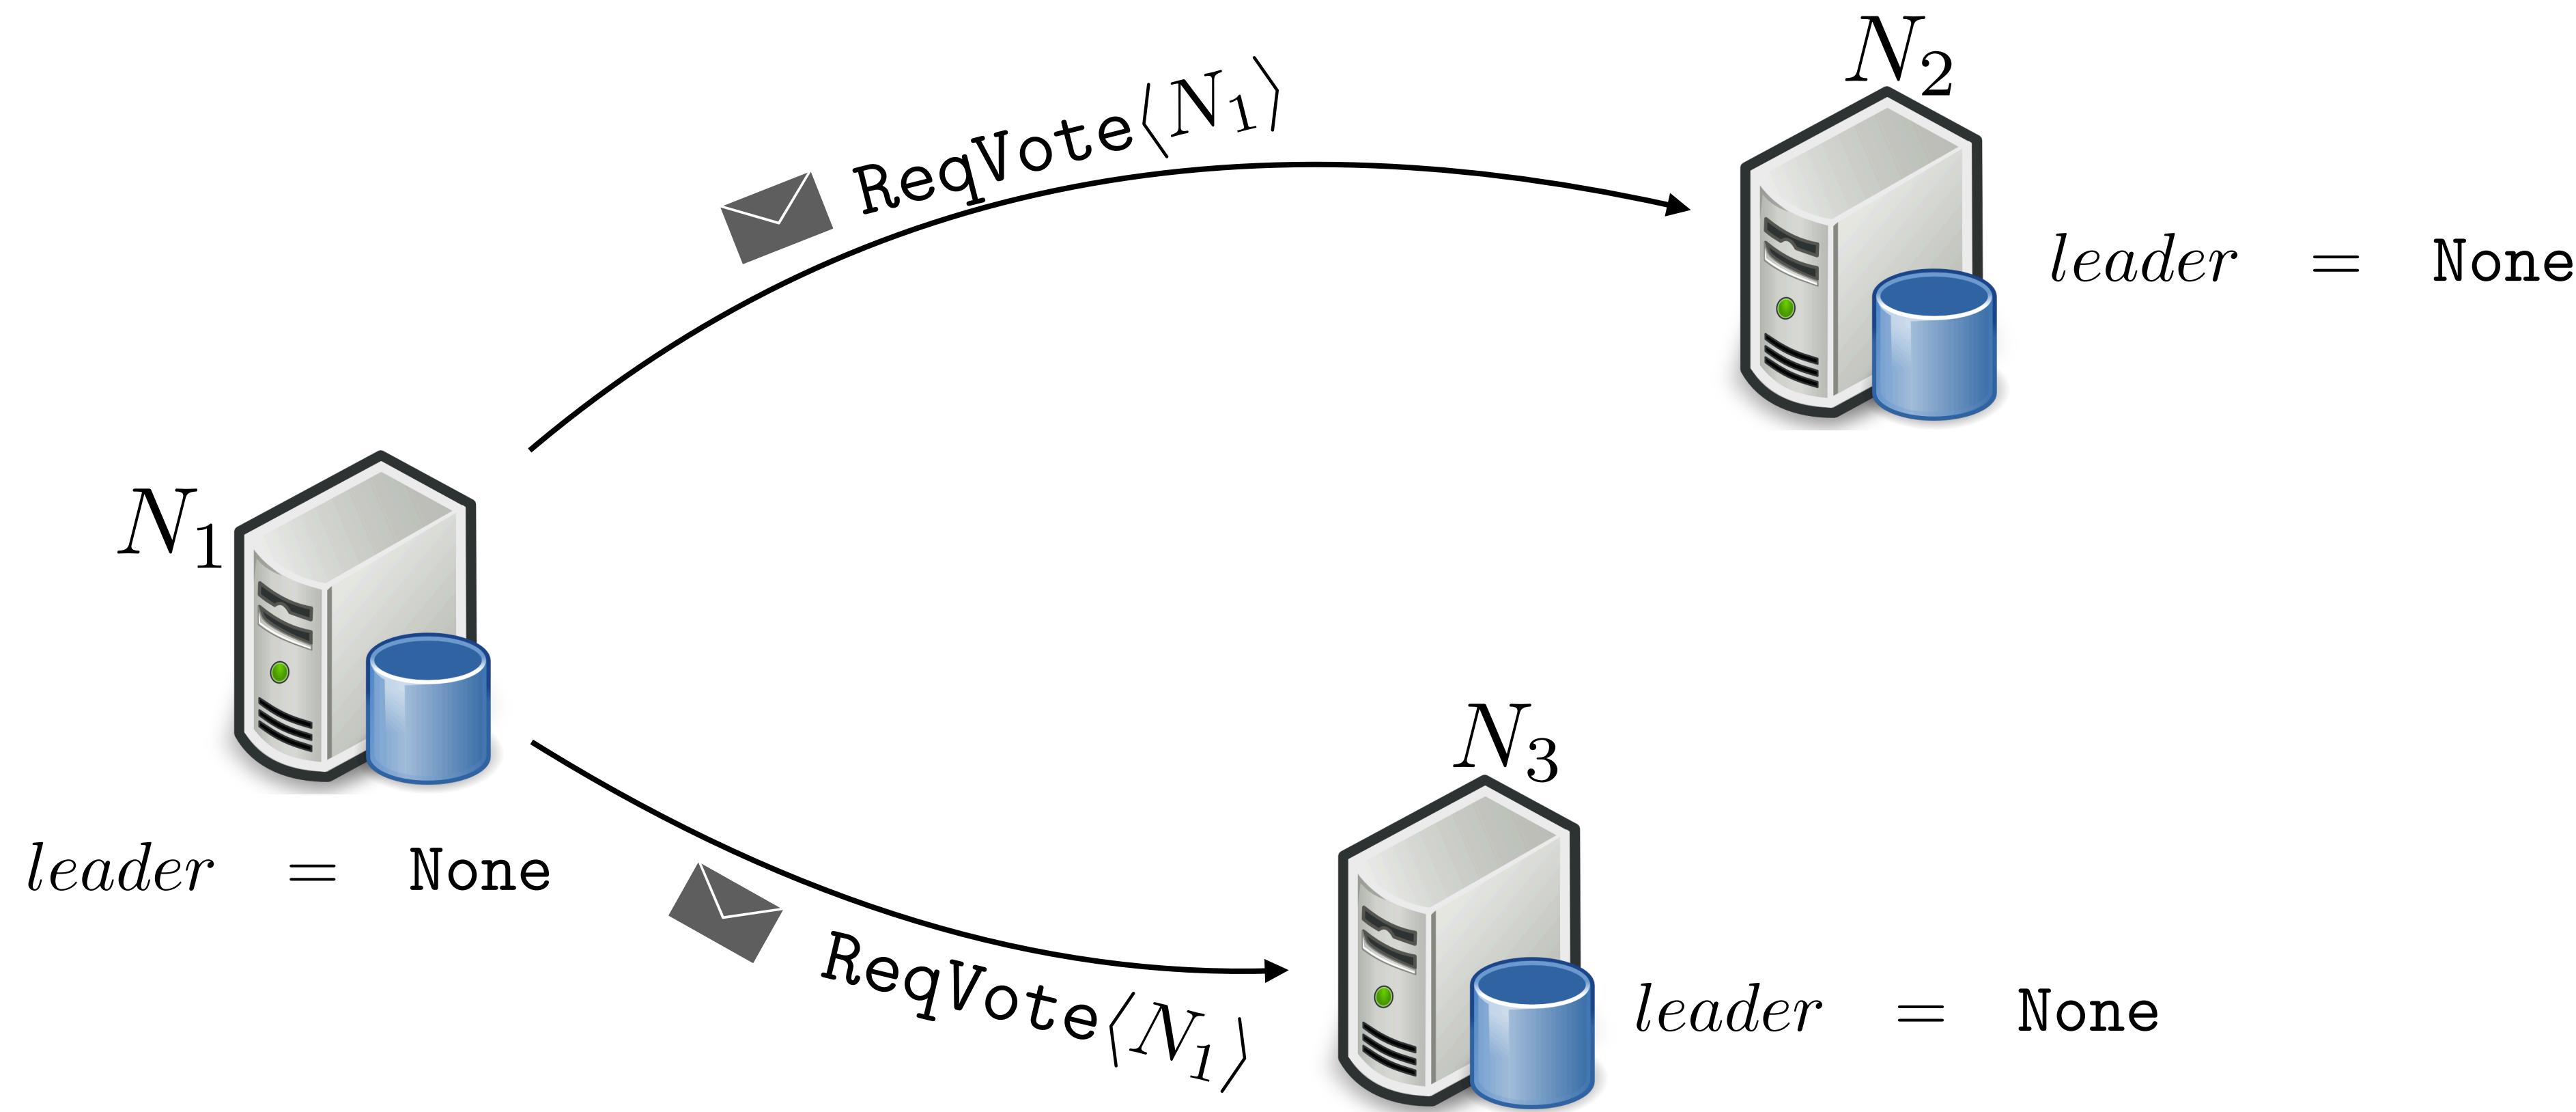
# Key Questions

Can we avoid strong consistency / consensus in practice?

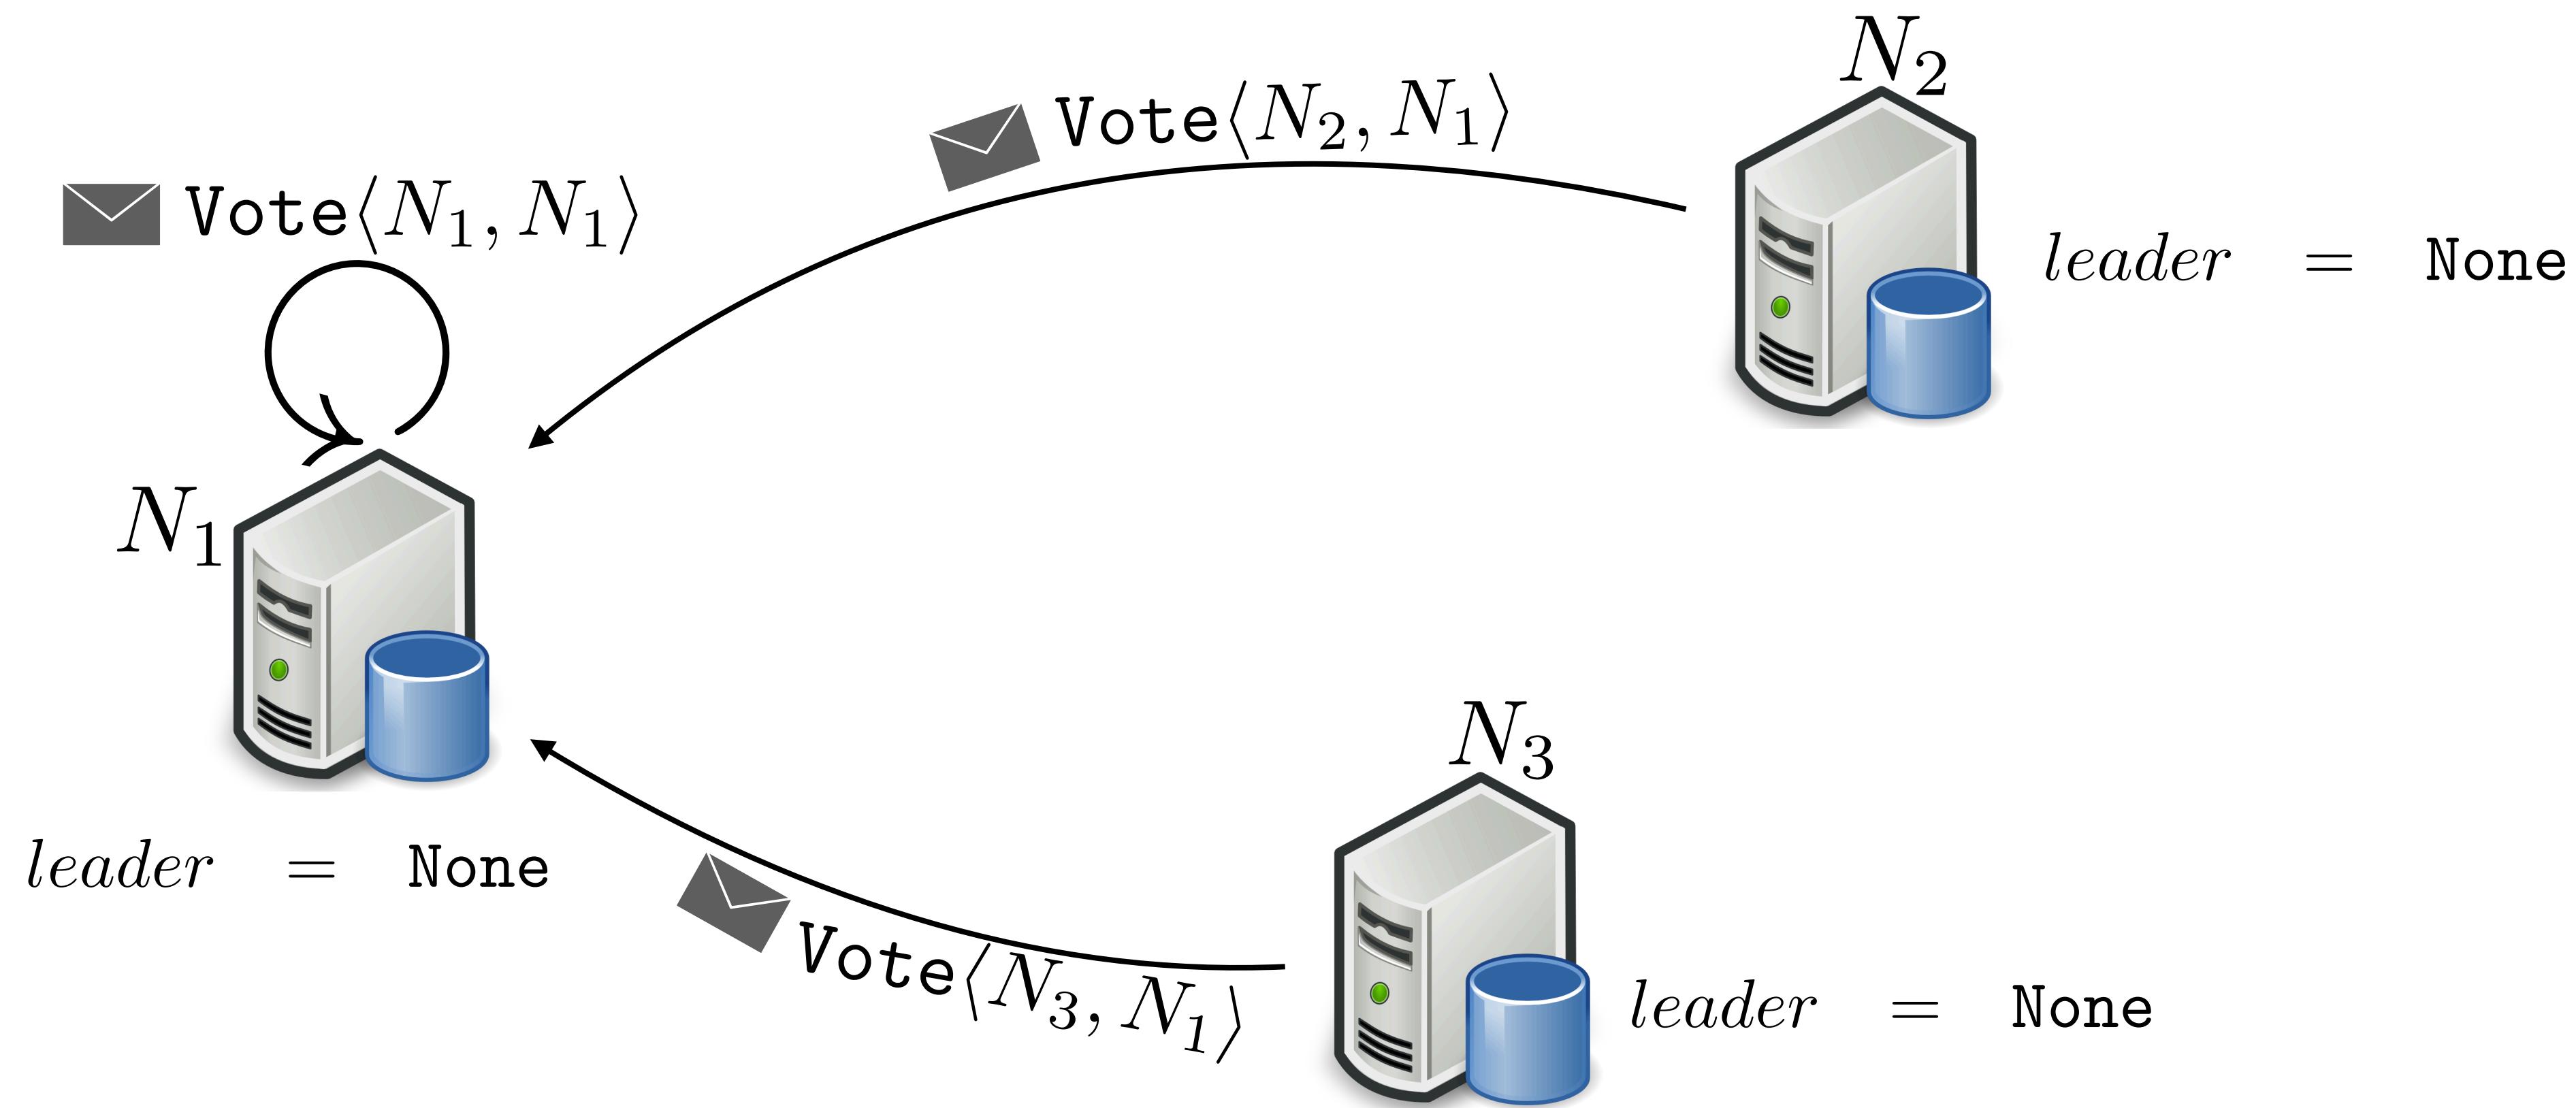
- Negatively affect the *availability* (CAP theorem).
- Extremely complex  $\Rightarrow$  Error-prone implementations

If not, can we make formal verification scale to implementations?

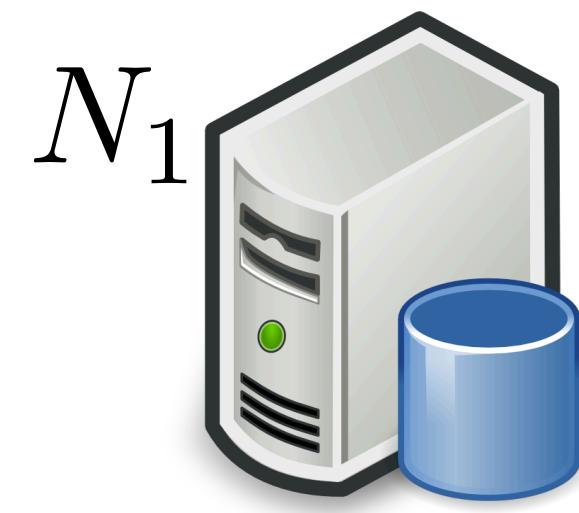
# A Simple Consensus Protocol: Leader Election



# Leader Election



# Leader Election

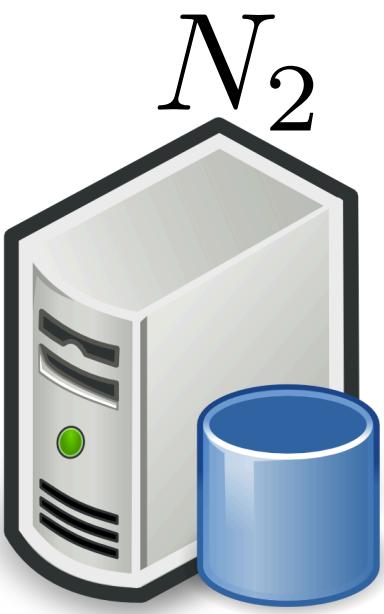


$leader$  = None  
 $votes$  =  $\{N_1, N_2, N_3\}$   
 $hasVoted$  = true

Quorum

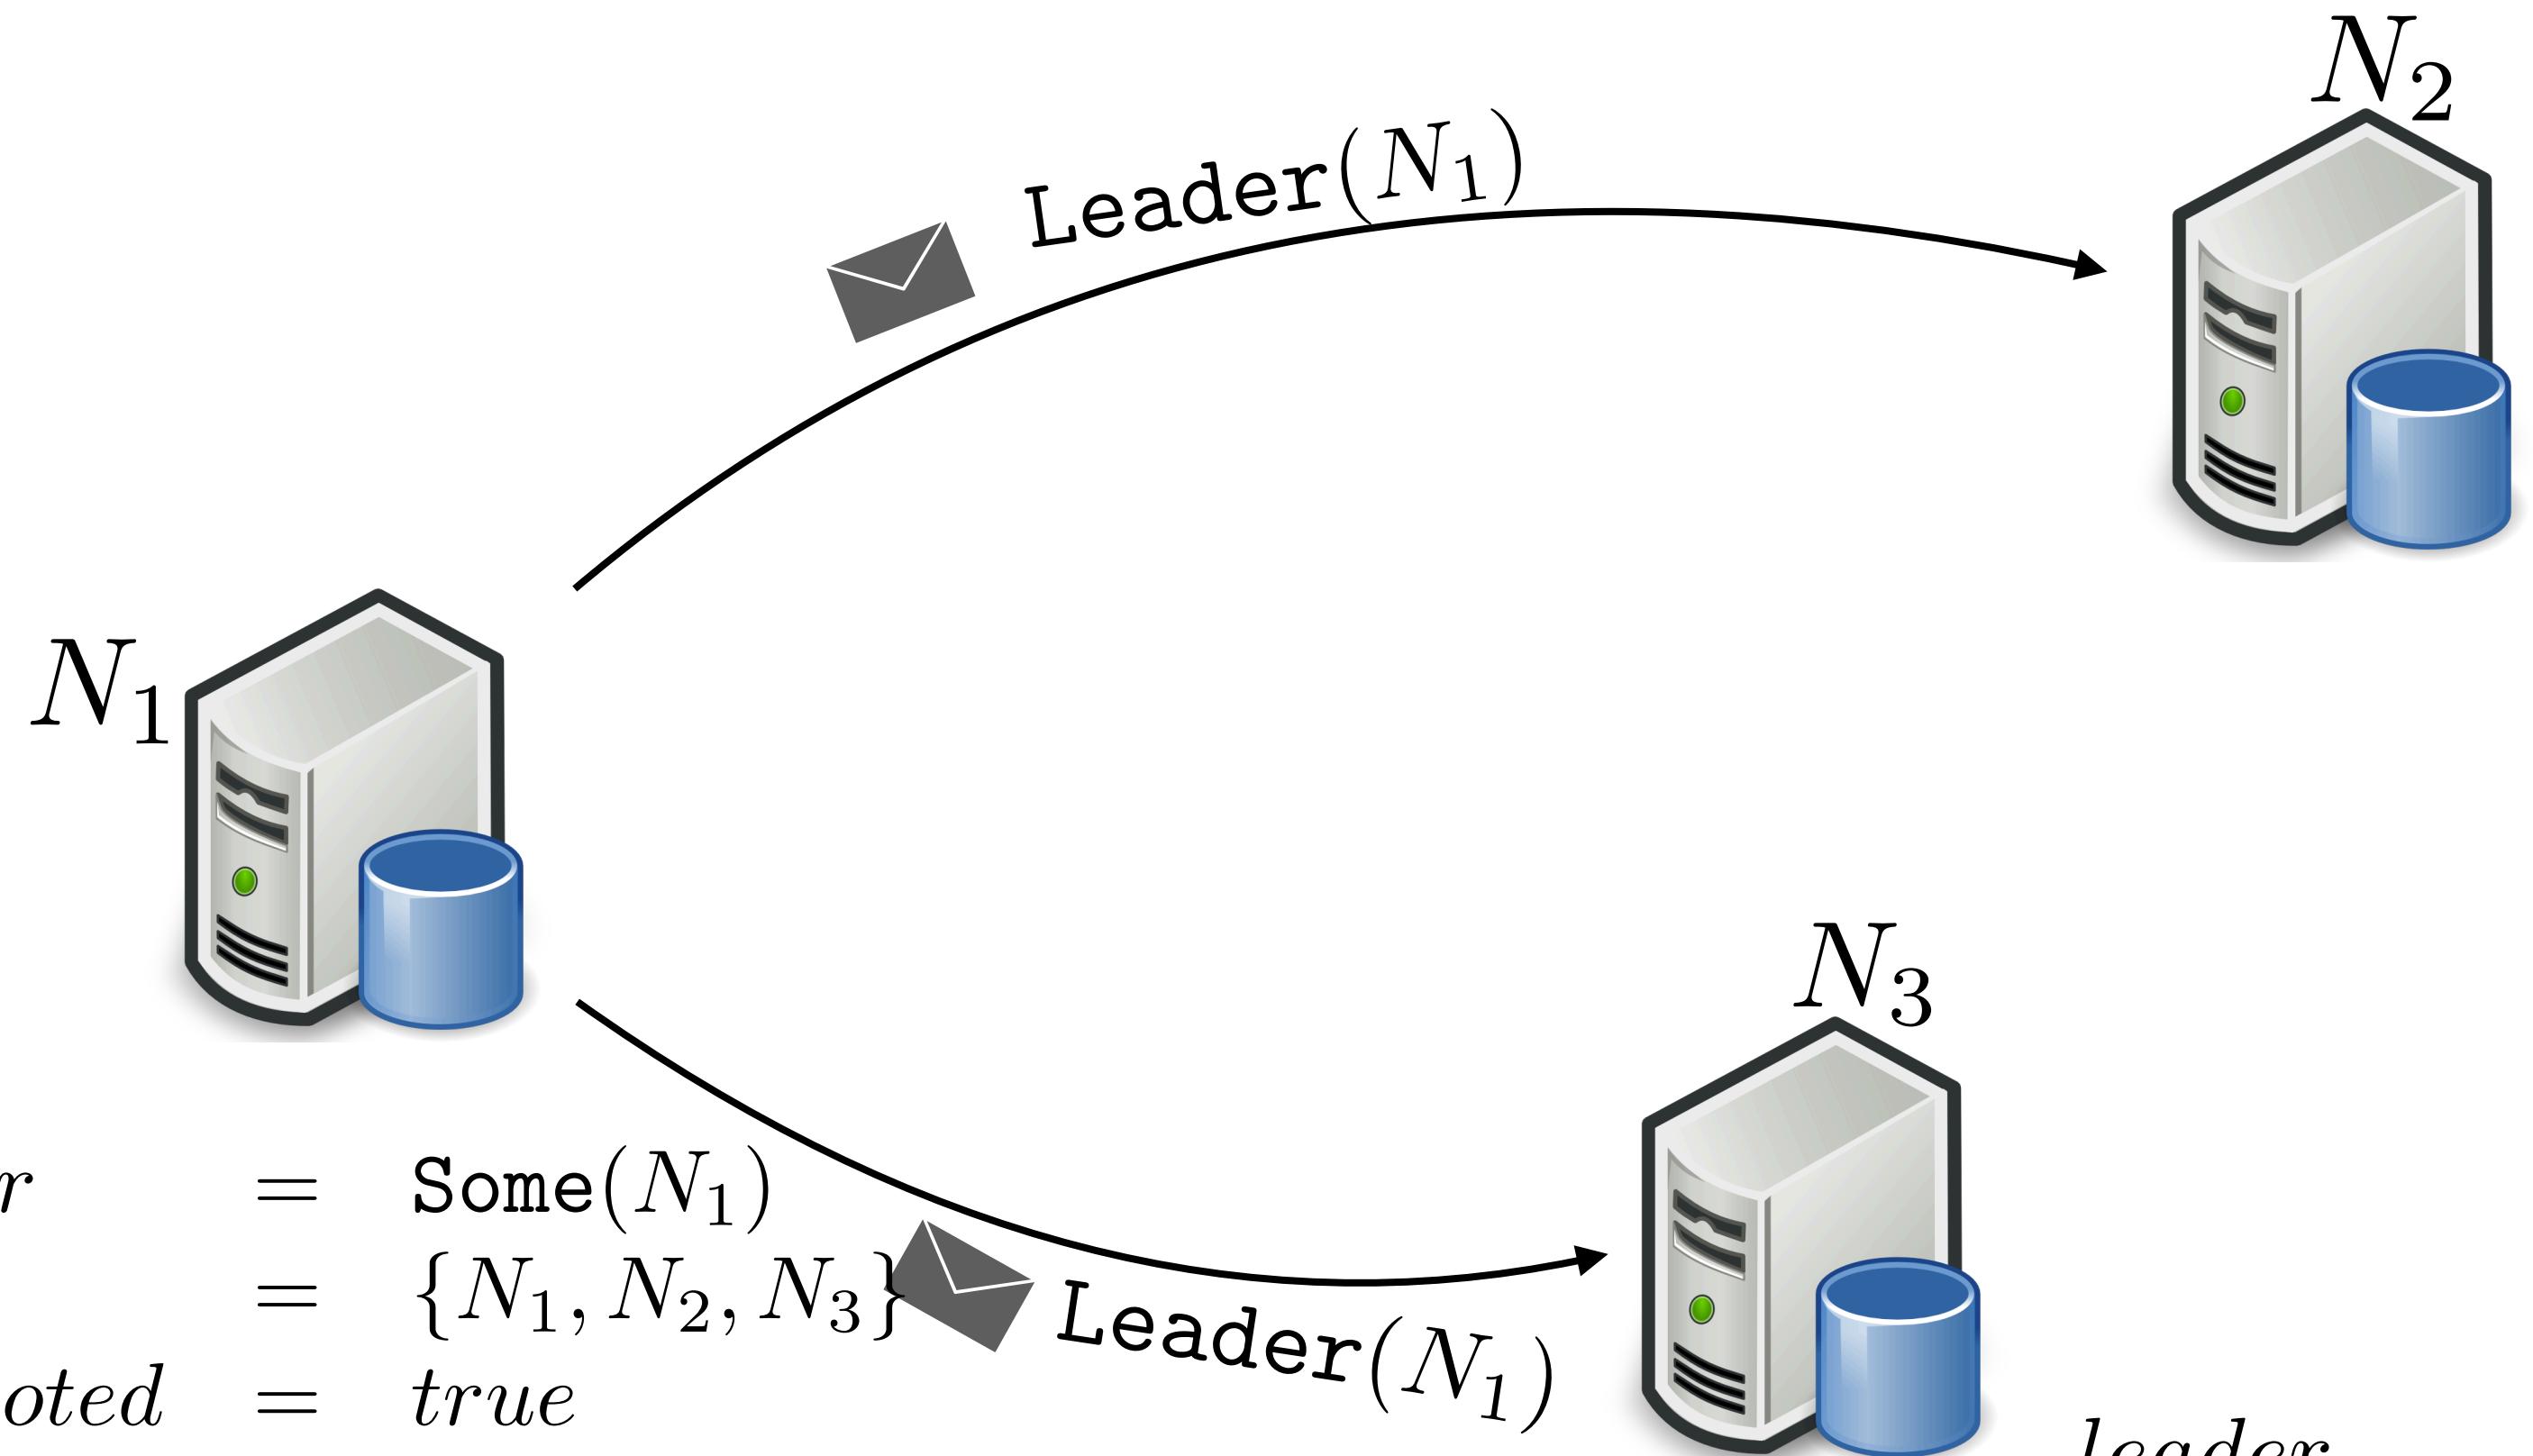


$leader$  = None  
 $votes$  = {}  
 $hasVoted$  = true



$leader$  = None  
 $votes$  = {}  
 $hasVoted$  = true

# Leader Election



$leader = \text{Some}(N_1)$

$votes = \{N_1, N_2, N_3\}$

$hasVoted = \text{true}$

$leader = \text{None}$

$votes = \{\}$

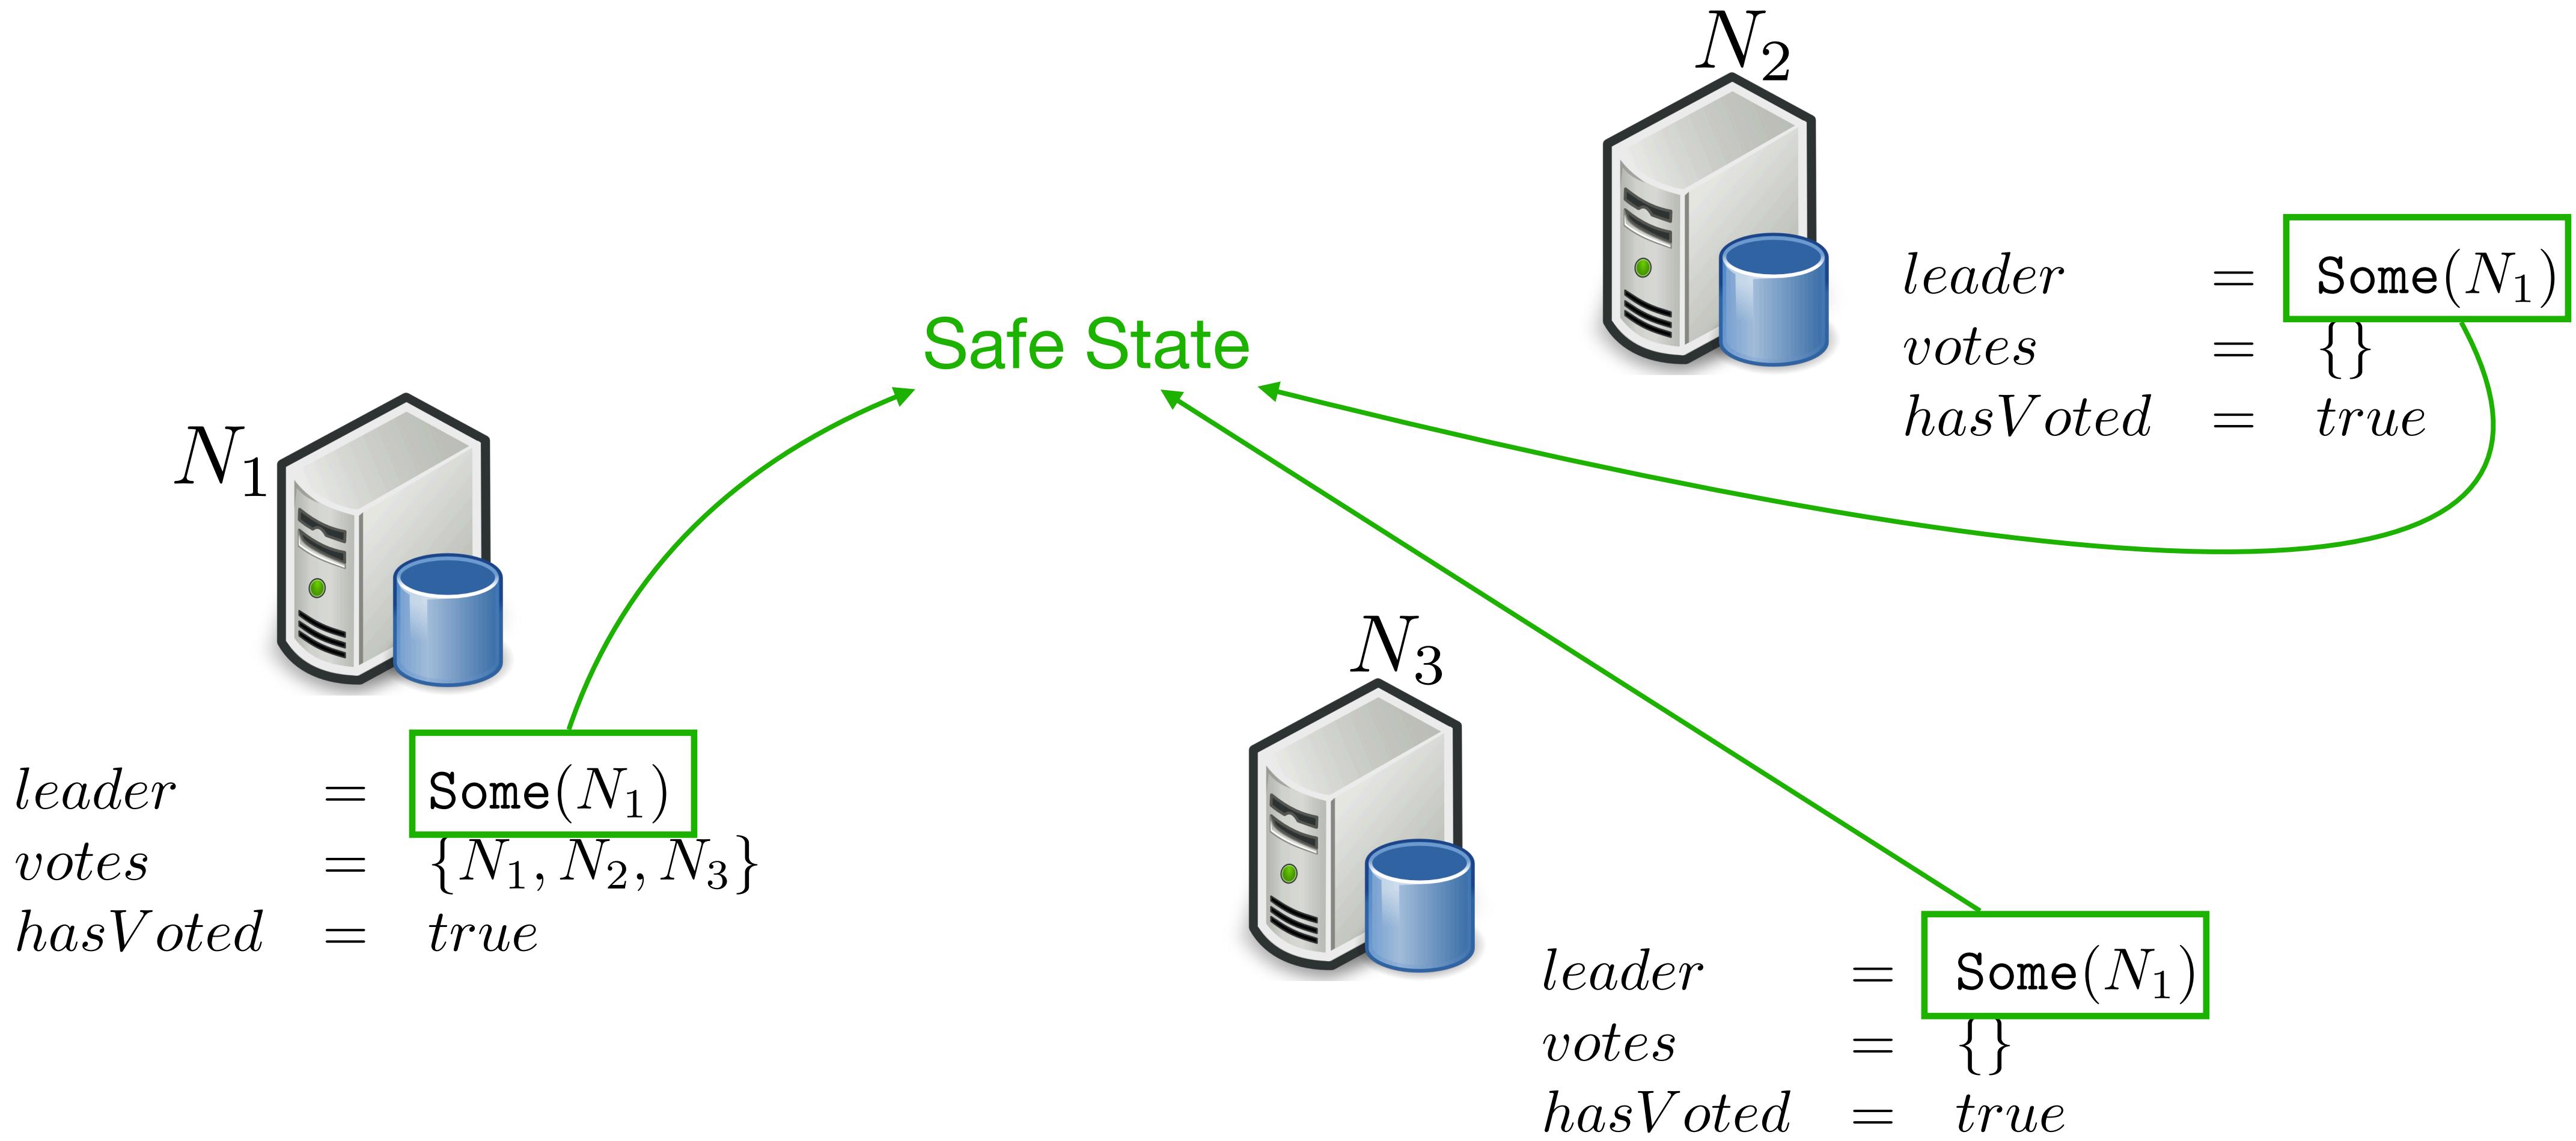
$hasVoted = \text{true}$

$leader = \text{None}$

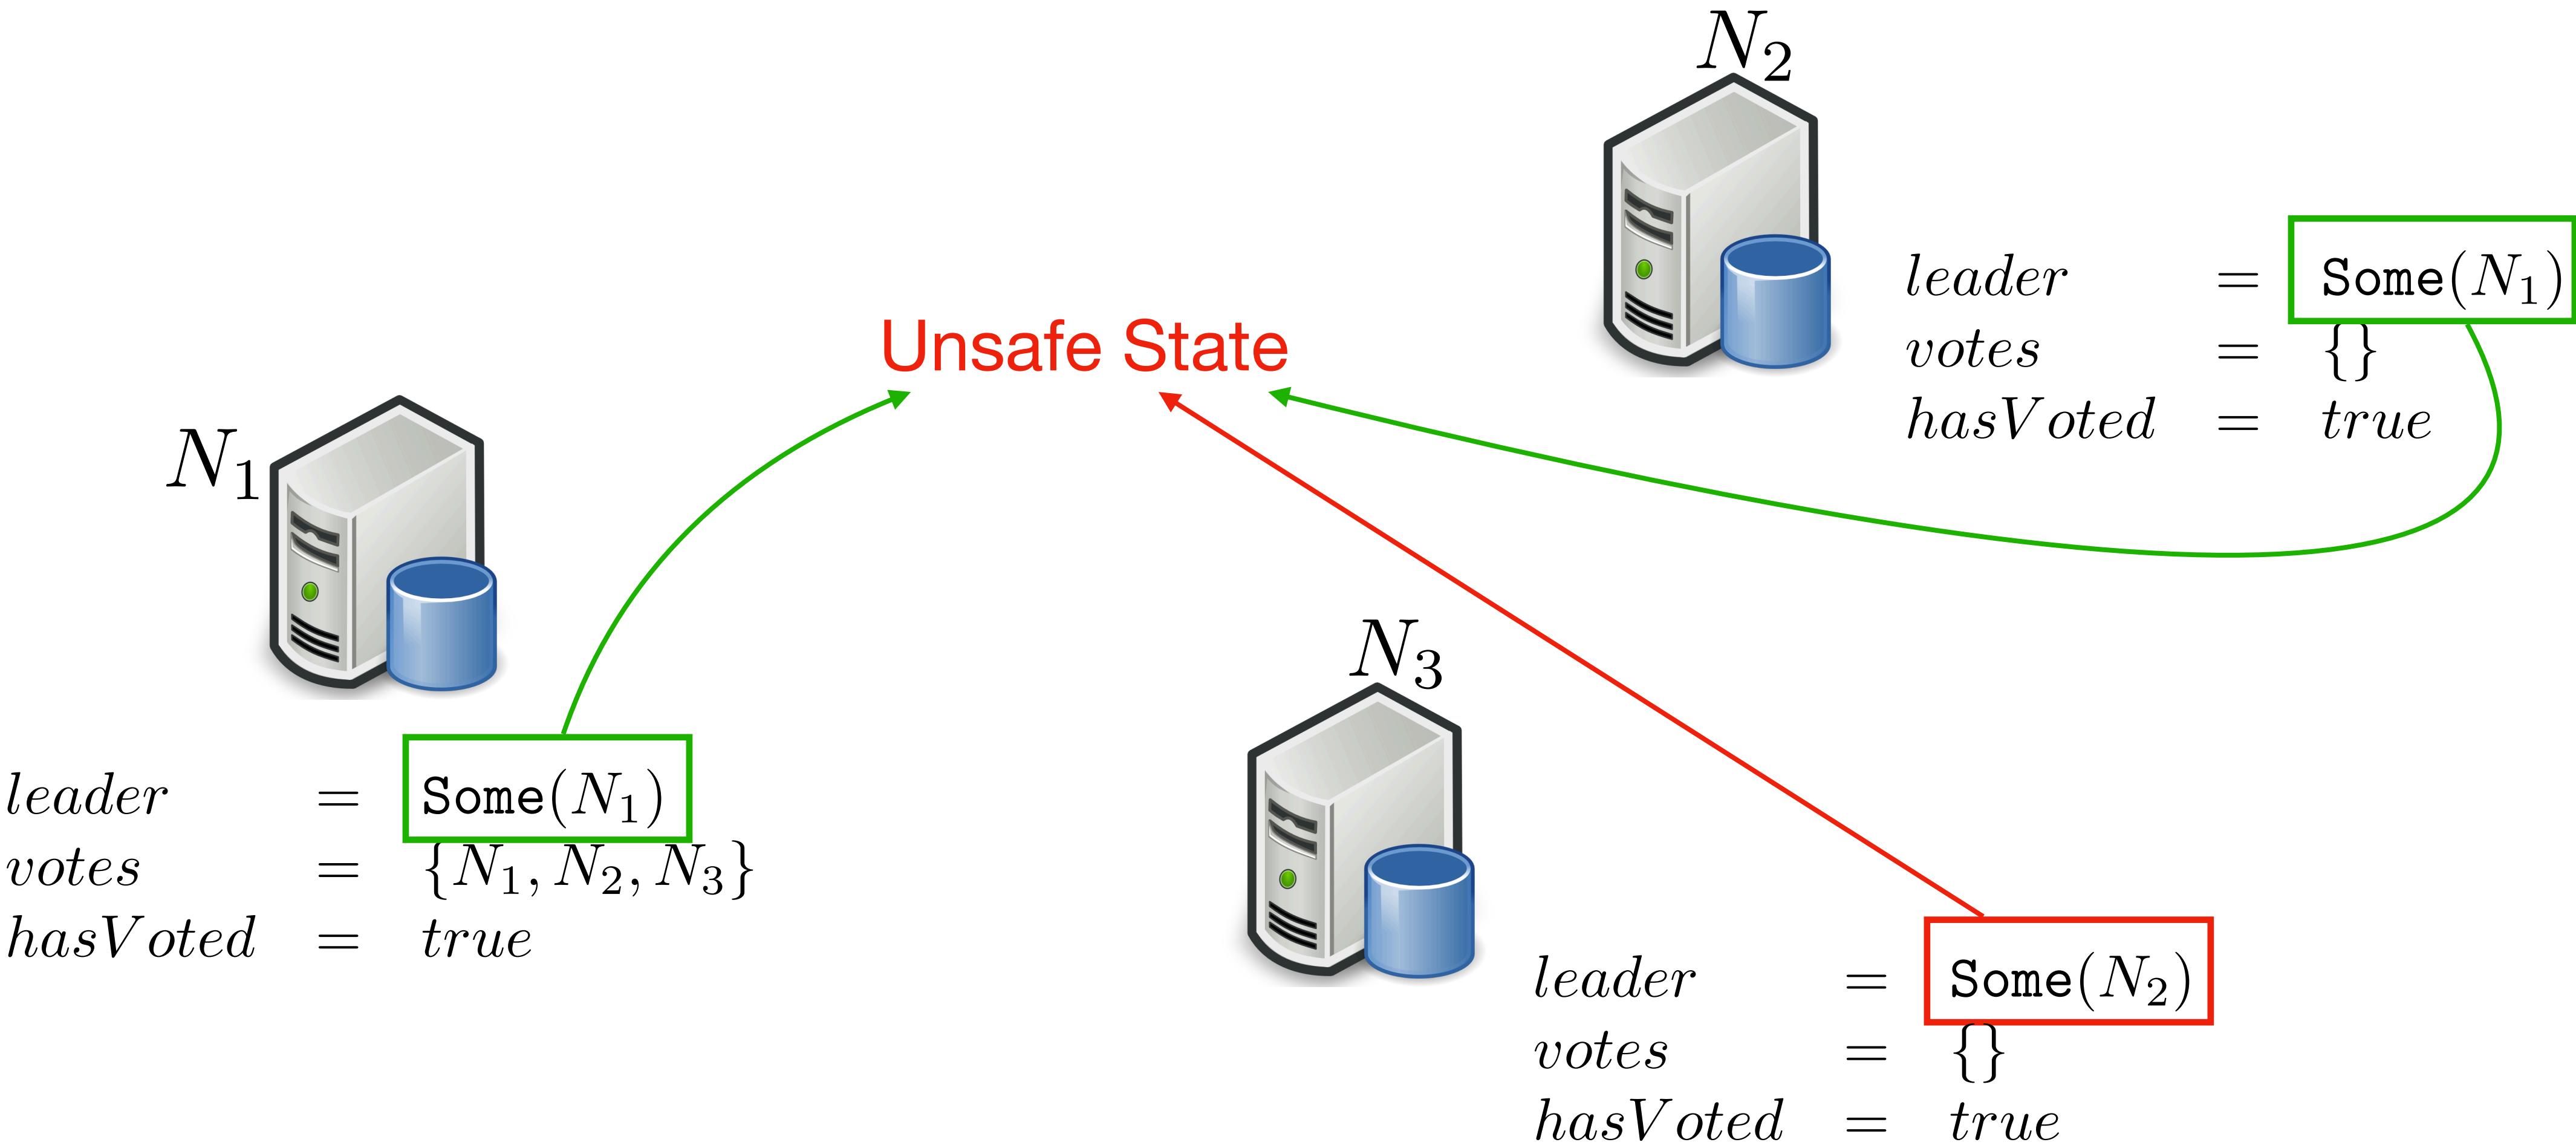
$votes = \{\}$

$hasVoted = \text{true}$

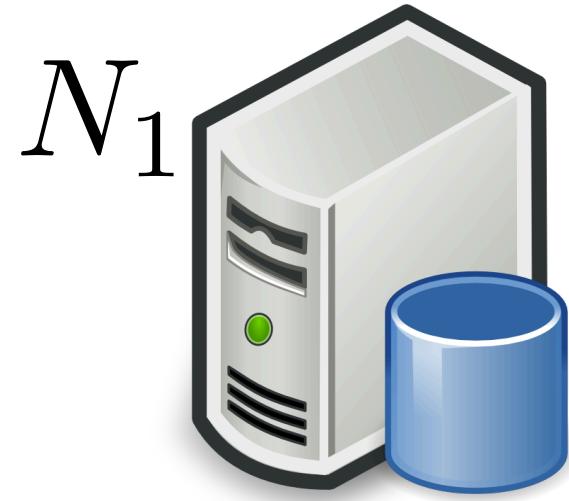
# Leader Election



# Leader Election



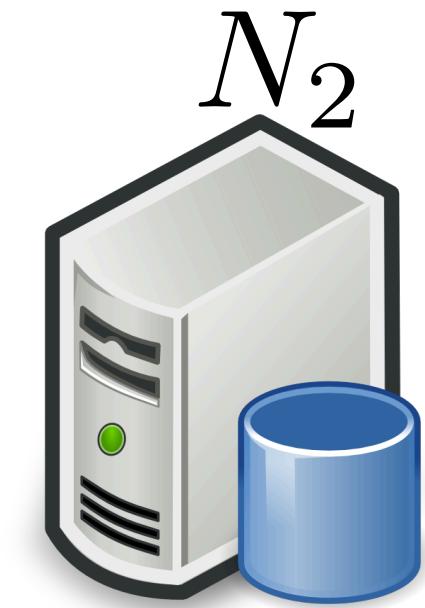
# Leader Election: Safety Spec



*leader* = Some( $N_1$ )  
*votes* =  $\{N_1, N_2, N_3\}$   
*hasVoted* = true



*leader* = Some( $N_2$ )  
*votes* = {}  
*hasVoted* = true



*leader* = Some( $N_1$ )  
*votes* = {}  
*hasVoted* = true

$$\phi = \forall(N_1, N_2, N_3, N_4). \text{leader}[N_1] = \text{Some}(N_2) \wedge \text{leader}[N_3] = \text{Some}(N_4) \Rightarrow N_2 = N_4$$

# Proving Safety of Leader Election: Inductive Invariant

$$\mathbb{I} = \bigwedge_{i=1}^6 \varphi_i$$

$\varphi_1 = \forall N. \text{Leader}\langle N \rangle \in msgs \Rightarrow votes[N] \in Quorum$

1. A Leader message exists only if the sender received a quorum of vote

$\varphi_2 = \forall N_1, N_2. N_1 \in votes[N_2] \Rightarrow \text{Vote}\langle N_1, N_2 \rangle \in msgs$

2. A node counts a vote only if the corresponding Vote message exists

$\varphi_3 = \forall N_1, N_2. \text{Vote}\langle N_1, N_2 \rangle \in msgs \Rightarrow hasVoted[N_1] = true$

3. If a node has voted then its *hasVoted* flag is set

$\varphi_4 = \forall N_1, N_2, N_3. \text{Vote}\langle N_1, N_2 \rangle \in msgs \wedge \text{Vote}\langle N_1, N_3 \rangle \in msgs \Rightarrow N_2 = N_3$

4. A node only ever casts one vote

$\varphi_5 = \forall N_1, N_2. leader[N_1] = N_2 \Rightarrow \text{Leader}\langle N_2 \rangle \in msgs$

5. A node sets a leader only if the corresponding Leader message exists

$\varphi_6 = \forall N_1, N_2. \text{Leader}\langle N_1 \rangle \in msgs \wedge \text{Leader}\langle N_2 \rangle \in msgs \Rightarrow N_1 = N_2$

6. A Leader message is only ever broadcast for one node

# Inductive Invariant $\Rightarrow$ Undecidability

$$\varphi_2 = \boxed{\forall} N_1, N_2. N_1 \in votes[N_2] \Rightarrow \underbrace{\text{Vote}\langle N_1, N_2 \rangle}_{\text{Is actually written as:}} \in msgs$$
$$\exists m. m \in msgs \wedge \text{typeOf}(m) = \text{Vote} \wedge \text{sender}(m) = N_1 \wedge \text{receiver}(m) = N_2$$

$\forall^* \exists^*$  Quantifier alternation is undecidable.

2

# Undecidability $\Rightarrow$ Manual Proofs

## PROBLEM: Too MUCH WORK!

- ▶ Hawblitzel et al. 2015 *IronFleet: Proving Practical Distributed Systems Correct.*  
“...developing IronFleet and applying it to two real systems required approximately 3.7 person-years”
- ▶ Wilcox et al. 2015 *Verdi: A Framework for Implementing and Formally Verifying Distributed Systems.*
- ▶ Woos et al. 2016 *Planning for Change in a Formal Verification of the Raft Consensus Protocol.*  
“...required iteratively discovering and proving 90 system invariants”  
“...530 lines of code and 50,000 lines of proof.”
- ▶ Padon et al. 2017 *Paxos made EPR: Decidable Reasoning about Distributed Protocols.*
- ▶ Taube et al. 2018 *Modularity for Decidability of Deductive Verification with Applications to Distributed Systems.*  
“...took approximately 3 person-months ...300 [lines] of invariants and ghost code.”
- ▶ Gleissenthall et al. 2019. *Pretend Synchrony: Synchronous Verification of Asynchronous Distributed Systems.*

How do we avoid this?

# Network is the root of complexity

$$\phi = \forall(N_1, N_2, N_3, N_4). \text{leader}[N_1] = \text{Some}(N_2) \wedge \text{leader}[N_3] = \text{Some}(N_4) \Rightarrow N_2 = N_4$$

$$\varphi_1 = \forall N. \text{Leader}\langle N \rangle \in \boxed{\text{msgs}} \Rightarrow \text{votes}[N] \in \text{Quorum}$$

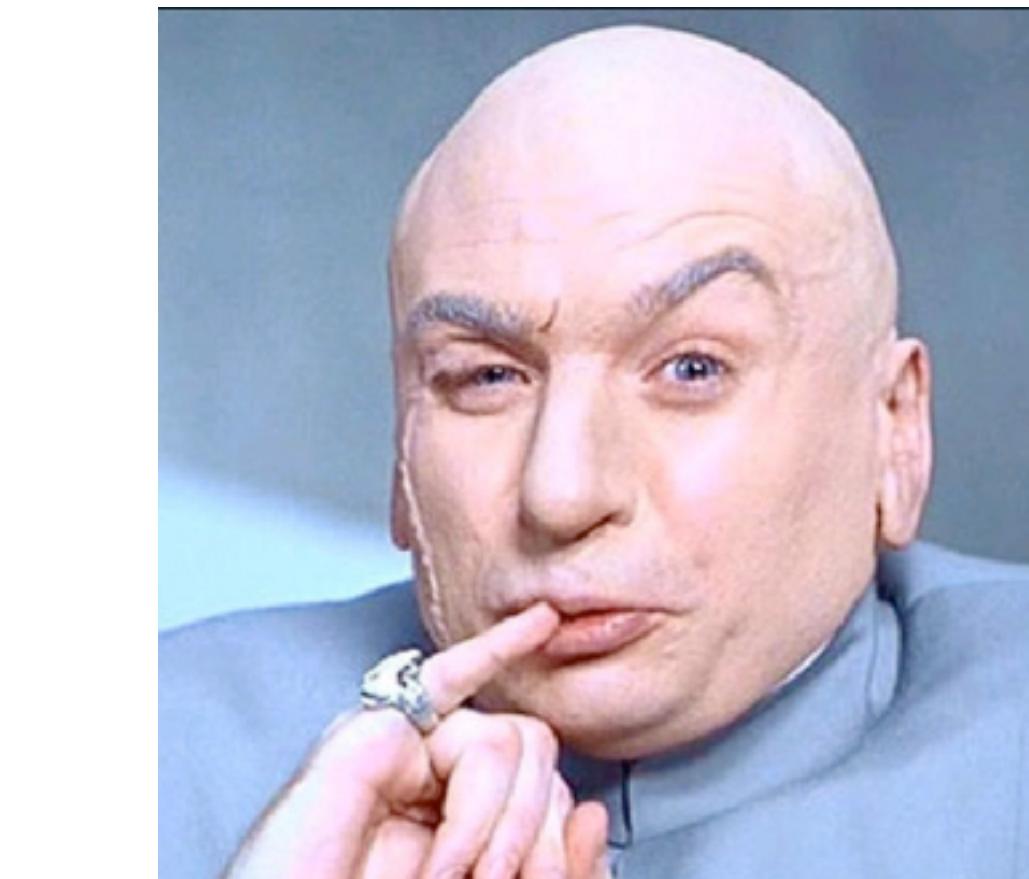
$$\varphi_2 = \forall N_1, N_2. N_1 \in \text{votes}[N_2] \Rightarrow \text{Vote}\langle N_1, N_2 \rangle \in \boxed{\text{msgs}}$$

$$\varphi_3 = \forall N_1, N_2. \text{Vote}\langle N_1, N_2 \rangle \in \boxed{\text{msgs}} \Rightarrow \text{hasVoted}[N_1] = \text{true}$$

$$\begin{aligned} \varphi_4 = \forall N_1, N_2, N_3. \text{Vote}\langle N_1, N_2 \rangle \in \text{msgs} \wedge \text{Vote}\langle N_1, N_3 \rangle \in \boxed{\text{msgs}} \\ \Rightarrow N_2 = N_3 \end{aligned}$$

$$\varphi_5 = \forall N_1, N_2. \text{leader}[N_1] = N_2 \Rightarrow \text{Leader}\langle N_2 \rangle \in \boxed{\text{msgs}}$$

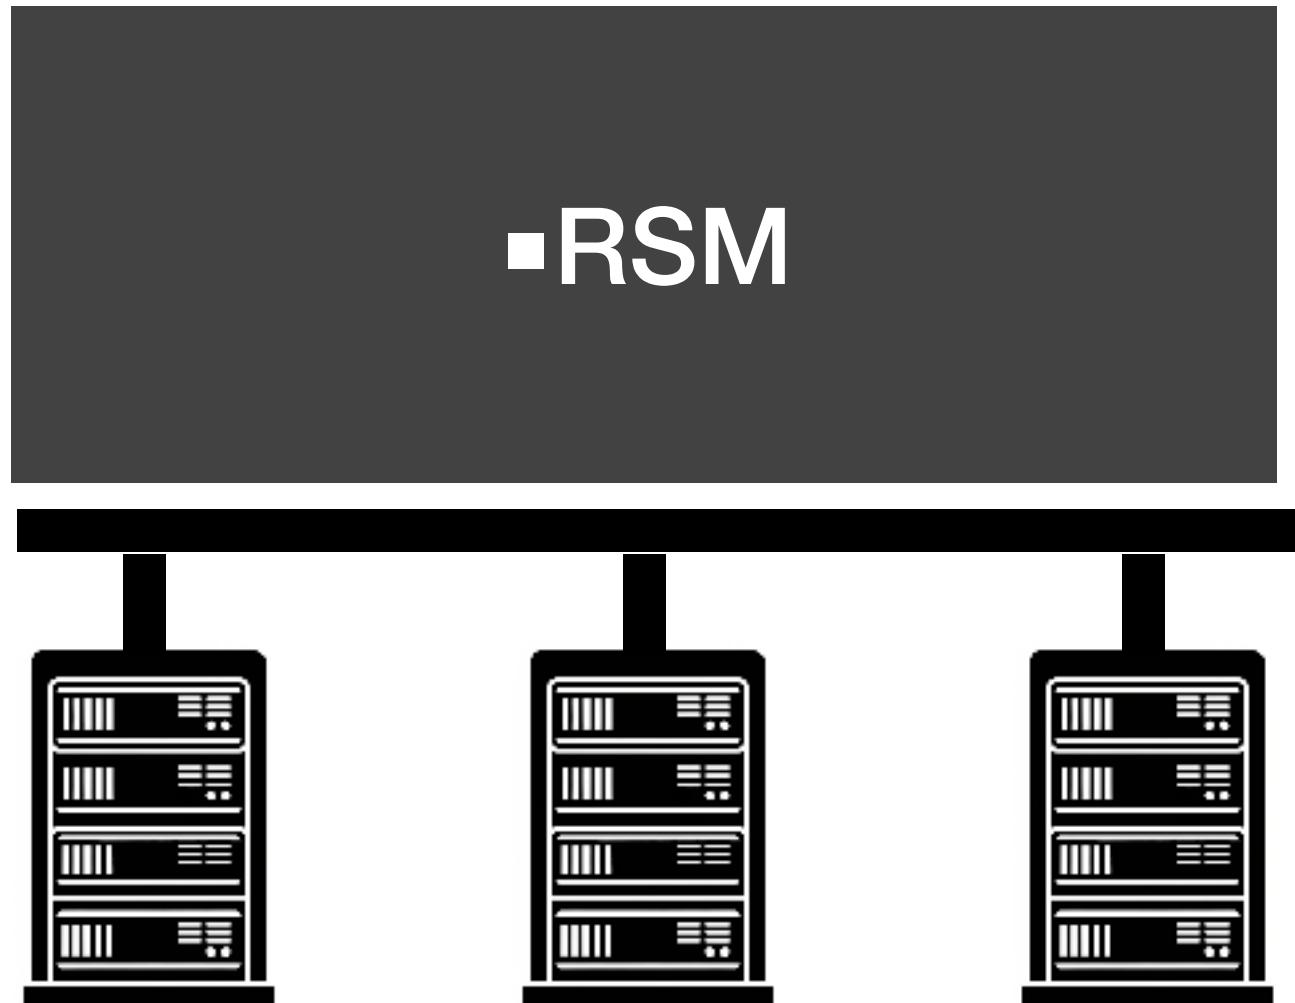
$$\varphi_6 = \forall N_1, N_2. \text{Leader}\langle N_1 \rangle \in \text{msgs} \wedge \text{Leader}\langle N_2 \rangle \in \boxed{\text{msgs}} \Rightarrow N_1 = N_2$$



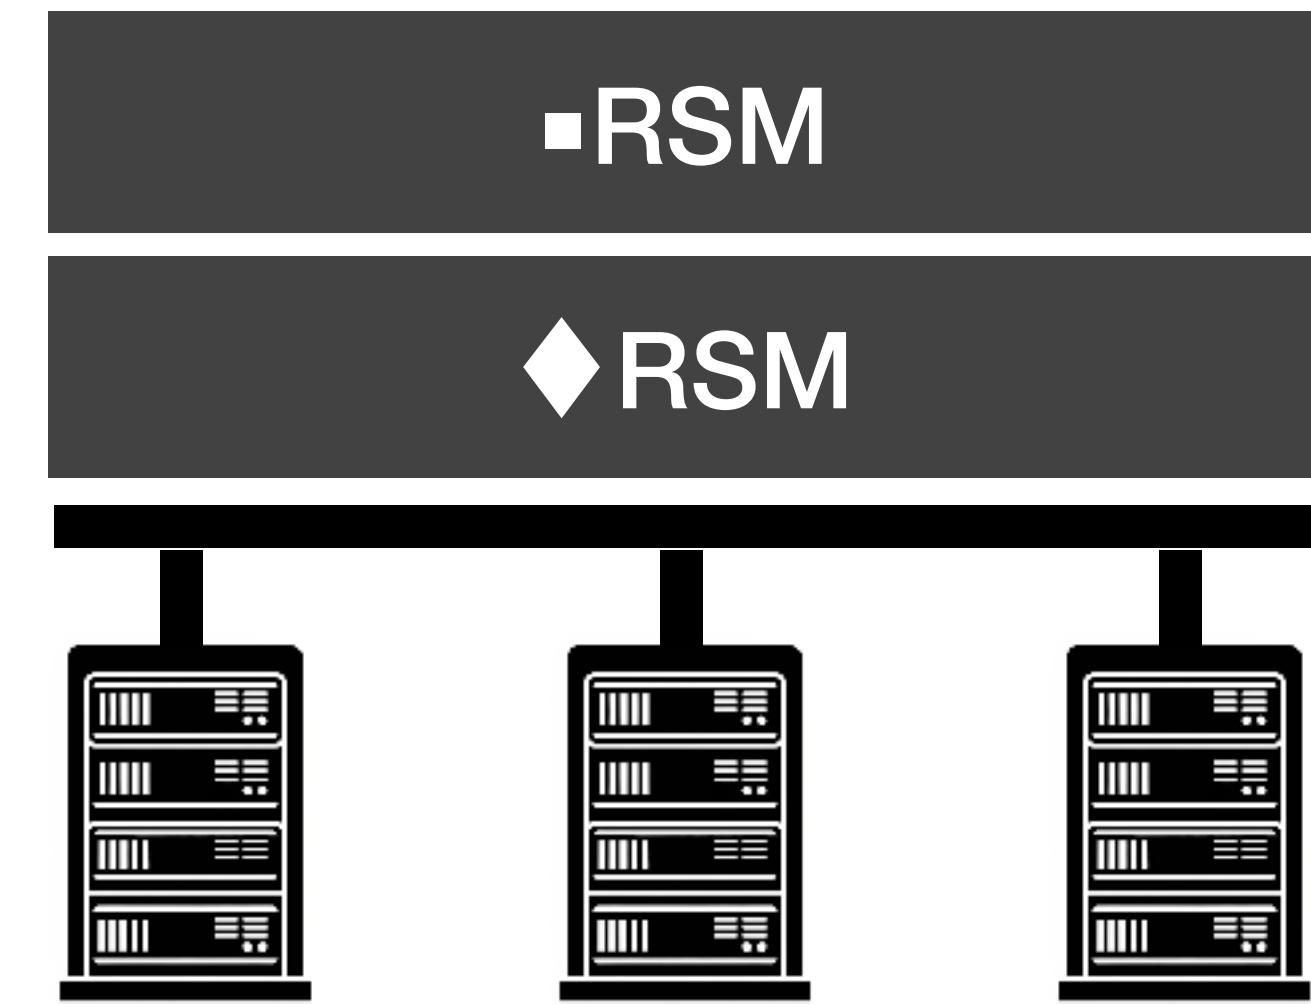
# Network is the root of complexity

*m~~o~~gs*

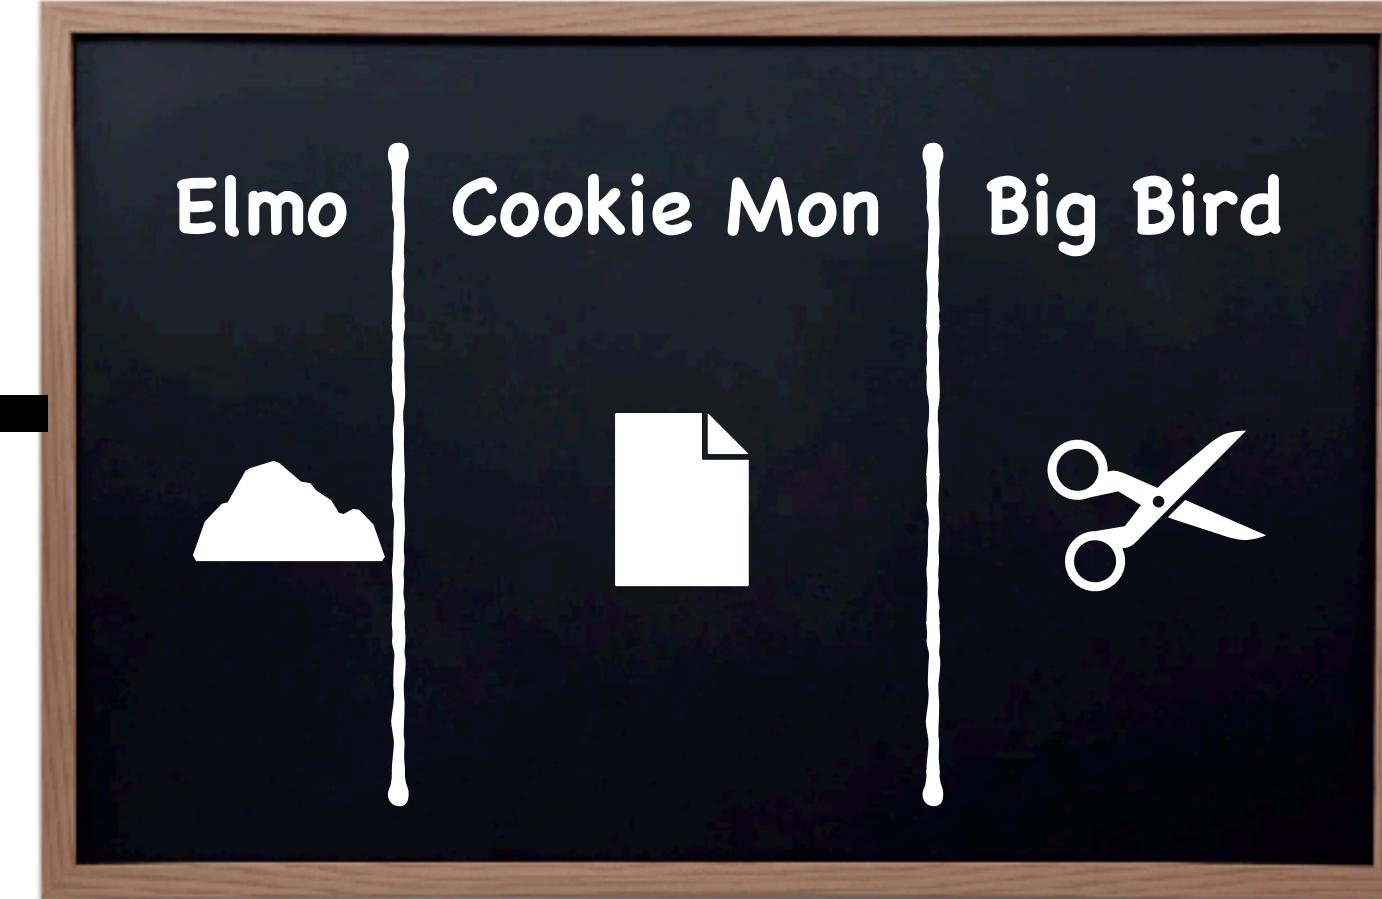
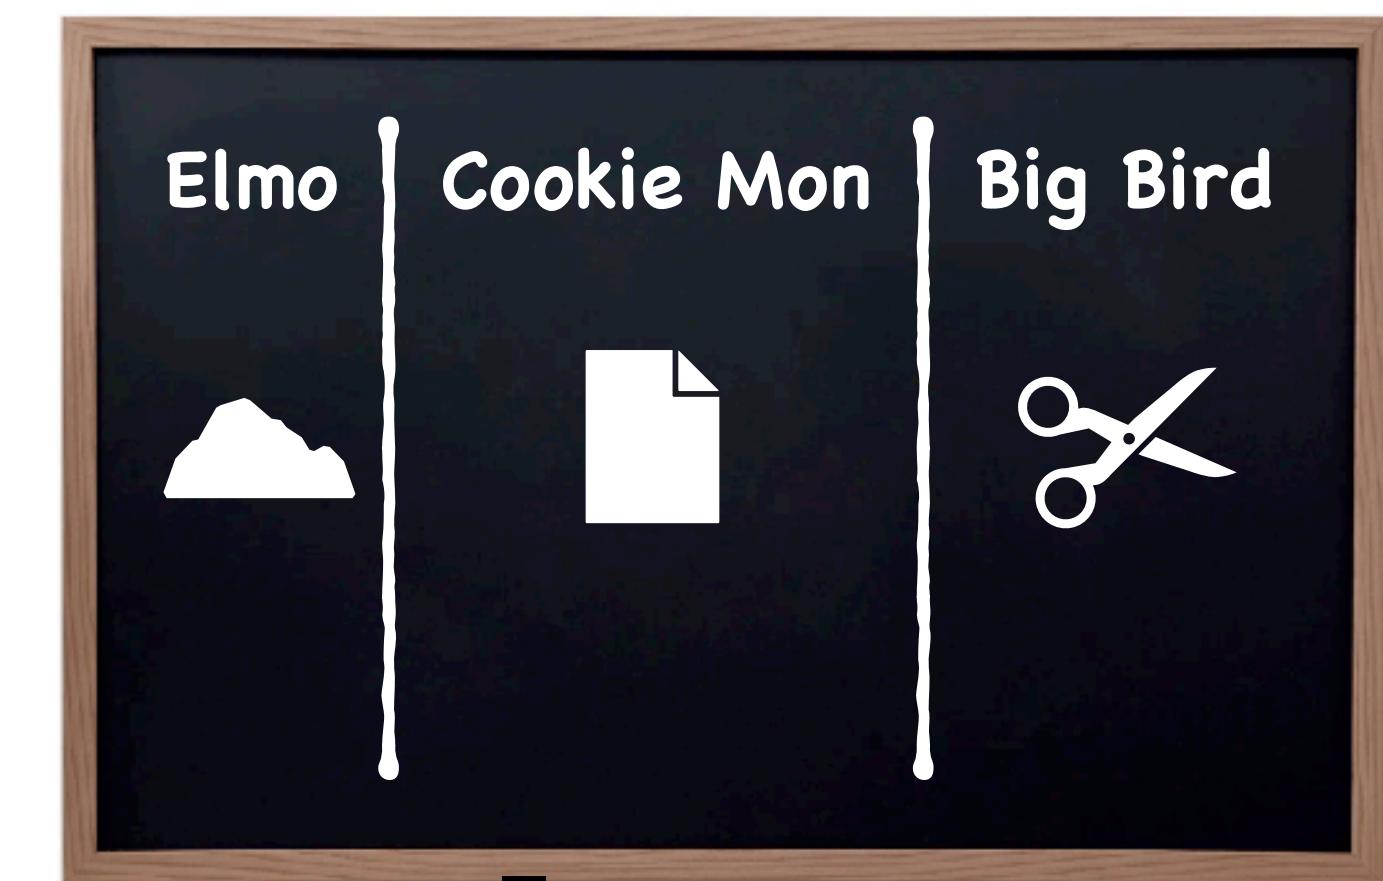
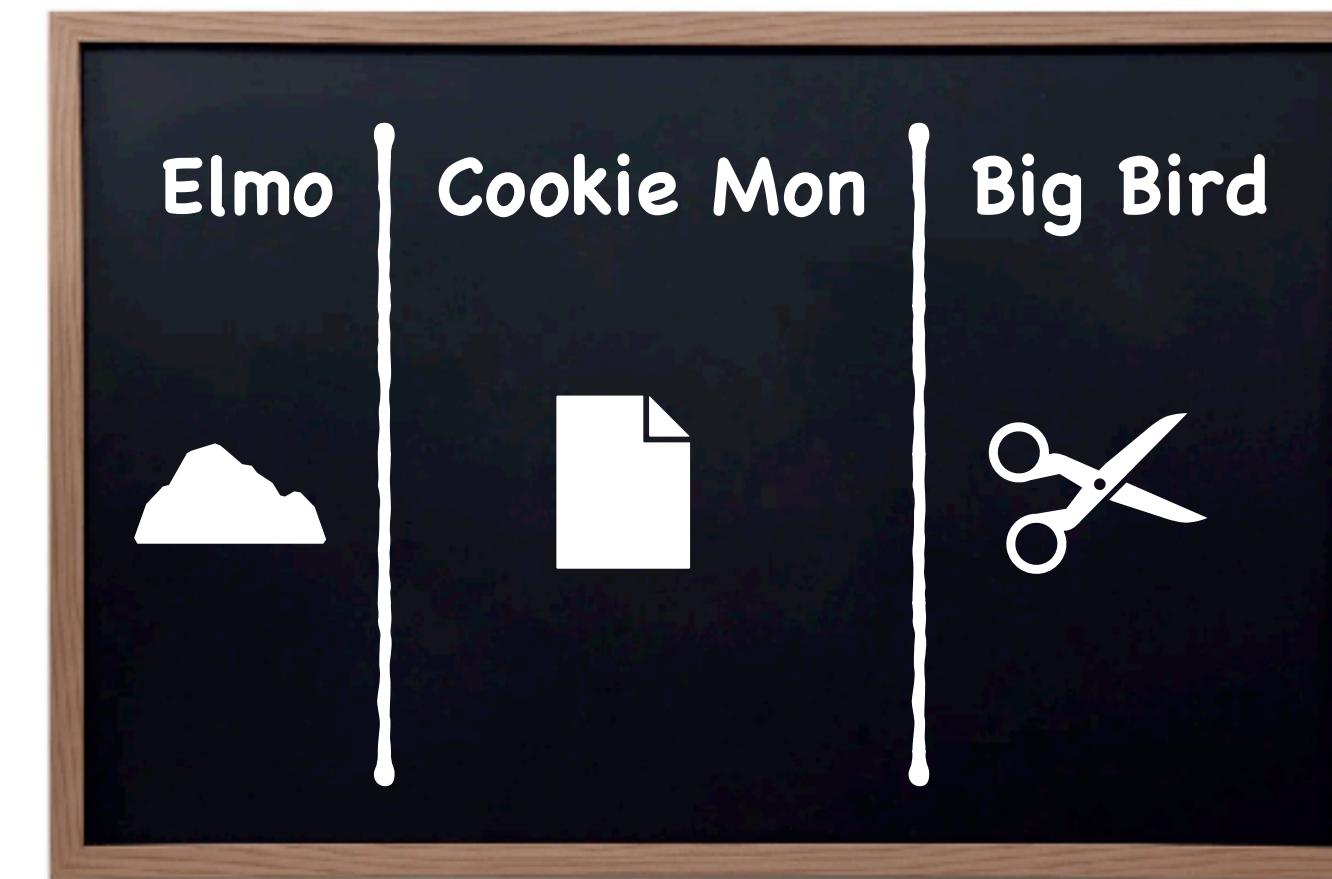
Instead of the conventional *monolithic* approach...



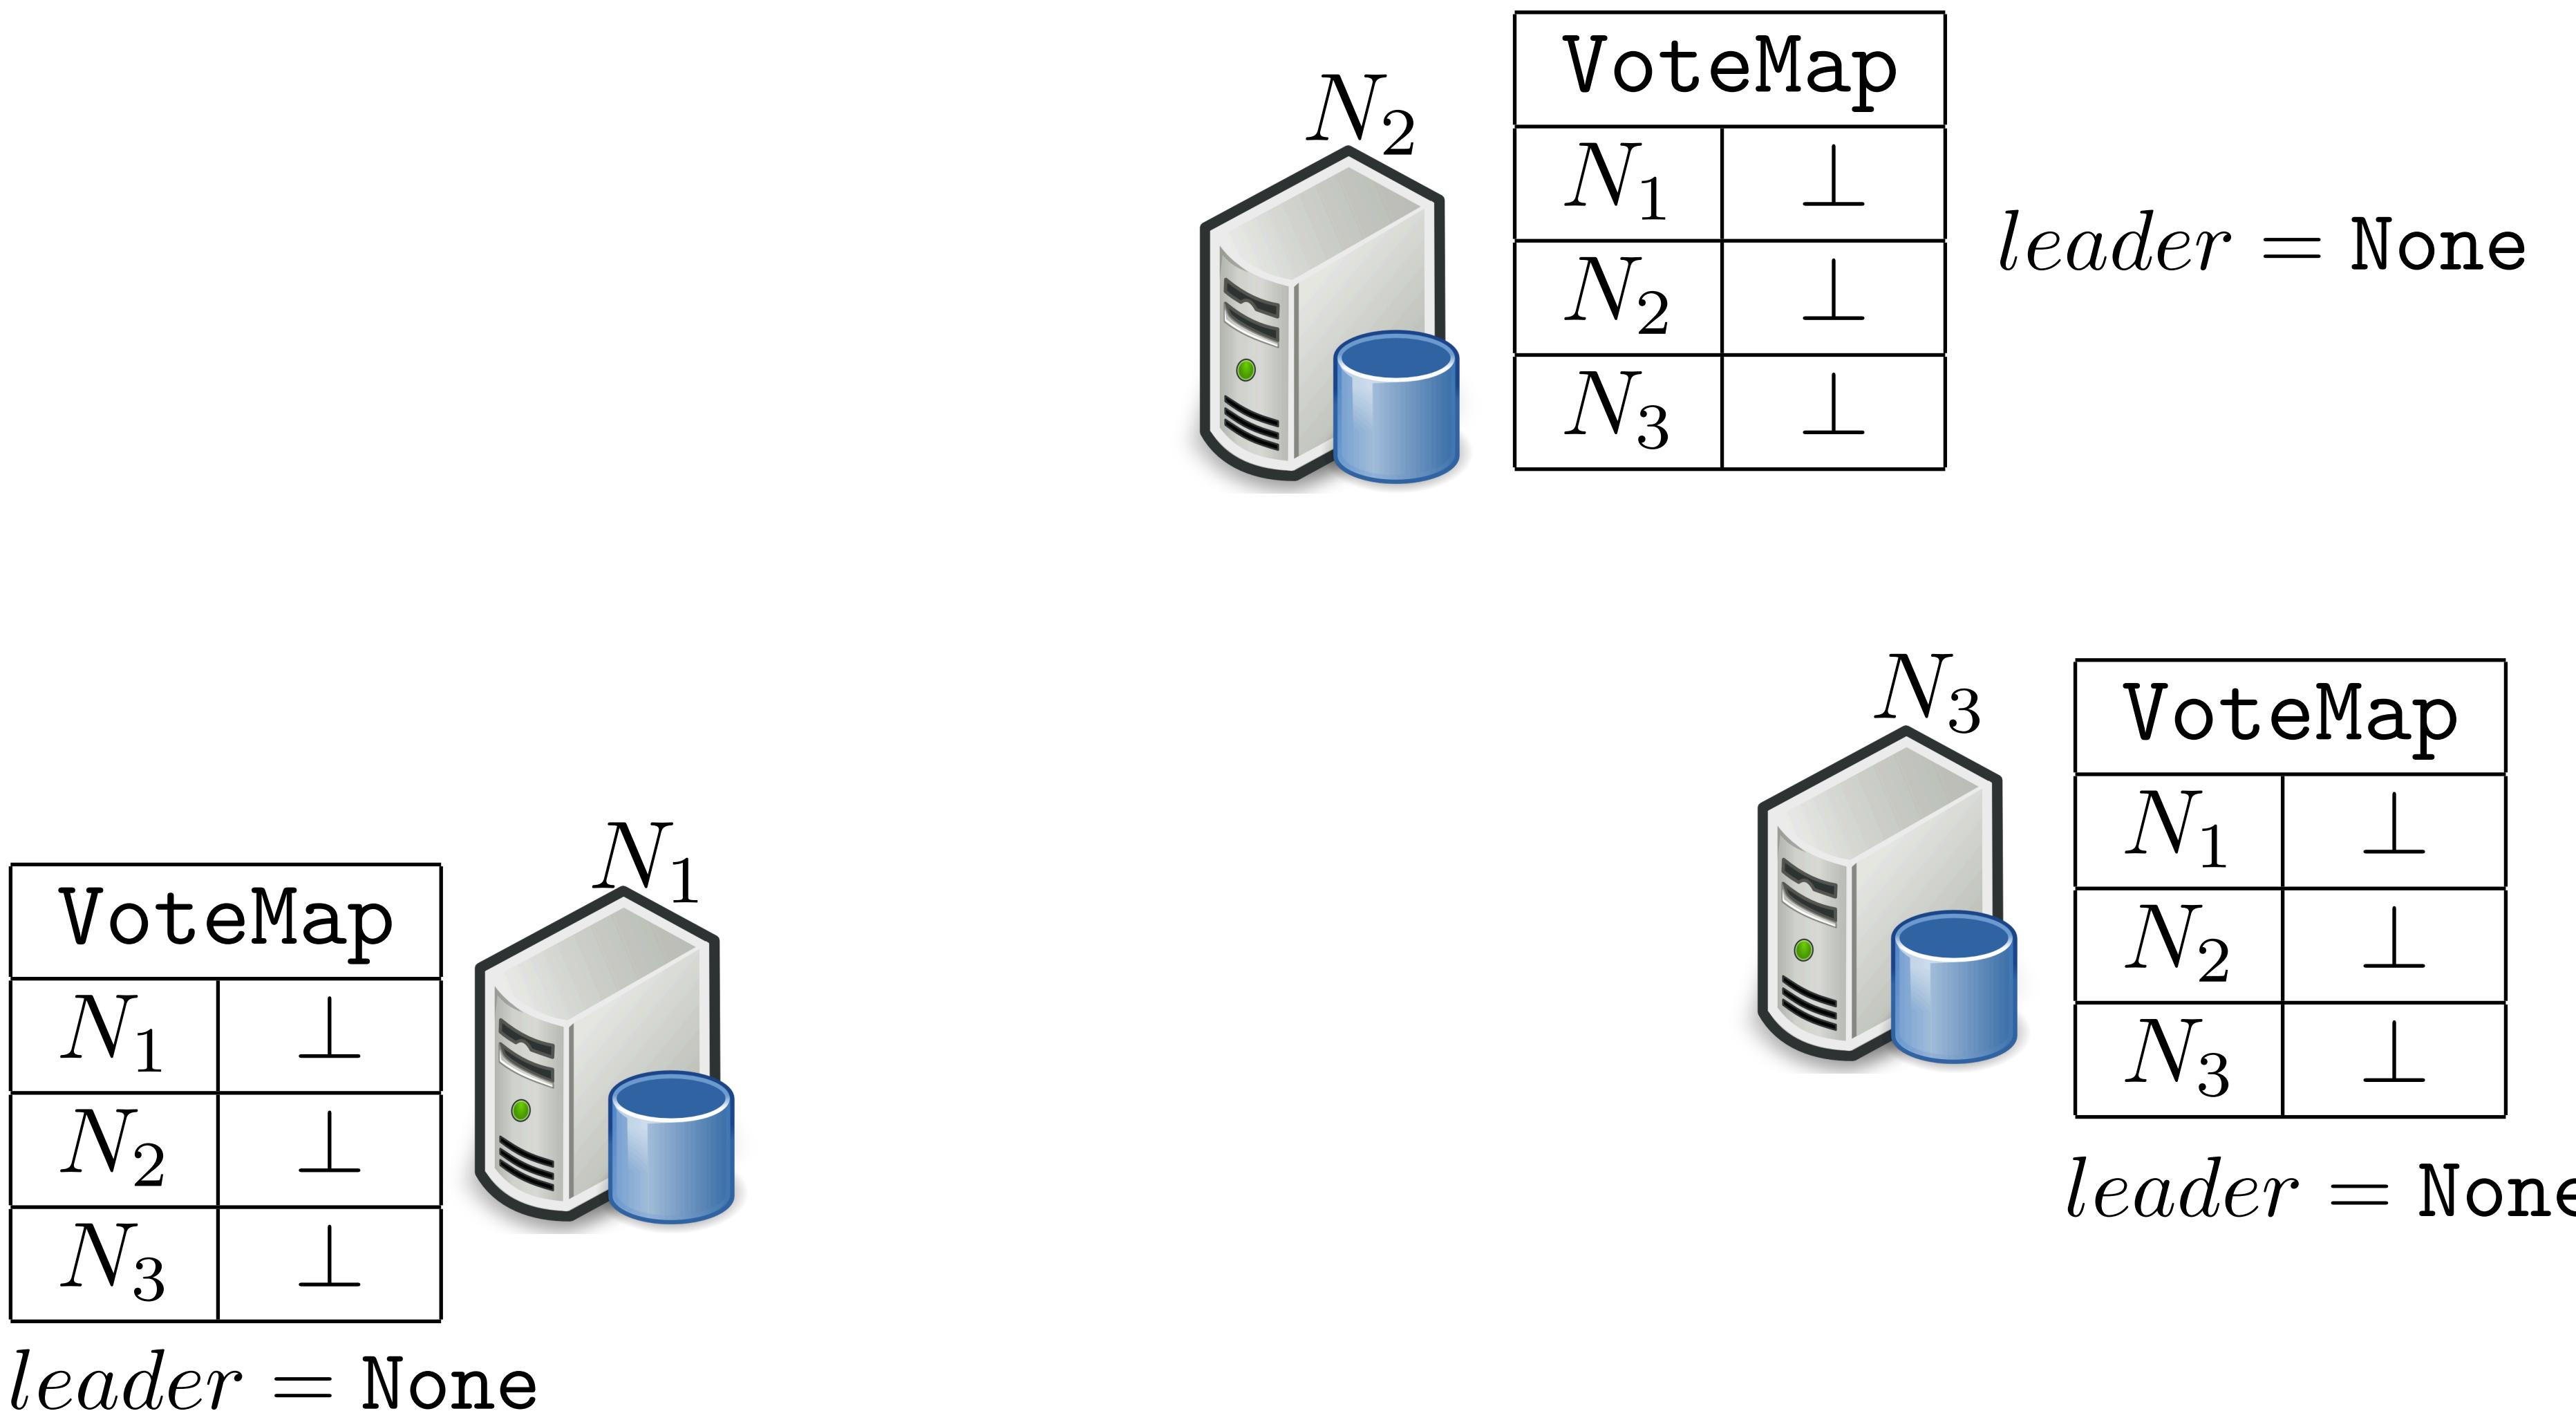
Let's try a *modular* approach of building ■RSMs



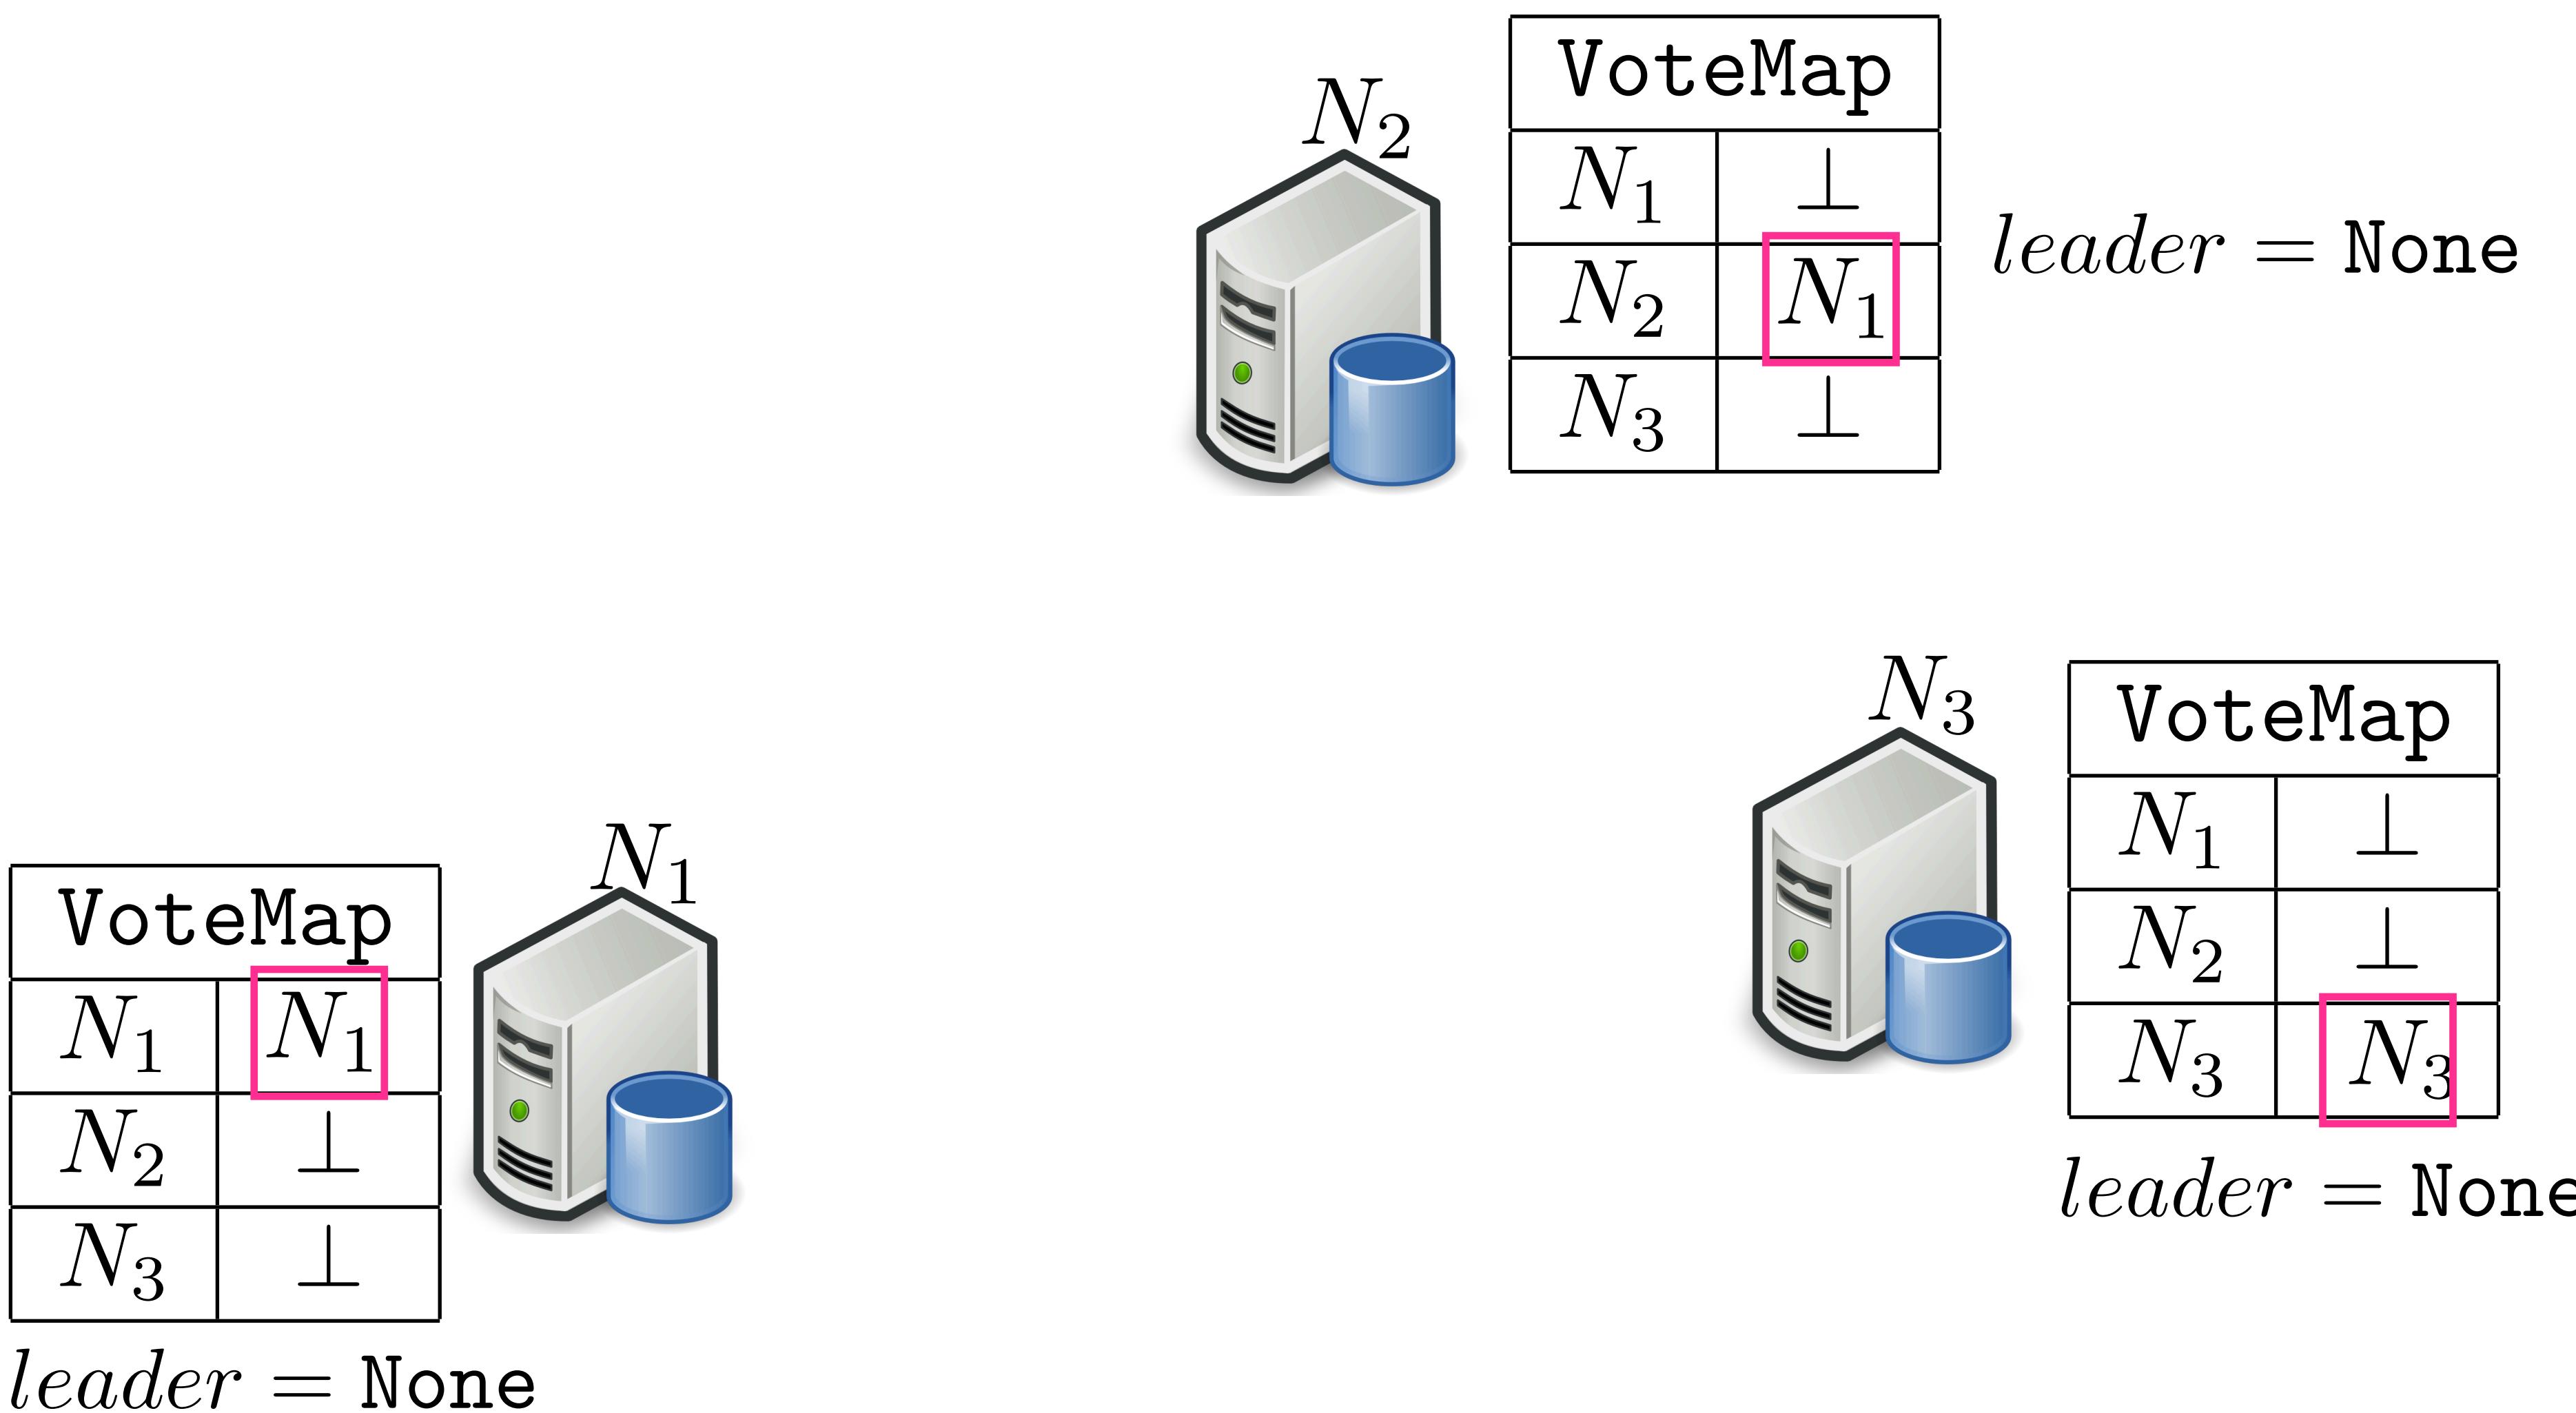
# Working with ♦RSM



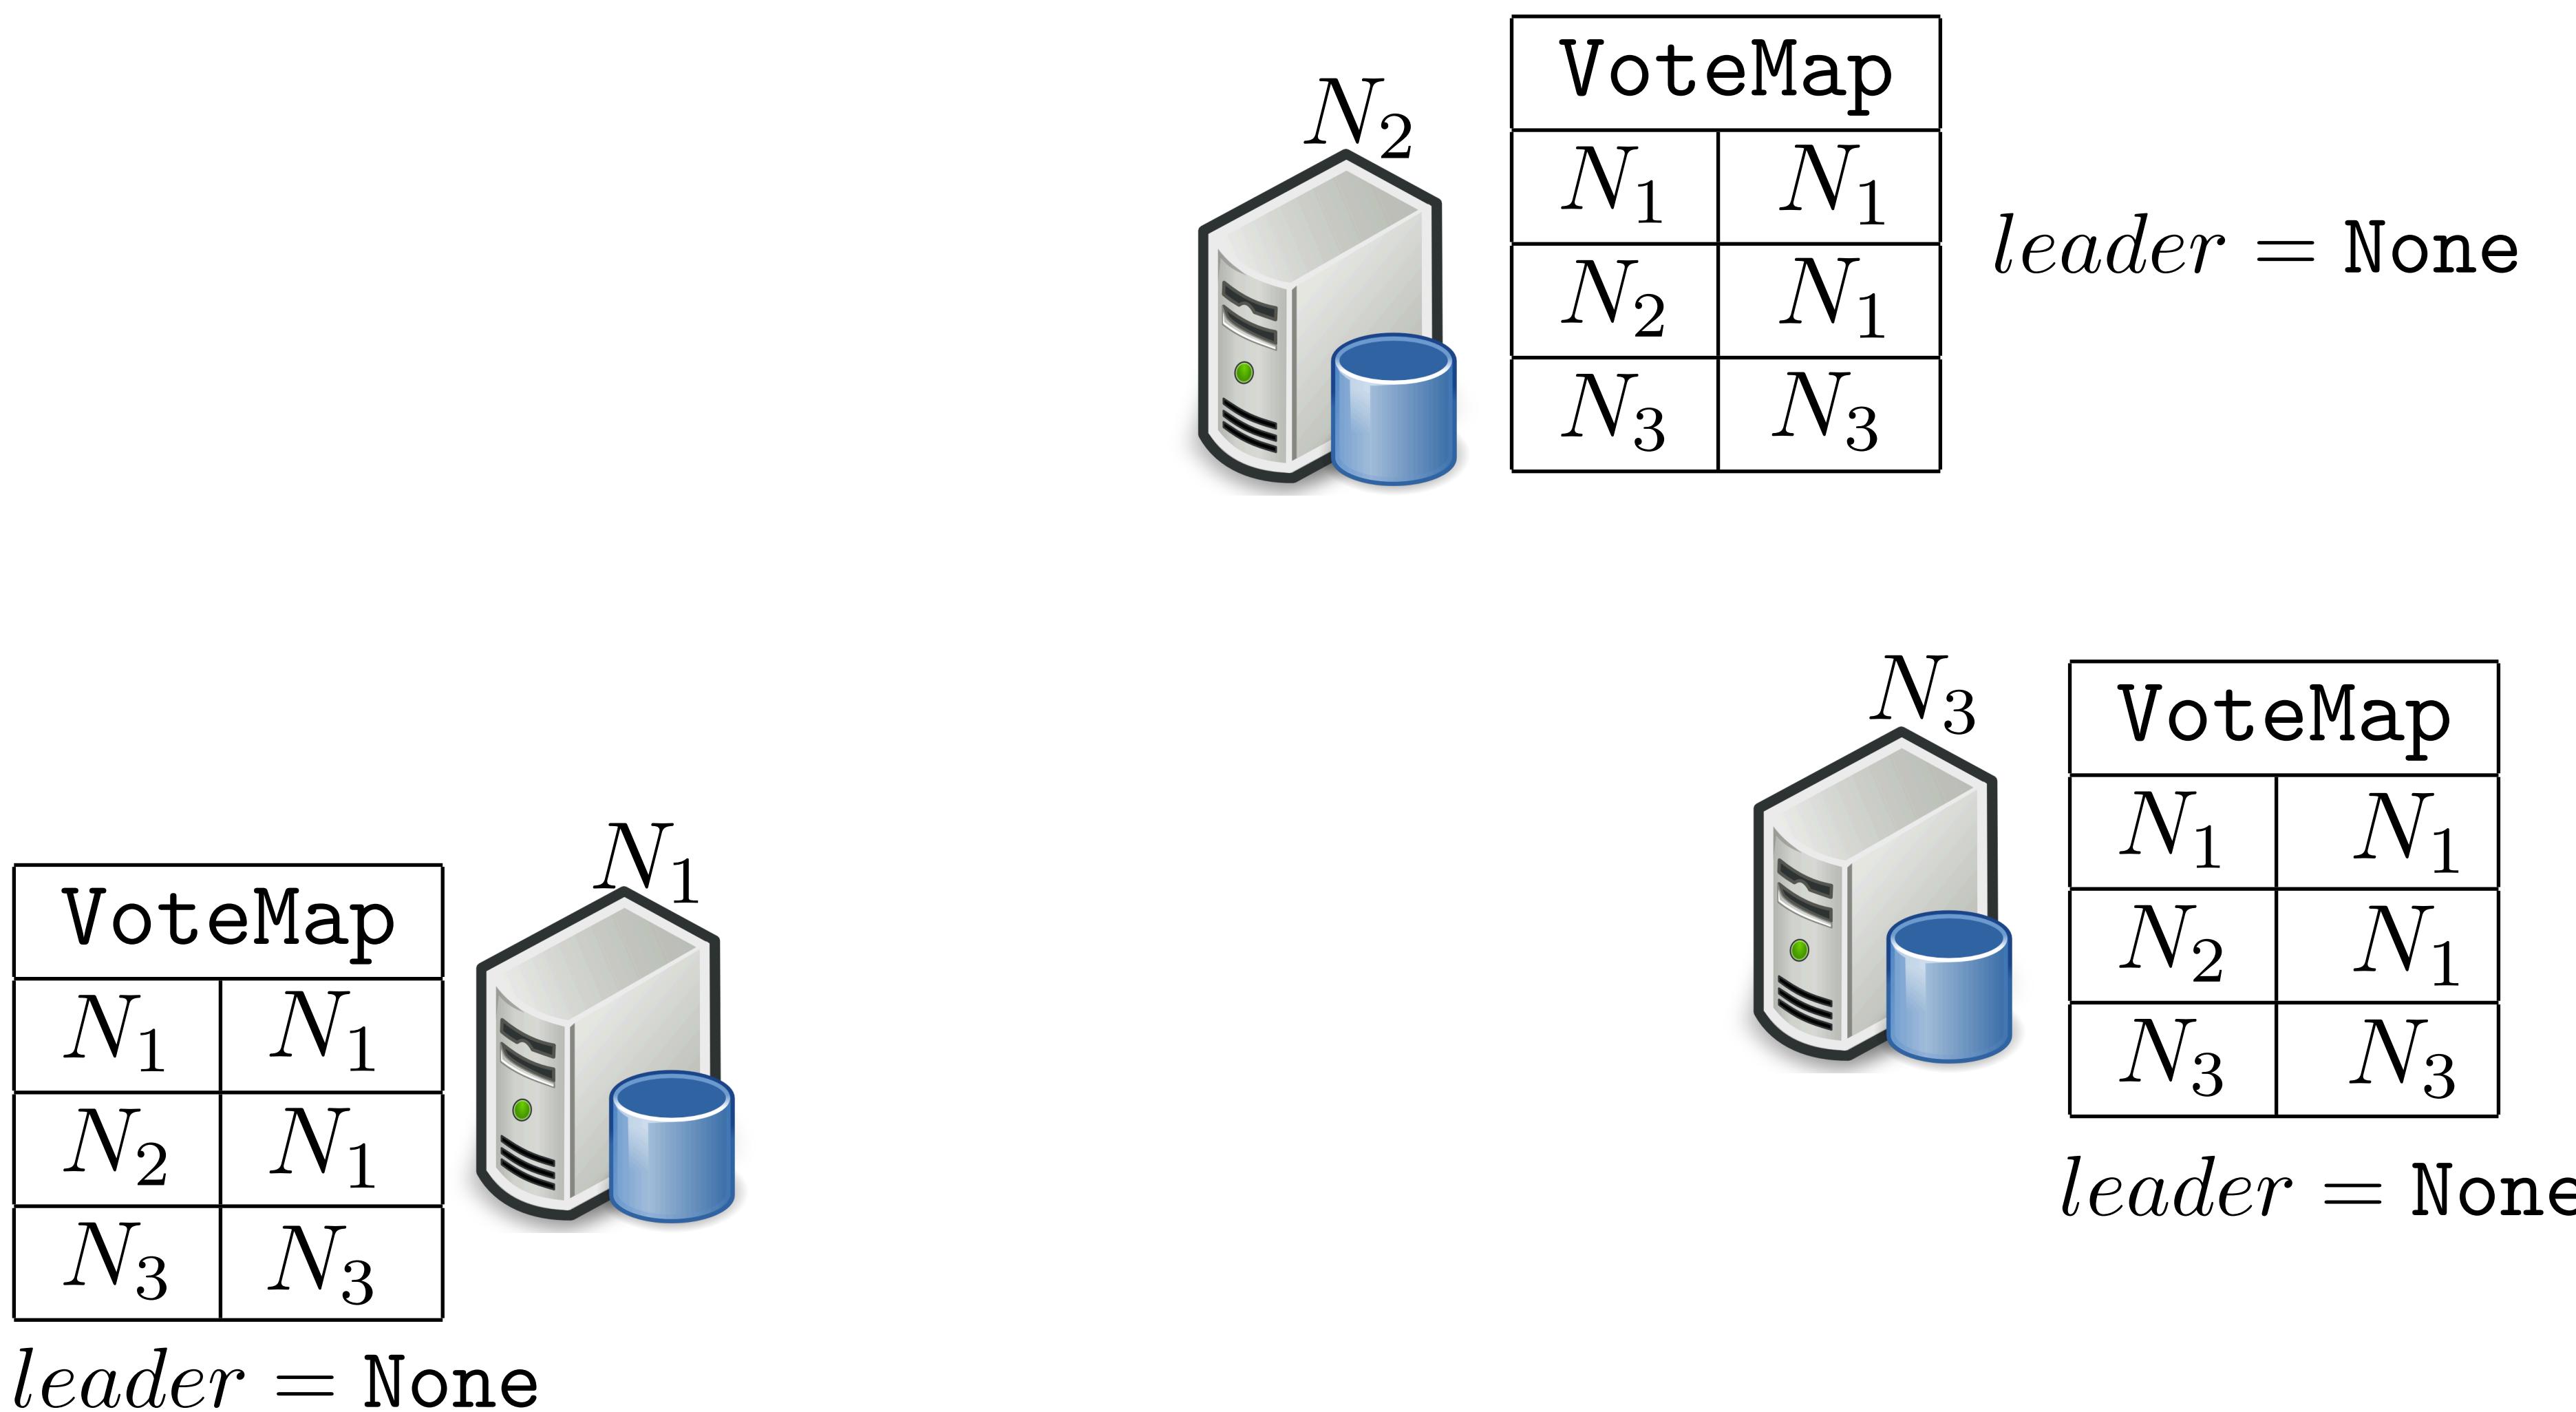
# Leader Election on ♦RSM



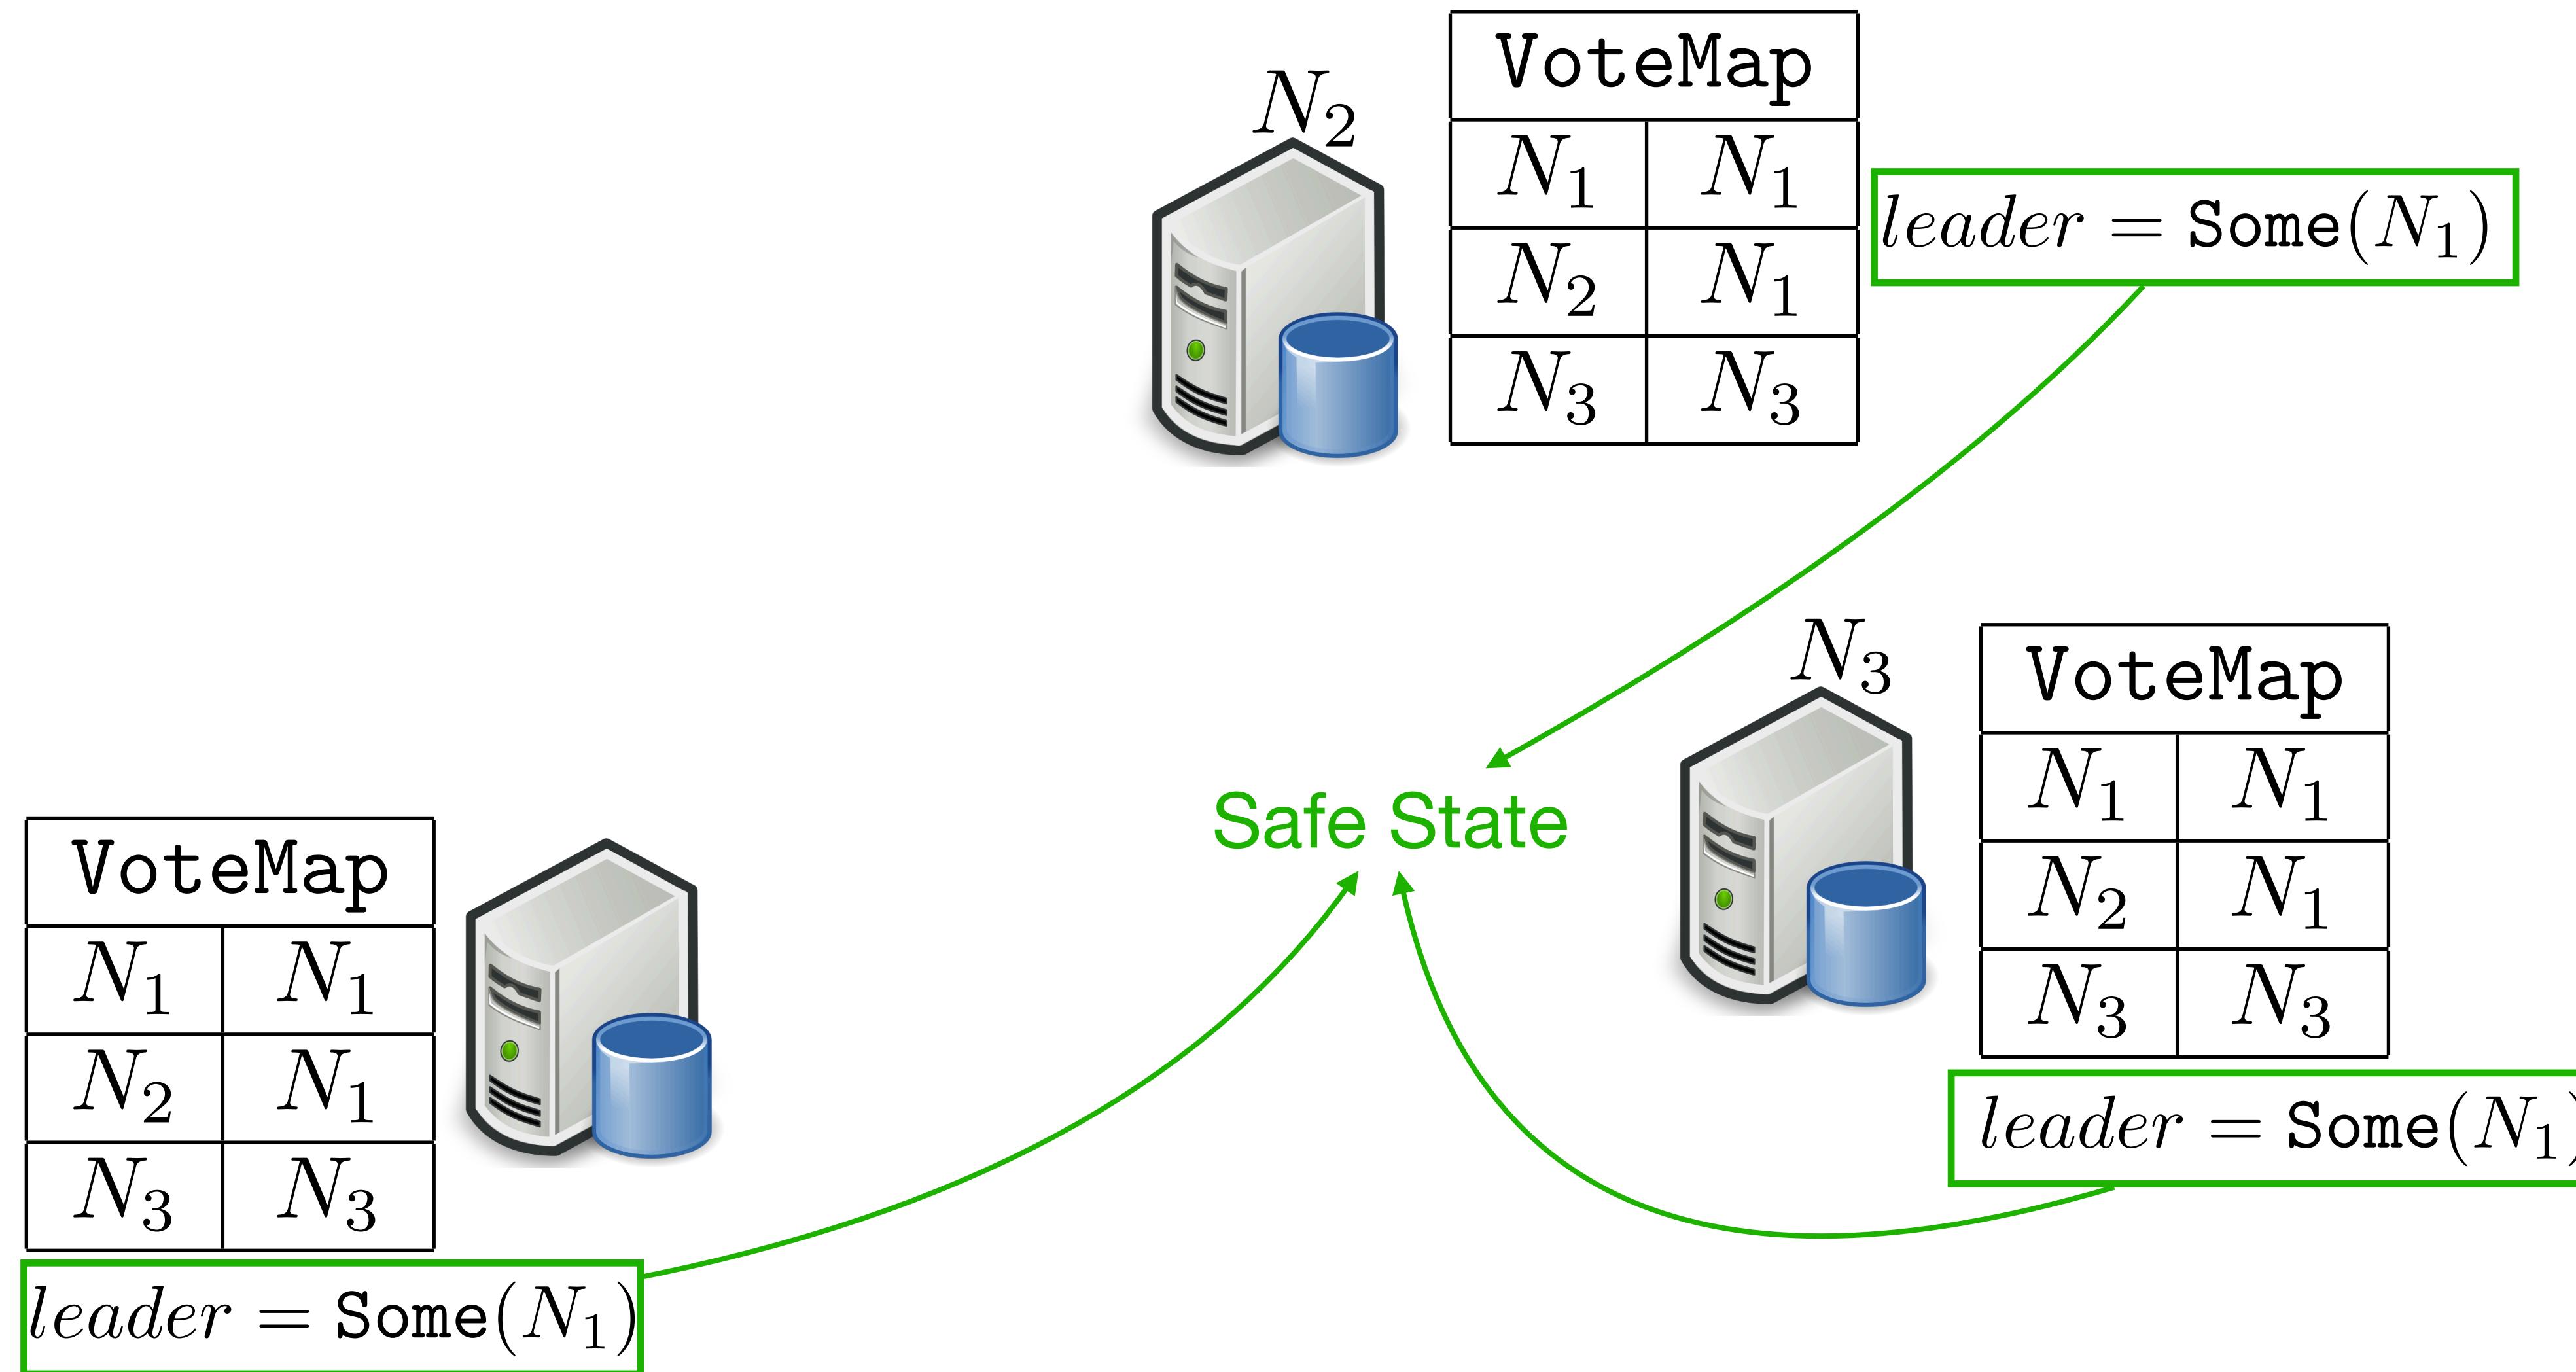
# Leader Election on ♦RSM



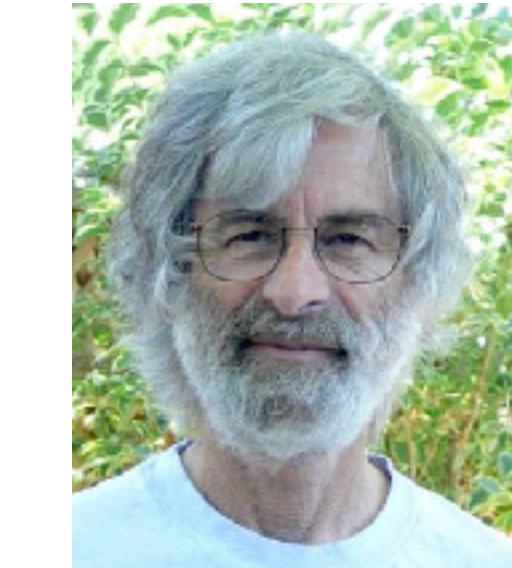
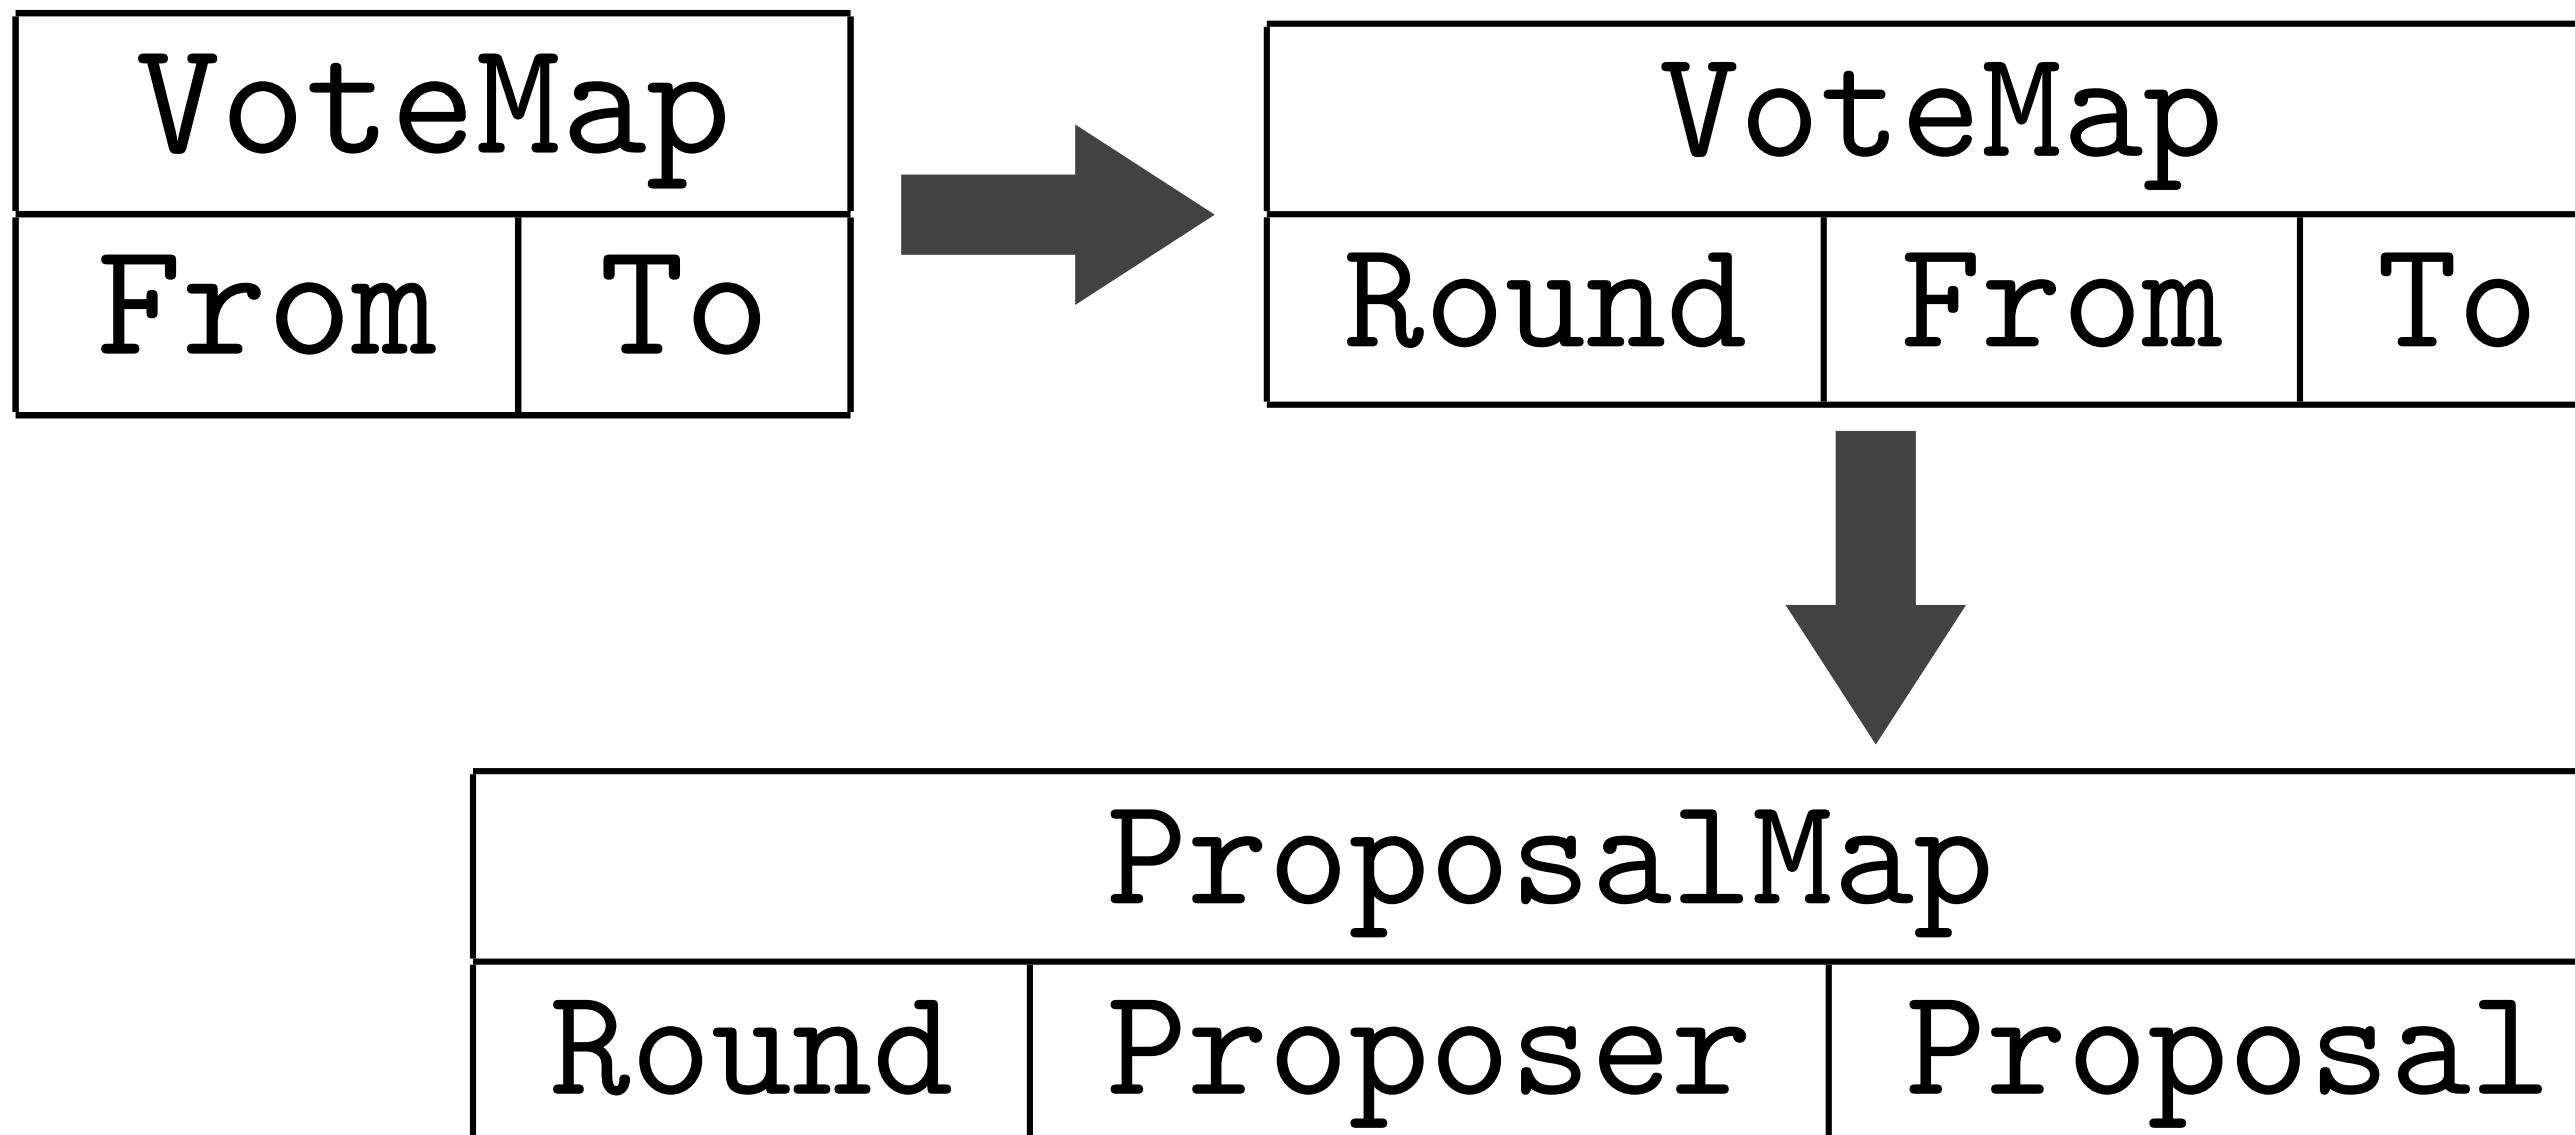
# Leader Election on ♦RSM



# Leader Election on ♦RSM



# Generalized Leader Election on ♦RSM



Paxos Made Simple

Leslie Lamport

01 Nov 2001



Wednesday, May 21, 2025 at 3:52 PM



MS

○ Marc Shapiro -- at work <marc.shapiro@mit.edu>  
To: Gowtham Kaki

[External email - use caution]

Thanks for your talk today. It was inspiring. Thanks to you, I finally understood Paxos!  
Your explanation was so simple.

Marc

# A Proof System for ♦RSM-based Consensus

**THEOREM 4.3 (SUP VERIFICATION).** *A replicated update system RS has stable update precondition  $\Phi^s$  if the following verification conditions hold:*

## Initiation Safety

$$\forall N \in \text{Nid}_{\text{RS}}. \quad \forall E \in \text{Effect}_{\text{RS}}. \quad \forall S \in \text{AState}_{\text{RS}}. \quad \forall V_1, V_2 \in \text{Val}_{\text{RS}}. \quad \forall f^a \in \text{Action}. \\ f^a(N, V_1, S) = \langle V_2, E \rangle \implies \langle N, E, S \rangle \in \Phi^s$$

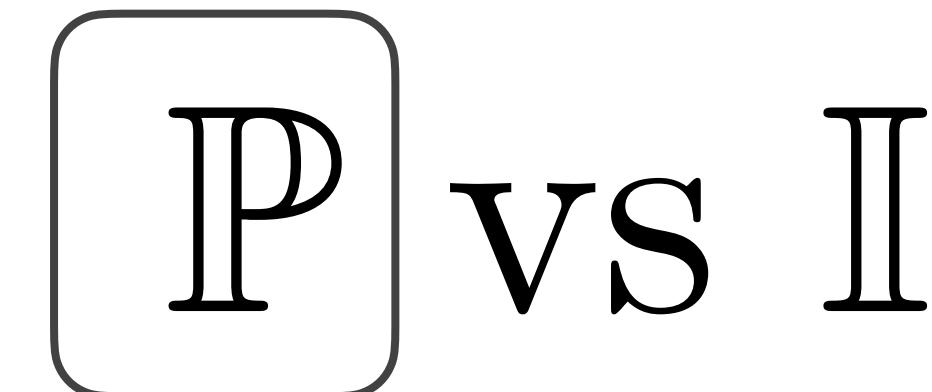
## Racing-State Stability

$$\forall N_1, N_2 \in \text{Nid}_{\text{RS}}. \quad \forall E_1, E_2 \in \text{Effect}_{\text{RS}}. \quad \forall S, S' \in \text{AState}_{\text{RS}}. \\ \langle N_1, E_1, S \rangle \in \Phi^s \wedge \langle N_2, E_2, S \rangle \in \Phi^s \wedge f_{\text{RS}}^h(E_2, S) = S' \implies \langle N_1, E_1, S' \rangle \in \Phi^s$$

## Racing-State Commutativity

$$\forall N_1, N_2 \in \text{Nid}_{\text{RS}}. \quad \forall E_1, E_2 \in \text{Effect}_{\text{RS}}. \quad \forall S, S'_1, S''_1, S'_2, S''_2 \in \text{AState}_{\text{RS}}. \\ (\langle N_1, E_1, S \rangle \in \Phi^s \wedge \langle N_2, E_2, S \rangle \in \Phi^s \wedge f_{\text{RS}}^h(E_1, S) = S'_1 \\ \wedge f_{\text{RS}}^h(E_2, S'_1) = S''_1 \wedge f_{\text{RS}}^h(E_2, S) = S'_2 \wedge f_{\text{RS}}^h(E_1, S'_1) = S''_2) \implies S''_1 = S''_2$$

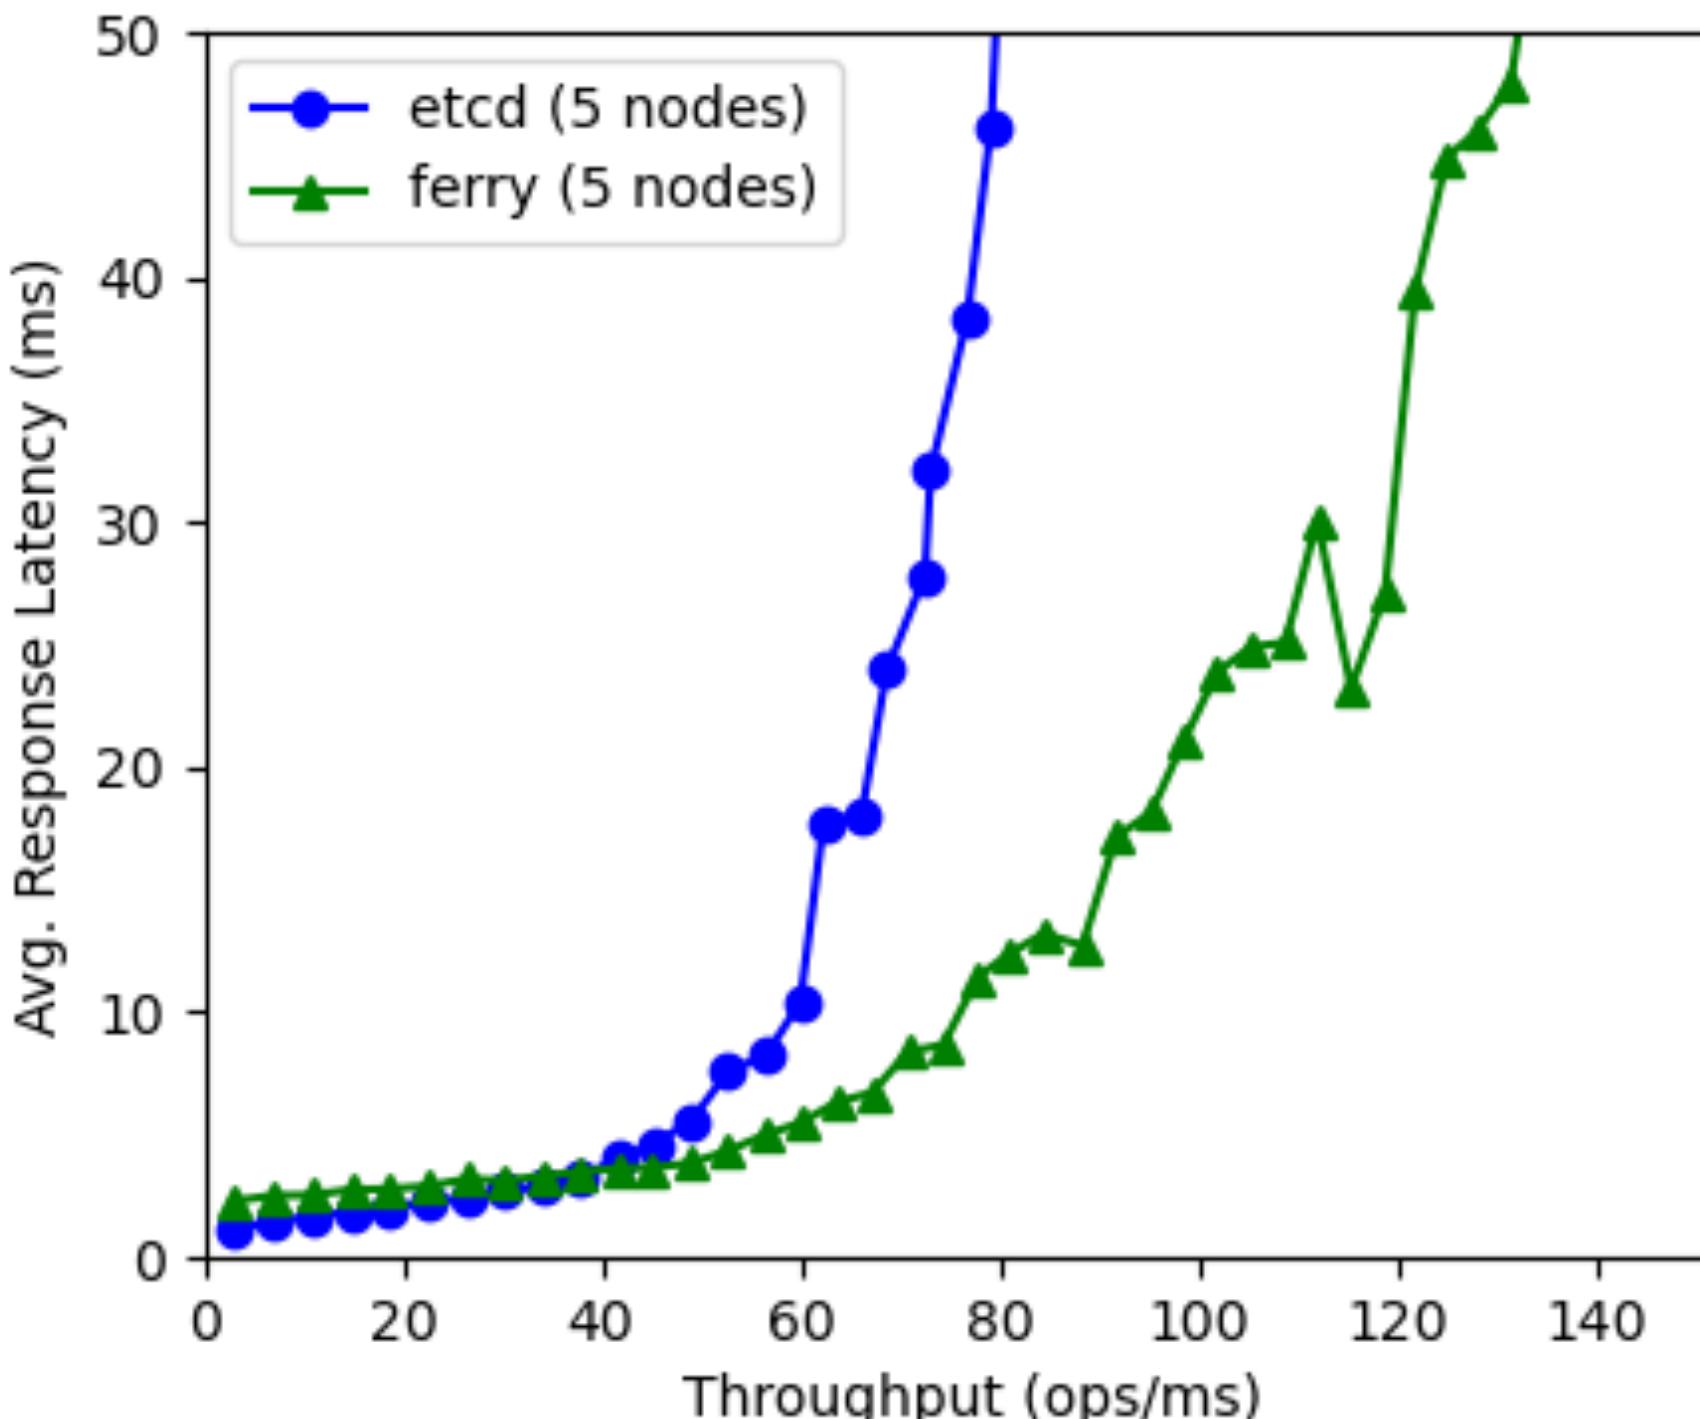
Uses Strong Update Preconditions  
instead of Inductive Invariants



**THEOREM 4.8 (DECIDABILITY FRIENDLY).** *The Def. 4.6 conditions are decidable as long as the relations  $R^a$ ,  $R^h$ ,  $R^d$ , and  $\Phi^s$ , taken together, are in the Extended EPR logic fragment.*

# Super-V

- A Haskell library to implement and formally verify distributed protocols.
- Case study: Raft log replication protocol.
  - Implementation: ~400 LOC.
  - Specification ( $P$ ): ~110 LOC. No quantifier alternation. Extended EPR.
- Automatically verified using CVC5 in under 50mins on Intel Core I5 laptop.



- Performs similarly or better than etcd-raft!



**Bolt-On Strong Consistency: Specification, Implementation, and Verification**

NICHOLAS V. LEWCHENKO, University of Colorado Boulder, USA

GOWTHAM KAKI, University of Colorado Boulder, USA

BOR-YUH EVAN CHANG\*, University of Colorado Boulder, USA and Amazon, USA

[OOPSLA 2025]

## 2. Improving Generality: From *safety* to *security*

### Contributors



Kirby Linvill



Eric Wustrow



Evan Chang

### Tools



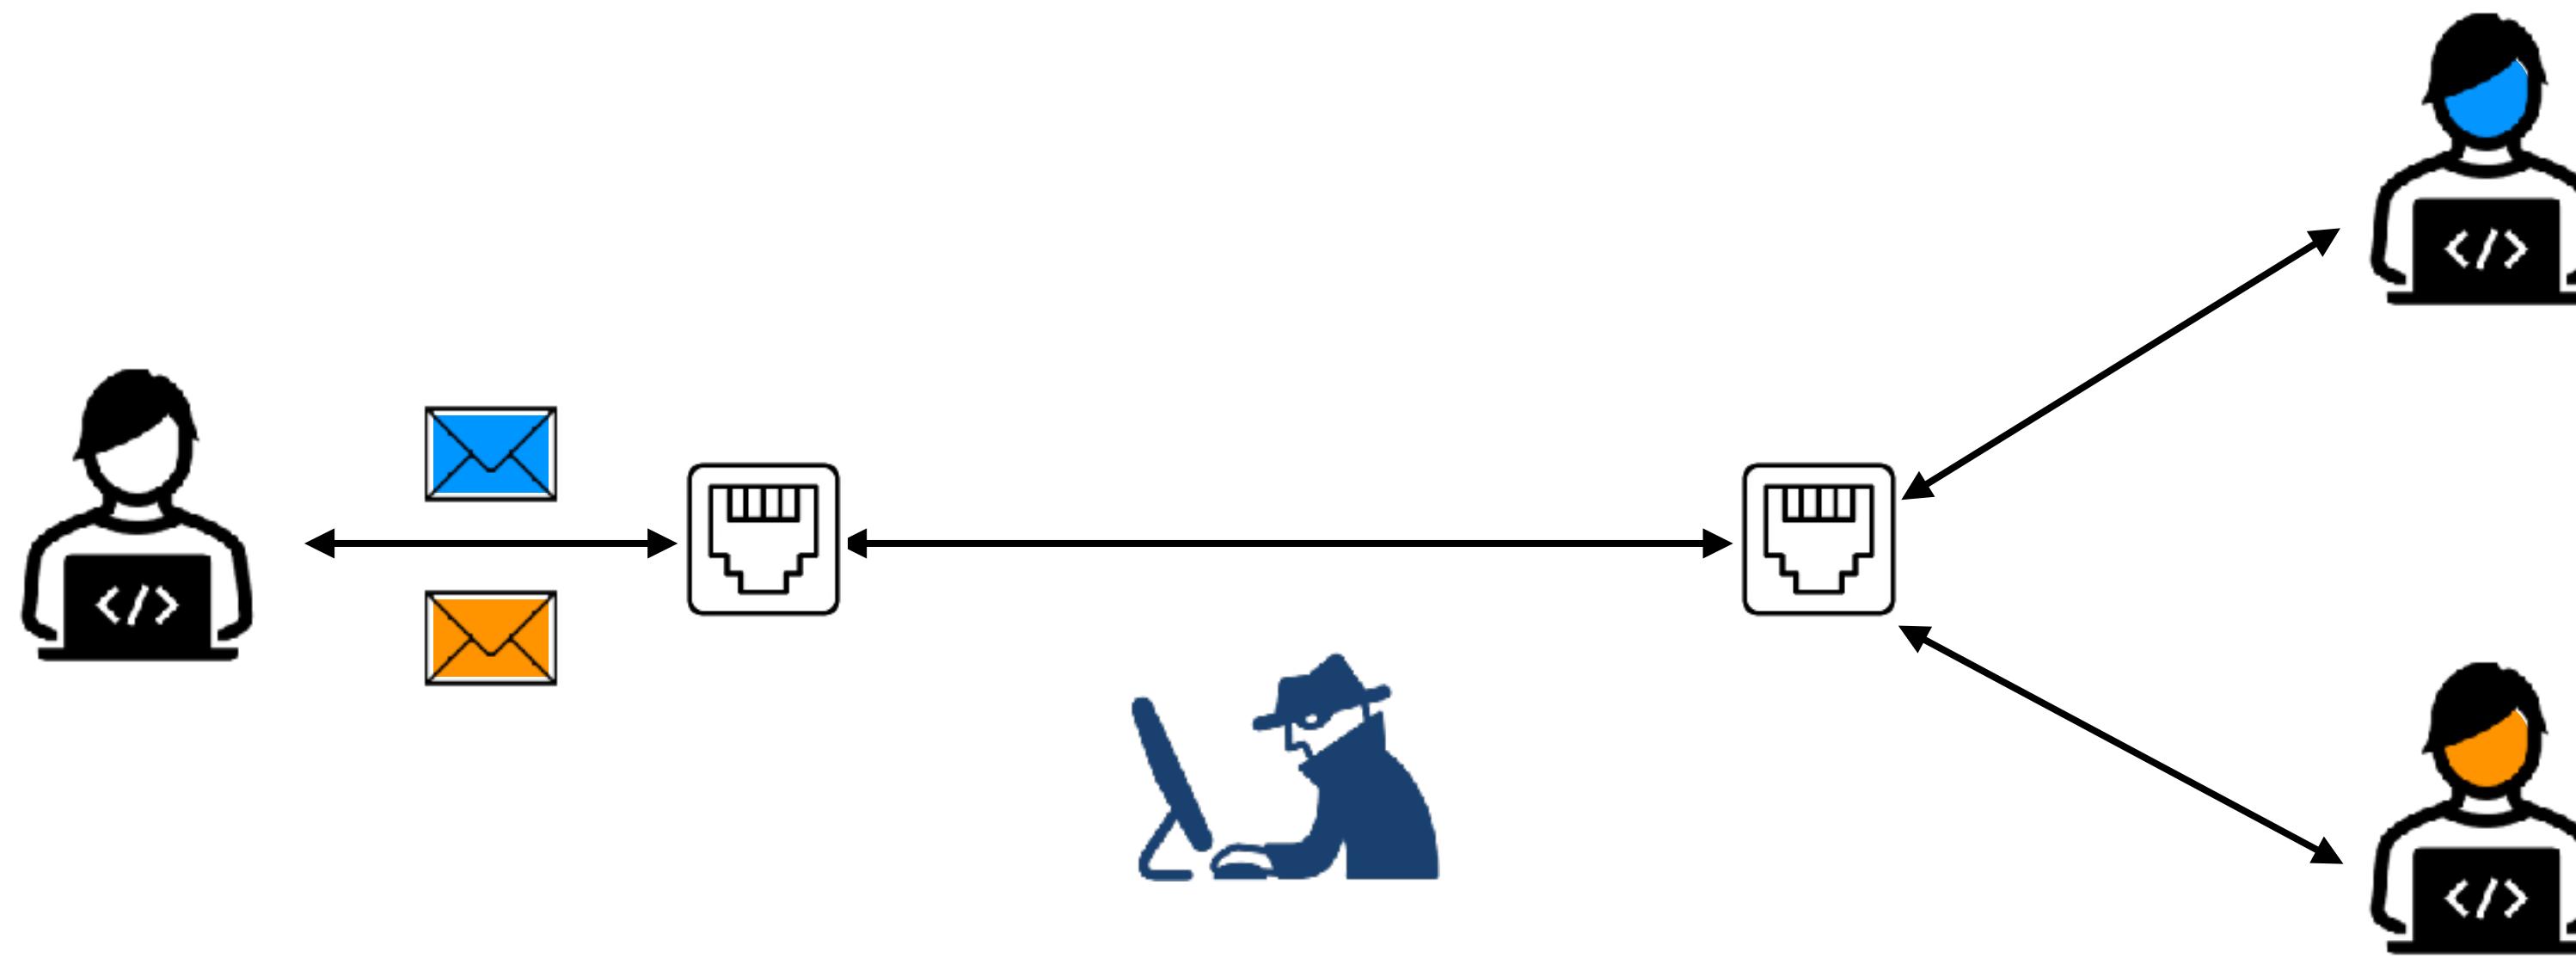
**Waldo**

<https://github.com/cuplv/Waldo>

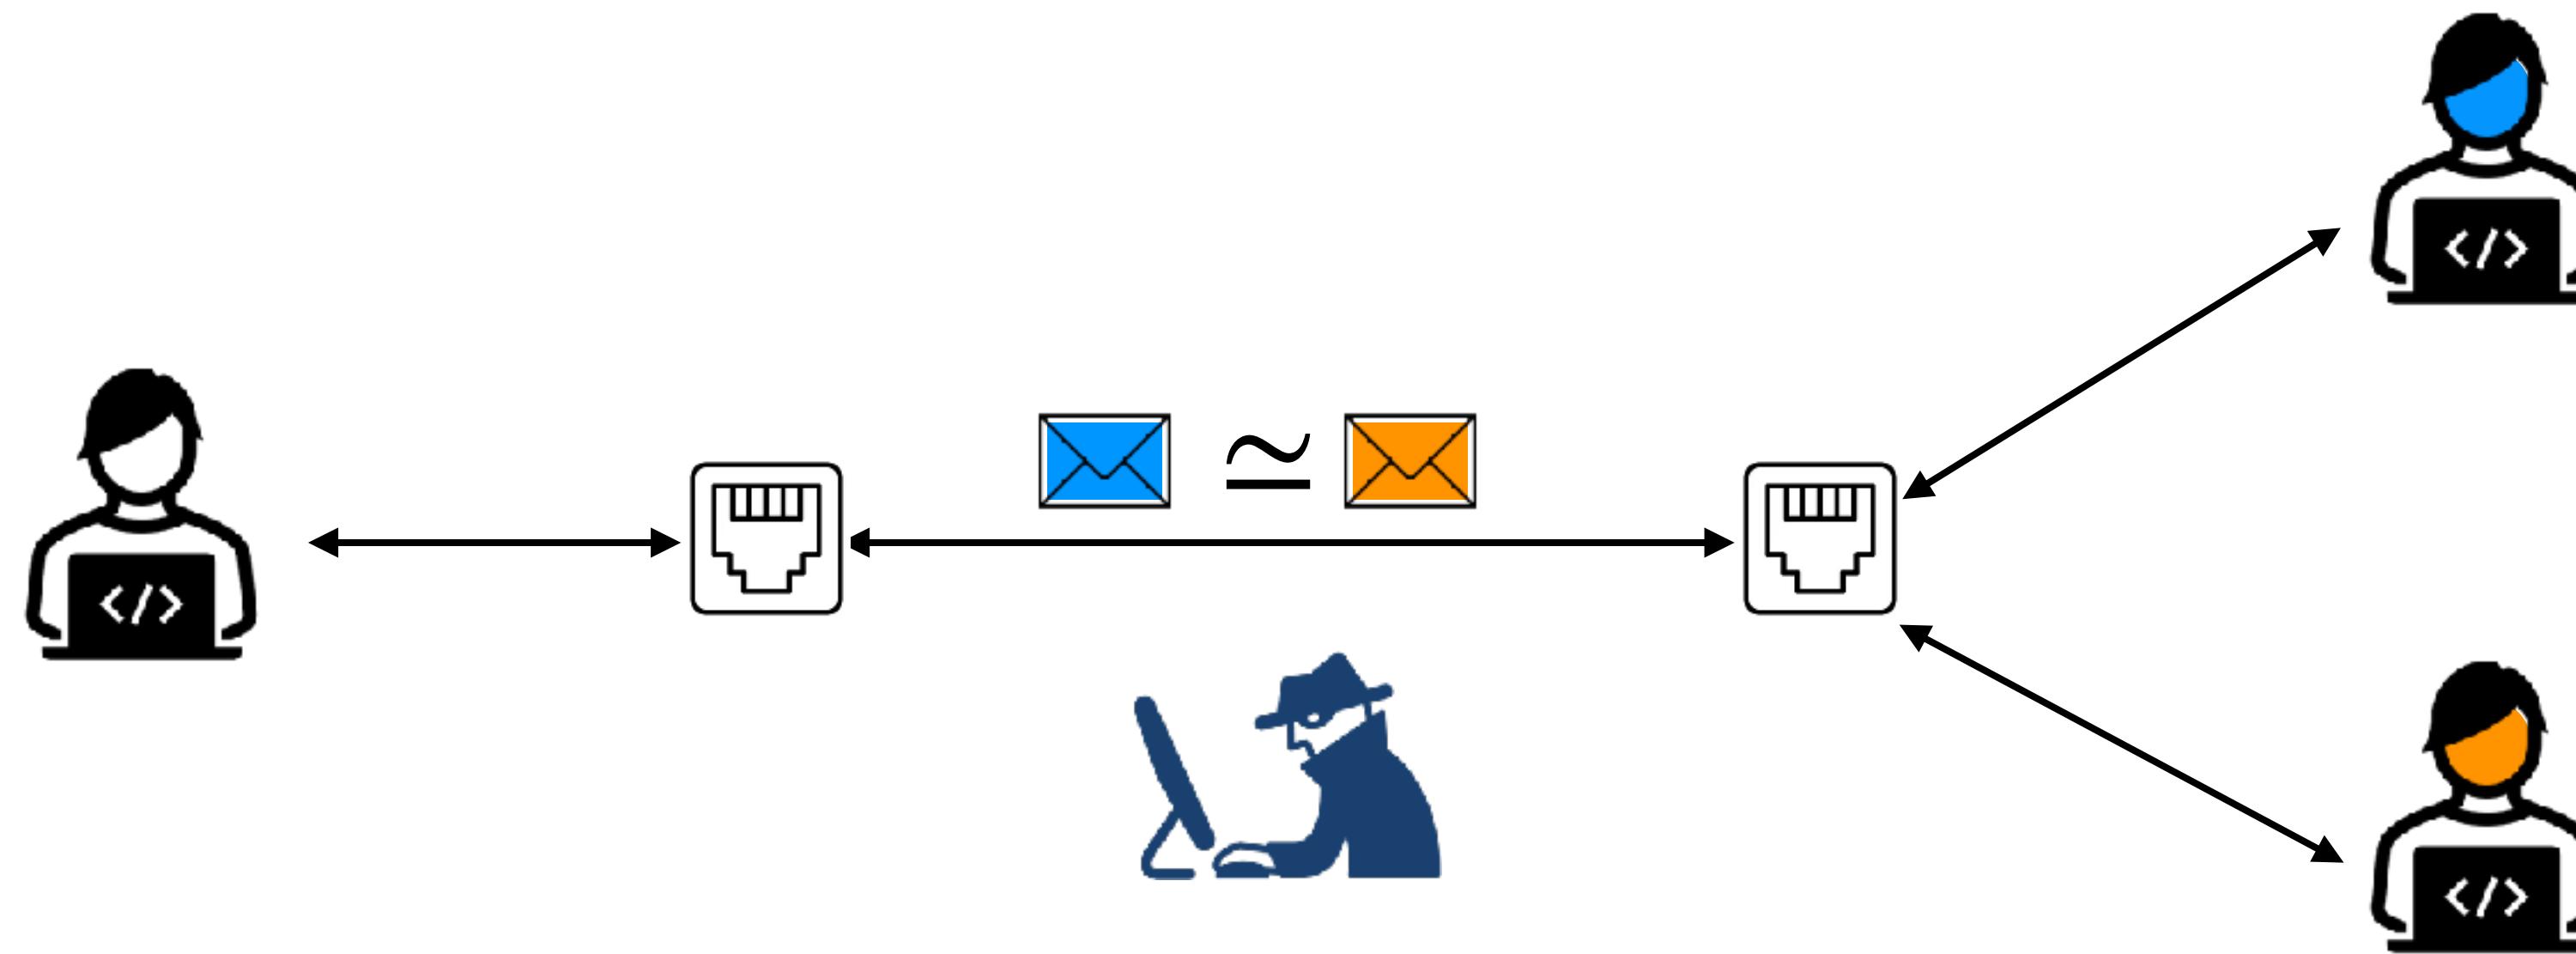
### Publications

- Kirby Linvill, Gowtham Kaki, and Eric Wustrow. **Verifying Indistinguishability of Privacy-Preserving Protocols.** Proceedings of the ACM on Programming Languages (PACMPL), issue **OOPSLA2 2023**, Article No.: 273, Pages 1442 - 1469.
- Kirby Linvill, Gowtham Kaki, and Bor-Yuh Evan Chang. **A Bisimulation Proof Framework for Symbolic Active Indistinguishability.** Pre-print.

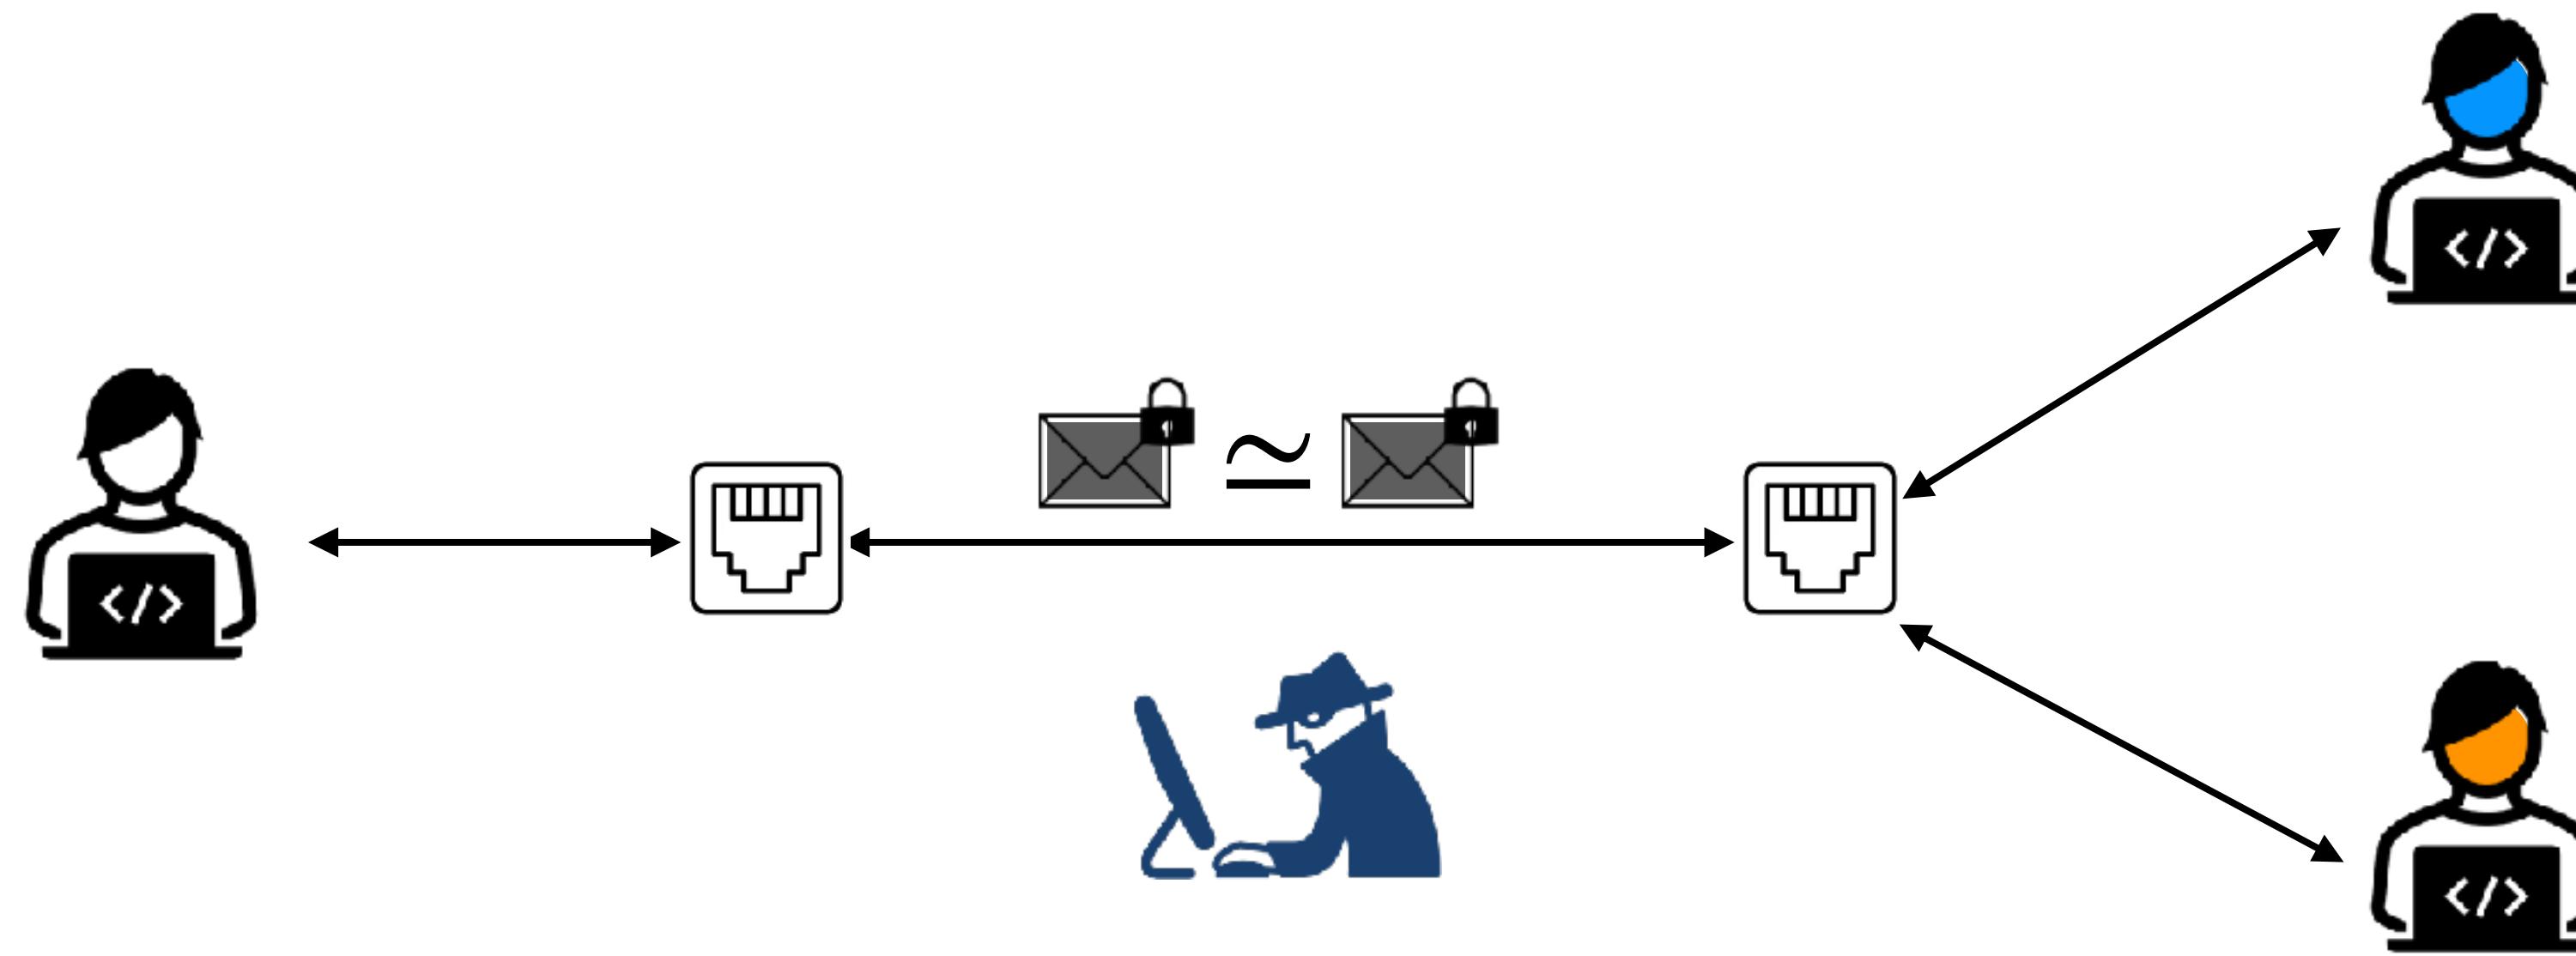
# Privacy on Public Networks



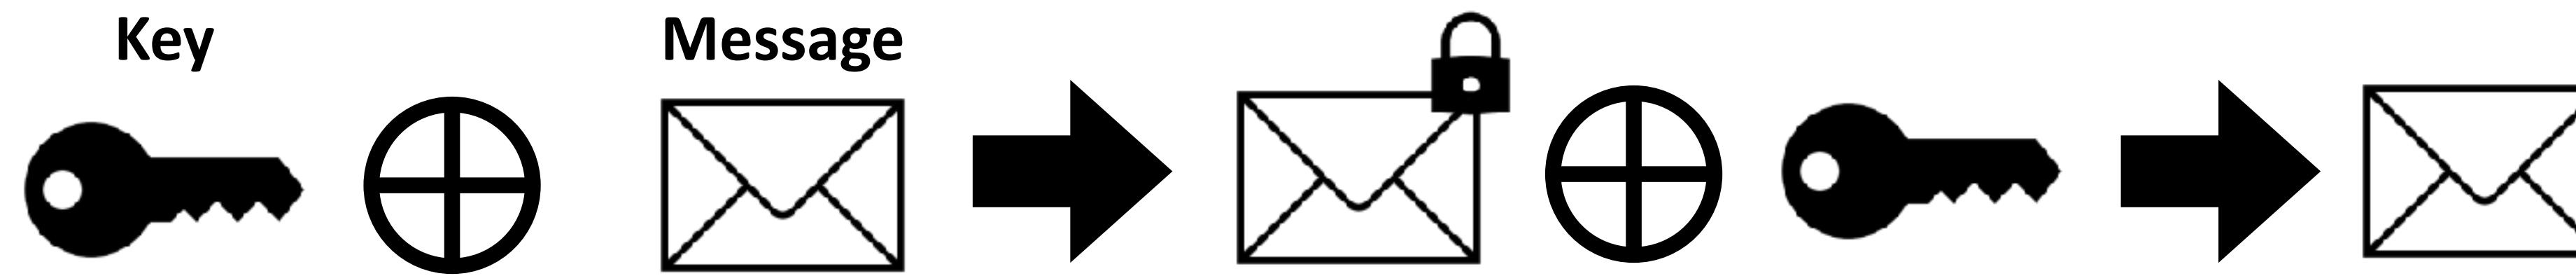
# Privacy on Public Networks



# Privacy on Public Networks



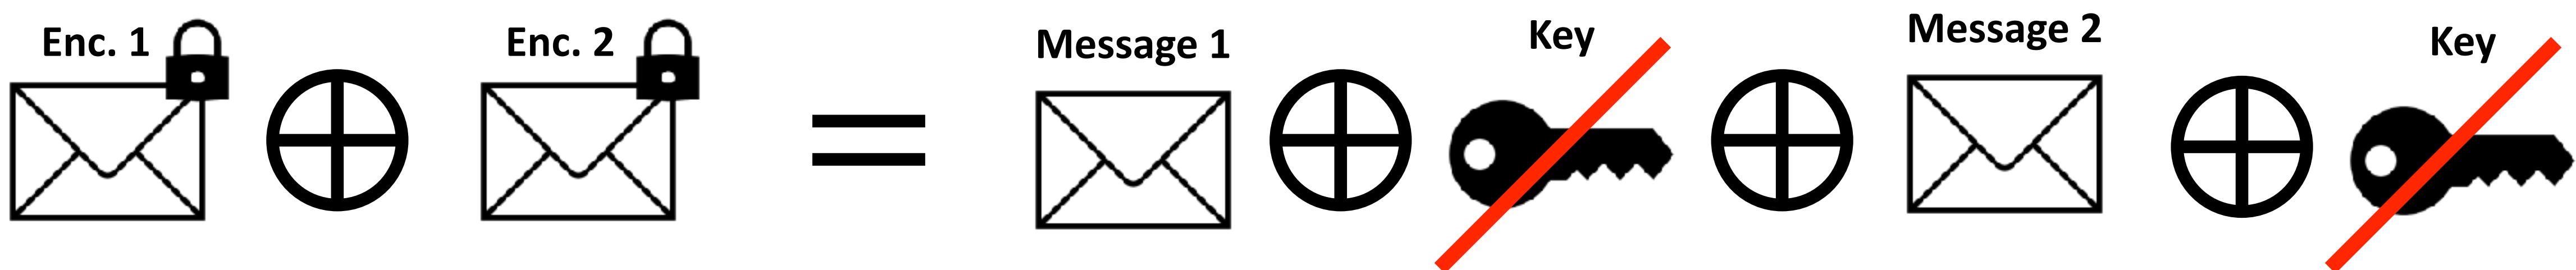
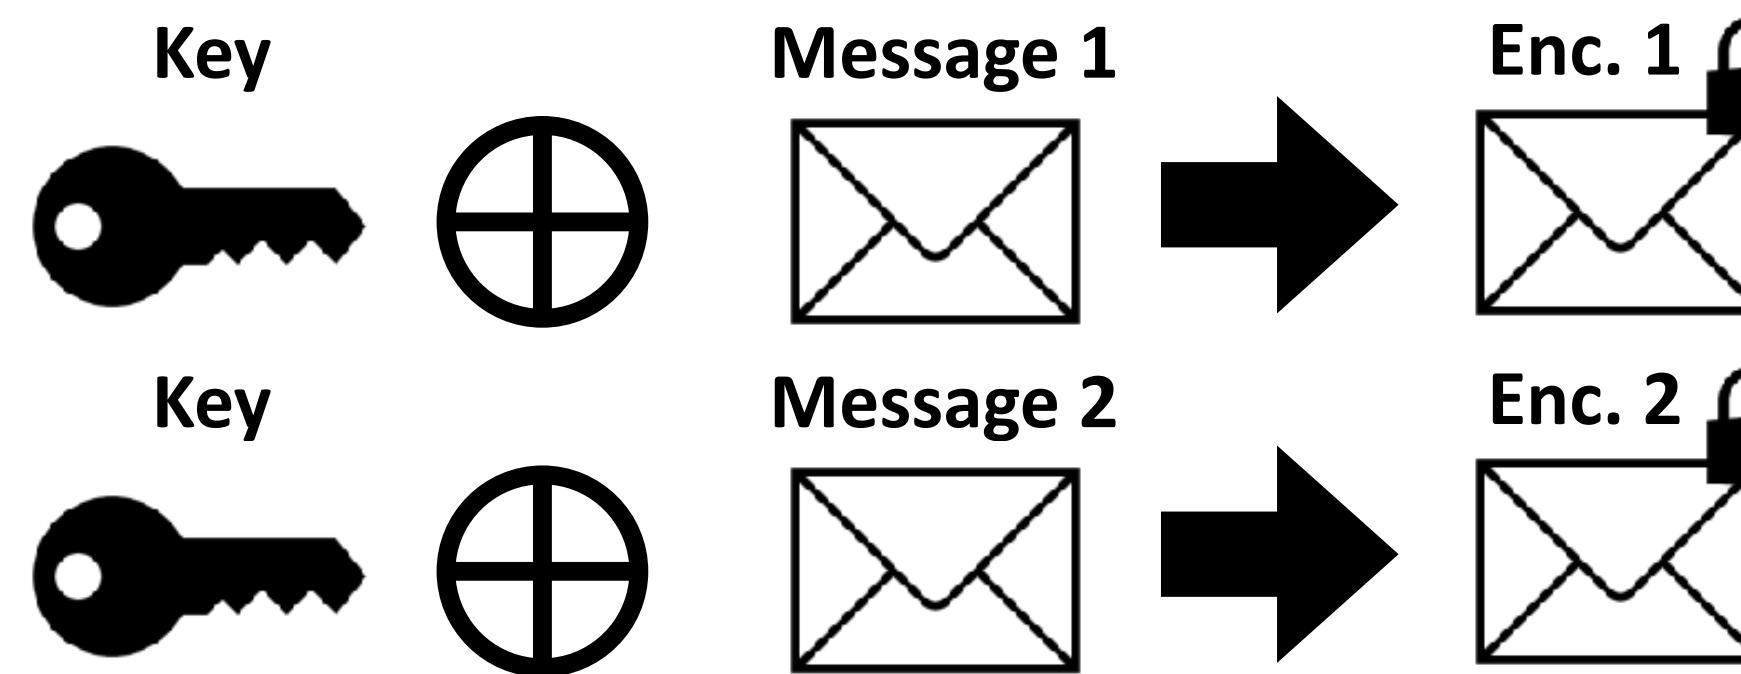
# One-Time Pad (OTP) Encryption



1. Generate a key such that  $\text{length(key)} = \text{length(message)}$
2. Exclusive Or (XOR) the key and message to encrypt
3. XOR the key and message to decrypt

# Improper One-Time Pad (OTP) Encryption

What if we re-use the key?



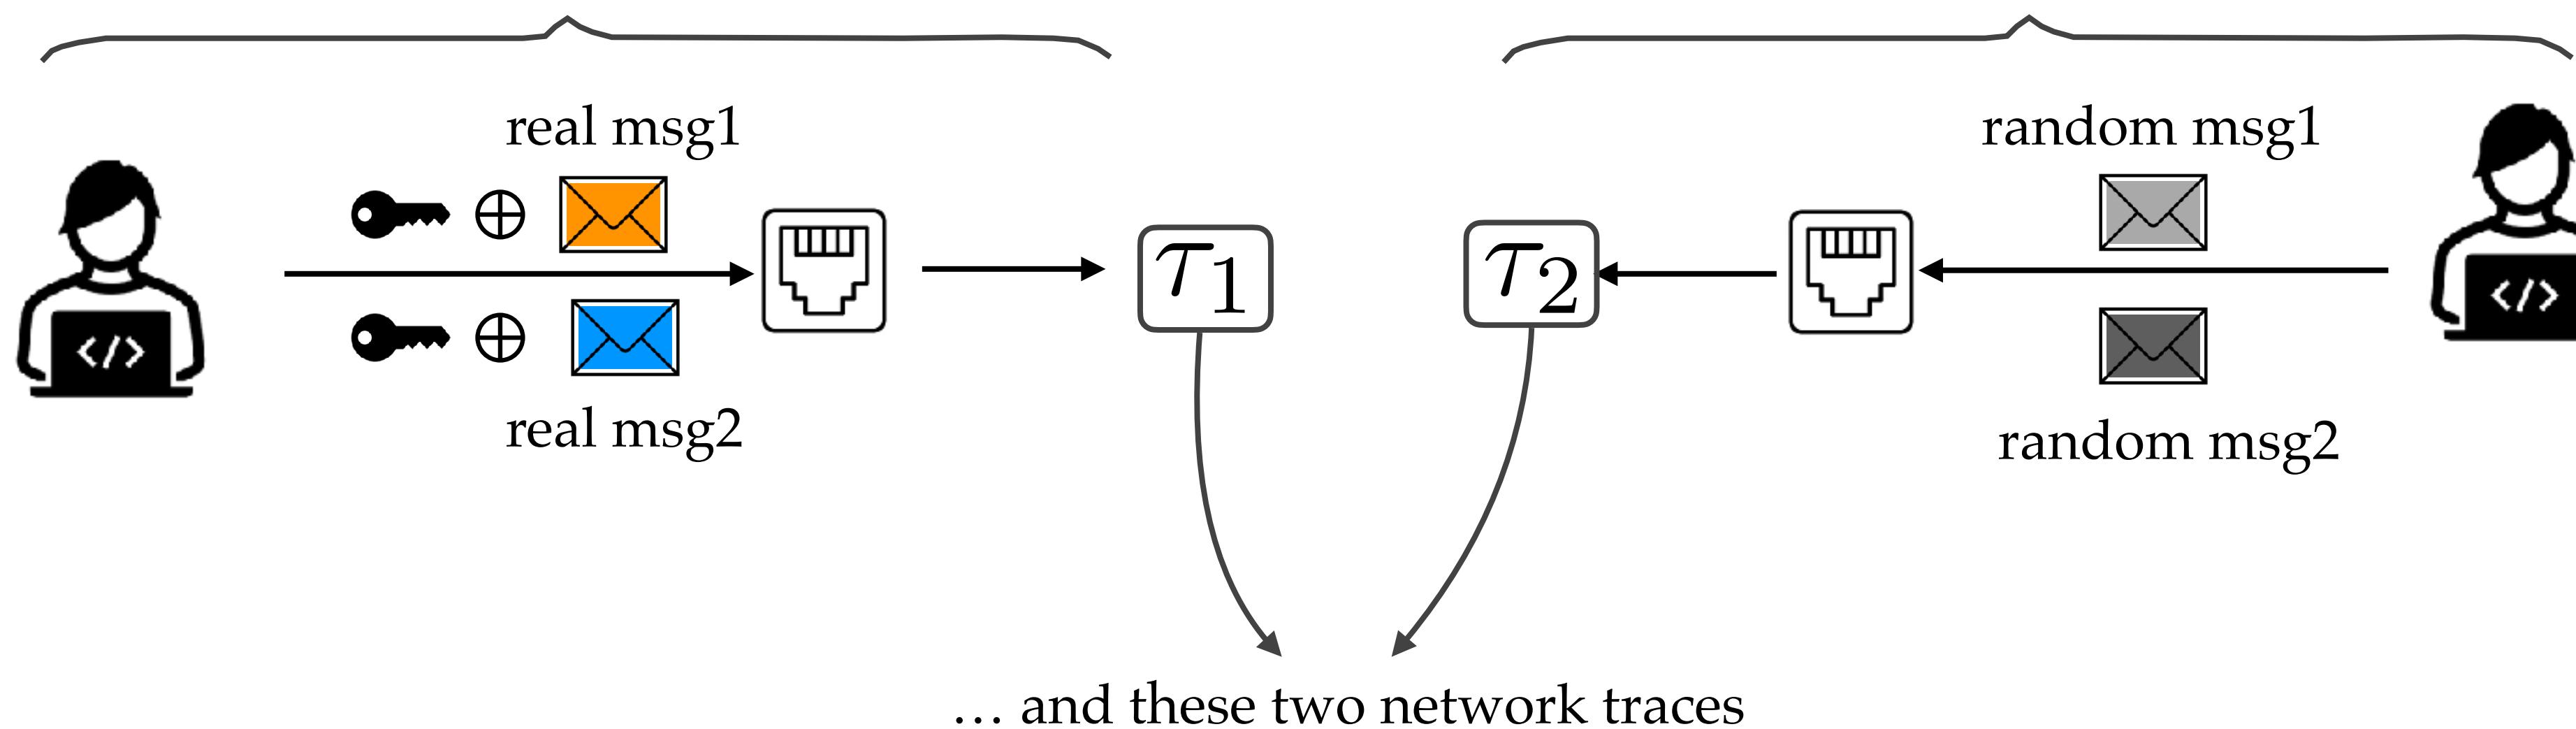
Since  $x \oplus x = 0$ ,  $x \oplus 0 = x$ , and  $\oplus$  is commutative

*This leaks which bits of Message 1 and Message 2 differ*

# Improper One-Time Pad (OTP) Encryption



Attacker can distinguish between these two scenarios ...



# Symbolic Model of Cryptography

An axiomatization of cryptography to enable symbolic reasoning

- Plaintexts and ciphertexts are uninterpreted symbols.
- Encryption and decryption are uninterpreted functions related by:

$$\forall m. \forall k_1, k_2. \text{decrypt}(\text{encrypt}(m, k_1), k_2) = m \Leftrightarrow k_1 = k_2$$

- Decryption is the only operation defined on ciphertexts.
- The only way to decrypt a ciphertext is by knowing the key, in which case the corresponding plaintext can be *deduced*.
- A pair of symbols,  $S_1$  and  $S_2$ , are *distinguishable* iff there exists a deduction possible on one but not the other. Otherwise they are *indistinguishable*. Likewise with traces of symbols.

In the symbolic model, OTP is a perfect encryption primitive.

⇒

We can “prove” in the symbolic model that  $\tau_1$  and  $\tau_2$  are indistinguishable

# Probabilistic Model of Cryptography

## Verifying Indistinguishability of Privacy-Preserving Protocols

KIRBY LINVILL, University of Colorado Boulder, USA

GOWTHAM KAKI, University of Colorado Boulder, USA

ERIC WUSTROW, University of Colorado Boulder, USA

[OOPSLA 2023]

- Introduces probabilistic model of cryptography.
- Crypto primitives are modeled as functions that sample bitstrings from pre-defined distributions.
- E.g., a perfect block cipher with block size L can be modeled as:

$$c \xleftarrow{\$} \{0, 1\}^L$$

- We can now reason about the (symbolic) probability of generating a trace.
- Two encrypted network traces,  $\tau_1$  and  $\tau_2$ , are indistinguishable iff they are equally likely to have been generated by a protocol execution.
- Alternatively, two protocol executions,  $\pi_1$  and  $\pi_2$ , are indistinguishable iff they are equally likely to generate any given encrypted trace  $\tau$ .

$$\Pr[\pi_1 \Downarrow \tau] = \Pr[\pi_2 \Downarrow \tau]$$

We adapted Probabilistic Relational Hoare Logic (PRHL; Barthe et al, POPL 2009) to prove verification conditions of this form ...

# Probabilistic Model of Cryptography

## Verifying Indistinguishability of Privacy-Preserving Protocols

KIRBY LINVILL, University of Colorado Boulder, USA

GOWTHAM KAKI, University of Colorado Boulder, USA

ERIC WUSTROW, University of Colorado Boulder, USA

[OOPSLA 2023]

$$\Pr[\pi_1 \Downarrow \tau] = \Pr[\pi_2 \Downarrow \tau]$$



Waldo

<https://github.com/cuplv/Waldo>

We adapted Probabilistic Relational Hoare Logic (PRHL; Barthe *et al*, POPL 2009) to prove verification conditions of this form ...

... and built an SMT-aided proof automation framework in F\*.

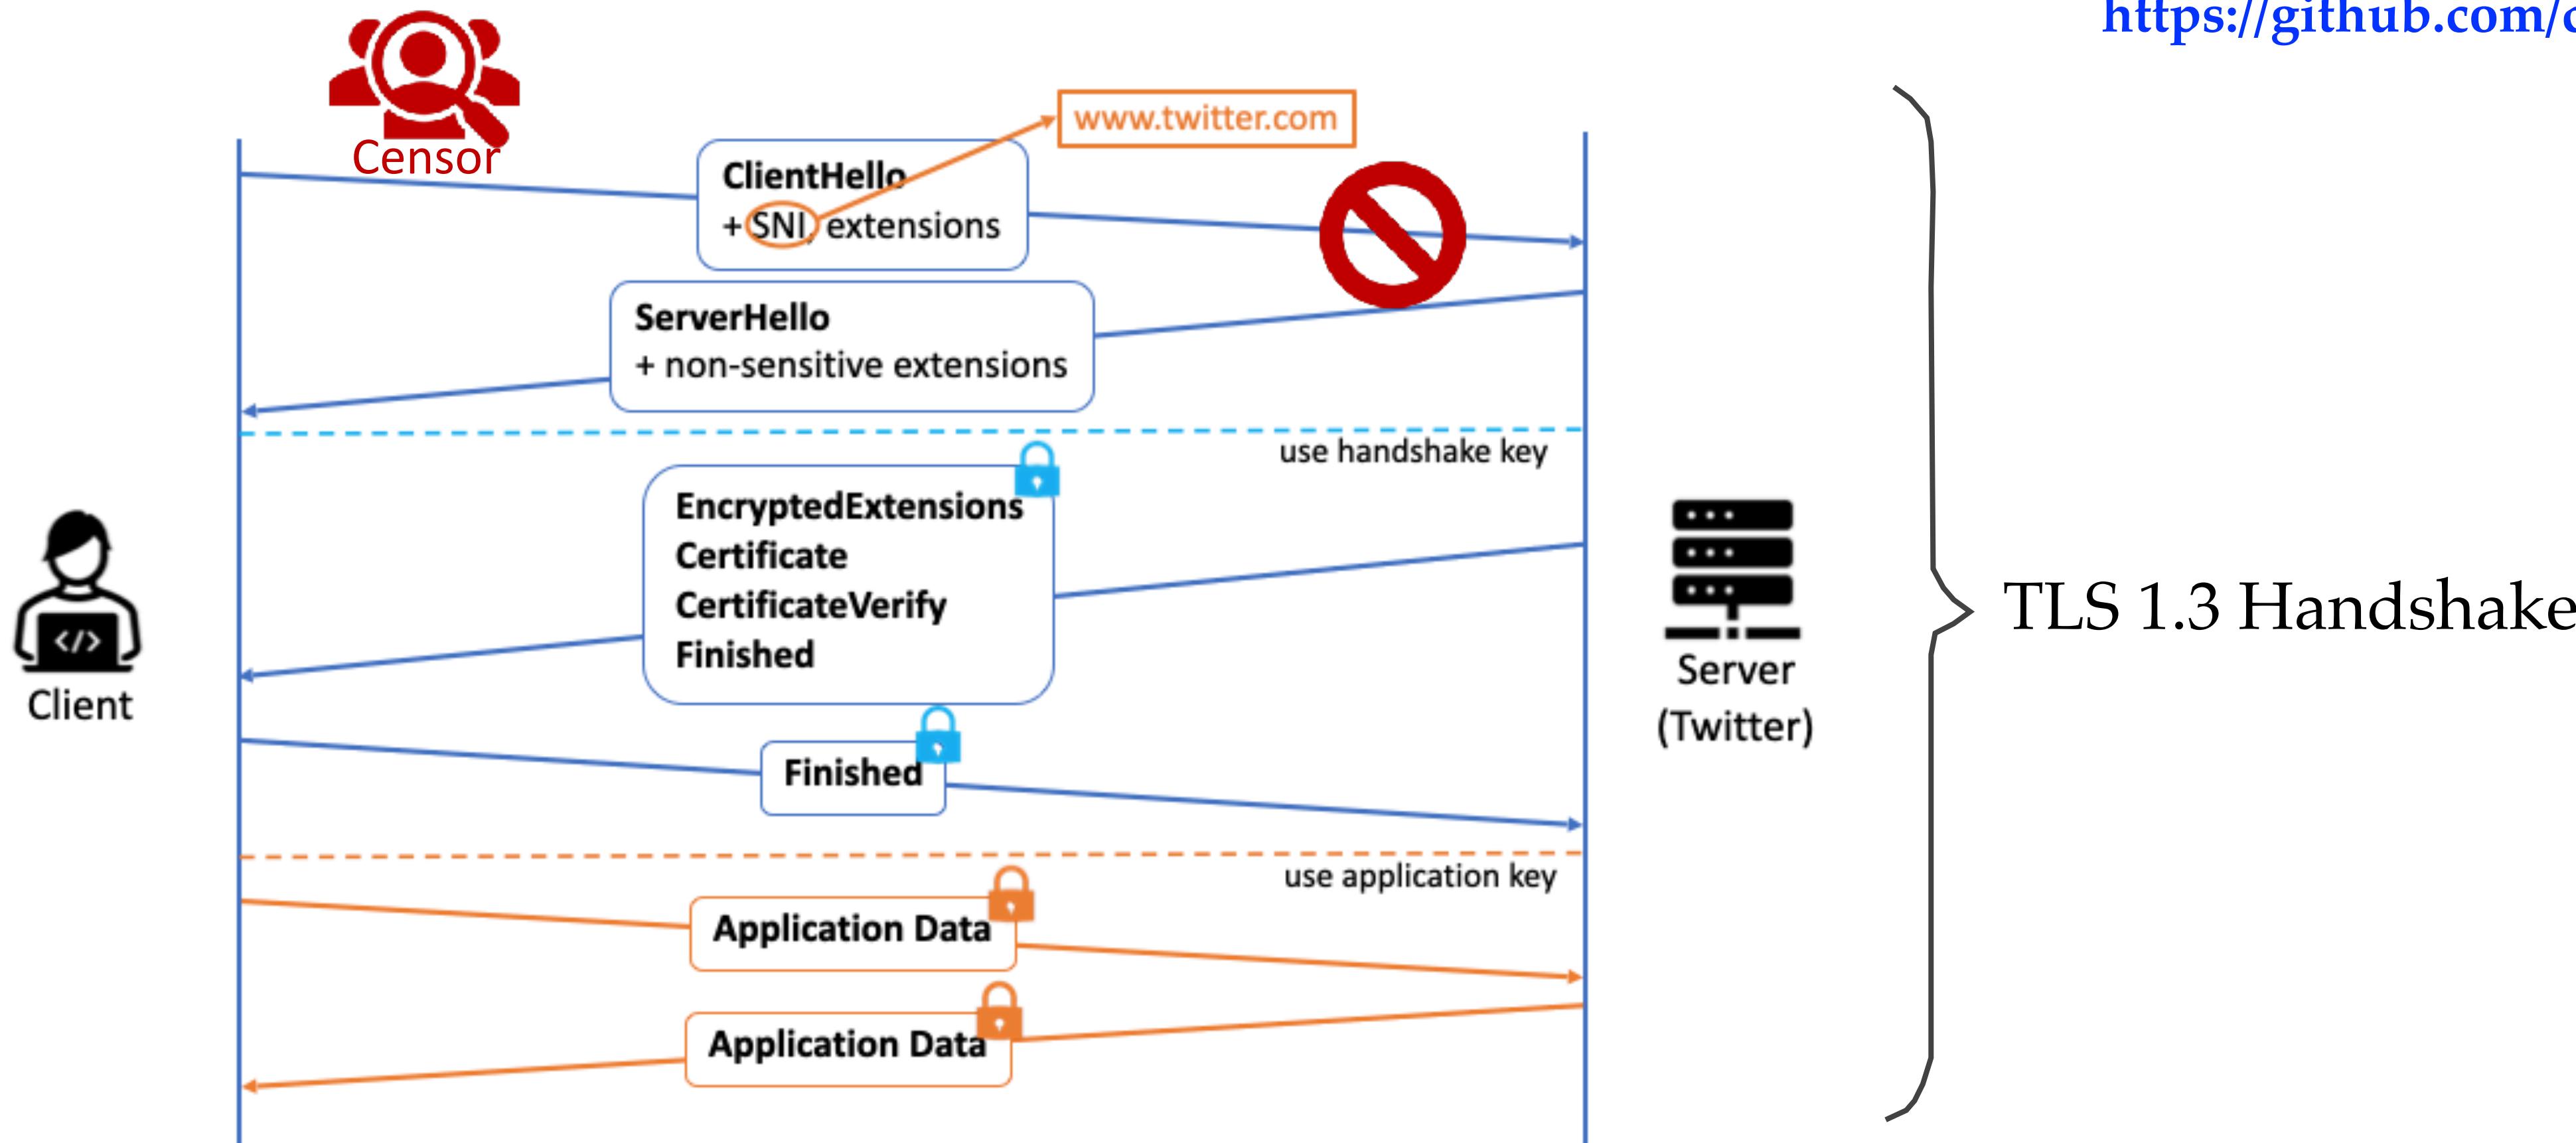
# Waldo

- A framework to verify privacy guarantees of communication protocols
- Practical impact: verification of TLS ECH Extension



**Waldo**

<https://github.com/cuplv/Waldo>

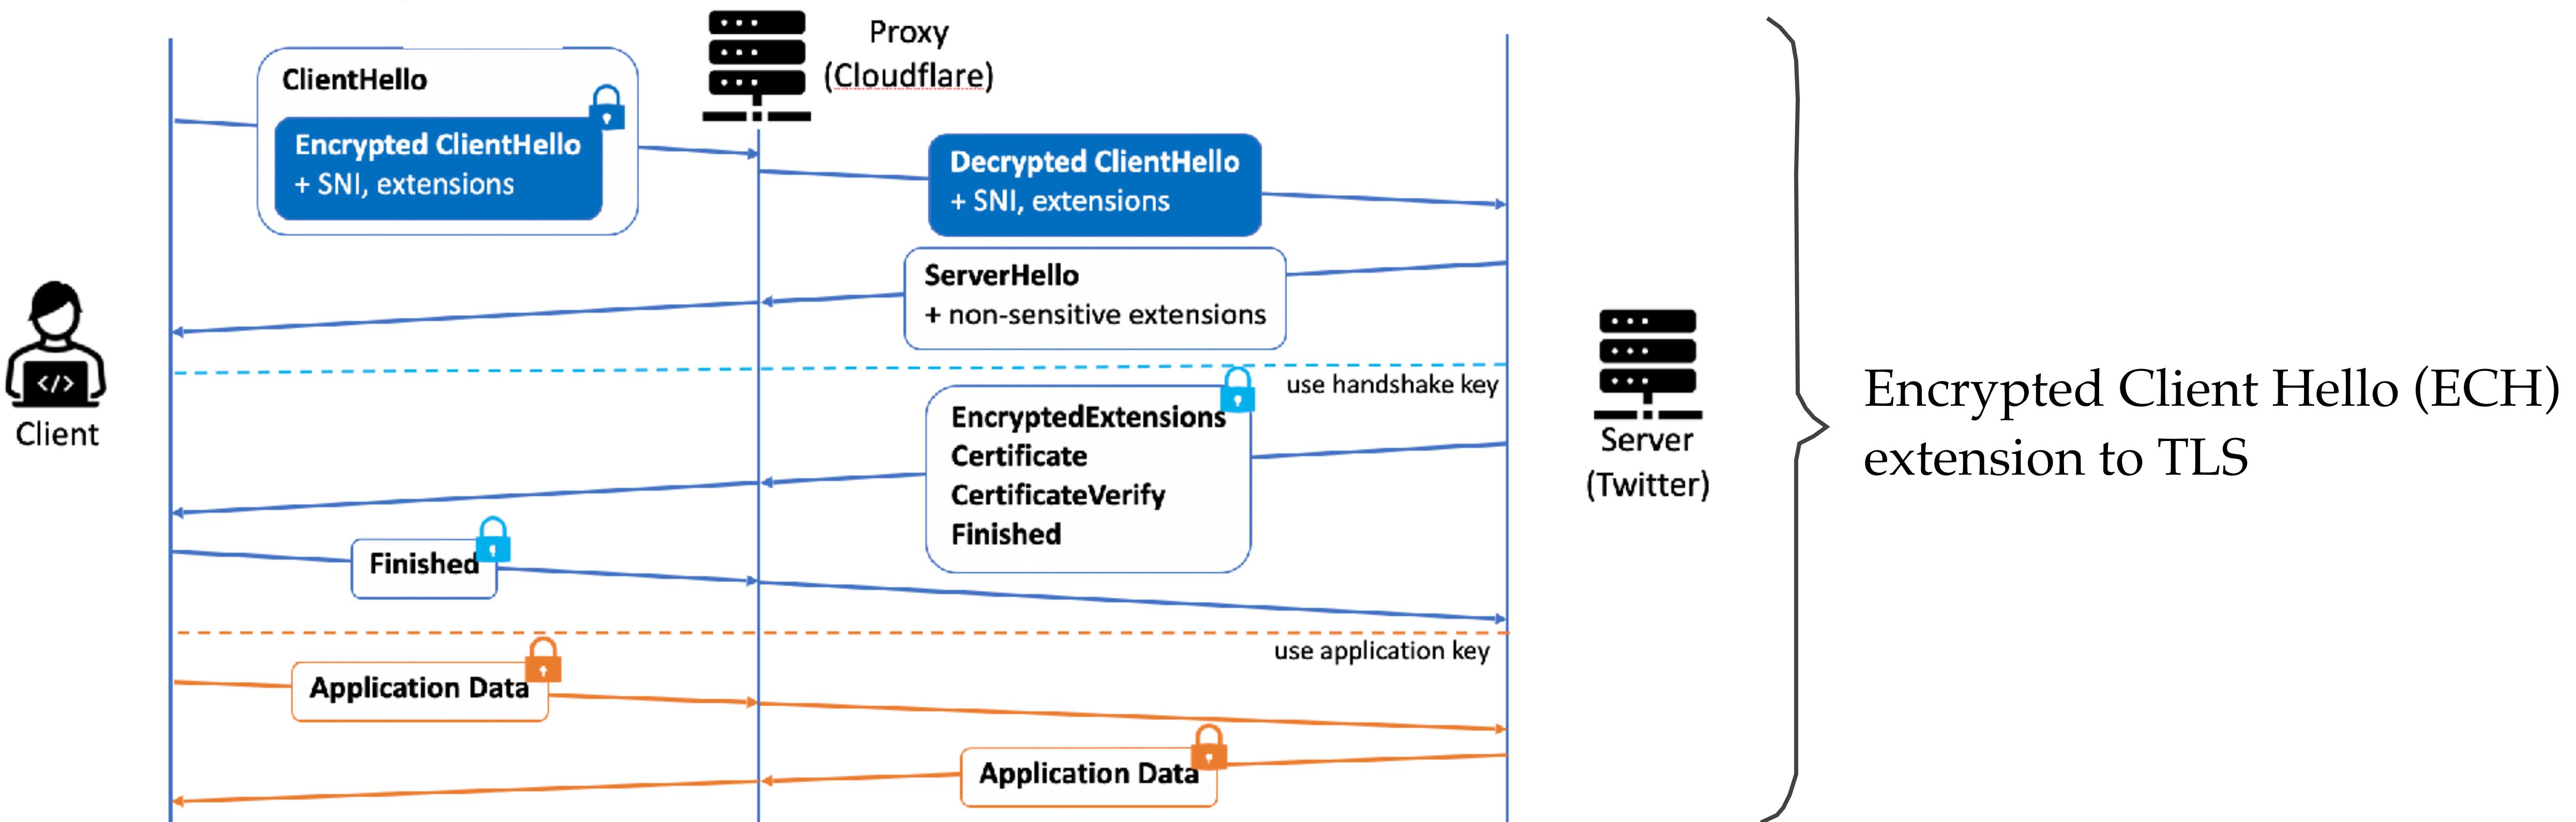


# Waldo

- A framework to verify privacy guarantees of communication protocols
- Practical impact: verification of TLS ECH Extension



<https://github.com/cuplv/Waldo>



# 3. Improving Usability: From *experts* to *programmers*

## Contributors



Nick Lewchenko



Christian Fontenot



Shawn Meier



Evan Chang

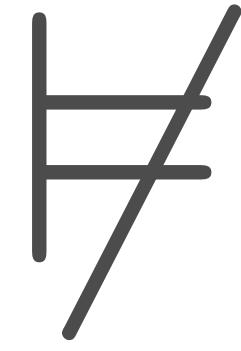
## Tools



Raven



Historia



DisProve

<https://github.com/cuplv/ravencheck>

<https://github.com/cuplv/Historia>

<https://github.com/cuplv/disprove>

## Publications

- Nicholas Lewchenko, Gowtham Kaki, and Bor-Yuh Evan Chang. **Effectively-Propositional Higher-Order Functional Programming.** Under editorial review for ACM on Programming Languages (PACMPL), issue **OOPSLA1 2026**.
- Shawn Meier, Sergio Mover, Gowtham Kaki, and Bor-Yuh Evan Chang. **Historia: Refuting Callback Reachability with Message-History Logics.** Proceedings of the ACM on Programming Languages, Volume 7, Issue **OOPSLA2 2023**. Article No.: 289, Pages 1905 - 193.
- Christian Fontenot, Gowtham Kaki, and Bor-Yuh Evan Chang. **Dissprove: Goal-Directed Verification of Parametrized Affine Actor Systems.** Under submission to POPL 2026.

# Artifacts of Formal Verification

- Specifications
- Invariants
- Proofs

# Artifacts of Formal Verification

- Specifications

More often than not:

- Intuitively known to programmers
- Can be recovered from the asserts

- Invariants

- Notoriously hard to discern.
- Quantified invariant inference is an open problem.

- Proofs

Encoding in a decidable logic guarantees proof automation.

Is it possible to circumvent the need for invariants?

- An (inductive) invariant characterizes every *reachable* state of the system.
- What if we reason about *unreachable* states instead?

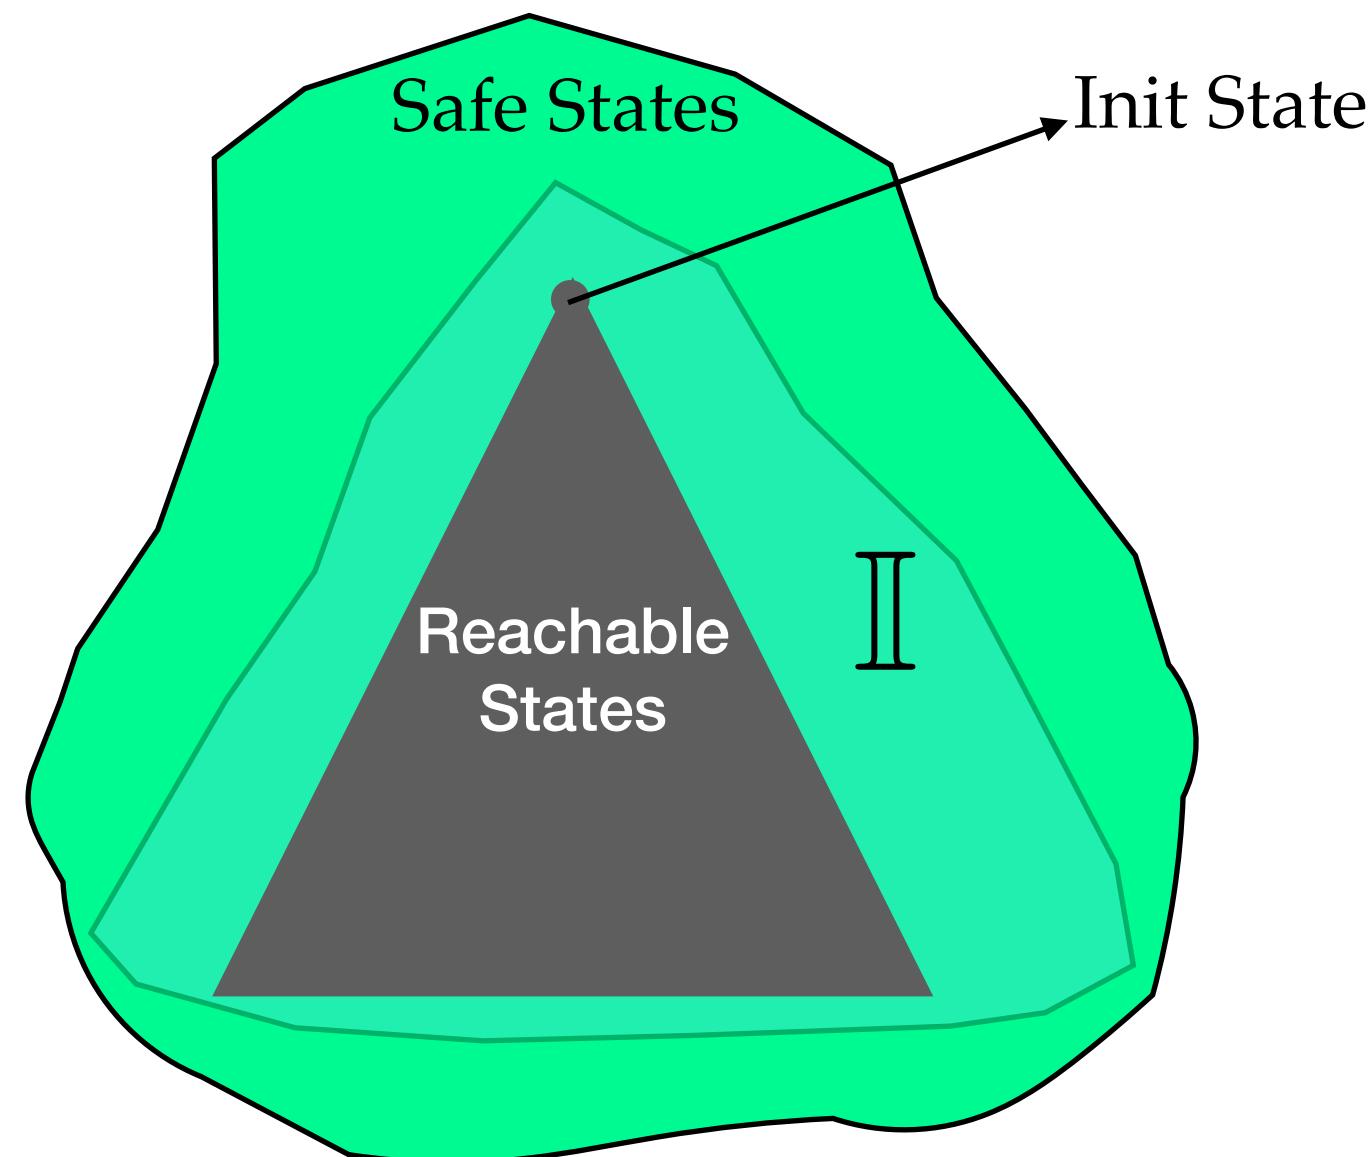
```
event eInitialize :  
    (Id: int, nodes: map[int, Node]);  
event eVote : int;  
event eLeader : int;  
  
machine Node {  
    var my_id: int;  
    var my_votes: int;  
    var my_leader: int;  
  
    on eInitialize do (int) {  
        ...  
        my_leader = -1;  
    }  
    ...  
    on eLeader do (id: int) {  
        assert(my_leader == -1 ||  
               my_leader == id);  
        my_leader = id;  
    }  
}
```

An implementation of Leader Election protocol in P language (Desai et al, PLDI'13)

# Goal-Directed Verification

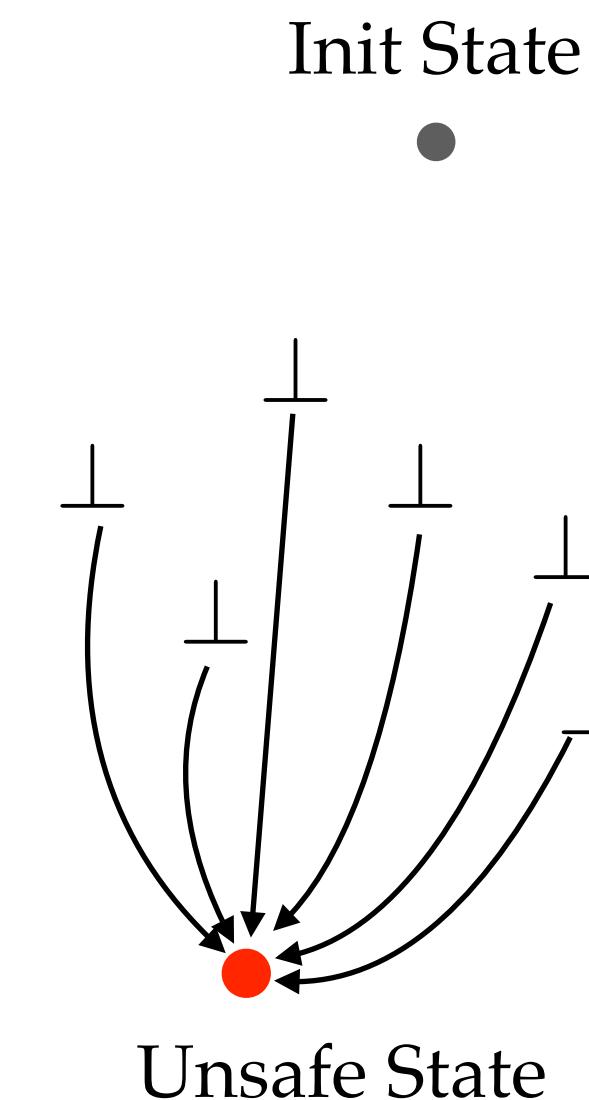
- An (inductive) invariant characterizes every *reachable* state of the system.
- What if we reason about *unreachable* states instead?

Regular (Forwards) Verification



All states reachable from the initial state are safe

Goal-Directed Verification



Unsafe state is unreachable from the init state

# Goal-Directed Verification in Practice

**Historia: Refuting Callback Reachability with Message-History Logics**

SHAWN MEIER, University of Colorado Boulder, USA

SERGIO MOVER, LIX, École Polytechnique, CNRS, Institut Polytechnique de Paris, France

GOWTHAM KAKI, University of Colorado Boulder, USA

BOR-YUH EVAN CHANG\*, University of Colorado Boulder & Amazon, USA

[OOPSLA 2023]

<https://github.com/cuplv/Historia>

- A goal-directed formal verification engine for Android apps.
- Proved the absence of a specific class of callback ordering-driven bugs in 47 open-source apps with 2M lines of code *without* requiring explicit inductive invariants.



Historia

**Dissprove: Goal-Directed Verification of Parametrized Affine Actor Systems**

ANONYMOUS AUTHOR(S)

[\[Under submission to POPL 2026\]](#)



DisProve

<https://github.com/cuplv/disprove>

- Adapts goal-directed formal verification to distributed protocols.
- Proved the safety of leader election protocol with *unbounded* number of participants *without* needing inductive invariants or update preconditions!
- Guaranteed to terminate on distributed algorithms composed of affine actors.

Key bottleneck: computational complexity!

Focus of future work.

# Conclusion

- Programming distributed systems is hard due to the sheer combinatorial complexity of many things that can go wrong.
- Formal verification can help ensure fail-safe operation and resilience to adversarial attacks.
- There exist several impediments to widespread adoption of formal methods in distributed systems development. Three main ones are poor scalability, insufficient generality, and steep learning curve.
- My research addresses these impediments by developing novel programming abstractions and automated reasoning techniques. The goal is to make provably-safe distributed systems development accessible to mainstream application developers.
- Research results have been published in open-access journals and implemented in several open-source tools.

## Acknowledgements



**Computer Science**  
COLLEGE OF ENGINEERING AND APPLIED SCIENCE

**Research & Innovation Office**

**Engineering Education and AI-Augmented Learning IRT**  
COLLEGE OF ENGINEERING AND APPLIED SCIENCE

**Google**